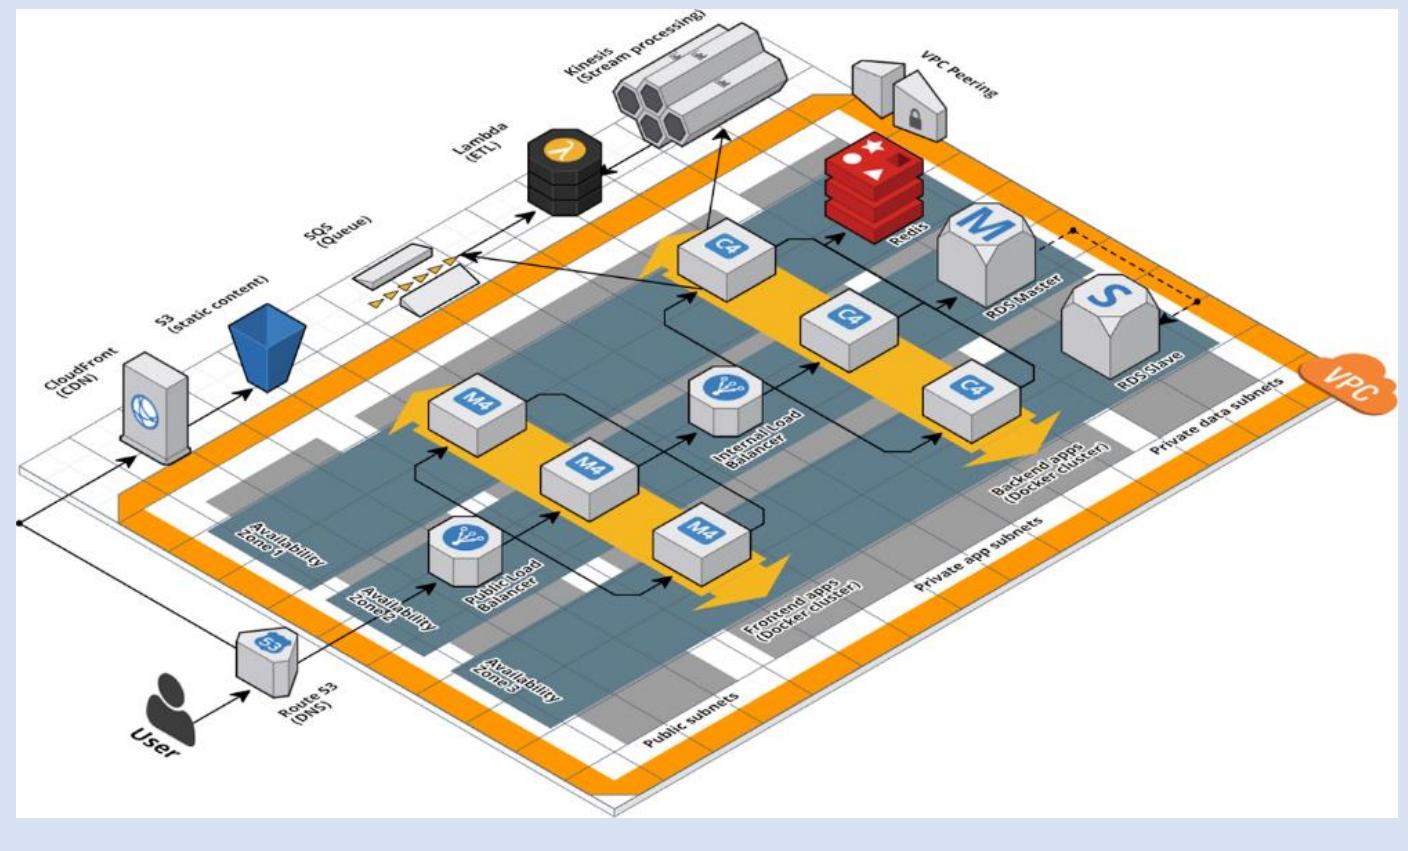


The goal of this document is to present the details related with Terraform. How to use it using a step-by-step approach. Terraform is a solution to create cloud environments using code instead of using the cloud providers User Interfaces. Therefore, the management of the cloud resources takes less effort and the creation, or update, process is faster.

The following figure taken the main literature reference in this field: *Terraform Up & Running, Writing Infrastructure as Code, Yevgeniy Brikman, 3<sup>rd</sup> edition (2022)* exemplifies a complex infrastructure that could be created using Terraform.



## Contents

A.	Terraform installation procedure .....	3
B.	Obtaining the access for sandbox or AWS academy learner lab.....	4
C.	Deploying a single EC2 instance .....	6
D.	Deploying a single EC2 instance and uploading files .....	7
E.	Sending commands during boot of a single EC2 instance.....	9
F.	Change the listener of a Kafka broker deployed with terraform .....	11
G.	Adding more Kafka brokers to cluster, and configuring each one individually.....	13
H.	How to show the dependencies in terraform? .....	16

I.	Sharing terraform state between resources using AWS S3.....	17
J.	Creating Lambda function and API Gateway trigger using S3 .....	22
K.	Kong and Konga deployed in AWS EC2 using Terraform.....	25
L.	Camunda Engine deployed in AWS EC2 using Terraform .....	28
M.	Quarkus project deployed in AWS EC2 using Terraform .....	30
N.	Ollama deployed in AWS EC2 using Terraform .....	32
	References.....	34

## A. Terraform installation procedure

Use the URL <https://developer.hashicorp.com/terraform/downloads> and choose the package appropriate for your operating system and then follow the indications.

For better Terraform code rendering, you can optionally, install the Terraform extension for VSCode.

The screenshot shows the HashiCorp Terraform extension page in the VS Code Marketplace. The extension is version v2.25.2, developed by HashiCorp, with 2,407,473 installs and a 4.5-star rating from 165 reviews. It provides syntax highlighting and autocompletion for Terraform. The extension is currently enabled globally. Below the main card, there are tabs for Details, Feature Contributions, Changelog, and Runtime Status. A note says to read the Troubleshooting Guide for answers to common questions. The Features section lists various capabilities like IntelliSense, Syntax validation, and Terraform commands. The IntelliSense and Autocomplete section explains that it supports code completion, parameter info, quick info, and member lists. The More Info section provides publishing details: Published on 12/02/2018 at 08:31:09, Last released on 15/12/2022 at 18:25:23, and Identifier hashicorp.terraform. On the right side, there are categories for Programming Languages, Linters, and Formatters, as well as links to Extension Resources like Marketplace, Repository, License, and HashiCorp.

## B. Obtaining the access for sandbox or AWS academy learner lab

- B.1. Choose your learner lab, then start it.

The screenshot shows the AWS Academy Cloud Foundations – Sandbox environment overview. On the left, there's a sidebar with icons for Conta, Módulos, Fóruns, Notas, Cursos, Calendário, Caixa de entrada, Histórico, and Ajuda. The main content area has tabs for Details, AWS, Start Lab, End Lab, 2:35, Instructions, and Actions. Under Actions, 'Show' is selected. A terminal window on the right shows the command 'ddd\_v1\_v\_y10\_1052360@runweb71380:~\$'. The page title is 'ACFv2BR-34198 > Módulos > Sandbox > Sandbox Environment'. The content includes sections for 'AWS Academy Cloud Foundations – Sandbox' and 'Environment Overview', stating that the sandbox provides an environment for ad-hoc exploration of AWS services. It also mentions a region restriction to the 'us-east-1' Region.

- B.2. When the lab is started, select Details, and then Show:

The screenshot shows the 'Details' tab of the AWS CLI interface. The 'Show' button is highlighted. Other buttons include AWS, Start Lab, End Lab, 2:34, Instructions, and Actions. Below the buttons, a terminal window shows the command '2360@runweb71380:~\$'. The overall interface is similar to the one shown in the previous screenshot.

- B.3. A similar access configuration is available for you (they are both examples that could be available, the first one is secretKey/accessKey, and the second is secretKey/accessKey/token). You will have to adapt your terraform calls accordingly with the provided AWS access.

The screenshot shows the 'Cloud Access' configuration page. It includes sections for 'AWS CLI' (with a 'Show' button), 'Cloud Labs' (showing session time, start date, end date, and accumulated lab time), 'ips' (public and private IP addresses), 'SSH key' (with 'Show', 'Download PEM', and 'Download PPK' buttons), 'AWS SSO' (with a 'Download URL' button), and a table for 'SecretKey', 'BastionHost', 'Region', and 'AccessKey'. The 'SecretKey' row is redacted with a blue bar.

SecretKey	[REDACTED]
BastionHost	18.234.191.90
Region	us-east-1
AccessKey	AKIAQPVWLKEA366JUOFG

The screenshot shows the AWS Learner Lab interface. On the left, there's a sidebar with icons for Account, Courses, Calendar, Inbox, History, and Help. The main area has a top navigation bar with 'ALLv1-38078 > Modules > Learner Lab > Learner Lab' and a status bar showing 'Used \$0 of \$100', '03:17', and various lab control buttons.

**Cloud Access:**

```
[default]
aws_access_key_id=ASIA5
aws_secret_access_key=B8
aws_session_token=FwoGZx
steYrhYfMsolDYdxhnglRk8
EdyjdJrOBASusvbx2u+j26
0d16j8NuC01vdQxe6oJ7CJZ
K3NnZs+huswkh-fmPf3SP5w
aq1gZsopr3TOabf1UBR+Bukq
```

**Cloud Labs:**

Remaining session time: 03:17:28(198 minutes)  
Session started at: 2023-01-30T03:08:38-0800  
Session to end at: 2023-01-30T07:08:38-0800  
Accumulated lab time: 00:42:00 (42 minutes)

No running instance

SSH key [Show](#) [Download PEM](#) [Download PPK](#)  
AWS SSO [Download URL](#)

AWSAccountID	188216847906
Region	us-east-1

## C. Deploying a single EC2 instance

- C.1. Configure the AWS access inside the terraform script file. For that, create the following script, replacing the `access_key`, `secret_key`, `token`, `region`, the `ami`, and the `instance_type` accordingly with your needs. Save the file in text format with a `.tf` extension.

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region = "us-east-1"
  access_key  = "YOUR_ACCESS_KEY"
  secret_key   = "YOUR_SECRET_KEY"
  token        = "YOUR_TOKEN"
}

resource "aws_instance" "example" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.micro"

  tags = {
    Name = "terraform-example"
  }
}
```

- C.2. Deploy the code with the following commands. You should use only one `.tf` file in each directory. A local configuration is created and then reflected in the cloud environment:

```
terraform init
terraform apply
```

or sequentially:

```
terraform init && terraform apply
```

- C.3. You can verify the configuration of the current configuration using the following command:

```
terraform show
```

- C.4. Clean up when you're done:

```
terraform destroy
```

*Hint: to facilitate the terraform command you can also use the -auto-approve option*

## D. Deploying a single EC2 instance and uploading files

D.1. As previous, configure your AWS access keys as environment variables, and then, create the following .tf file, adapting the **region**, the **ami**, and the **instance\_type** accordingly with you needs.

Notice that **access\_key**, **secret\_key** and **token** could be directly written in the source code.

Also get available the kafka tgz file and test.pem on your local path.

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }

  provider "aws" {
    region      = "us-east-1"
    access_key  = "YOUR_ACCESS_KEY"
    secret_key  = "YOUR_SECRET_KEY"
    token       = "YOUR_TOKEN"
  }

  resource "aws_instance" "exampleFileTransfer" {
    ami           = "ami-045269a1f5c90a6a0"
    instance_type = "t2.micro"
    vpc_security_group_ids = [aws_security_group.instance.id]
    key_name     = "vokey"
    tags = {
      Name = "terraform-example"
    }

    connection {
      type     = "ssh"
      user    = "ec2-user"
      private_key = file("test.pem")
      host    = "${self.public_dns}"
    }

    provisioner "file" {
      source      = "kafka_2.13-3.9.0.tgz"
      destination = "/home/ec2-user/kafka_2.13-3.9.0.tgz"
    }
  }

  resource "aws_security_group" "instance" {
    name = var.security_group_name
    ingress {
      from_port   = 22
      to_port     = 22
      protocol    = "tcp"
      cidr_blocks = ["0.0.0.0/0"]
    }
  }

  variable "security_group_name" {
    description = "The name of the security group"
    type        = string
    default     = "terraform-example-instance"
  }
}
```

Where **key\_name** is obtained from the name of the keypairs available directly from the AWS console:

The screenshot shows the AWS EC2 Key Pairs page. At the top, there's a search bar and a 'Create key pair' button. Below that, a table lists one key pair: 'vokey' (rsa type, created on 2023/01/24 15:37 GMT+0), with a fingerprint listed as 'base0:aa:be:42:8f:54:ea:e0:70:99:6d:32:...'. The left sidebar has links for EC2 Dashboard, EC2 Global View, Events, and Tags.

The provisioner "file" tag can be repeated as many times as needed, e.g.:

```
provisioner "file" {
  source      = "kafka_2.13-3.6.2.tgz"
  destination = "/home/ec2-user/kafka_2.13-3.6.2.tgz"
}

provisioner "file" {
  source      = "/root/zookeeper/apache-zookeeper-3.8.4-bin.tar.gz"
  destination = "/home/ec2-user/apache-zookeeper-3.8.4-bin.tar.gz"
}
```

Take attention to the **security resource** that allows SSH connections to your instance. If any other ingress is needed, you can add more in the same .tf file.

Also, the connection provides the **connection** configuration to enable the file upload (it uses a scp mechanism).

Also take note that provisioner is last action taken.

D.2. After changing for your configuration. Try it. Deploy the code:

```
terraform init
terraform apply
```

or

```
terraform init && terraform apply
```

D.3. Check if the file is uploaded correctly using a SSH connection directly to the EC2 instance:

```
ssh -i test.pem ec2-user@YOUR_DNS_NAME
```

*To remember: Use the following command to set the permissions of your private key file so that only you can read it.*

```
chmod u=rwx,g=,o= myKeyAWS.pem
```

D.4. Clean up when you're done:

```
terraform destroy
```

## E. Sending commands during boot of a single EC2 instance

User\_data is a shell script or cloud-init directive that will be executed during instance creation only.

The result of the `user_data` execution can be seen in EC2 Console (select the Instance, click Actions -> Monitor and troubleshoot -> Get system log) and you can find its execution log on the EC2 Instance itself (typically in `/var/log/cloud-init*.log`), both of which are useful for debugging, and neither of which is available with provisioners.

E.1. As previous, configure your AWS access keys as environment variables, and then, create the following .tf file, adapting the `region`, the `ami`, and the `instance_type` accordingly with you needs.

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access_key  = "YOUR_ACCESS_KEY"
  secret_key  = "YOUR_SECRET_KEY"
  token       = "YOUR_TOKEN"
}

resource "aws_instance" "exampleKAFKA" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.small"
  vpc_security_group_ids = [aws_security_group.instance.id]
  key_name     = "vokey"

  user_data = "${file("creation.sh")}"

  user data replace on change = true

  tags = {
    Name = "terraform-example-kafka"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port   = 2181
    to_port     = 2181
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port   = 9092
    to_port     = 9092
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port     = 0
    to_port       = 0
    protocol     = "-1"
    cidr_blocks  = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }
}
```

```

variable "security group name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-example-instance"
}

```

Where the script `user_data = "${file("creation.sh")}"` is a local file that could contain some work that you need to perform.

For instance, installing and starting ZooKeeper and Kafka, automatically!

```

#!/bin/bash
echo "Starting..."
cd
sudo wget https://dlcdn.apache.org/zookeeper/zookeeper-3.9.3/apache-zookeeper-3.9.3-bin.tar.gz
sudo tar -zxf apache-zookeeper-3.9.3-bin.tar.gz
sudo mv apache-zookeeper-3.9.3-bin /usr/local/zookeeper
sudo mkdir -p /var/lib/zookeeper
echo "tickTime=2000
dataDir=/var/lib/zookeeper
clientPort=2181" > /usr/local/zookeeper/conf/zoo.cfg

sudo yum -y install java-17-amazon-corretto-devel.x86_64

sudo /usr/local/zookeeper/bin/zkServer.sh start

sudo wget https://downloads.apache.org/kafka/3.9.0/kafka_2.13-3.9.0.tgz
sudo tar -zxf kafka_2.13-3.9.0.tgz
sudo mv kafka_2.13-3.9.0 /usr/local/kafka
sudo mkdir /tmp/kafka-logs

sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon /usr/local/kafka/config/server.properties

echo "Finished."

```

The **egress** tag allows the EC2 instance to outbound to another machine in internet.

Recall that a provisioner is the last action taken. That is why the binary files for Zookeeper and Kafka are being downloaded instead of being received by a provisioner.

If needed, the logging of the EC2 instance provisioning can be consulted at `/var/log/cloud-init-output.log`

## F. Change the listener of a Kafka broker deployed with terraform

F.1. Use the previous terraform file for deploying kafka broker in AWS.

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access_key  = "YOUR_ACCESS_KEY"
  secret_key  = "YOUR_SECRET_KEY"
  token       = "YOUR_TOKEN"
}

resource "aws_instance" "exampleKAFKAConfigured" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.small"
  vpc_security_group_ids = [aws_security_group.instance.id]
  key_name     = "vockey"

  user_data = "${file("creation.sh")}"

  user_data_replace_on_change = true

  tags = {
    Name = "terraform-example-kafka2"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port   = 2181
    to_port     = 2181
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port   = 9092
    to_port     = 9092
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port     = 0
    to_port       = 0
    protocol      = "-1"
    cidr_blocks   = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-example-instance"
}
```

F.2. Then, the script for provisioning the kafka broker need to be adapted with the new configuration for the listener at file /usr/local/kafka/config/server.properties. For that the name of the EC2 instance need to be known. The instance metadata service available at <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-metadata.html> could be used for that purpose.

```
#!/bin/bash

echo "Starting..."
cd
sudo wget https://dlcdn.apache.org/zookeeper/zookeeper-3.9.3/apache-zookeeper-3.9.3-bin.tar.gz
sudo tar -zxf apache-zookeeper-3.9.3-bin.tar.gz
sudo mv apache-zookeeper-3.9.3-bin /usr/local/zookeeper
sudo mkdir -p /var/lib/zookeeper
echo "tickTime=2000
dataDir=/var/lib/zookeeper
clientPort=2181" > /usr/local/zookeeper/conf/zoo.cfg

sudo yum -y install java-17-amazon-corretto-devel.x86_64
sudo /usr/local/zookeeper/bin/zkServer.sh start

sudo wget https://downloads.apache.org/kafka/3.9.0/kafka_2.13-3.9.0.tgz
sudo tar -zxf kafka_2.13-3.9.0.tgz
sudo mv kafka_2.13-3.9.0 /usr/local/kafka
sudo mkdir /tmp/kafka-logs

ip=`curl http://169.254.169.254/latest/meta-data/public-hostname`
sudo sed -i "s/#listeners=PLAINTEXT:\:\/\/:9092/listeners=PLAINTEXT:\:\/\/$ip:9092/g" /usr/local/kafka/config/server.properties

# due to AWS network stablishment process, check if 60 seconds is enough for your situation
(sleep 120 && sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon /usr/local/kafka/config/server.properties )&
echo "Finished."
```

## G. Adding more Kafka brokers to cluster, and configuring each one individually

Solving this problem in small solutions:

- **First solution:** simply add the count element to your aws instance. When you execute the command "terraform apply", those correspondingly number of similar instances will be created. It's a fast and easy way of provisioning multiple EC2 instances. However, all the internal detailed configurations of the installed software need to be done manually per EC2 instance: which at scale could be unfeasible.

```
resource "aws_instance" "exampleCLUSTER" {
  ami                  = "ami-045269a1f5c90a6a0"
  instance_type        = "t2.small"

  count = 4

  vpc_security_group_ids  = [aws_security_group.instance.id]
  key_name                = "vockekey"

  user_data = "${file("creation.sh")}"

  user_data_replace_on_change = true

  tags = {
    Name = "terraform-example-kafka"
  }
}
```

- **Second solution:** being able to provision multiple EC2 instances, and at the same time configure all the internal detailed configurations of each one of the EC2 instances. It has the advantage of accelerating the manual configurations by automation, and at scale. The drawback is the complexity and the lack of global configuration, i.e., configuration related with all the EC2 instances for all the EC2 instances.

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access_key  = "YOUR_ACCESS_KEY"
  secret_key  = "YOUR_SECRET_KEY"
  token       = "YOUR_TOKEN"
}

variable "nBroker" {
  description = "number of brokers"
  type        = number
  default     = 3
}

resource "aws_instance" "exampleCluster" {
  ami                  = "ami-045269a1f5c90a6a0"
  instance_type        = "t2.small"
  count               = var.nBroker
  vpc_security_group_ids  = [aws_security_group.instance.id]
  key_name             = "vockekey"

  user_data      = base64encode(templatefile("creation.sh", {
    idBroker = "${count.index}"
    totalBrokers = var.nBroker
  }))

  user_data_replace_on_change = true

  tags = {
```

```

        Name = "terraform-example-kafka.${count.index}"
    }

}

output "publicdnslist" {
    value = "${formatlist("%v", aws_instance.exampleCluster.*.public_dns)}"
}

resource "aws security group" "instance" {
    name = var.security_group_name
    ingress {
        from_port    = 22
        to_port      = 22
        protocol     = "tcp"
        cidr_blocks = ["0.0.0.0/0"]
    }
    ingress {
        from_port    = 2181
        to_port      = 2181
        protocol     = "tcp"
        cidr_blocks = ["0.0.0.0/0"]
    }
    ingress {
        from_port    = 2888
        to_port      = 2888
        protocol     = "tcp"
        cidr_blocks = ["0.0.0.0/0"]
    }
    ingress {
        from_port    = 3888
        to_port      = 3888
        protocol     = "tcp"
        cidr_blocks = ["0.0.0.0/0"]
    }
    ingress {
        from_port    = 9092
        to_port      = 9092
        protocol     = "tcp"
        cidr_blocks = ["0.0.0.0/0"]
    }
    egress {
        from_port      = 0
        to_port        = 0
        protocol       = "-1"
        cidr_blocks   = ["0.0.0.0/0"]
        ipv6_cidr_blocks = [":/:0"]
    }
}

variable "security_group_name" {
    description = "The name of the security group"
    type         = string
    default      = "terraform-example-instances"
}

```

```

#!/bin/bash

echo "Starting..."
cd
sudo wget https://dlcdn.apache.org/zookeeper/zookeeper-3.9.3/apache-zookeeper-3.9.3-bin.tar.gz
sudo tar -zxf apache-zookeeper-3.9.3-bin.tar.gz
sudo mv apache-zookeeper-3.9.3-bin /usr/local/zookeeper
sudo mkdir -p /var/lib/zookeeper
echo "tickTime=2000
dataDir=/var/lib/zookeeper
clientPort=2181
maxClientCnxns=60
initLimit=10
syncLimit=5" > /usr/local/zookeeper/conf/zoo.cfg

sudo yum -y install java-17-amazon-corretto-devel.x86_64
echo ${idBroker} > /var/lib/zookeeper/myid
sudo /usr/local/zookeeper/bin/zkServer.sh start
sudo wget https://downloads.apache.org/kafka/3.9.0/kafka_2.13-3.9.0.tgz
sudo tar -zxf kafka_2.13-3.9.0.tgz

```

```

sudo mv kafka_2.13-3.9.0 /usr/local/kafka
sudo mkdir /tmp/kafka-logs

ip=`curl http://169.254.169.254/latest/meta-data/public-hostname`
sudo sed -i "s/#listeners=PLAINTEXT:\:\/\/:9092/listeners=PLAINTEXT:\:\/\/$ip:9092/g"
/usr/local/kafka/config/server.properties

sudo sed -i "s/broker.id=0/broker.id=${idBroker}/g" /usr/local/kafka/config/server.properties

sudo sed -i "s/offsets.topic.replication.factor=1/offsets.topic.replication.factor=${totalBrokers}/g"
/usr/local/kafka/config/server.properties
sudo sed -i
"s/transaction.state.log.replication.factor=1/transaction.state.log.replication.factor=${totalBrokers}/g"
/usr/local/kafka/config/server.properties
sudo sed -i "s/transaction.state.log.min_isr=1/transaction.state.log.min_isr=${totalBrokers}/g"
/usr/local/kafka/config/server.properties

# due to AWS network establishment process, check if 30 seconds is enough for your situation
(sleep 30 && sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon
/usr/local/kafka/config/server.properties )&

echo "Finished."

```

As explained in tutorial P4 - Distributed Kafka service, two global configurations are still missing in all the EC2 instances. Login in each machine and setup the following configuration:

At /usr/local/kafka/config/server.properties

```

#zookeeper connectivity (one per EC2 VM of this cluster)
zookeeper.connect=ec2-54-90-57-82.compute-1.amazonaws.com:2181,ec2-54-173-171-63.compute-
1.amazonaws.com:2181,ec2-54-236-47-54.compute-1.amazonaws.com:2181

```

And at /usr/local/zookeeper/conf/zoo.cfg

```

server.1=ec2-54-90-57-82.compute-1.amazonaws.com:2888:3888
server.2=ec2-54-173-171-63.compute-1.amazonaws.com:2888:3888
server.3=ec2-54-236-47-54.compute-1.amazonaws.com:2888:3888

```

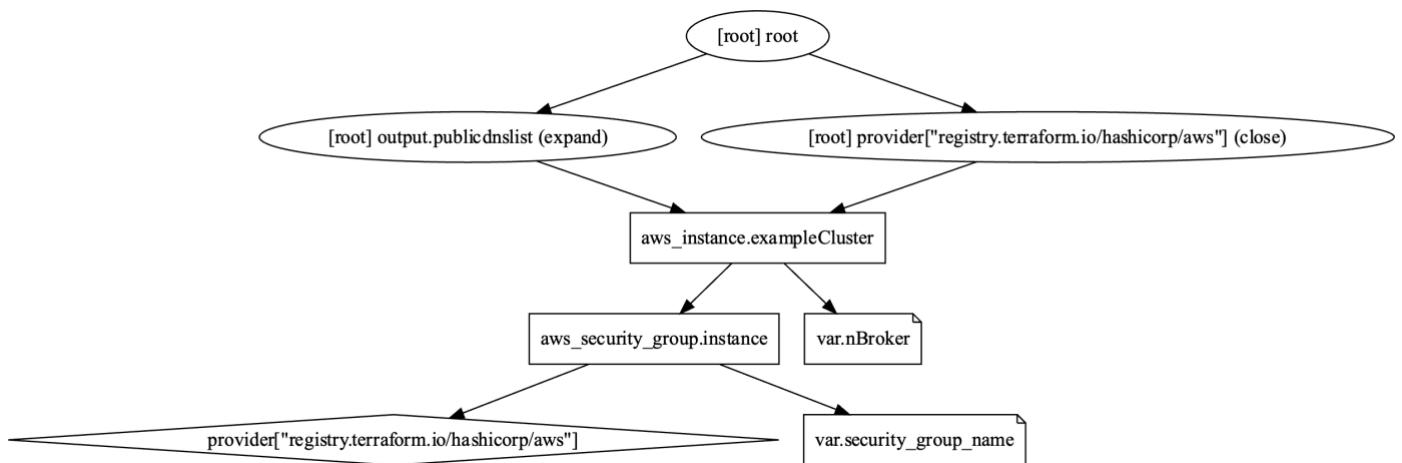
## H. How to show the dependencies in terraform?

To show the dependencies of your deployment environment use the following command:

```
terraform graph -draw-cycles
```

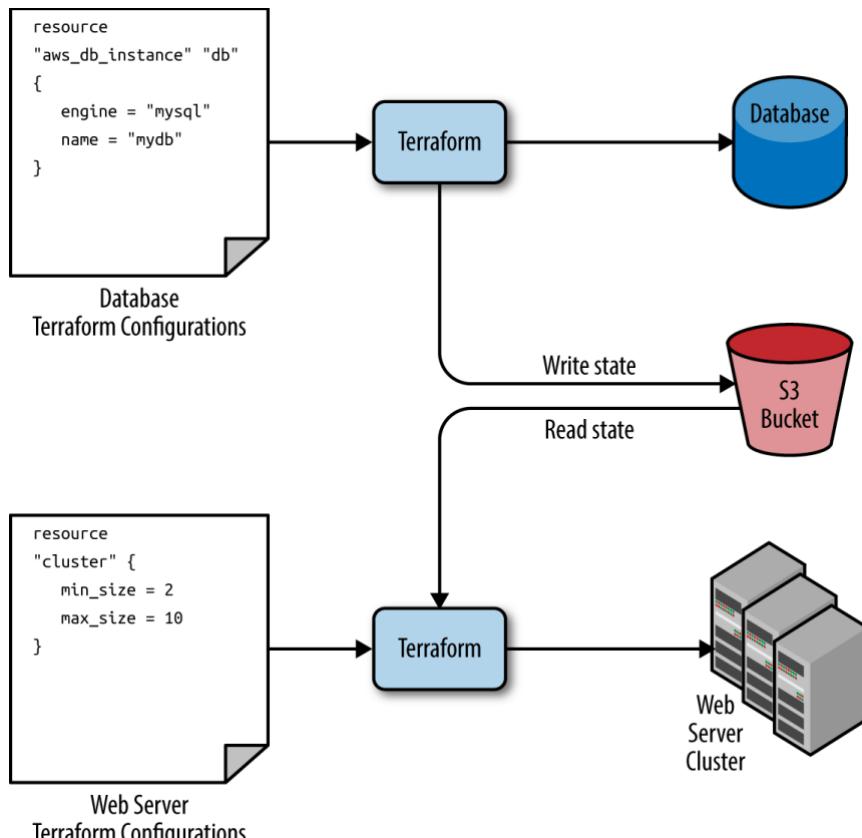
Then copy+paste the output for the graphviz, e.g., on <https://dreampuf.github.io/GraphvizOnline>.

Here is an example of you obtain:



## I. Sharing terraform state between resources using AWS S3

The terraform state sharing allows the provisioning of an environment where the posterior resources depend on the formers, as depicted in the next figure. It could be done using a cloud storage or a local storage. S3 is an AWS service to store data.



Source: Brikman, Y. (2022). *Terraform: Up and Running*. "O'Reilly Media, Inc."

Notice that to write the state in the storing an output should be defined in the terraform script file, i.e., to store locally the state of terraform:

```
output "publicdnslist" {
  value = "${formatlist("%v", aws_instance.exampleCluster.*.public_dns)}"
}
```

Then, you can list all the outputs, e.g.:

```
terraform state list
```

returning:

```
aws_instance.exampleCluster[0]
aws_instance.exampleCluster[1]
aws_instance.exampleCluster[2]
aws_security_group.instance
```

Or just show one of them:

```
terraform state show 'aws_instance.exampleCluster[0]'
```

In Terraform, the state is a representation of the configuration of your infrastructure that is stored locally or remotely. Accessing the state locally involves using the terraform state command.

I.1. Choose the AWS S3 service, then, click on “create bucket”:

The screenshot shows the AWS S3 Buckets page. At the top, a green banner indicates "Successfully created bucket 'terraform-s3-2025-01-15'". Below this, there's an "Account snapshot" section with a link to "View details". Under "General purpose buckets", there is one entry: "General purpose buckets (1) All AWS Regions". A "Create bucket" button is visible at the top right of this section.

Give it a unique name:

The screenshot shows the "Create bucket" wizard. In the "General configuration" step, the "Bucket type" is set to "General purpose". The "Bucket name" is "terraform-s3-2025-01-15". In the "Object Ownership" step, "ACLs disabled (recommended)" is selected. In the "Block Public Access settings for this bucket" step, "Block all public access" is checked. In the "Bucket Versioning" step, "Disable" is selected. The bottom of the screen shows standard AWS navigation links like CloudShell, Feedback, and a footer with copyright information.

I.2. Terraform state will be stored in the S3 cloud environment. For that, firstly, create in a separate directory, a terraform file to create an AWS RDS (*e.g.*: mySQL cloud database) and store remotely the name of the DB.

```
terraform {
  backend "s3" {
    bucket = "terraform-s3-2025-01-15"
```

```

key= "stage/data-stores/mysql/terraform.tfstate"
region = "us-east-1"
access_key  = "YOUR_ACCESS_KEY"
secret_key  = "YOUR_SECRET_KEY"
token       = "YOUR_TOKEN"
}
required_version = ">= 1.0.0, < 2.0.0"

required providers {
  aws = {
    source  = "hashicorp/aws"
    version = "~> 4.0"
  }
}

provider "aws" {
  region = "us-east-1"
  access_key  = "YOUR_ACCESS_KEY"
  secret_key  = "YOUR_SECRET_KEY"
  token       = "YOUR_TOKEN"
}

variable "db_username" {
  description = "The username for the database"
  type        = string
  sensitive   = true
  default     = "teste"
}

variable "db_password" {
  description = "The password for the database"
  type        = string
  sensitive   = true
  default     = "testeteste"
}

variable "db_name" {
  description = "The name to use for the database"
  type        = string
  default     = "example_database_stage"
}

resource "aws_db_instance" "example" {
  identifier_prefix  = "terraform-up-and-running"
  engine            = "mysql"
  allocated_storage = 20
  instance_class    = "db.t4g.micro"
  skip_final_snapshot = true
  publicly_accessible = true
  vpc_security_group_ids = [aws_security_group.rds.id]
  db_name           = var.db_name

  username = var.db_username
  password = var.db_password
}

output "address" {
  value      = aws_db_instance.example.address
  description = "Connect to the database at this endpoint"
}

output "port" {
  value      = aws_db_instance.example.port
  description = "The port the database is listening on"
}

resource "aws_security_group" "rds" {
  name = var.security_group_name
  ingress {
    from_port    = 3306
    to_port      = 3306
    protocol     = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
  egress {
    from_port      = 0
    to_port        = 0
    protocol       = "-1"
    cidr_blocks   = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }
}

```

```

}

variable "security group name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-rds-instance"
}

```

Execute the usual command:

```
terraform init && terraform apply
```

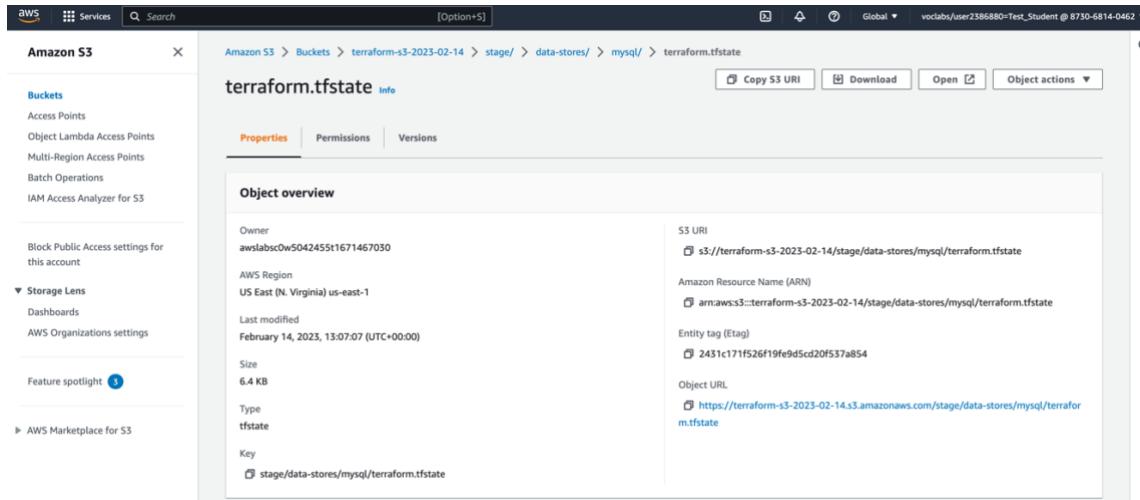
You can now check the outputs were written correctly to terraform state using the following commands:

```
terraform state list
```

```
terraform state show 'aws_db_instance.example'
```

```
terraform output
```

Moreover, if you navigate until:



You'll see that the terrafrom state of the recently created mysql database is stored there.

- I.3. Now, in another directory, create a terraform file that can create an AWS EC2 instance writing the previous database address and port to a text file inside the EC2. The address and port are read from the state file of the previous step. With this mechanism you can create resources that are configured with parameters from the previous created resources. Facilitating the automation provisioning process!

```

terraform {
  required version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }

  data "terraform_remote_state" "db" {
    backend = "s3"
    config = {
      bucket = "terraform-s3-2025-01-15"
    }
  }
}

```

```

key= "stage/data-stores/mysql/terraform.tfstate"
region = "us-east-1"
access_key  = "YOUR_ACCESS_KEY"
secret_key  = "YOUR_SECRET_KEY"
token       = "YOUR_TOKEN"
}

provider "aws" {
  region = "us-east-1"
  access_key  = "YOUR_ACCESS_KEY"
  secret_key  = "YOUR_SECRET_KEY"
  token       = "YOUR_TOKEN"
}

resource "aws instance" "exampleCluster" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.small"
  vpc_security_group_ids = [aws security group.instance.id]
  key_name      = "vockey"

  user_data     = base64encode(templatefile("creation.sh", {
    address = data.terraform_remote_state.db.outputs.address
    port   = data.terraform_remote_state.db.outputs.port
  }))
}

user_data_replace_on_change = true

tags = {
  Name = "terraform-example-sharestate"
}
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port  = 22
    to_port    = 22
    protocol   = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port     = 0
    to_port       = 0
    protocol     = "-1"
    cidr_blocks  = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-example-instance2"
}

```

And creation.sh with the following source code:

```

#!/bin/bash
cd

echo ${address} > address.txt
echo ${port} > port.txt

```

Execute the usual command:

```
terraform init && terraform apply
```

Login in the newly created EC2 instance and verify in the root directory that two new files are now created with the content provided by DB configuration.

## J. Creating Lambda function and API Gateway trigger using S3

J.1. To create a lambda function, written in JAVA 8, that you previously compiled locally, you need to transfer the .jar file to S3, then create the AWS lambda function resource, and finally, add an API Gateway trigger for that lambda function.

```
terraform {
  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.15.0"
    }
    random = {
      source  = "hashicorp/random"
      version = "~> 3.1.0"
    }
    archive = {
      source  = "hashicorp/archive"
      version = "~> 2.2.0"
    }
  }
  required_version = "~> 1.2"
}

provider "aws" {
  region = "us-east-1"
  access_key  =
  secret_key  =
  token =
}

## S3 PART
resource "aws_s3_bucket" "lambda_bucket" {
  bucket = "terraform-lambda-2025-01-16"
}

resource "aws_s3_object" "lambda_hello_world" {
  bucket = aws_s3_bucket.lambda_bucket.id
  key    = "lambda-java-example-0.0.1-SNAPSHOT.jar"
  source = "${path.module}/lambda-java-example-0.0.1-SNAPSHOT.jar"
}

## LAMBDA FUNCTION PART
resource "aws_lambda_function" "hello_world" {
  function_name = "HelloWorld"
  s3_bucket    = aws_s3_bucket.lambda_bucket.id
  s3_key       = aws_s3_object.lambda_hello_world.key
  runtime      = "java11"
  handler      = "example.Hello::handleRequest"
  role         = "arn:aws:iam::8730681408695462:role/LabRole"
}

resource "aws_cloudwatch_log_group" "hello_world" {
  name = "/aws/lambda/${aws_lambda_function.hello_world.function_name}"
  retention_in_days = 30
}

### API GATEWAY PART
resource "aws_api_gateway_rest_api" "example" {
  name      = "ServerlessExample"
  description = "Terraform Serverless Application Example"
}

resource "aws_api_gateway_resource" "proxy" {
  rest_api_id = "${aws_api_gateway_rest_api.example.id}"
  parent_id   = "${aws_api_gateway_rest_api.example.root_resource_id}"
  path_part   = "{proxy+}"
  path_part   = "helloworldpath"
}

resource "aws_api_gateway_method" "proxy" {
  rest_api_id  = "${aws_api_gateway_rest_api.example.id}"
  resource_id  = "${aws_api_gateway_resource.proxy.id}"
  http_method  = "ANY"
  authorization = "NONE"
}

resource "aws_api_gateway_integration" "lambda" {
  rest_api_id = "${aws_api_gateway_rest_api.example.id}"
}
```

```

resource_id = "${aws_api_gateway_method.proxy.resource_id}"
http_method = "${aws_api_gateway_method.proxy.http_method}"

integration http method = "POST"
type          = "AWS_PROXY"
uri           = "${aws_lambda_function.hello_world.invoke_arn}"
}

resource "aws api gateway deployment" "example" {
depends_on = [
  aws_api_gateway_integration.lambda,
]
rest_api_id = "${aws_api_gateway_rest_api.example.id}"
stage_name  = "test"
}

resource "aws_lambda_permission" "apigw" {
statement_id  = "AllowAPIGatewayInvoke"
action        = "lambda:InvokeFunction"
function_name = "${aws_lambda_function.hello_world.function_name}"
principal     = "apigateway.amazonaws.com"

# The /*/* portion grants access from any method on any resource
# within the API Gateway "REST API".
source_arn = "${aws_api_gateway_rest_api.example.execution_arn}/*/*"
}

# Output value definitions
output "lambda_bucket_name" {
description = "Name of the S3 bucket used to store function code."
value = aws_s3_bucket.lambda_bucket.id
}

output "function_name" {
description = "Name of the Lambda function."
value = aws_lambda_function.hello_world.function_name
}

output "base_url" {
value = "${aws_api_gateway_deployment.example.invoke_url}"
}

```

Where the field role of resource "aws\_lambda\_function":

```
role = "arn:aws:iam::8730681408695462:role/LabRole"
```

could be consulted on your AWS account, in the IAM service, or replacing your AWSAccountId number directly in the bold part of the field:

The screenshot shows two side-by-side interfaces. On the left is the AWS IAM 'LabRole' configuration page, which includes sections for Summary, Permissions, and CloudWatch Metrics. It lists several AWS managed policies attached to the role. On the right is the 'Cloud Access' interface, showing session details like '03:11', 'Used \$2.9 of \$100', and a 'Cloud Access' section containing AWS CLI commands.

J.2. To test your lambda function, externally from any machine, you need to connect to AWS API Gateway, as explained in the Lambda-AWS tutorial:

```
curl -i -H "Content-Type: Application/json" --data "@body.json" -X POST https://<yourURL>/test/helloworldpath
```

And the result obtained is the following:

```
HTTP/2 200
content-type: application/json
content-length: 43
date: Wed, 15 Feb 2023 16:36:51 GMT
x-amzn-requestid: f2027a23-803d-42d3-8537-e2fe256d7480
x-amz-apigw-id: AY6FiH8XoAMFbvA=
x-custom-header: my custom header value
x-amzn-trace-id: Root=1-63ed0a23-652880af6750972f7523767e; Sampled=0
x-cache: Miss from cloudfront
via: 1.1 cb4f40303e252a22c4df5918669814ac.cloudfront.net (CloudFront)
x-amz-cf-pop: LIS50-C1
x-amz-cf-id: CAQtv416vFuEzDsQv7tj56smEc7s4ZkcEffM3VobsGSKnB-1H2y4g==

{"message": "Hello Integracao Empresarial!"}
```

## K. Kong and Konga deployed in AWS EC2 using Terraform

K.1. To install Kong you will need to follow the indications as explained in the Kong tutorial. They are now adapted to be hyperautomated using Terraform.

Where the .tf file can include the following instructions:

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region     = "us-east-1"
  access_key = ""
  secret_key = ""
  token     = ""
}

resource "aws_instance" "exampleInstallKong" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.small"
  vpc_security_group_ids = [aws_security_group.instance.id]
  key_name      = "vockey"

  user_data = "${file("deploy.sh")}"

  user_data_replace_on_change = true

  tags = {
    Name = "terraform-example-Kong"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }

  egress {
    from_port      = 0
    to_port        = 0
    protocol       = "-1"
    cidr_blocks   = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-kong-instance"
}

output "address" {
  value      = aws_instance.exampleInstallKong.public_dns
  description = "Connect to the database at this endpoint"
}
```

And the correspondingly script with the following:

```
#!/bin/bash
echo "Starting..."
```

```

sudo yum install -y docker
sudo service docker start
sudo docker network create kong-net
sudo docker run -d --name kong-database \
--network=kong-net \
-p 5432:5432 \
-e "POSTGRES_USER=kong" \
-e "POSTGRES_DB=kong" \
-e "POSTGRES_PASSWORD=kongpass" \
postgres:13

sudo docker run --rm --network=kong-net \
-e "KONG_DATABASE=postgres" \
-e "KONG_PG_HOST=kong-database" \
-e "KONG_PG_PASSWORD=kongpass" \
-e "KONG_PASSWORD=test" \
kong/kong-gateway:3.9.0.0 kong migrations bootstrap

sudo docker run -d --name kong-gateway \
--network=kong-net \
-e "KONG_DATABASE=postgres" \
-e "KONG_PG_HOST=kong-database" \
-e "KONG_PG_USER=kong" \
-e "KONG_PG_PASSWORD=kongpass" \
-e "KONG_PROXY_ACCESS_LOG=/dev/stdout" \
-e "KONG_ADMIN_ACCESS_LOG=/dev/stdout" \
-e "KONG_PROXY_ERROR_LOG=/dev/stderr" \
-e "KONG_ADMIN_ERROR_LOG=/dev/stderr" \
-e "KONG_ADMIN_LISTEN=0.0.0.0:8001, 0.0.0.0:8444 ssl" \
-e "KONG_ADMIN_GUI_URL=http://localhost:8002" \
-e KONG_LICENSE_DATA \
-p 8000:8000 \
-p 8443:8443 \
-p 8001:8001 \
-p 8002:8002 \
-p 8445:8445 \
-p 8003:8003 \
-p 8004:8004 \
-p 127.0.0.1:8444:8444 \
kong/kong-gateway:3.9.0.0

echo "Finished."

```

K.2. To install the UI Konga again you can follow the indications as explained in the Kong tutorial. They are now adapted to be hyperautomated using Terraform.

Where the .tf file can include the following instructions:

```

terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access_key  = ""
  secret_key  = ""
  token       = ""
}

resource "aws_instance" "exampleInstallKonga" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.small"
  vpc_security_group_ids = [aws_security_group.instance.id]
  key_name     = "vockey"
}

```

```

user_data = "${file("deploy.sh")}"
user data replace on change = true

tags = {
  Name = "terraform-example-Konga"
}
}

resource "aws_security_group" "instance" {
  name = var.security_group_name
  ingress {
    from_port    = 0
    to_port      = 0
    protocol     = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }
  egress {
    from_port      = 0
    to_port        = 0
    protocol       = "-1"
    cidr_blocks   = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-konga-instance"
}

output "address" {
  value      = aws_instance.exampleInstallKonga.public dns
  description = "Connect to the database at this endpoint"
}

```

And the correspondingly script with the following:

```

#!/bin/bash
echo "Starting..."

sudo yum install -y docker
sudo service docker start
sudo docker pull pantsel/konga
sudo docker run -d --name konga -p 1337:1337 pantsel/konga
echo "Finished."

```

## L. Camunda Engine deployed in AWS EC2 using Terraform

L.1. To install Camunda Engine you will need to follow the indications as explained in the Camunda tutorial. They are now adapted to be hyperautomated using Terraform.

Where the .tf file can include the following instructions:

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region     = "us-east-1"
  access_key = ""
  secret_key = ""
  token     = ""
}

resource "aws_instance" "exampleInstallCamundaEngine" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.small"
  vpc_security_group_ids = [aws_security_group.instance.id]
  key_name      = "vockey"

  user_data = "${file("deploy.sh")}"

  user_data_replace_on_change = true

  tags = {
    Name = "terraform-example-Camunda"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }

  egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-Camunda-instance"
}

output "address" {
  value      = aws_instance.exampleInstallCamundaEngine.public_dns
  description = "Connect to the database at this endpoint"
}
```

And the correspondingly script with the following:

```
#!/bin/bash
echo "Starting..."
sudo yum update -y
sudo yum install -y docker
sudo service docker start
sudo docker pull camunda/camunda-bpm-platform:latest
sudo docker run -d --name camunda -p 8080:8080 camunda/camunda-bpm-platform:latest
echo "Finished."
```

## M. Quarkus project deployed in AWS EC2 using Terraform

M.1. To deploy a previous created and compiled Quarkus project you will need to have the image previously pushed to git, Then, use the following files for reference. Where the .tf file can include the following instructions:

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access_key  = "YOUR ACES KEY"
  secret_key  = "YOUR SECRET KEY"
  token       = "YOUR TOKEN"
}

resource "aws_instance" "exampleDeployQuarkus" {
  ami           = "ami-0e7290665643979b5" # Amazon Linux ARM AMI built by Amazon Web Services - FOR DOCKER image COMPATIBILITY if compiled previously on ARM
  instance_type = "t4g.nano"
  vpc_security_group_ids = [aws_security_group.instance.id]
  key_name     = "vockey"

  user_data = "${file("quarkus.sh")}"
  user_data_replace_on_change = true

  tags = {
    Name = "terraform-deploy-QuarkusProject"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name
  ingress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }
  egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-Quarkus-instance"
}

output "address" {
  value      = aws_instance.exampleDeployQuarkus.public_dns
  description = "Address of the Quarkus EC2 machine"
}
```

And the correspondingly script with the following:

```
#!/bin/bash
echo "Starting..."
sudo yum install -y docker
sudo service docker start
sudo docker login -u "YOUR DOCKER USERNAME" -p "YOUR DOCKER PASSWORD"
sudo docker pull YOUR DOCKER USERNAME/tryout1:1.0.0-SNAPSHOT
sudo docker run -d --name tryout2 -p 9000:9000 YOUR DOCKER USERNAME/tryout1:1.0.0-SNAPSHOT
echo "Finished."
```

Docker images built for ARM architecture may not work on AMD architecture machines. This is because ARM and AMD processors use different instruction sets and have different hardware architectures.

For this example, the AMI identifier and instance type were extracted from the launching instance AWS GUI.

The screenshot shows the AWS EC2 'Launch an instance' wizard. The 'Summary' section is highlighted with a red box. It includes the following details:

- Software Image (AMI)**: Amazon Linux 2 LTS Arm64 Kernel 5.10 AMI 2.0.20250108.0 arm64 HVM gp2
- Virtual server type (instance type)**: t4g.nano
- Firewall (security group)**: New security group
- Storage (volumes)**: 1 volume(s) - 8 GiB

The 'Free tier' information is also visible:

**Free tier:** In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million IOs, 1 GB of snapshots, and 100 GB of bandwidth to the internet.

Other sections shown include 'Name and tags', 'Application and OS Images (Amazon Machine Image)', 'Amazon Machine Image (AMI)', 'Architecture', 'Instance type', and 'Additional costs apply for AMIs with pre-installed software'.

## N. Ollama deployed in AWS EC2 using Terraform

N.1. To install Ollama in a AWS EC2 instance using Terraform follows the following the .tf file to create the needed resources:

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access_key  = "YOUR AWS ACCESS KEY"
  secret_key  = "YOUR AWS SECRET KEY"
  token       = "YOUR AWS TOKEN"
}

resource "aws_instance" "exampleOllamaConfiguration" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.medium"
  count         = 1
  vpc_security_group_ids = [aws_security_group.instance.id]

  root_block_device {
    volume_size = 24 # In Gb
  }

  key_name  = "vokey"
  user_data = "${file("creation.sh")}"

  user_data_replace_on_change = true

  tags = {
    Name = "terraform-example-ollama"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port   = 11434
    to_port     = 11434
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-ollama-example-instance"
}

output "address" {
  value      = aws_instance.exampleOllamaConfiguration[*].public_dns
  description = "Address of the Kafka EC2 machine with ollama"
}
```

N.2. And the following bash script to be executed after EC2 resource creation by the following script creation.sh:

```
#!/bin/bash
cd
sudo yum update -y

sudo curl -fsSL https://ollama.com/install.sh | sh

export HOME=$HOME:/usr/local/bin
sudo sed -i "s/\"[Install\"]/Environment=\"OLLAMA_HOST=0.0.0.0:11434\"\\n\[Install\]/g"
/etc/systemd/system/ollama.service

sudo systemctl enable ollama
sudo systemctl start ollama

ollama pull llama3.2
```

*Hint 1:* Additionally, if inside the EC2 instance you would like to check the models loaded use the following command:

```
[ec2-user@ip-172-31-23-142 ~]$ ollama list
NAME           ID          SIZE      MODIFIED
llama3.2:1b   baf6a787fdff   1.3 GB    7 minutes ago
```

*Hint 2:* Additionally, if inside the EC2 instance you would like to check the processors loaded use the following command:

```
[ec2-user@ip-172-31-23-142 ~]$ ollama ps
NAME           ID          SIZE      PROCESSOR      UNTIL
llama3.2:1b   baf6a787fdff   2.2 GB    100% CPU      4 minutes from now
```

*Hint 3:* Additionally, a request example can be tested with the following remote curl command:

```
curl http://<EC2 DNS NAME>:11434/api/generate -d '{
  "model": "llama3.2",
  "prompt": "Why is the sky blue?",
  "stream": false
}'
```

The response obtained is similar with the following:

```
{"model":"llama3.2","created_at":"2025-02-12T15:18:11.983219976Z","response":"The sky appears blue because of a phenomenon called scattering, which occurs when sunlight interacts with the tiny molecules of gases in the Earth's atmosphere.\n\nHere's a simplified explanation:\n1. **Sunlight**: When the sun shines, it emits white light, which contains all the colors of the visible spectrum (red, orange, yellow, green, blue, indigo, and violet).\n2. **Atmospheric scattering**: As sunlight enters the Earth's atmosphere, it encounters tiny molecules of gases such as nitrogen (N2) and oxygen (O2). These molecules are much smaller than the wavelength of light.\n3. **Scattering**: When sunlight hits these tiny molecules, it scatters in all directions. However, not all wavelengths scatter equally. The shorter (blue) wavelengths are scattered more than the longer (red) wavelengths by the smaller gas molecules.\n4. **Blue light dominates**: As a result of this scattering effect, the blue light is dispersed throughout the atmosphere, while the red and orange light continue to travel in a straight line.\n5. **Our eyes perceive the sky as blue**: From our perspective on the surface of the Earth, we see the scattered blue light coming from all directions, which gives the sky its blue appearance.\n\nThis phenomenon is more pronounced during the daytime when the sun is overhead, and the scattering effect is greater. During sunrise and sunset, the sun's rays have to travel through more of the atmosphere, scattering off even more molecules and creating the warm hues we see.\n\nSo, in short, the sky appears blue because of the scattering of sunlight by tiny molecules in the Earth's atmosphere, which favors shorter wavelengths (like blue light) over longer wavelengths (like red light).","done":true,"done_reason":"stop","context":128006,9125,128007,271,38766,1303,33025,2696,25,6790,220,2366,18,271,128009,128006,882,128007,271,10445,374,279,13180,6437,30,128009,128006,78191,128007,271,791,13180,8111,6437,1606,315,264,25885,2663,72916,11,902,13980,994,40120,84261,449,279,13987,35715,315,45612,304,279,9420,596,16975,382,8586,596,264,44899,16540,1473,16,13,3146,31192,4238,96618,3277,279,7160,65880,11,433,73880,4251,3177,11,902,5727,682,279,8146,315,279,9621,20326,320,1171,11,19087,11,14071,11,6307,11,6437,11,1280,7992,11,323,80836,4390,17,13,3146,1688,8801,33349,72916,96618,1666,40120,29933,279,9420,596,16975,11,433,35006,13987,35715,315,45612,1778,439,47503,320,4517,8,323,24463,320,46,17,570,4314,35715,527,1790,9333,1109,279,46406,315,3177,627,18,13,3146,3407,31436,96618,3277,40120,13280,1521,13987,35715,11,433,1156,10385,304,682,18445,13,4452,11,539,682,93959,45577,18813,13,578,24210,320,12481,8,93959,527,38067,810,1109,279,5129,320,1171,8,93959,555,279,9333,6962,35715,627,19,13,3146,10544,3177,83978,96618,1666,264,1121,315,420,72916,2515,11,279,6437,3177,374,77810,6957,279,16975,11,1418,279,2579,323,19087,3177,3136,311,5944,304,264,7833,1584,627,20,13,3146,8140,6548,45493,279,13180,439,6437,96618,5659,1057,13356,389,279,7479,315,279,9420,11,584,1518,279,38067,6437,3177,5108,505,682,18445,11,902,6835,279,13180,1202,6437,11341,382,2028,25885,374,810,38617,2391,279,62182,994,279,7160,374,32115,11,323,279,72916,2515,374,7191,13,12220,64919,323,44084,11,279,7160,596,45220,617,311,5944,1555,810,315,279,16975,11,72916,1022,1524,810,35715,323,6968,279,8369,82757,584,1518,382,4516,11,304,2875,11,279,13180,8111,6437,1606,315,279,72916,315,40120,555,13987,35715,304,279,9420,596,16975,11,902,54947,24210,93959,320,4908,6437,3177,8,927,5129,93959,320,4908,2579,3177,570],"total_duration":72783017580,"load_duration":3846465248,"prompt_eval_count":31,"prompt_eval_duration":3108000000,"eval_count":344,"eval_duration":65784000000}
```

## References

URLs with other resources

- Terraform Up & Running, Writing Infrastructure as Code, Yevgeniy Brikman, 3<sup>rd</sup> edition (2022).  
<https://www.terraformupandrunning.com/>
  - The github with all the source code from the book is at: <https://github.com/brikis98/terraform-up-and-running-code>
- <https://developer.hashicorp.com/terraform/downloads>
- <https://developer.hashicorp.com/terraform/tutorials/aws-get-started>
- <https://stackoverflow.com/questions/41596412/how-to-use-terraform-output-as-input-variable-of-another-terraform-template>
- <https://developer.hashicorp.com/terraform/tutorials/aws/lambda-api-gateway>
- <https://docs.aws.amazon.com/lambda/latest/dg/lambda-runtimes.html>
- <https://registry.terraform.io/providers/hashicorp/aws/2.34.0/docs/guides/serverless-with-aws-lambda-and-api-gateway>
- Ollama API documentation: <https://github.com/ollama/ollama/blob/main/docs/api.md>
- Ollama FAQs documentation: <https://github.com/ollama/ollama/blob/main/docs/faq.md>
- Ollama readme page <https://github.com/ollama/ollama?tab=readme-ov-file>