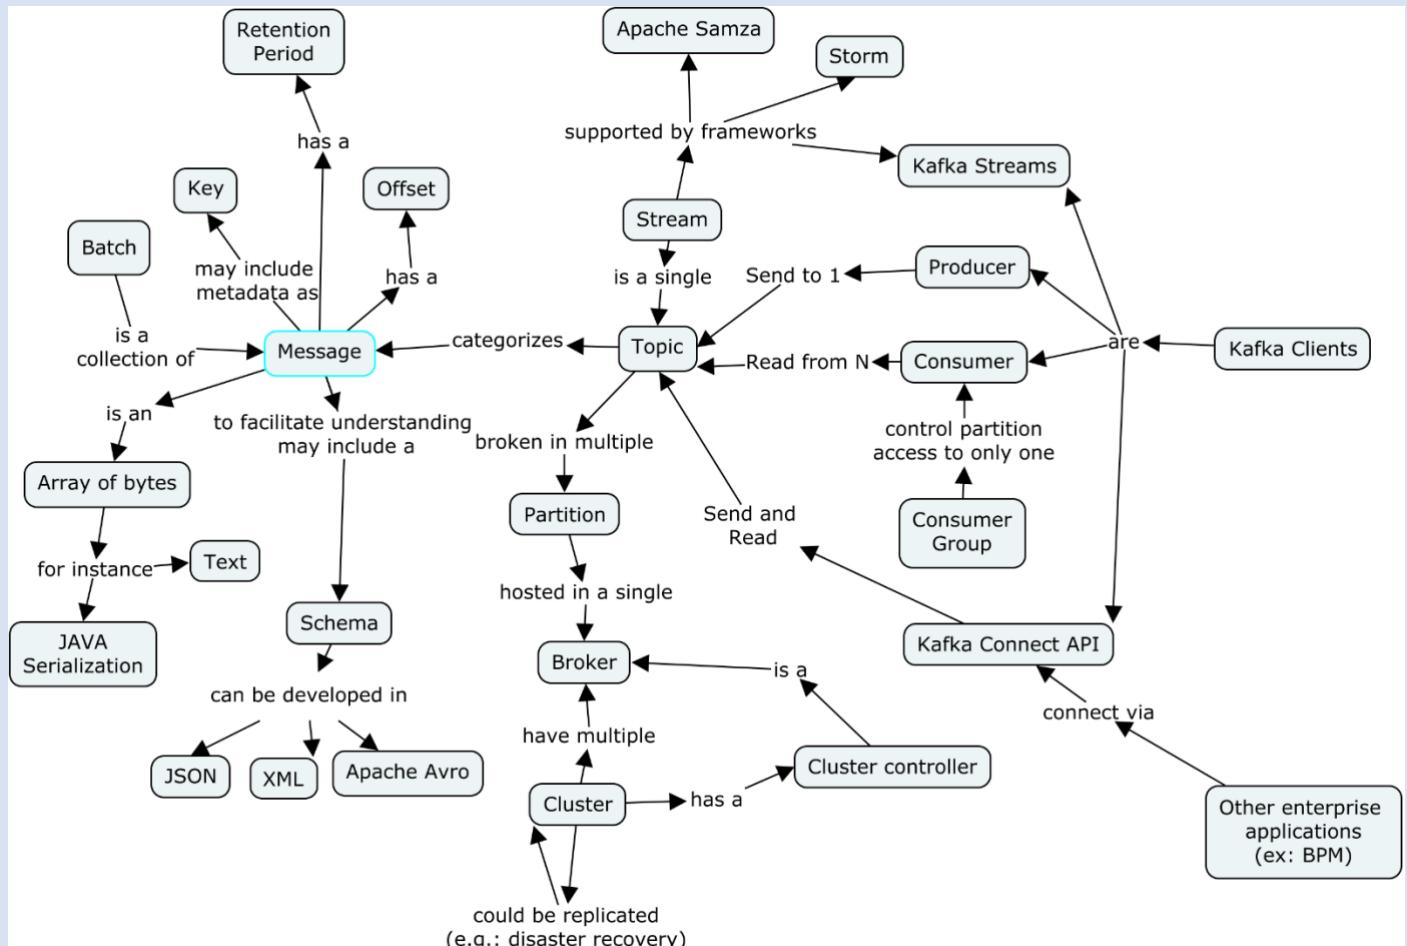


The goal of this document is to show how to operate the basics of a remote Kafka service: installation, starting, accessing, testing, stopping and backup. Remote installation step by step is presented, where the remote installation requires an Amazon AWS account. The following concepts are contained in Kafka:



The following contents is presented in this document.

Contents

A. Creating AWS Academy account and AWS private key	3
B. Creating and launching an AWS EC2 instance	6
C. Access the AWS EC2 instance using PuTTY (for windows only)	10
D. Access the AWS EC2 instance using ssh (for linux/macOS) and scp (for linux/macOS)	12
E. Access the AWS EC2 instance using FileZilla (for windows and macOS)	14

F.	Installing manually Kafka in the AWS EC2 instance.....	15
G.	Testing Kafka locally in the AWS EC2 instance.....	17
H.	Testing Kafka in the AWS EC2 instance remotely	18
I.	Stopping the AWS EC2 instance	21
	Other references.....	22

A. Creating AWS Academy account and AWS private key

The goal of this section is to activate your AWS account and create your private key that will allow you to access your EC2 instance.

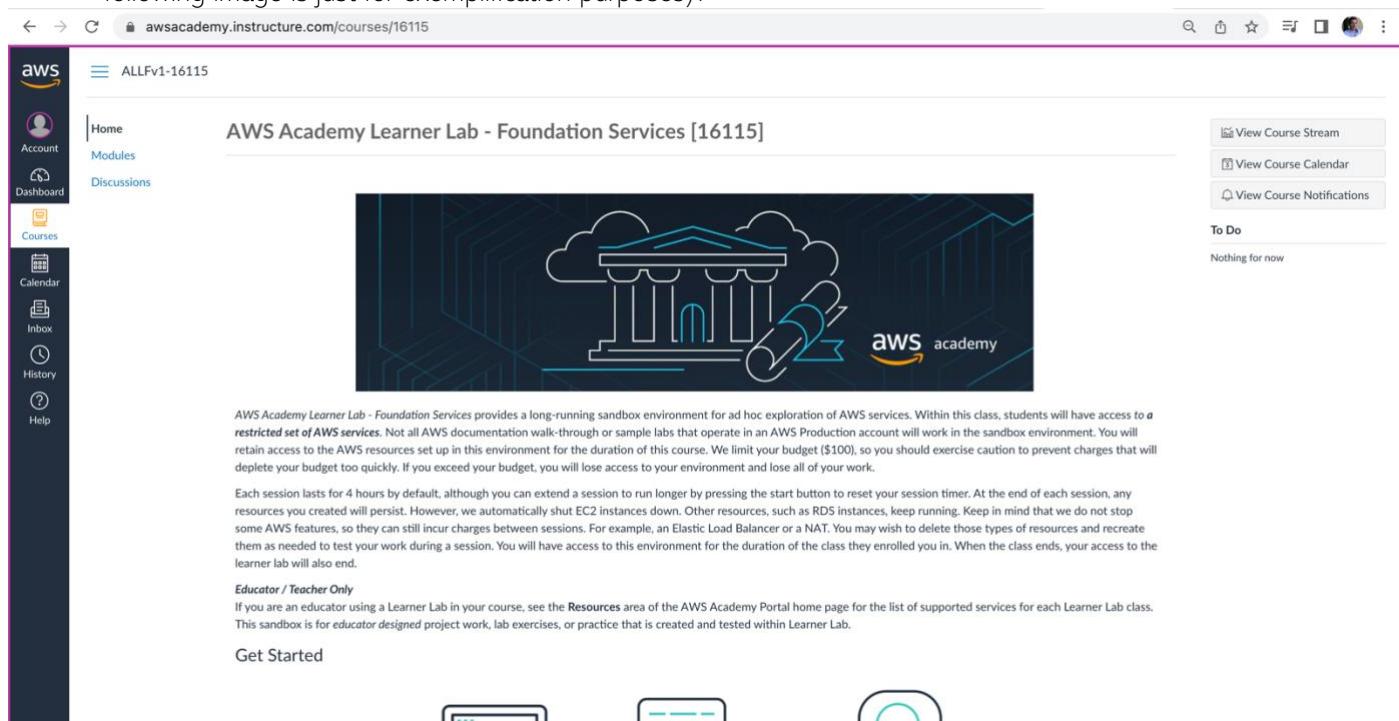
- A.O. Create an amazon AWS Academy account using the invitation email sent by Amazon Academy on the faculties request.

While you are in the process of account creation be careful with your graduation date – define it, at least, accordingly with the end of semester date.

Important remark: the usage of cloud resources implies costs. An EC2 AWS instance is a machine stored in an Amazon AWS facility that is using the usual computational resources: CPU, memory, storage, communications, energy, physical space, among others. This is a reality that any organization faces daily.

Therefore, you should be aware that while your AWS instance is started, your AWS account also starts to be billed.
We absolutely recommend you to always stop your instances when you are not working with them. Always manage your resources wisely!

- A.1. After registration, accept the terms, and then you can access your account and have a similar canvas (the following image is just for exemplification purposes).



The screenshot shows the AWS Academy Learner Lab - Foundation Services [16115] course page. The left sidebar has a dark theme with icons for Account, Courses, Calendar, Inbox, History, and Help. The main content area has a light background. At the top, there's a navigation bar with a search icon, a star icon, and other course-related links. Below the navigation is a large banner featuring a stylized illustration of clouds and a building, with the AWS academy logo. To the right of the banner are three buttons: 'View Course Stream', 'View Course Calendar', and 'View Course Notifications'. A 'To Do' section below them says 'Nothing for now'. The central part of the page contains descriptive text about the learner lab, including a note about restricted AWS services and budget limits. It also mentions session duration and educator/teacher access. At the bottom, there are three small icons representing different AWS services: a server, a dashed line, and a rainbow arc.

- A.2. Select the available course (the following image is just for exemplification purposes):

A.3. Select "Modules" and then "Learner Lab – Foundational Services"

A.4. Click "Start Lab" on the right side

A.5. Then, when the indicator is **green**, click "AWS Details" to access AWS Management Console on the right side. As indicated by the red boxes below. Store the returned key file in your computer for later use. You can choose .pem to use with OpenSSH and/or .ppk to use with PuTTY. If you are unsure store the keys in both formats on your computer.

A.6. Then, click "AWS" to access AWS Management Console on the left side. As indicated by the red box below.

The screenshot shows the AWS Academy interface. On the left, there's a sidebar with icons for Account, Courses, Calendar, Inbox, History, and Help. The main area has a breadcrumb trail: ALLFv1-16115 > Modules > Learner Lab Foundation Services > Learner Lab - Foundational Services. A red box highlights the "AWS" link in the top navigation bar. To the right, there's a terminal window showing a command prompt and a "Learner Lab - Foundational Level" section with various links.

A.7. Then, click "AWS Details" to consult all the credentials to access your AWS area.

The screenshot shows the AWS Academy interface. The sidebar includes Account, Courses, Calendar, Inbox, History, and Help. The main area has a breadcrumb trail: awsacademy.instructure.com/courses/107105/modules/items/10016408. A red box highlights the "AWS Details" link in the top navigation bar. The "AWS Details" section displays AWS CLI credentials (access key ID, secret access key, session token) and session details (start time, end time, accumulated lab time). It also shows AWS Account ID and Region information.

A.8. In the left upper corner, select "Services -> Compute (EC2)". You can also search for EC2.

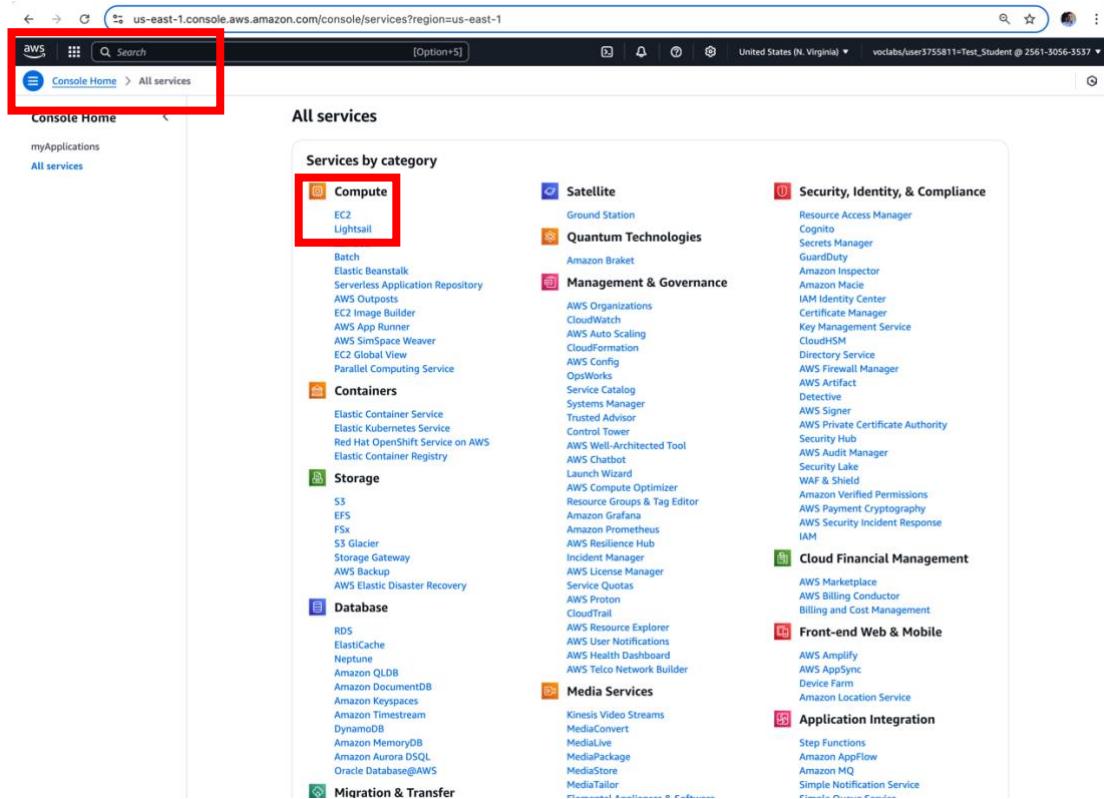
The screenshot shows the AWS Management Console. The left sidebar has a "Services" dropdown menu. A red box highlights the "Compute" option under the EC2 category. The main pane lists various AWS services like Blockchain, Analytics, Customer Engagement, etc., with the EC2 service highlighted.

You will obtain the AWS EC2 dashboard. All the options are vertically organized in the left side.

B. Creating and launching an AWS EC2 instance

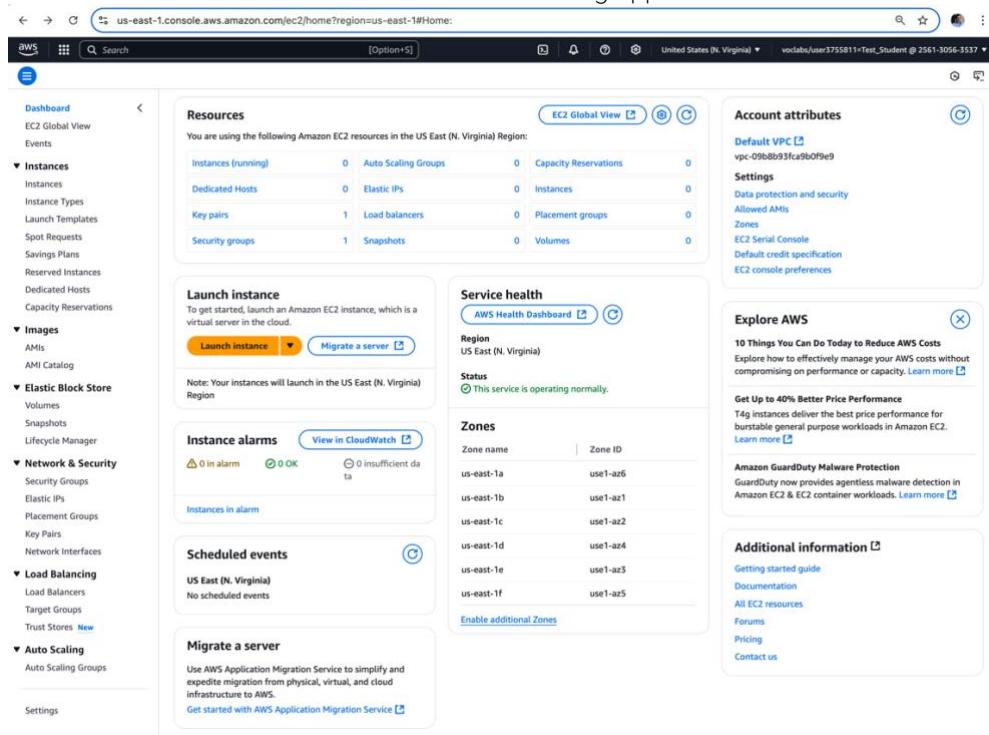
The goal of this section is to create your first EC2 instance using the available Amazon Machine Images (AMI) available. This requires the previous execution of section A.

- B.1. Login in your account, then access the AWS console, and in the left upper corner, select “A11 Services -> Compute (EC2)”.



The screenshot shows the AWS Console interface. In the top navigation bar, the URL is "us-east-1.console.aws.amazon.com/console/services?region=us-east-1". The top right shows "United States (N. Virginia)" and a user profile. Below the navigation bar is the "Console Home" link and the "All services" link. The main area is titled "All services" and contains a sidebar titled "Services by category". The "Compute" category is highlighted with a red box. Other categories listed include Satellite, Quantum Technologies, Management & Governance, Storage, Database, Migration & Transfer, Security, Identity, & Compliance, Cloud Financial Management, Front-end Web & Mobile, Media Services, Application Integration, and CloudWatch Metrics.

- B.2. Select “Instances -> Instances”. The following appearance will be shown.



The screenshot shows the AWS EC2 Instances page. The left sidebar includes links for Dashboard, Instances, Images, Elastic Block Store, Network & Security, Load Balancing, and Auto Scaling. The main content area has sections for Resources (listing instances, dedicated hosts, key pairs, and security groups), Launch instance (with a prominent red box around the "Launch instance" button), Service health (AWS Health Dashboard), Zones (listing availability zones like us-east-1a through us-east-1f), and Additional information (links to Getting started guide, Documentation, All EC2 resources, Forums, Pricing, and Contact us).

B.3. Select "Launch Instance" and a similar interface will be presented as shown below. The goal is to define the Amazon Machine Image (AMI) that will contain all the software configuration required to launch your instance.

The screenshot shows the AWS EC2 'Launch an instance' wizard. The top navigation bar includes links for 'Search', '[Option+S]', 'United States (N. Virginia)', and 'voclabs/user3755811=Test_Student @ 2561-3056-3537'. Below the navigation is a breadcrumb trail: 'EC2 > Instances > Launch an instance'.

Name and tags

Name: e.g. My Web Server [Add additional tags](#)

Application and OS Images (Amazon Machine Image)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below.

[Search our full catalog including 1000s of application and OS images](#)

Quick Start

Quick start options include: Amazon Linux, macOS, Ubuntu, Windows, Red Hat, SUSE Linux, and Debian.

Amazon Machine Image (AMI)

Selected: Amazon Linux 2023 AMI
AMI ID: ami-05576a079321f21f8
Virtualization: hvm ENA enabled: true Root device type: ebs [Free tier eligible](#)

Description

Amazon Linux 2023 is a modern, general purpose Linux-based OS that comes with 5 years of long term support. It is optimized for AWS and designed to provide a secure, stable and high-performance execution environment to develop and run your cloud applications.

Amazon Linux 2023 AMI 2023.6.20250107.0 x86_64 HVM kernel-6.1.10-100.1.1.amzn2.0.1-1.1.1

Instance type

Selected: t2.micro [Free tier eligible](#)

Family: t2 1 vCPU 1 GiB Memory Current generation: true
On-Demand Windows base pricing: 0.0162 USD per Hour
On-Demand Ubuntu Pro base pricing: 0.0134 USD per Hour On-Demand SUSE base pricing: 0.0116 USD per Hour
On-Demand RHEL base pricing: 0.026 USD per Hour On-Demand Linux base pricing: 0.0116 USD per Hour

[All generations](#) [Compare instance types](#)

Summary

Number of instances: 1

Software Image (AMI)
Amazon Linux 2023 AMI 2023.6.2... [read more](#)
ami-05576a079321f21f8

Virtual server type (instance type)
t2.micro

Firewall (security group)
New security group

Storage (volumes)
1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the internet.

[Cancel](#) [Launch instance](#) [Preview code](#)

B.4. Use the following configurations:

Step 1 – Choose an Amazon Machine Image (AMI)	select "Amazon Linux 2 AMI (HVM) - Kernel 5.10, SSD Volume Type" (available on January 2025) with "64-bit (x86)"
Step 2 – Choose an Instance Type	select "t2.micro"
Step 3 – Configure Instance Details	check if number of instances = 1 Keep all the remaining options as default
Step 4 – Storage	keep the default size
Step 5 – Key Pair (login)	Key Pair name = vockey
Step 6 – Add Tags	Key = application Value = kafka
Step 7 – Configure Security Group <i>(it could be configured later in the instance console or</i>	Select "Create a new security group" Security group name: launch-Kafka Select "Add rule" Type = "Custom TCP Rule" Protocol = TCP Port range = 2181

<p><i>since the creation using terraform)</i></p>	<p>Source = Anywhere (0.0.0.0, ::/0) Description = zookeeper Select "Add rule" Type = "Custom TCP Rule" Protocol = TCP Port range = 9092 Source = Anywhere (0.0.0.0, ::/0) Description = kafka</p> <p>Select "Add rule" Type = "SSH" Protocol = TCP Port range = 22 Source = Anywhere (0.0.0.0, ::/0) Description = remote access</p>
<p>Step 8 – Review Instance Launch</p>	

B.5. Press "Launch" in the bottom of the page



Your AWS account starts to be billed at this point.

B.6. Go to EC2 dashboard as explained in A.1., and you will find an appearance similar with the following image.

Instances (1/1) Info	
<input checked="" type="checkbox"/> Name <input type="text" value="teste kafka"/>	<input type="checkbox"/> Instance ID <input type="text" value="i-064b2df1db33ddcde"/>
<input type="checkbox"/> Instance state <input checked="" type="radio"/> Running <input type="radio"/> Stopped <input type="radio"/> Pending <input type="radio"/> Terminating <input type="radio"/> Shutting down	
<input type="checkbox"/> Instance type <input type="text" value="t2.micro"/>	
<input type="checkbox"/> Status check <input checked="" type="radio"/> 2/2 checks passed <input type="radio"/> 1/2 checks passed <input type="radio"/> 0/2 checks passed	
<input type="checkbox"/> Alarm status <input checked="" type="radio"/> No alarms <input type="radio"/> 1 alarm	
<input type="checkbox"/> Availability Zone <input type="text" value="us-east-1d"/>	

i-064b2df1db33ddcde (teste kafka)

Details		Status and alarms	Monitoring	Security	Networking	Storage	Tags
Instance summary							
Instance ID	<input type="text" value="i-064b2df1db33ddcde"/>						
Public IPv4 address	<input type="text" value="52.23.182.107 open address"/>						
Private IP4 address	<input type="text" value="172.31.16.172"/>						
IPv6 address	<input type="text" value="–"/>						
Instance state	<input checked="" type="radio"/> Running <input type="radio"/> Stopped <input type="radio"/> Pending <input type="radio"/> Terminating <input type="radio"/> Shutting down						
Hostname type	<input type="text" value="IP name: ip-172-31-16-172.ec2.internal"/>						
Answer private resource DNS name	<input type="text" value="IPV4 (A)"/>						
Auto-assigned IP address	<input type="text" value="52.23.182.107 [Public IP]"/>						
IAM Role	<input type="text" value="–"/>						
Subnet ID	<input type="text" value="subnet-0fc8666a60fe9f5f"/>						
Instance ARN	<input type="text" value="arn:aws:ec2:us-east-1:256130563537:instance/i-064b2df1db33ddcde"/>						
IMDSv2	<input type="checkbox"/> Required						
Operator	<input type="text" value="–"/>						
Instance details							
AMI ID	<input type="text" value="ami-0454e52560c7f5c55"/>						
Monitoring	<input type="checkbox"/> disabled						
Platform details	<input type="text" value="Linux/UNIX"/>						

You can check the IPv4 and v6 public address, and the state of the instance. In this example it is *running*.

B.7. Stop the instance, checking the row with the desired instance and then pressing "Instance state -> Stop". Wait for the instance State to change to stopped.

B.8. To allow a proper future execution of Kafka, change the type of instance, by choosing the following options, and then choose **t2.small**.

The screenshot shows the AWS EC2 Instances page. On the left, there's a sidebar with various navigation links like EC2 Dashboard, Events, Tags, Reports, Limits, Instances, Launch Templates, Spot Requests, Reserved Instances, Dedicated Hosts, Scheduled Instances, Capacity Reservations, Images, AMIs, Bundle Tasks, Elastic Block Store, Volumes, Snapshots, Lifecycle Manager, Network & Security, Security Groups, Elastic IPs, Placement Groups, Key Pairs, and Network Interfaces. The main area has tabs for Launch Instance, Connect, and Actions. Under Actions, a dropdown menu is open with options: Connect, Get Windows Password, Create Template From Instance, Launch More Like This, Instance Settings (which is currently selected), Add/Edit Tags, Attach to Auto Scaling Group, Attach/Replace IAM Role, Change Instance Type, Change Termination Protection, View/Change User Data, Change Shutdown Behavior, Change T2/T3 Unlimited, Get System Log, Get Instance Snapshot, Modify Instance Placement, and Modify Capacity Reservation Settings. Below this, there's a detailed view of an instance: Instance ID: i-03a6db7323ba18361, Private IP: 172.31.30.91. The instance state is stopped. The instance type is t2.small. It has one elastic IP assigned: us-east-1a. The AMI ID is amazon-ami-hvm-2018.03.0.20181129-x86_64-gp2 (ami-0080e4c5bc079760e). The owner is 032896652389. The launch time was February 25, 2019 at 1:13:28 PM UTC. The termination protection is False. The monitoring status is basic. The alarm status is None. On the right, there's a large table with columns: Availability Zone, Instance State, Status Checks, Alarm Status, Public DNS (IPv4), IPv4 Public IP, IPv6 IPs, Key Name, and Mon. There are three rows in the table, each corresponding to an instance. At the bottom of the page, there are links for Feedback, English (US), and footer links for © 2008 - 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved., Privacy Policy, and Terms of Use.

C. Access the AWS EC2 instance using PuTTY (for windows only)

The goal of this section is to describe how your AWS EC2 instance can be accessed using an SSH client. It will be useful to install, configure and manage the Kafka server.

- C.1. Install PuTTY from <https://www.chiark.greenend.org.uk/~sgtatham/putty/latest.html>
- C.2. (*optional – only if the key is a .pem file*) Convert your AWS private key created in section A, using PuTTYgen. For that, execute PuTTYgen from start menu.
- C.3. (*optional – only if the key is a .pem file*) Then, in type of key to generate choose: RSA
- C.4. (*optional – only if the key is a .pem file*) Press “Load” and locate your AWS private key stored in step A.4. (with .pem file extension)
- C.5. (*optional – only if the key is a .pem file*) Press “Save private key” and choose “Yes” when asking to save without passphrase
- C.6. (*optional – only if the key is a .pem file*) Store with the same name as the AWS private key (with .ppk file extension)
- C.7. Start PuTTY from start menu
- C.8. Fill hostname with “user_name@public_dns_name”, where user_name = ec2-user and public_dns_name= <as presented in B.7.>



Public DNS name is changed whenever an AWS EC2 instance reboot is performed.

- C.9. Choose SSH connection type and check if port = 22
- C.10. You can configure PuTTY to automatically send "keepalive" data at regular intervals to keep the session active. This is useful to avoid disconnecting the instance from session inactivity. In the Category panel, choose Connection and enter the required range in the Seconds between keepalives field. For example, if the session disconnects after 10 minutes of inactivity, enter 180 to configure PuTTY to send keepalive data every 3 minutes.
- C.11. In the Category panel, choose Connection, choose SSH, choose Auth, select “Browse” and choose the .ppk file from step C.6.



Start your AWS instance.

- C.12. In the Category panel, choose Session, and then select “Open”
- C.13. Confirm the certificate in the dialog box. You will be presented with the following similar screen:

```
ec2-user@ip-172-31-10-11: ~
Using username "ec2-user".
Authenticating with public key "imported-openssh-key"
Last login: Thu Dec 27 11:43:38 2018 from
[ec2-user@ip-172-31-10-11 ~]$
```

- C.14. Update your AWS instance with the command:

```
sudo yum update
```

- C.15. Check JDK packages available with the command:

```
sudo yum search "java-17"
```

- C.16. Execute the command (or another equivalent package)

```
sudo yum install java-17-amazon-corretto-devel.x86_64
```

C.17. Check if JAVA is currently version 17 with the command:

```
java -version
```

If any other version is referred you may remove it, e.g., first check the installed java versions with the command:

```
yum list installed |grep java
```

and then remove the ones that you don't need with the command

```
sudo yum remove java-1.7.0-openjdk
```

C.18. Execute the command:

```
exit
```



Stop your AWS instance.

D. Access the AWS EC2 instance using ssh (for linux/macOS) and scp (for linux/macOS)

The goal of this section is to describe how your AWS EC2 instance can be accessed using an SSH client in linux. It will be useful to install, configure and manage the Kafka server.



Start your AWS instance.

Public DNS name is changed whenever an AWS EC2 instance reboot is performed.

D.1. In a new terminal session, change directories to the location of the private key file that you created when you created the AWS instance, for instance:

```
[oracle@soabpm-vm ~]$ cd /media/sf_SharedFolderForOracleSOA
```

C.19. Use the following command to set the permissions of your private key file so that only you can read it.

```
[oracle@soabpm-vn sf SharedFolderForOracleSOA]$ chmod u=rwx,g=,o= myKeyAWS.pem
```

C.20. Use the ssh command to connect to the instance. You specify the private key (.pem) file and user_name@public_dns_name. For example, if you used Amazon Linux 2 or the Amazon Linux AMI, the user name is ec2-user. In

```
[oracle@soabpm-vm sf_SharedFolderForOracleSOA]$ ssh -i myKeyAWS.pem ec2-user@YOURIP  
The authenticity of host 'X.X.X.X (X.X.X.X)' can't be established.  
RSA key fingerprint is 25:1c:ef:ee:c1:87:61:e9:ea:f6:e3:45:1e:3a:63:73.  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added 'X.X.X.X' (RSA) to the list of known hosts.  
Last login: Fri Feb  8 14:59:07 2019 from Y.Y.Y.Y
```

- | (- | - /) Amazon Linux AMI

```
https://aws.amazon.com/amazon-linux-ami/2018.03-release-notes/
5 package(s) needed for security, out of 6 available
Run "sudo yum update" to apply all updates.
ZooKeeper JMX enabled by default
Using config: /usr/local/zookeeper/bin/../conf/zoo.cfg
Starting zookeeper ... already running as process 11211.
[ec2-user@ip-X-X-X-X ~]$
```

C.21. Update your AWS instance with the command:

```
sudo yum update
```

C.22. Check JDK packages available with the command:

```
sudo yum search "java-17"
```

C.23. Execute the command (or another equivalent package)

```
sudo yum install java-17-amazon-corretto-devel.x86_64
```

C.24. Check if JAVA is currently version 17 with the command:

```
java -version
```

If any other version is referred you may remove it, e.g., first check the installed java versions with the command:

```
yum list installed |grep java
```

and then remove the ones that you don't need with the command

```
sudo yum remove java-1.7.0-openjdk
```

C.25. Execute the command:

exit

C.26. To transfer files from your machine to your AWS instance, use the scp application in linux. For instance:

P1-Kafka-in-AWSAcademy-v3.0.docx

```
scp -i my-key-pair.pem /path/SampleFile.txt ec2-user@c2-198-51-100-1.compute-1.amazonaws.com:~
```

C.27. To transfer files from your AWS instance to your machine, use the scp application in linux. For instance:

```
scp -i my-key-pair.pem ec2-user@ec2-198-51-100-1.compute-1.amazonaws.com:~/SampleFile.txt ~/SampleFile2.txt
```

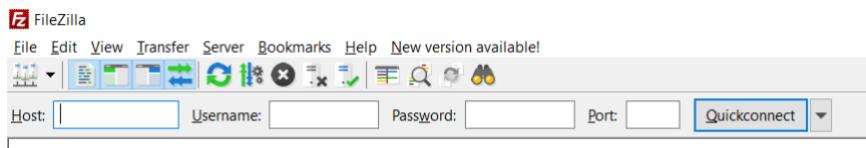


Stop your AWS instance.

E. Access the AWS EC2 instance using FileZilla (for windows and macOS)

The goal of this section is to describe how your AWS EC2 instance can be accessed using an SFTP client. It will be useful to transfer files between your computer and AWS EC2 instance and back.

- E.1. Install FileZilla from <https://filezilla-project.org/download.php?type=client>
- E.2. Open FileZilla.
- E.3. Go to the FileZilla Edit -> settings, and on the left, click SFTP.
- E.4. Add a new private key. (Your .ppk key)
- E.5. If you are using a .pem key you must convert it accordingly with steps C.2. to C.6., otherwise it will not work
- E.6. At the top in the Quickconnect bar (similar with the image below), put your Public DNS in the host, ec2-user, port 22 (Port 22 is SFTP rather than FTP, AWS will kick back FTP.), and NO PASSWORD.



Start your AWS instance.

- E.7. Click Quickconnect.
- E.8. File transfer to AWS EC2 instance is only a matter of dragging files from the left to right side.



Stop your AWS instance.

F. Installing manually Kafka in the AWS EC2 instance

The goal of this section is to describe the required steps to install the Zookeeper and Kafka servers in your AWS EC2 instance.



Start your AWS instance.

F.1. Access your AWS EC2 instance using previous steps of this tutorial.

F.2. To change for your home directory, type the command:

```
cd
```

F.3. Download the Kafka binary version from <https://kafka.apache.org/downloads>. Recommended version is kafka_2.13-3.9.0.tgz (in January 2025).

```
wget https://downloads.apache.org/kafka/3.9.0/kafka_2.13-3.9.0.tgz
```

F.4. Download the ZooKeeper binary version from <https://zookeeper.apache.org/releases.html>. Recommended version is apache-zookeeper-3.9.3-bin.tar.gz (in January 2025).

```
wget https://dlcdn.apache.org/zookeeper/zookeeper-3.9.3/apache-zookeeper-3.9.3-bin.tar.gz
```

F.5. Then, type the command:

```
ls -ltr
```

to check if you have both kafka and zookeeper files available.

F.6. Update your AWS instance with the command:

```
sudo yum update
```

F.7. Execute the command (or another equivalent package)

```
sudo yum install java-17-amazon-corretto-devel.x86_64
```

F.8. To extract the zookeeper files, type the command:

```
tar -zxf apache-zookeeper-3.9.3-bin.tar.gz
```

F.9. Move the new zookeeper directory to system typing the command:

```
sudo mv apache-zookeeper-3.9.3-bin /usr/local/zookeeper
```

F.10. Create a new directory using the command:

```
sudo mkdir -p /var/lib/zookeeper
```

F.11. Create the baseline zookeeper server configuration with the command:

```
cat > /usr/local/zookeeper/conf/zoo.cfg << EOF
```

and then write the following content to the file directly in the command line:

```
tickTime=2000  
dataDir=/var/lib/zookeeper  
clientPort=2181  
EOF
```

F.12. Start ZooKeeper server with the command:

```
sudo /usr/local/zookeeper/bin/zkServer.sh start
```

F.13. To extract the kafka server, type the command:

```
tar -zxf kafka_2.13-3.9.0.tgz
```

F.14. Move the new kafka directory to system typing the command:

```
sudo mv kafka_2.13-3.9.0 /usr/local/kafka
```

F.15. Create a new directory using the command:

```
sudo mkdir /tmp/kafka-logs
```

F.16. Start Kafka with the command:

```
sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon /usr/local/kafka/config/server.properties
```

Hint 1: if you are short in the instance memory reduce the JVM usage with the command:

```
export KAFKA_HEAP_OPTS="-Xmx256M -Xms256M"
```

Hint 2: you can always check if Zookeeper and Kafka servers are started correctly with the command:

```
ps -ef |grep java
```



Stop your AWS instance.

Exercise 1: To automate the start of Zookeeper and Kafka servers with instance login write both starting commands in the file:

```
.bash_profile
```

G. Testing Kafka locally in the AWS EC2 instance

The goal of this section is to test if the previous installation in section E. was executed successfully. The production and consumption of Kafka messages are tested. It is considered that in this environment the servers and the clients are installed in the same machine.



Start your AWS instance.

G.1. Access your AWS EC2 instance using the steps C.7. to C.12.

G.2. First, install telnet package in your AWS EC2 instance typing the command:

```
sudo yum install telnet.x86_64
```

G.3. Test ZooKeeper installation typing the command:

```
telnet localhost 2181
```

then write the command:

```
srvr
```

and it is expected to receive the following output:

```
[ec2-user@ ~]$ telnet localhost 2181
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
srvr
Zookeeper version: 3.4.12-e5259e437540f349646870ea94dc2658c4e44b3b, built on 03/27/2018 03:55 GMT
Latency min/avg/max: 0/0/0
Received: 1
Sent: 0
Connections: 1
Outstanding: 0
Zxid: 0x0
Mode: standalone
Node count: 4
Connection closed by foreign host.
```

G.4. Test Kafka installation, with the creation of a new topic named "test", typing the command:

```
sudo /usr/local/kafka/bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --
-partitions 1 --topic test
```

G.5. Type the command to check if the topic "test" is created successfully:

```
sudo /usr/local/kafka/bin/kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic test
```

G.6. Type the command to produce messages to the topic "test":

```
sudo /usr/local/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic test
```

And write some text messages, one per line. Then, to finish press Ctrl+D.

G.7. Type the command to consume messages from the topic "test":

```
sudo /usr/local/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic test --from-
beginning
```

All the sent messages from topic "test" will be presented. Then, to finish press Ctrl+C.

Exercise 2: Try creating different PuTTY sessions for producer and consumer and check the run-time production and consumption of messages.



Stop your AWS instance.

H. Testing Kafka in the AWS EC2 instance remotely

The goal of this section is to test if the previous installation in section E. was executed successfully. The production and consumption of Kafka messages are tested. It is considered that in this environment the servers and the clients are installed in different machines.

This section requires the previous installation of Kafka, as described in section E.



Start your AWS instance.

- H.0. Access your AWS EC2 instance using the steps C.7. to C.12.
H.1. From your local docker container, from local macOS, local linux, or other AWS EC2 linux environment, check if you can connect to the AWS EC2 instance Zookeeper server using the following command:

```
telnet YourPublicIP 2181
```

remember that *YourPublicIP* can be found in the AWS EC2 dashboard as presented below:

The screenshot shows the AWS EC2 Instances page. A single instance named 'i-0e47631114bff84b1' is listed as 'Running'. The 'Public IPv4 address' (34.203.227.188) and 'Public IPv4 DNS' (ec2-34-203-227-188.compute-1.amazonaws.com) are highlighted with red boxes. The left sidebar shows various AWS services like EC2 Dashboard, Events, Tags, and Instances (with one item listed). The bottom of the page shows detailed instance information for 'Instance: i-0e47631114bff84b1', including platform (Amazon Linux Inferred), AMI ID (ami-0947d2ba12ee1ff5), and monitoring status (disabled).

After that, send the command:

```
SRVR
```

You should be able to see an interaction similar with the following:

```
root@29f9117a8312:~# telnet [REDACTED] 2181
Trying [REDACTED] ...
Connected to [REDACTED]
Escape character is '^J'.
srvr
Zookeeper version: 3.4.12-e5259e437540f349646870ea94dc2658c4e44b3b, built on 03/27/2018 03:55 GMT
Latency min/avg/max: 0/0/7
Received: 6819
Sent: 6822
Connections: 2
Outstanding: 0
Zxid: 0xee
Mode: standalone
Node count: 133
Connection closed by foreign host.
```

Hint 3: Public IP or Public DNS Name could be.

- H.2. Using the FileZilla, download the file /usr/local/kafka/config/server.properties from your AWS EC2 instance.
- H.3. Edit the previous file adding the following uncommented text. Where *Your Public DNS Name* can be found in the AWS EC2 dashboard as presented above.

```
...
#####
# Socket Server Settings #####
#
# The address the socket server listens on. It will get the value returned from
# java.net.InetAddress.getCanonicalHostName() if not configured.
# FORMAT:
#   listeners = listener_name://host_name:port
# EXAMPLE:
#   listeners = PLAINTEXT://your.host.name:9092
#listeners=PLAINTEXT://:9092
listeners=PLAINTEXT://Your Public DNS Name:9092
...
```

Hint 4: Public DNS name is changed whenever an AWS EC2 instance reboot is performed, therefore reconfiguration of this file could be needed often.

- H.4. Save the file and upload, using FileZilla, for your AWS EC2 instance in /usr/local/kafka/config/server.properties.

Hint 5 regarding H.2., H.3. and H.4.: another option to change the content of the server.properties file is using the vi application directly in the AWS EC2 environment.

- H.5. Stop Kafka server using the command:

```
sudo /usr/local/kafka/bin/kafka-server-stop.sh -daemon /usr/local/kafka/config/server.properties
```

- H.6. Check that Kafka server is stopped with:

```
ps -ef |grep java|grep kafka
```

- H.7. Restart Kafka with the command, as in F.15.:

```
sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon /usr/local/kafka/config/server.properties
```

- H.8. To test the connection, from your local docker container, from local macOS, local linux, or other AWS EC2 linux environment, type the command to **check if the topic "test" is created successfully in the AWS EC2 instance**:

```
/usr/local/kafka/bin/kafka-topics.sh --bootstrap-server YourPublicDNSName:9092 --describe --topic test
```

- H.9. You are now able to communicate with the AWS EC2 instance. Any other command from section G. can now be tested.

Exercise 3: Create an environment where the messages are produced from one AWS EC2 machine, sent to Kafka server in AWS EC2 instance, and then consumed by another AWS EC2 machine. Check if the flow of messages is executing correctly. Screenshot of the solution:

The screenshot shows three terminal windows on an AWS EC2 instance. A red arrow points from the first window to the third window.

- Terminal 1 (Left):** Shows the creation of a Kafka topic named "EXERCISE3".

```
root@29f9117a8312:~# /usr/local/kafka/bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1 --topic EXERCISE3
Creating topic "EXERCISE3".
root@29f9117a8312:~# /usr/local/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic EXERCISE3
>first message
>second message
>third message
>final message
>
```
- Terminal 2 (Middle):** Shows the output of the Kafka log file "000.log" with four messages: "first message", "second message", "third message", and "final message".

```
[ec2-user@ip-172-31-10-10 ~] $ tail -f /tmp/kafka-logs/EXERCISE3-0/000.log
first message
second message
third message
final message
```
- Terminal 3 (Right):** Shows the consumption of the Kafka topic "EXERCISE3" using the Kafka console consumer.

```
root@29f9117a8312:~# /usr/local/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic EXERCISE3 --from-beginning
from-beginning
first message
second message
third message
final message
```



Stop your AWS instance.

I. Stopping the AWS EC2 instance

The goal of this section is to describe how an AWS EC2 instance can be stopped. This section requires the execution of section B.

- I.1. Go to AWS EC2 dashboard as explained in A.1., and you will find an appearance similar with the following image.

The screenshot shows the AWS EC2 Instances page. On the left, there is a navigation sidebar with various services like EC2 Dashboard, Instances, Network & Security, Load Balancing, and Auto Scaling. The main area shows a table of instances. One row is selected, and its 'Instance state' cell is highlighted with a red box, showing 'Running'. The table includes columns for Name, Instance ID, Instance type, Status check, Alarm status, Availability Zone, Public IPv4 DNS, Public IPv4 IP, and Elastic IP. Below the table, a detailed view of the selected instance ('i-0e47631114bff84b1') is shown. The 'Details' tab is active, displaying sections for Instance summary, Instance details, and Info. The Instance summary section shows the instance ID (i-0e47631114bff84b1), instance state (Running), instance type (t2.micro), and IAM Role. It also displays a warning message about user permissions. The Instance details section provides information such as Platform (Amazon Linux Inferred), AMI ID, AMI name, AMI location, AMI launch index, and Credit specification. The Info section shows monitoring status, termination protection, lifecycle, key pair name, and kernel ID.

- I.2. Stop the instance, checking the row with the desired instance and then pressing "Instance state -> Stop instance state". Wait for the instance State to change to stopped.

Other references

- Connecting to Your Linux Instance Using SSH,
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/AccessingInstancesLinux.html>
- To create a Maven project in eclipse: <https://www.tech-recipes.com/rx/39279/create-a-new-maven-project-in-eclipse/>
- About Kafka group-ids: <https://stackoverflow.com/questions/35561110/can-multiple-kafka-consumers-read-same-message-from-the-partition>
- Where to download apache maven <https://maven.apache.org/download.cgi>
- How to install apache maven <https://maven.apache.org/install.html>