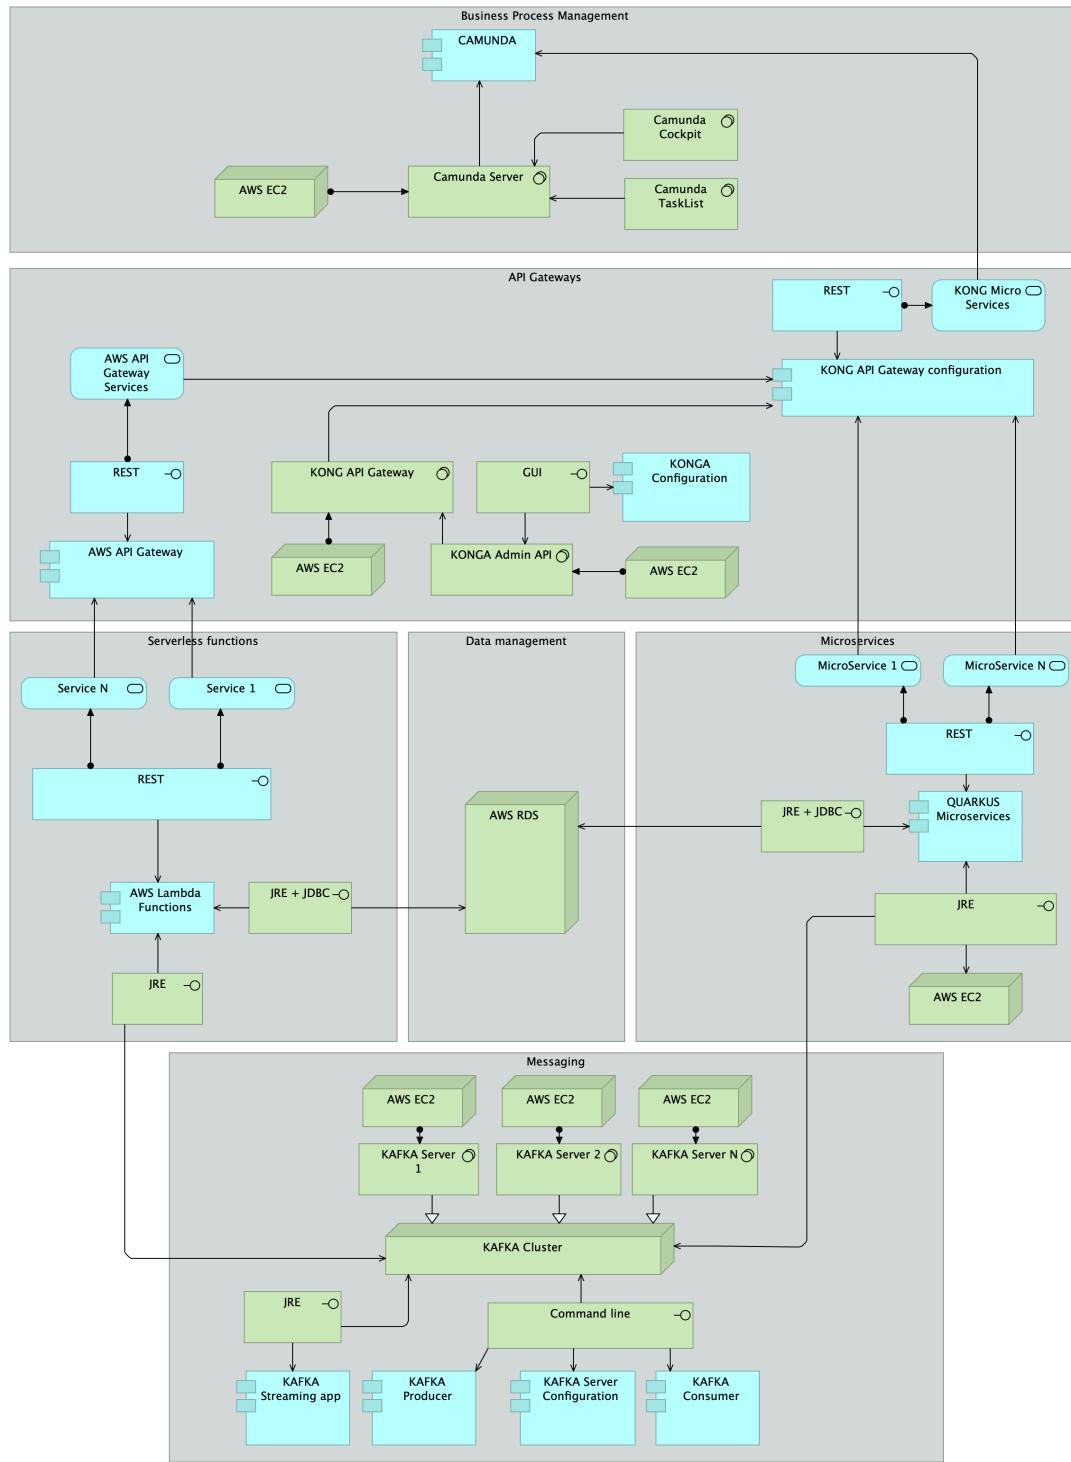


The following tutorials describe, in detail, the expected technological implementation using examples. The target integration architecture is overviewed in the following ArchiMate model (Application and Technological layers).



ArchiMate used elements:

Application Component An application component represents an encapsulation of application functionality aligned to implementation structure, which is modular and replaceable.

Application Interface An application interface represents a point of access where application services are made available to a user, another application component, or a node.

Application Service An application service represents an explicitly defined exposed application behavior.

Node A node represents a computational or physical resource that hosts, manipulates, or interacts with other computational or physical resources.

System Software A system software represents software that provides or contributes to an environment for storing, executing, and using software or data deployed within it.

Technology Interface A technology interface represents a point of access where technology services offered by a node can be accessed.

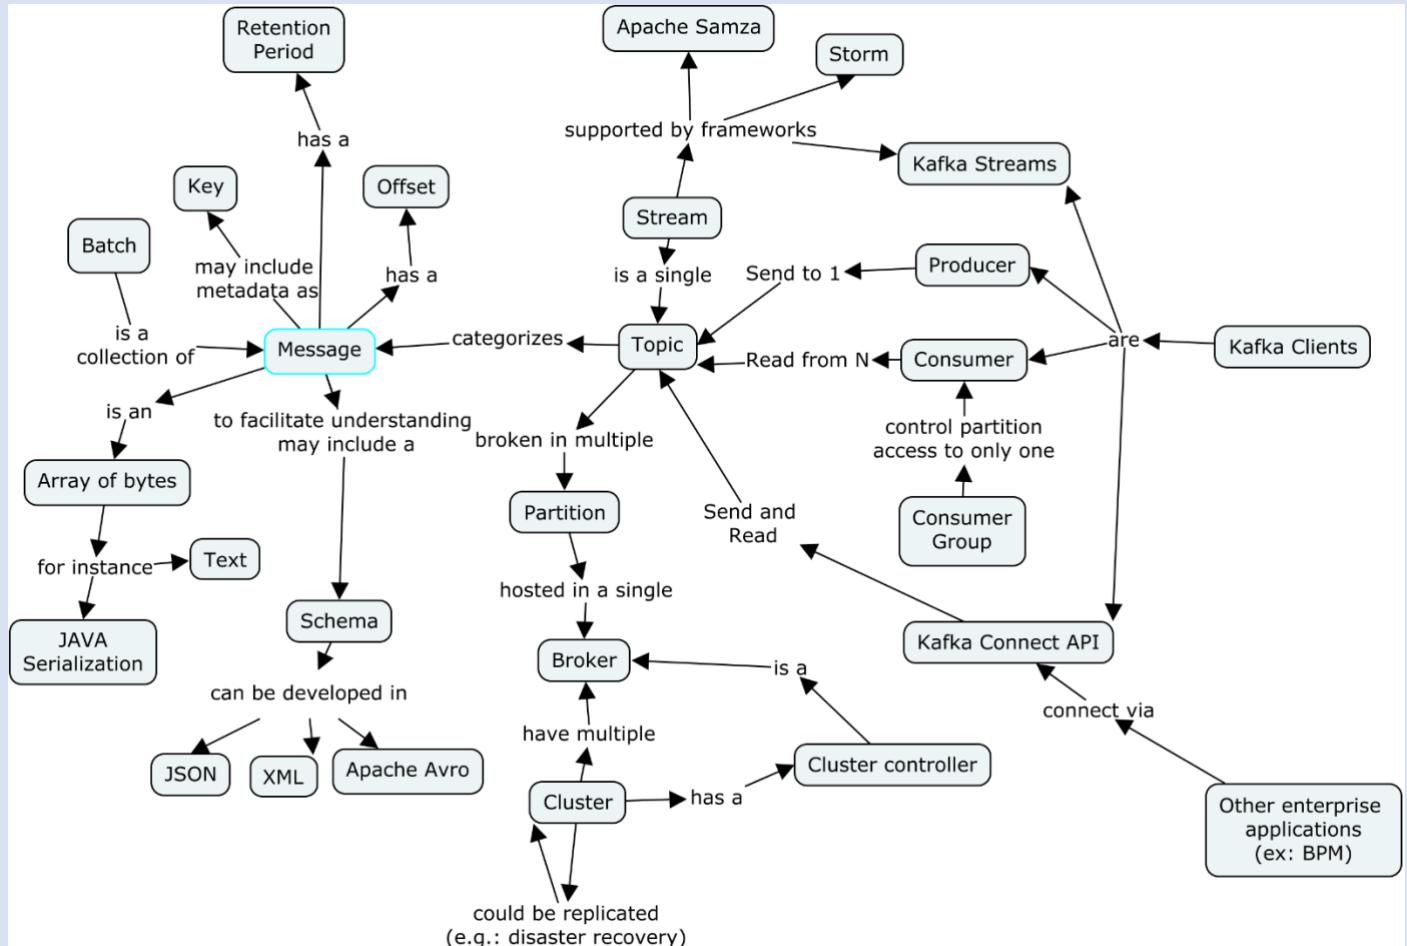
→ The serving relationship represents that an element provides its functionality to another element.

→ The assignment relationship represents the allocation of responsibility, performance of behavior, storage, or execution.

→ The specialization relationship represents that an element is a particular kind of another element.

P1-Kafka-in-AWSAcademy-v3.0	2
P2-HyperAutomation-TERRAFORM-v1.3	24
P3-Kafka-Distributed-v1.5	61
P4-Kafka-Streams-AWS-v1.1	69
P5-RDS-database-v1.1	82
P6-Quarkus_v1.9	87
P7-BPMN-CAMUNDA-v1.11	113
P8-Lambda-AWS-v1.2	149
P9-KONG-v0.8	161

The goal of this document is to show how to operate the basics of a remote Kafka service: installation, starting, accessing, testing, stopping and backup. Remote installation step by step is presented, where the remote installation requires an Amazon AWS account. The following concepts are contained in Kafka:



The following contents is presented in this document.

Contents

A. Creating AWS Academy account and AWS private key	3
B. Creating and launching an AWS EC2 instance	6
C. Access the AWS EC2 instance using PuTTY (for windows only)	10
D. Access the AWS EC2 instance using ssh (for linux/macOS) and scp (for linux/macOS)	12
E. Access the AWS EC2 instance using FileZilla (for windows and macOS)	14

F.	Installing manually Kafka in the AWS EC2 instance.....	15
G.	Testing Kafka locally in the AWS EC2 instance.....	17
H.	Testing Kafka in the AWS EC2 instance remotely	18
I.	Stopping the AWS EC2 instance	21
	Other references.....	22

A. Creating AWS Academy account and AWS private key

The goal of this section is to activate your AWS account and create your private key that will allow you to access your EC2 instance.

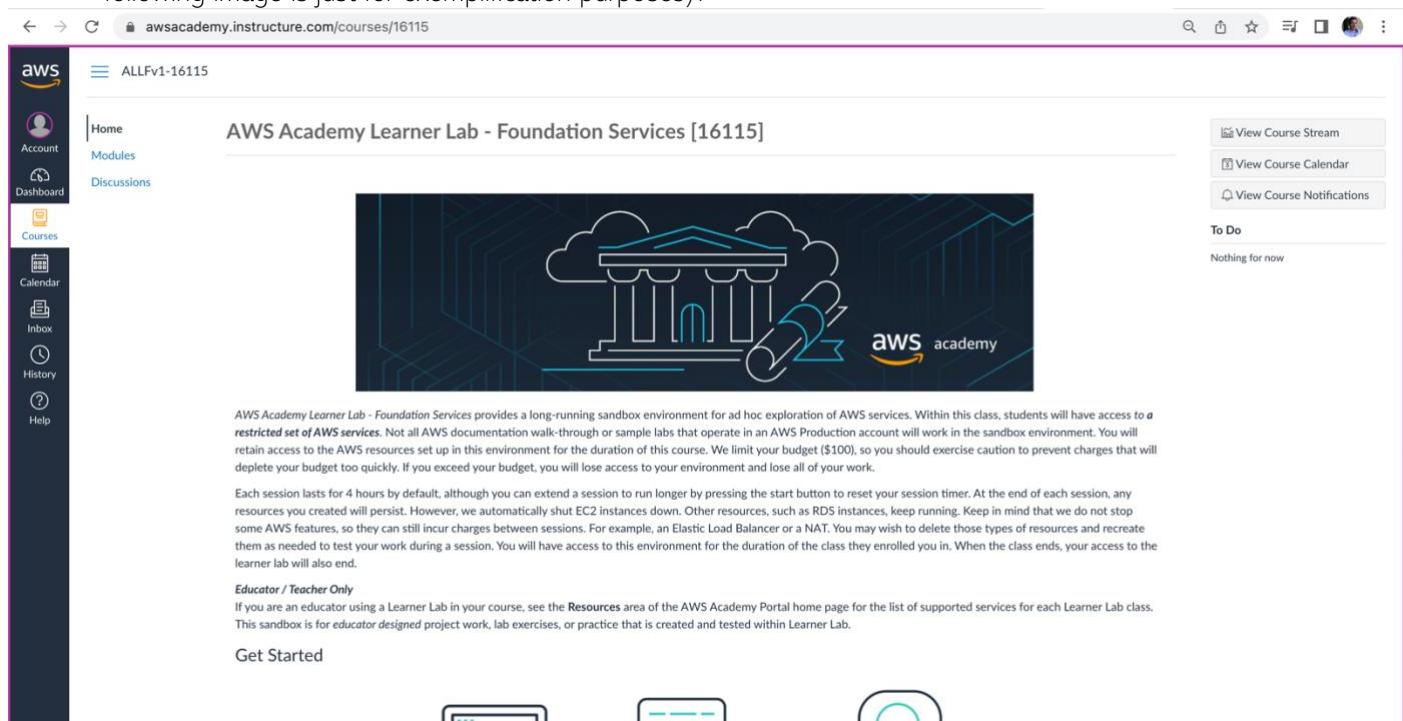
- A.O. Create an amazon AWS Academy account using the invitation email sent by Amazon Academy on the faculties request.

While you are in the process of account creation be careful with your graduation date – define it, at least, accordingly with the end of semester date.

Important remark: the usage of cloud resources implies costs. An EC2 AWS instance is a machine stored in an Amazon AWS facility that is using the usual computational resources: CPU, memory, storage, communications, energy, physical space, among others. This is a reality that any organization faces daily.

Therefore, you should be aware that while your AWS instance is started, your AWS account also starts to be billed.
We absolutely recommend you to always stop your instances when you are not working with them. Always manage your resources wisely!

- A.1. After registration, accept the terms, and then you can access your account and have a similar canvas (the following image is just for exemplification purposes).



The screenshot shows the AWS Academy Learner Lab - Foundation Services [16115] course page. The left sidebar has a dark theme with icons for Account, Courses, Calendar, Inbox, History, and Help. The main content area has a light background. At the top, there's a navigation bar with a search icon, a star icon, and other course-related links. Below the navigation is a large banner featuring a stylized illustration of clouds and a bridge, with the AWS academy logo. The banner text reads: "AWS Academy Learner Lab - Foundation Services provides a long-running sandbox environment for ad hoc exploration of AWS services. Within this class, students will have access to a restricted set of AWS services. Not all AWS documentation walk-through or sample labs that operate in an AWS Production account will work in the sandbox environment. You will retain access to the AWS resources set up in this environment for the duration of this course. We limit your budget (\$100), so you should exercise caution to prevent charges that will deplete your budget too quickly. If you exceed your budget, you will lose access to your environment and lose all of your work." Below the banner, there's a detailed description of the session length, resource persistence, and budget limits. A section for "Educator / Teacher Only" is present, followed by a "Get Started" button and three small decorative icons at the bottom.

- A.2. Select the available course (the following image is just for exemplification purposes):

A.3. Select "Modules" and then "Learner Lab – Foundational Services"

A.4. Click "Start Lab" on the right side

A.5. Then, when the indicator is **green**, click "AWS Details" to access AWS Management Console on the right side. As indicated by the red boxes below. Store the returned key file in your computer for later use. You can choose .pem to use with OpenSSH and/or .ppk to use with PuTTY. If you are unsure store the keys in both formats on your computer.

A.6. Then, click "AWS" to access AWS Management Console on the left side. As indicated by the red box below.

The screenshot shows the AWS Academy interface. On the left, there's a sidebar with icons for Account, Courses, Calendar, Inbox, History, and Help. The main area has a breadcrumb trail: ALLFv1-16115 > Modules > Learner Lab Foundation Services > Learner Lab - Foundational Services. A red box highlights the "AWS" link in the top navigation bar. To the right, there's a terminal window showing a command prompt and some text, followed by a "Learner Lab - Foundational Level" section with various links.

A.7. Then, click "AWS Details" to consult all the credentials to access your AWS area.

The screenshot shows the AWS Academy interface. The sidebar includes Account, Courses, Calendar, Inbox, History, and Help. The main area has a breadcrumb trail: awsacademy.instructure.com/courses/107105/modules/items/10016408. A red box highlights the "AWS Details" link in the top navigation bar. The "AWS Details" section displays AWS CLI credentials (access key ID, secret access key, session token) and session details like start and end times, accumulated lab time, and region information.

A.8. In the left upper corner, select "Services -> Compute (EC2)". You can also search for EC2.

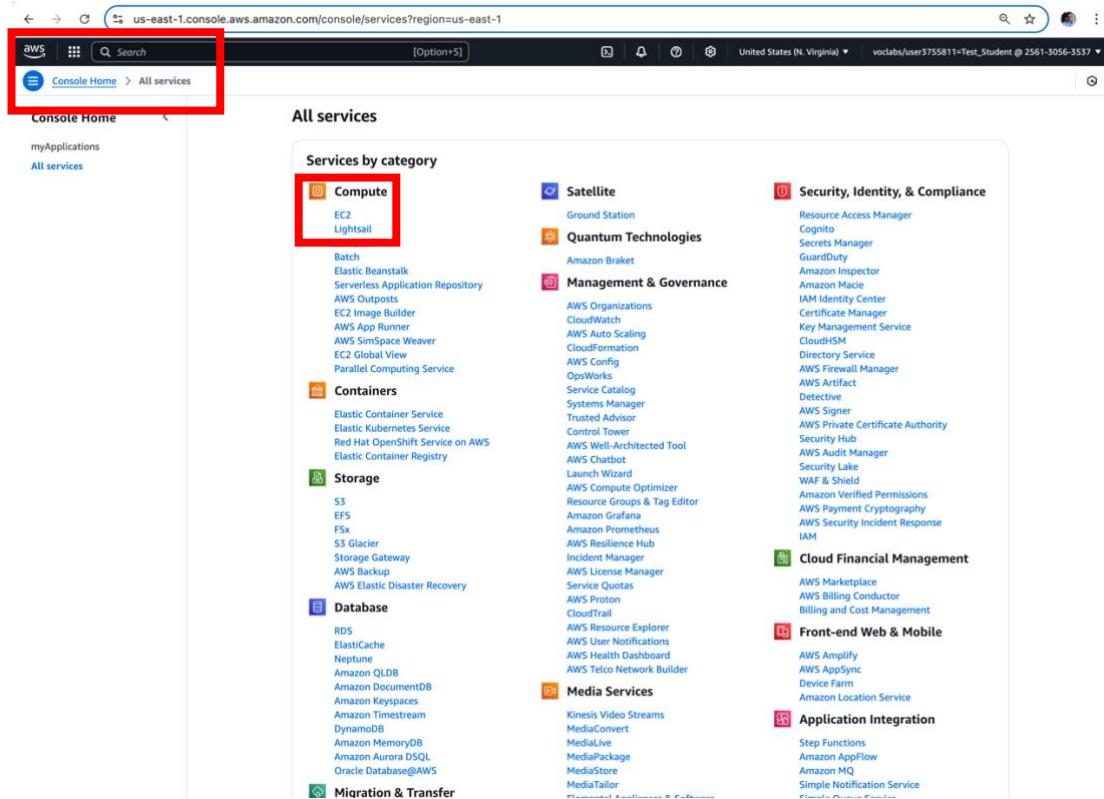
The screenshot shows the AWS Management Console. The left sidebar has a "Services" dropdown menu with "Compute" selected, indicated by a red box. Other services listed include EC2, Lambda, EKS, ECR, Storage, Database, and Media Services. The main pane shows a grid of AWS services, with the "Compute" section highlighted in red.

You will obtain the AWS EC2 dashboard. All the options are vertically organized in the left side.

B. Creating and launching an AWS EC2 instance

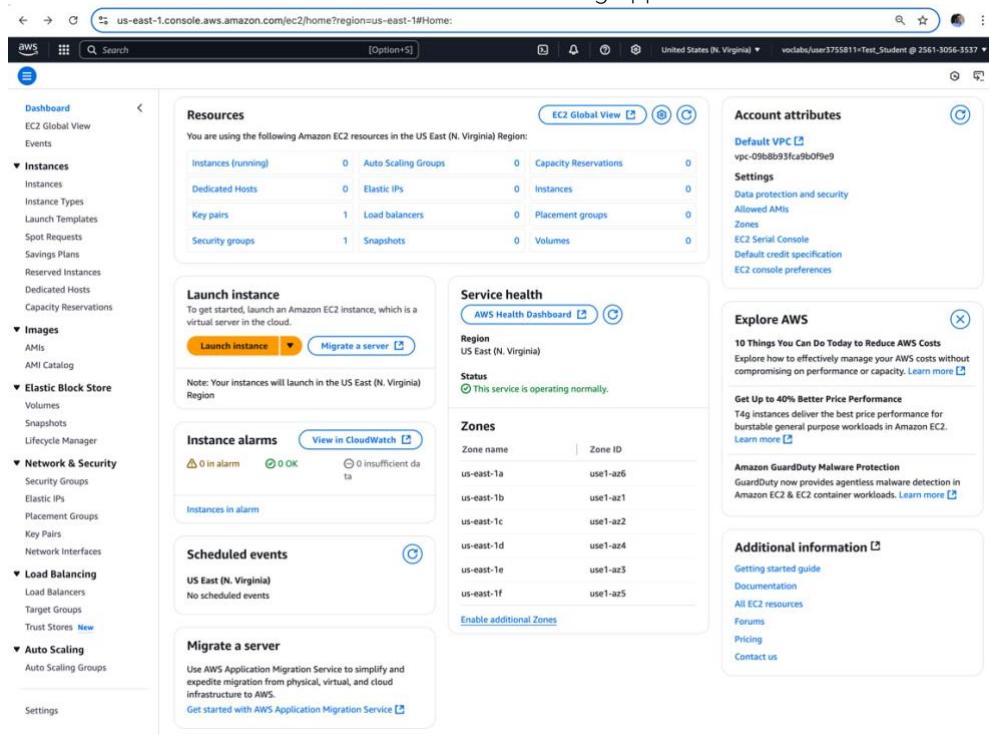
The goal of this section is to create your first EC2 instance using the available Amazon Machine Images (AMI) available. This requires the previous execution of section A.

- B.1. Login in your account, then access the AWS console, and in the left upper corner, select “A11 Services -> Compute (EC2)”.



The screenshot shows the AWS Console homepage with a red box highlighting the 'Compute' category under 'Services by category'. Other categories like 'Storage', 'Database', and 'Migration & Transfer' are also visible. The 'Compute' category contains links for EC2, Lightsail, and other services.

- B.2. Select “Instances -> Instances”. The following appearance will be shown.



The screenshot shows the AWS EC2 Instances page. A red box highlights the 'Launch instance' button. The page displays various EC2 resources and settings, including Auto Scaling Groups, Capacity Reservations, and Account attributes.

B.3. Select "Launch Instance" and a similar interface will be presented as shown below. The goal is to define the Amazon Machine Image (AMI) that will contain all the software configuration required to launch your instance.

The screenshot shows the AWS EC2 'Launch an instance' wizard. The top navigation bar includes links for 'Search', '[Option+S]', 'United States (N. Virginia)', and 'voclabs/user3755811=Test_Student @ 2561-3056-3537'. Below the navigation is a breadcrumb trail: 'EC2 > Instances > Launch an instance'.

Name and tags

Name: e.g. My Web Server [Add additional tags](#)

Application and OS Images (Amazon Machine Image)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below.

[Search our full catalog including 1000s of application and OS images](#)

Quick Start

Quick start options include: Amazon Linux, macOS, Ubuntu, Windows, Red Hat, SUSE Linux, and Debian.

Amazon Machine Image (AMI)

Selected: Amazon Linux 2023 AMI
AMI ID: ami-05576a079321f21f8
Virtualization: hvm ENA enabled: true Root device type: ebs [Free tier eligible](#)

Description

Amazon Linux 2023 is a modern, general purpose Linux-based OS that comes with 5 years of long term support. It is optimized for AWS and designed to provide a secure, stable and high-performance execution environment to develop and run your cloud applications.

Amazon Linux 2023 AMI 2023.6.20250107.0 x86_64 HVM kernel-6.1.10-100.1.1.amzn2.0.1-1.1.1

Instance type

Selected: t2.micro [Free tier eligible](#)

Family: t2 1 vCPU 1 GiB Memory Current generation: true
On-Demand Windows base pricing: 0.0162 USD per Hour
On-Demand Ubuntu Pro base pricing: 0.0134 USD per Hour On-Demand SUSE base pricing: 0.0116 USD per Hour
On-Demand RHEL base pricing: 0.026 USD per Hour On-Demand Linux base pricing: 0.0116 USD per Hour

[All generations](#) [Compare instance types](#)

Summary

Number of instances: 1

Software Image (AMI)
Amazon Linux 2023 AMI 2023.6.2... [read more](#)
ami-05576a079321f21f8

Virtual server type (instance type)
t2.micro

Firewall (security group)
New security group

Storage (volumes)
1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the internet.

[Cancel](#) [Launch instance](#) [Preview code](#)

B.4. Use the following configurations:

Step 1 – Choose an Amazon Machine Image (AMI)	select "Amazon Linux 2 AMI (HVM) - Kernel 5.10, SSD Volume Type" (available on January 2025) with "64-bit (x86)"
Step 2 – Choose an Instance Type	select "t2.micro"
Step 3 – Configure Instance Details	check if number of instances = 1 Keep all the remaining options as default
Step 4 – Storage	keep the default size
Step 5 – Key Pair (login)	Key Pair name = vockey
Step 6 – Add Tags	Key = application Value = kafka
Step 7 – Configure Security Group <i>(it could be configured later in the instance console or</i>	Select "Create a new security group" Security group name: launch-Kafka Select "Add rule" Type = "Custom TCP Rule" Protocol = TCP Port range = 2181

<p><i>since the creation using terraform)</i></p>	<p>Source = Anywhere (0.0.0.0, ::/0) Description = zookeeper Select "Add rule" Type = "Custom TCP Rule" Protocol = TCP Port range = 9092 Source = Anywhere (0.0.0.0, ::/0) Description = kafka</p> <p>Select "Add rule" Type = "SSH" Protocol = TCP Port range = 22 Source = Anywhere (0.0.0.0, ::/0) Description = remote access</p>
<p>Step 8 – Review Instance Launch</p>	

B.5. Press "Launch" in the bottom of the page



Your AWS account starts to be billed at this point.

B.6. Go to EC2 dashboard as explained in A.1., and you will find an appearance similar with the following image.

The screenshot shows the AWS EC2 Instances dashboard. On the left, there's a navigation sidebar with options like Dashboard, EC2 Global View, Events, Instances (selected), Images, Elastic Block Store, Network & Security, Load Balancing, and Auto Scaling. The main area displays a table titled 'Instances (1/1) Info' with one row. The row contains the instance name 'teste kafka', instance ID 'i-064b2df1db33ddcde', and the 'Instance state' column which shows 'Running' with a green status indicator. A red box highlights this 'Running' status. Below the table, there's a detailed view for the instance 'i-064b2df1db33ddcde (teste kafka)'. This view includes sections for Instance summary, IAM Role, IMDSv2, Operator, Instance details (AMI ID: ami-0454e52560c7f5c55), and Platform details (Linux/UNIX). The Public IPv4 address listed is 52.23.182.107.

You can check the IPv4 and v6 public address, and the state of the instance. In this example it is *running*.

B.7. Stop the instance, checking the row with the desired instance and then pressing "Instance state -> Stop". Wait for the instance State to change to stopped.

B.8. To allow a proper future execution of Kafka, change the type of instance, by choosing the following options, and then choose **t2.small**.

The screenshot shows the AWS EC2 Instances page. On the left, there's a sidebar with various navigation links like EC2 Dashboard, Events, Tags, Reports, Limits, Instances, Launch Templates, Spot Requests, Reserved Instances, Dedicated Hosts, Scheduled Instances, Capacity Reservations, Images, AMIs, Bundle Tasks, Elastic Block Store, Volumes, Snapshots, Lifecycle Manager, Network & Security, Security Groups, Elastic IPs, Placement Groups, Key Pairs, and Network Interfaces. The main area has tabs for Launch Instance, Connect, and Actions. Under Actions, a dropdown menu is open with options: Connect, Get Windows Password, Create Template From Instance, Launch More Like This, Instance Settings (which is currently selected), Add/Edit Tags, Attach to Auto Scaling Group, Attach/Replace IAM Role, Change Instance Type, Change Termination Protection, View/Change User Data, Change Shutdown Behavior, Change T2/T3 Unlimited, Get System Log, Get Instance Snapshot, Modify Instance Placement, and Modify Capacity Reservation Settings. Below this, there's a detailed view of an instance: Instance ID: i-03a6db7323ba18361, Private IP: 172.31.30.91. The instance state is stopped. The instance type is t2.small. It has one elastic IP assigned: us-east-1a. The AMI ID is amazon-ami-hvm-2018.03.0.20181129-x86_64-gp2 (ami-0080e4c5bc079760e). The owner is 062896652389. The launch time was February 25, 2019 at 1:13:28 PM UTC. The termination protection is False. The monitoring status is basic. There are no status checks or alarms. The instance is located in availability zone east-1b. The public DNS is ip-172-31-30-91.ec2.internal and the private DNS is 172.31.30.91. The secondary private IP is 172.31.30.91. The VPC ID is vpc-9f1cb283, Subnet ID is subnet-0d9a2747, and Network interfaces are eth0. The source/dest check is True, and T2/T3 Unlimited is Disabled. The EBS-optimized is False, Root device type is ebs, and the root device is /dev/xvda. There are no block devices, elastic graphics ID, or elastic inference accelerator ID. At the bottom, there are links for Feedback, English (US), and footer links for © 2008 - 2019, Amazon Web Services, Inc. or its affiliates. All rights reserved., Privacy Policy, and Terms of Use.

C. Access the AWS EC2 instance using PuTTY (for windows only)

The goal of this section is to describe how your AWS EC2 instance can be accessed using an SSH client. It will be useful to install, configure and manage the Kafka server.

- C.1. Install PuTTY from <https://www.chiark.greenend.org.uk/~sgtatham/putty/latest.html>
- C.2. (*optional – only if the key is a .pem file*) Convert your AWS private key created in section A, using PuTTYgen. For that, execute PuTTYgen from start menu.
- C.3. (*optional – only if the key is a .pem file*) Then, in type of key to generate choose: RSA
- C.4. (*optional – only if the key is a .pem file*) Press “Load” and locate your AWS private key stored in step A.4. (with .pem file extension)
- C.5. (*optional – only if the key is a .pem file*) Press “Save private key” and choose “Yes” when asking to save without passphrase
- C.6. (*optional – only if the key is a .pem file*) Store with the same name as the AWS private key (with .ppk file extension)
- C.7. Start PuTTY from start menu
- C.8. Fill hostname with “user_name@public_dns_name”, where user_name = ec2-user and public_dns_name= <as presented in B.7.>



Public DNS name is changed whenever an AWS EC2 instance reboot is performed.

- C.9. Choose SSH connection type and check if port = 22
- C.10. You can configure PuTTY to automatically send "keepalive" data at regular intervals to keep the session active. This is useful to avoid disconnecting the instance from session inactivity. In the Category panel, choose Connection and enter the required range in the Seconds between keepalives field. For example, if the session disconnects after 10 minutes of inactivity, enter 180 to configure PuTTY to send keepalive data every 3 minutes.
- C.11. In the Category panel, choose Connection, choose SSH, choose Auth, select “Browse” and choose the .ppk file from step C.6.



Start your AWS instance.

- C.12. In the Category panel, choose Session, and then select “Open”
- C.13. Confirm the certificate in the dialog box. You will be presented with the following similar screen:

```
Using username "ec2-user".
Authenticating with public key "imported-openssh-key"
Last login: Thu Dec 27 11:43:38 2018 from
[ec2-user@ip-10-0-2-11 ~]$
```

- C.14. Update your AWS instance with the command:

```
sudo yum update
```

- C.15. Check JDK packages available with the command:

```
sudo yum search "java-17"
```

- C.16. Execute the command (or another equivalent package)

```
sudo yum install java-17-amazon-corretto-devel.x86_64
```

C.17. Check if JAVA is currently version 17 with the command:

```
java -version
```

If any other version is referred you may remove it, e.g., first check the installed java versions with the command:

```
yum list installed |grep java
```

and then remove the ones that you don't need with the command

```
sudo yum remove java-1.7.0-openjdk
```

C.18. Execute the command:

```
exit
```



Stop your AWS instance.

D. Access the AWS EC2 instance using ssh (for linux/macOS) and scp (for linux/macOS)

The goal of this section is to describe how your AWS EC2 instance can be accessed using an SSH client in linux. It will be useful to install, configure and manage the Kafka server.



Start your AWS instance.

Public DNS name is changed whenever an AWS EC2 instance reboot is performed.

D.1. In a new terminal session, change directories to the location of the private key file that you created when you created the AWS instance, for instance:

```
[oracle@soabpm-vm ~]$ cd /media/sf_SharedFolderForOracleSOA
```

C.19. Use the following command to set the permissions of your private key file so that only you can read it.

```
[oracle@soabpm-vm sf_SharedFolderForOracleSOA]$ chmod u=rwx,g=,o= myKeyAWS.pem
```

C.20. Use the ssh command to connect to the instance. You specify the private key (.pem) file and user_name@public_dns_name. For example, if you used Amazon Linux 2 or the Amazon Linux AMI, the user name is ec2-user. In

```
[oracle@soabpm-vm sf_SharedFolderForOracleSOA]$ ssh -i myKeyAWS.pem ec2-user@YOURIP
The authenticity of host 'X.X.X.X (X.X.X.X)' can't be established.
RSA key fingerprint is 25:1c:ef:ee:c1:87:61:e9:ea:f6:e3:45:1e:3a:63:73.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'X.X.X.X' (RSA) to the list of known hosts.
Last login: Fri Feb  8 14:59:07 2019 from Y.Y.Y.Y

      _\   _ )   Amazon Linux AMI
     __| \__|_ |
```

https://aws.amazon.com/amazon-linux-ami/2018.03-release-notes/
5 package(s) needed for security, out of 6 available
Run "sudo yum update" to apply all updates.
ZooKeeper JMX enabled by default
Using config: /usr/local/zookeeper/bin/../conf/zoo.cfg
Starting zookeeper ... already running as process 11211.
[ec2-user@ip-X-X-X-X ~]\$

C.21. Update your AWS instance with the command:

```
sudo yum update
```

C.22. Check JDK packages available with the command:

```
sudo yum search "java-17"
```

C.23. Execute the command (or another equivalent package)

```
sudo yum install java-17-amazon-corretto-devel.x86_64
```

C.24. Check if JAVA is currently version 17 with the command:

```
java -version
```

If any other version is referred you may remove it, e.g., first check the installed java versions with the command:

```
yum list installed |grep java
```

and then remove the ones that you don't need with the command

```
sudo yum remove java-1.7.0-openjdk
```

C.25. Execute the command:

```
exit
```

C.26. To transfer files from your machine to your AWS instance, use the scp application in linux. For instance:

```
scp -i my-key-pair.pem /path/SampleFile.txt ec2-user@c2-198-51-100-1.compute-1.amazonaws.com:~
```

C.27. To transfer files from your AWS instance to your machine, use the scp application in linux. For instance:

```
scp -i my-key-pair.pem ec2-user@ec2-198-51-100-1.compute-1.amazonaws.com:~/SampleFile.txt ~/SampleFile2.txt
```

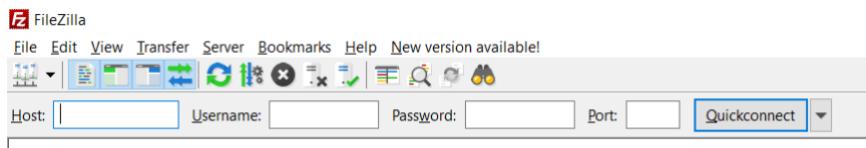


Stop your AWS instance.

E. Access the AWS EC2 instance using FileZilla (for windows and macOS)

The goal of this section is to describe how your AWS EC2 instance can be accessed using an SFTP client. It will be useful to transfer files between your computer and AWS EC2 instance and back.

- E.1. Install FileZilla from <https://filezilla-project.org/download.php?type=client>
- E.2. Open FileZilla.
- E.3. Go to the FileZilla Edit -> settings, and on the left, click SFTP.
- E.4. Add a new private key. (Your .ppk key)
- E.5. If you are using a .pem key you must convert it accordingly with steps C.2. to C.6., otherwise it will not work
- E.6. At the top in the Quickconnect bar (similar with the image below), put your Public DNS in the host, ec2-user, port 22 (Port 22 is SFTP rather than FTP, AWS will kick back FTP.), and NO PASSWORD.



Start your AWS instance.

- E.7. Click Quickconnect.
- E.8. File transfer to AWS EC2 instance is only a matter of dragging files from the left to right side.



Stop your AWS instance.

F. Installing manually Kafka in the AWS EC2 instance

The goal of this section is to describe the required steps to install the Zookeeper and Kafka servers in your AWS EC2 instance.



Start your AWS instance.

F.1. Access your AWS EC2 instance using previous steps of this tutorial.

F.2. To change for your home directory, type the command:

```
cd
```

F.3. Download the Kafka binary version from <https://kafka.apache.org/downloads>. Recommended version is kafka_2.13-3.9.0.tgz (in January 2025).

```
wget https://downloads.apache.org/kafka/3.9.0/kafka_2.13-3.9.0.tgz
```

F.4. Download the ZooKeeper binary version from <https://zookeeper.apache.org/releases.html>. Recommended version is apache-zookeeper-3.9.3-bin.tar.gz (in January 2025).

```
wget https://dlcdn.apache.org/zookeeper/zookeeper-3.9.3/apache-zookeeper-3.9.3-bin.tar.gz
```

F.5. Then, type the command:

```
ls -ltr
```

to check if you have both kafka and zookeeper files available.

F.6. Update your AWS instance with the command:

```
sudo yum update
```

F.7. Execute the command (or another equivalent package)

```
sudo yum install java-17-amazon-corretto-devel.x86_64
```

F.8. To extract the zookeeper files, type the command:

```
tar -zxf apache-zookeeper-3.9.3-bin.tar.gz
```

F.9. Move the new zookeeper directory to system typing the command:

```
sudo mv apache-zookeeper-3.9.3-bin /usr/local/zookeeper
```

F.10. Create a new directory using the command:

```
sudo mkdir -p /var/lib/zookeeper
```

F.11. Create the baseline zookeeper server configuration with the command:

```
cat > /usr/local/zookeeper/conf/zoo.cfg << EOF
```

and then write the following content to the file directly in the command line:

```
tickTime=2000  
dataDir=/var/lib/zookeeper  
clientPort=2181  
EOF
```

F.12. Start ZooKeeper server with the command:

```
sudo /usr/local/zookeeper/bin/zkServer.sh start
```

F.13. To extract the kafka server, type the command:

```
tar -zxf kafka_2.13-3.9.0.tgz
```

F.14. Move the new kafka directory to system typing the command:

```
sudo mv kafka_2.13-3.9.0 /usr/local/kafka
```

F.15. Create a new directory using the command:

```
sudo mkdir /tmp/kafka-logs
```

F.16. Start Kafka with the command:

```
sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon /usr/local/kafka/config/server.properties
```

Hint 1: if you are short in the instance memory reduce the JVM usage with the command:

```
export KAFKA_HEAP_OPTS="-Xmx256M -Xms256M"
```

Hint 2: you can always check if Zookeeper and Kafka servers are started correctly with the command:

```
ps -ef |grep java
```



Stop your AWS instance.

Exercise 1: To automate the start of Zookeeper and Kafka servers with instance login write both starting commands in the file:

```
.bash_profile
```

G. Testing Kafka locally in the AWS EC2 instance

The goal of this section is to test if the previous installation in section E. was executed successfully. The production and consumption of Kafka messages are tested. It is considered that in this environment the servers and the clients are installed in the same machine.



Start your AWS instance.

G.1. Access your AWS EC2 instance using the steps C.7. to C.12.

G.2. First, install telnet package in your AWS EC2 instance typing the command:

```
sudo yum install telnet.x86_64
```

G.3. Test ZooKeeper installation typing the command:

```
telnet localhost 2181
```

then write the command:

```
srvr
```

and it is expected to receive the following output:

```
[ec2-user@ ~]$ telnet localhost 2181
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
srvr
Zookeeper version: 3.4.12-e5259e437540f349646870ea94dc2658c4e44b3b, built on 03/27/2018 03:55 GMT
Latency min/avg/max: 0/0/0
Received: 1
Sent: 0
Connections: 1
Outstanding: 0
Zxid: 0x0
Mode: standalone
Node count: 4
Connection closed by foreign host.
```

G.4. Test Kafka installation, with the creation of a new topic named "test", typing the command:

```
sudo /usr/local/kafka/bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --
-partitions 1 --topic test
```

G.5. Type the command to check if the topic "test" is created successfully:

```
sudo /usr/local/kafka/bin/kafka-topics.sh --bootstrap-server localhost:9092 --describe --topic test
```

G.6. Type the command to produce messages to the topic "test":

```
sudo /usr/local/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic test
```

And write some text messages, one per line. Then, to finish press Ctrl+D.

G.7. Type the command to consume messages from the topic "test":

```
sudo /usr/local/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic test --from-
beginning
```

All the sent messages from topic "test" will be presented. Then, to finish press Ctrl+C.

Exercise 2: Try creating different PuTTY sessions for producer and consumer and check the run-time production and consumption of messages.



Stop your AWS instance.

H. Testing Kafka in the AWS EC2 instance remotely

The goal of this section is to test if the previous installation in section E. was executed successfully. The production and consumption of Kafka messages are tested. It is considered that in this environment the servers and the clients are installed in different machines.

This section requires the previous installation of Kafka, as described in section E.



Start your AWS instance.

- H.0. Access your AWS EC2 instance using the steps C.7. to C.12.
H.1. From your local docker container, from local macOS, local linux, or other AWS EC2 linux environment, check if you can connect to the AWS EC2 instance Zookeeper server using the following command:

```
telnet YourPublicIP 2181
```

remember that *YourPublicIP* can be found in the AWS EC2 dashboard as presented below:

The screenshot shows the AWS EC2 Instances page. A search bar at the top right contains the text "Successfully started i-0e47631114bff84b1". The main table has one row for an instance named "i-0e47631114bff84b1" which is "Running" on a "t2.micro" instance type. The Public IPv4 DNS is listed as "ec2-34-203-227-188.compute-1.amazonaws.com". Below the table, the "Instance: i-0e47631114bff84b1" details page is expanded. The "Details" tab is selected, showing the Public IPv4 address as "34.203.227.188" and the Public IPv4 DNS as "ec2-34-203-227-188.compute-1.amazonaws.com". Other tabs include Security, Networking, Storage, Status Checks, Monitoring, and Tags.

After that, send the command:

```
SRVR
```

You should be able to see an interaction similar with the following:

```
root@29f9117a8312:~# telnet [REDACTED] 2181
Trying [REDACTED] ...
Connected to [REDACTED]
Escape character is '^J'.
srvr
Zookeeper version: 3.4.12-e5259e437540f349646870ea94dc2658c4e44b3b, built on 03/27/2018 03:55 GMT
Latency min/avg/max: 0/0/7
Received: 6819
Sent: 6822
Connections: 2
Outstanding: 0
Zxid: 0xee
Mode: standalone
Node count: 133
Connection closed by foreign host.
```

Hint 3: Public IP or Public DNS Name could be.

- H.2. Using the FileZilla, download the file /usr/local/kafka/config/server.properties from your AWS EC2 instance.
- H.3. Edit the previous file adding the following uncommented text. Where *Your Public DNS Name* can be found in the AWS EC2 dashboard as presented above.

```
...
#####
# Socket Server Settings #####
#
# The address the socket server listens on. It will get the value returned from
# java.net.InetAddress.getCanonicalHostName() if not configured.
# FORMAT:
#   listeners = listener_name://host_name:port
# EXAMPLE:
#   listeners = PLAINTEXT://your.host.name:9092
#listeners=PLAINTEXT://:9092
listeners=PLAINTEXT://Your Public DNS Name:9092
...
```

Hint 4: Public DNS name is changed whenever an AWS EC2 instance reboot is performed, therefore reconfiguration of this file could be needed often.

- H.4. Save the file and upload, using FileZilla, for your AWS EC2 instance in /usr/local/kafka/config/server.properties.

Hint 5 regarding H.2., H.3. and H.4.: another option to change the content of the server.properties file is using the vi application directly in the AWS EC2 environment.

- H.5. Stop Kafka server using the command:

```
sudo /usr/local/kafka/bin/kafka-server-stop.sh -daemon /usr/local/kafka/config/server.properties
```

- H.6. Check that Kafka server is stopped with:

```
ps -ef |grep java|grep kafka
```

- H.7. Restart Kafka with the command, as in F.15.:

```
sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon /usr/local/kafka/config/server.properties
```

- H.8. To test the connection, from your local docker container, from local macOS, local linux, or other AWS EC2 linux environment, type the command to **check if the topic "test" is created successfully in the AWS EC2 instance**:

```
/usr/local/kafka/bin/kafka-topics.sh --bootstrap-server YourPublicDNSName:9092 --describe --topic test
```

- H.9. You are now able to communicate with the AWS EC2 instance. Any other command from section G. can now be tested.

Exercise 3: Create an environment where the messages are produced from one AWS EC2 machine, sent to Kafka server in AWS EC2 instance, and then consumed by another AWS EC2 machine. Check if the flow of messages is executing correctly. Screenshot of the solution:

The screenshot shows three terminal windows on an AWS EC2 instance. A red arrow points from the first window to the third window.

- Terminal 1 (Left):** Shows the creation of a Kafka topic named "EXERCISE3".

```
root@29f9117a8312:~# /usr/local/kafka/bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1 --topic EXERCISE3
Creating topic "EXERCISE3".
root@29f9117a8312:~# /usr/local/kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic EXERCISE3
>first message
>second message
>third message
>final message
>
```
- Terminal 2 (Middle):** Shows the output of the Kafka log for the "EXERCISE3-0" partition.

```
[ec2-user@ip-172-31-10-10 ~] $ tail -f /tmp/kafka-logs/EXERCISE3-0/000.log
EMA[...]&first messageE3V!Rg[pbg[pb
EMA[...]&second message<
EMA[...]&(second messageE3V!Rg[pbg[pb
EMA[...]&third messageE3V!Rg[pbg[pb
EMA[...]&final messageE3V!Rg[pbg[pb
```
- Terminal 3 (Right):** Shows the consumption of the messages from the Kafka topic.

```
root@29f9117a8312:~# /usr/local/kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic EXERCISE3 --from-beginning
from-beginning
first message
second message
third message
final message
```



Stop your AWS instance.

I. Stopping the AWS EC2 instance

The goal of this section is to describe how an AWS EC2 instance can be stopped. This section requires the execution of section B.

- I.1. Go to AWS EC2 dashboard as explained in A.1., and you will find an appearance similar with the following image.

The screenshot shows the AWS EC2 Instances page. On the left, there is a navigation sidebar with various services like EC2 Dashboard, Instances, Network & Security, Load Balancing, and Auto Scaling. The main area shows a table of instances. One row is selected, and its 'Instance state' cell is highlighted with a red box. The table includes columns for Name, Instance ID, Instance type, Status check, Alarm status, Availability Zone, Public IPv4 DNS, Public IPv4 IP, and Elastic IP. Below the table, a detailed view of the selected instance ('i-0e47631114bff84b1') is shown. The 'Details' tab is active, displaying sections for Instance summary, Instance details, and Info. The Instance summary section shows the instance ID (i-0e47631114bff84b1), instance state (Running), instance type (t2.micro), and IAM Role. It also displays a warning message about user permissions. The Instance details section provides information such as Platform (Amazon Linux Inferred), Platform details (Linux/UNIX), Launch time (Thu Jan 21 2021 17:39:17 GMT+0000), Stop-hibernate behavior (disabled), and State transition reason. The Info section shows monitoring status (disabled), termination protection (Disabled), lifecycle (normal), key pair name (ElPreparation2021), and kernel ID.

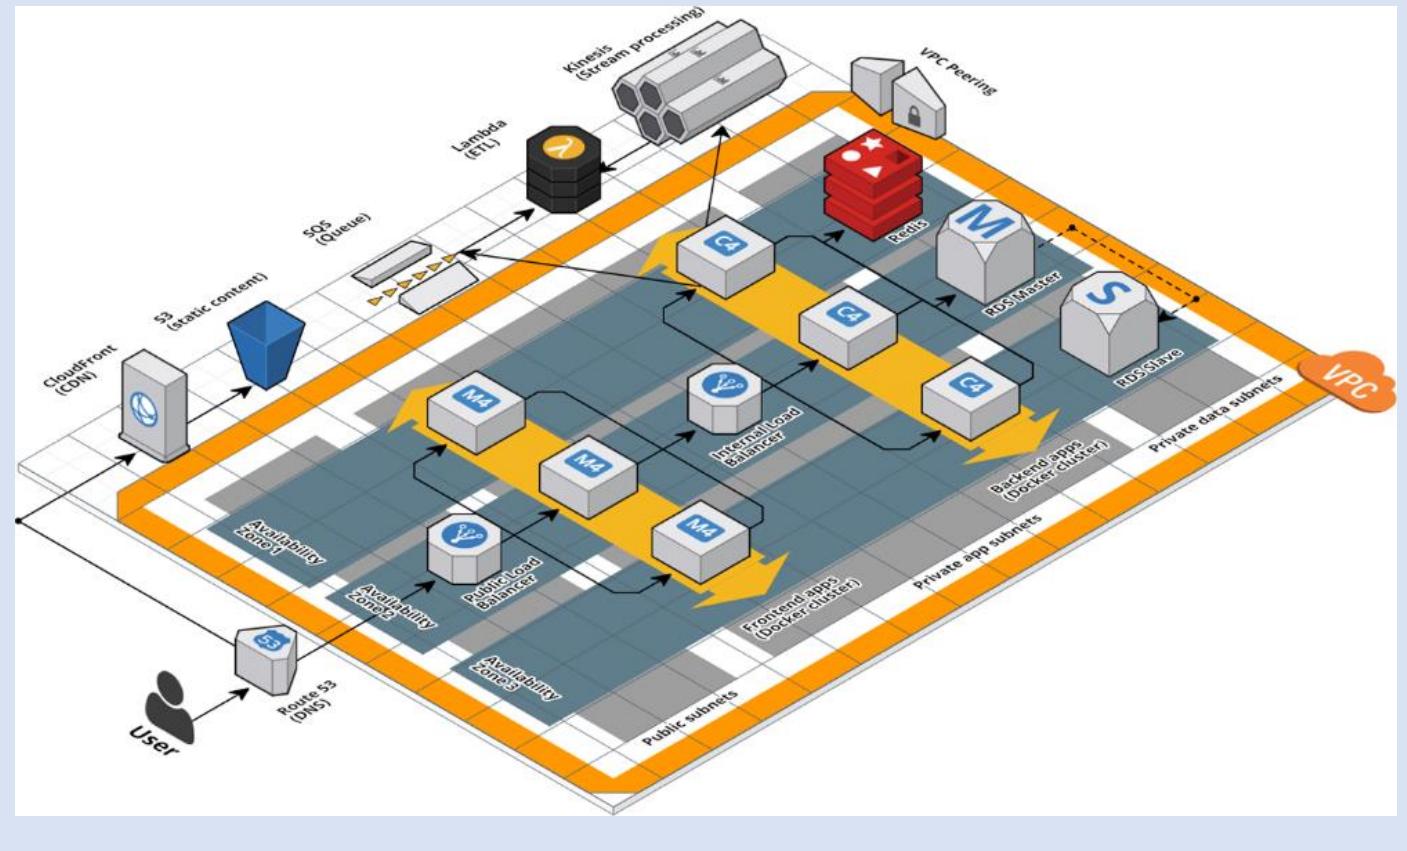
- I.2. Stop the instance, checking the row with the desired instance and then pressing "Instance state -> Stop instance state". Wait for the instance State to change to stopped.

Other references

- Connecting to Your Linux Instance Using SSH,
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/AccessingInstancesLinux.html>
- To create a Maven project in eclipse: <https://www.tech-recipes.com/rx/39279/create-a-new-maven-project-in-eclipse/>
- About Kafka group-ids: <https://stackoverflow.com/questions/35561110/can-multiple-kafka-consumers-read-same-message-from-the-partition>
- Where to download apache maven <https://maven.apache.org/download.cgi>
- How to install apache maven <https://maven.apache.org/install.html>

The goal of this document is to present the details related with Terraform. How to use it using a step-by-step approach. Terraform is a solution to create cloud environments using code instead of using the cloud providers User Interfaces. Therefore, the management of the cloud resources takes less effort and the creation, or update, process is faster.

The following figure taken the main literature reference in this field: *Terraform Up & Running, Writing Infrastructure as Code, Yevgeniy Brikman, 3rd edition (2022)*" exemplifies a complex infrastructure that could be created using Terraform.



Contents

A. Terraform installation procedure	3
B. Obtaining the access for sandbox or AWS academy learner lab	4
C. Deploying a single EC2 instance	6
D. Deploying a single EC2 instance and uploading files	7
E. Sending commands during boot of a single EC2 instance	9
F. Change the listener of a Kafka broker deployed with terraform	11

G.	Adding more Kafka brokers to cluster, and configuring each one individually	13
H.	How to show the dependencies in terraform?.....	16
I.	Sharing terraform state between resources using AWS S3	17
J.	Creating Lambda function and API Gateway trigger using S3.....	23
K.	Kong and Konga deployed in AWS EC2 using Terraform.....	27
L.	Camunda Engine deployed in AWS EC2 using Terraform	30
M.	Quarkus project deployed in AWS EC2 using Terraform	32
N.	Ollama deployed in AWS EC2 using Terraform.....	34
	References	37

A. Terraform installation procedure

Use the URL <https://developer.hashicorp.com/terraform/downloads> and choose the package appropriate for your operating system and then follow the indications.

For better Terraform code rendering, you can optionally, install the Terraform extension for VSCode.

The screenshot shows the HashiCorp Terraform extension page in the VS Code Marketplace. The extension is version v2.25.2, developed by HashiCorp, with 2,407,473 installs and a rating of 4.5 stars (165 reviews). It provides syntax highlighting and autocompletion for Terraform. The extension is currently enabled globally. The page includes tabs for Details, Feature Contributions, Changelog, and Runtime Status. The Details tab contains a troubleshooting guide and a list of features: IntelliSense, Syntax validation, Syntax highlighting, Code Navigation, Code Formatting, Code Snippets, Terraform Module Explorer, and Terraform commands. The Features section also describes IntelliSense and Autocomplete, stating that it provides code completion, parameter info, quick info, and member lists. A note mentions that Terraform constructs like resource and data, labels, blocks and attributes are auto completed both at the root of the document and inside other blocks. The extension is categorized under Programming Languages, Linters, and Formatters. Extension Resources include Marketplace, Repository, License, and HashiCorp. More Info shows the extension was published on 12/02/2016 at 08:31:09, last released on 15/12/2022 at 18:25:23, and has the identifier hashicorp.terraform.

B. Obtaining the access for sandbox or AWS academy learner lab

B.1. Choose your learner lab, then start it.

The screenshot shows the AWS Academy Cloud Foundations – Sandbox environment overview. The left sidebar includes links for Conta, Painel de controle, Cursos, Calendário, Caixa de entrada, Histórico, and Ajuda. The main content area displays the title "AWS Academy Cloud Foundations – Sandbox" and "Environment Overview". It states that the sandbox provides an environment for ad-hoc exploration of AWS services. A note indicates that the environment is cleaned up at the end of every session. A "Region restriction" section notes that service access is limited to the us-east-1 Region. On the right, there is a terminal window showing a bash prompt: `ddd_v1_w_yIO_1852360@runweb71380:~$`.

B.2. When the lab is started, select Details, and then Show:

The screenshot shows the AWS CLI Details page. The "Details" tab is selected. Below it, under "AWS:", there is a "Show" button. To the right of the "Show" button is a terminal window showing a bash prompt: `2360@runweb71380:~$`.

B.3. A similar access configuration is available for you (they are both examples that could be available, the first one is secretKey/accessKey, and the second is secretKey/accessKey/token). You will have to adapt your terraform calls accordingly with the provided AWS access.

The screenshot shows the Cloud Access page. Under "Cloud Labs", it displays session information: Remaining session time: 02:32:20(153 minutes), Session started at: 2023-01-23T05:33:06-0800, Session to end at: 2023-01-23T08:33:06-0800, and Accumulated lab time: 00:27:00 (27 minutes). It also lists network information: ips -- public:18.234.191.90, private:10.0.0.176. Under "AWS CLI:", there is a "Show" button. Below it, there are sections for "SSH key" (with "Show", "Download PEM", and "Download PPK" buttons) and "AWS SSO" (with a "Download URL" button). At the bottom, there is a table with the following data:

SecretKey	[REDACTED]
BastionHost	18.234.191.90
Region	us-east-1
AccessKey	AKIAQPVWLKEA366JUOFG

The screenshot shows the AWS Learner Lab interface. On the left, there's a sidebar with icons for Account, Courses, Dashboard, Calendar, Inbox, History, and Help. The main area has a breadcrumb navigation: ALLv1-38078 > Modules > Learner Lab > Learner Lab. At the top right, it says "Used \$0 of \$100", "03:17", and buttons for "Start Lab", "End Lab", "AWS Details", "Readme", and "Reset".

Cloud Access:

AWS CLI:
Copy and paste the following into `~/.aws/credentials`

```
[default]
aws_access_key_id=ASIA5
aws_secret_access_key=B8
aws_session_token=FwGZ2
steYrhYfMsolDYdxhnglRk8
8yjJrOBAsusVbx2u+j26
0d16j8huC01vdQxe6071CJ2
K3Nn2S+huswkh-fmPf3SP5w
aq1gZsopr3TOAbf1UBR+Bukq
```

Cloud Labs:

Remaining session time: 03:17:28 (198 minutes)
Session started at: 2023-01-30T03:08:38-0800
Session to end at: 2023-01-30T07:08:38-0800
Accumulated lab time: 00:42:00 (42 minutes)

No running instance

SSH key: Show | Download PEM | Download PPK

AWS SSO: Download URL

AWSAccountID	188216847906
Region	us-east-1

C. Deploying a single EC2 instance

- C.1. Configure the AWS access inside the terraform script file. For that, create the following script, replacing the **access_key**, **secret_key**, **token**, **region**, the **ami**, and the **instance_type** accordingly with your needs. Save the file in text format with a **.tf** extension.

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region = "us-east-1"
  access_key  = "YOUR_ACCESS_KEY"
  secret_key  = "YOUR_SECRET_KEY"
  token       = "YOUR_TOKEN"
}

resource "aws_instance" "example" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.micro"

  tags = {
    Name = "terraform-example"
  }
}
```

- C.2. Deploy the code with the following commands. You should use only one **.tf** file in each directory. A local configuration is created and then reflected in the cloud environment:

```
terraform init  
terraform apply
```

or sequentially:

```
terraform init && terraform apply
```

- C.3. You can verify the configuration of the current configuration using the following command:

```
terraform show
```

- C.4. Clean up when you're done:

```
terraform destroy
```

Hint: to facilitate the terraform command you can also use the -auto-approve option

D. Deploying a single EC2 instance and uploading files

D.1. As previous, configure your AWS access keys as environment variables, and then, create the following .tf file, adapting the **region**, the **ami**, and the **instance_type** accordingly with you needs.

Notice that **access_key**, **secret_key** and **token** could be directly written in the source code.

Also get available the kafka tgz file and test.pem on your local path.

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access_key  = "YOUR_ACCESS_KEY"
  secret_key  = "YOUR_SECRET_KEY"
  token       = "YOUR_TOKEN"
}

resource "aws_instance" "exampleFileTransfer" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.micro"
  vpc_security_group_ids = [aws_security_group.instance.id]
  key_name     = "vokey"
  tags = {
    Name = "terraform-example"
  }

  connection {
    type     = "ssh"
    user     = "ec2-user"
    private_key = file("test.pem")
    host     = "${self.public_dns}"
  }

  provisioner "file" {
    source      = "kafka_2.13-3.9.0.tgz"
    destination = "/home/ec2-user/kafka_2.13-3.9.0.tgz"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-example-instance"
}
```

Where **key_name** is obtained from the name of the keypairs available directly from the AWS console:

The screenshot shows the AWS Key Pairs page. At the top, there's a search bar and a 'Create key pair' button. Below that, a table lists one key pair: 'vokey' (rsa type, created on 2023/01/24 15:37 GMT+0), with a fingerprint listed as 'base0:aa:be:42:8f:54:ea:e0:70:99:6d:32:...'. The left sidebar has links for EC2 Dashboard, EC2 Global View, Events, and Tags.

The **provisioner "file"** tag can be repeated as many times as needed, e.g.:

```
provisioner "file" {
  source      = "kafka_2.13-3.6.2.tgz"
  destination = "/home/ec2-user/kafka_2.13-3.6.2.tgz"
}

provisioner "file" {
  source      = "/root/zookeeper/apache-zookeeper-3.8.4-bin.tar.gz"
  destination = "/home/ec2-user/apache-zookeeper-3.8.4-bin.tar.gz"
}
```

Take attention to the **security resource** that allows SSH connections to your instance. If any other ingress is needed, you can add more in the same .tf file.

Also, the connection provides the **connection** configuration to enable the file upload (it uses a scp mechanism).

Also take note that provisioner is last action taken.

D.2. After changing for your configuration. Try it. Deploy the code:

```
terraform init
terraform apply
```

or

```
terraform init && terraform apply
```

D.3. Check if the file is uploaded correctly using a SSH connection directly to the EC2 instance:

```
ssh -i test.pem ec2-user@YOUR_DNS_NAME
```

To remember: Use the following command to set the permissions of your private key file so that only you can read it.

```
chmod u=rwx,g=,o= myKeyAWS.pem
```

D.4. Clean up when you're done:

```
terraform destroy
```

E. Sending commands during boot of a single EC2 instance

User_data is a shell script or cloud-init directive that will be executed during instance creation only.

The result of the **user_data** execution can be seen in EC2 Console (select the Instance, click Actions -> Monitor and troubleshoot -> Get system log) and you can find its execution log on the EC2 Instance itself (typically in /var/log/cloud-init*.log), both of which are useful for debugging, and neither of which is available with provisioners.

E.1. As previous, configure your AWS access keys as environment variables, and then, create the following .tf file, adapting the **region**, the **ami**, and the **instance_type** accordingly with you needs.

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access key  = "YOUR_ACCESS_KEY"
  secret key  = "YOUR_SECRET_KEY"
  token       = "YOUR_TOKEN"
}

resource "aws_instance" "exampleKAFKA" {
  ami           = "ami-045269a1f5c90a6a0"
  instance type = "t2.small"
  vpc security group ids = [aws security group.instance.id]
  key_name     = "vockey"

  user_data = "${file("creation.sh")}"

  user_data_replace_on_change = true

  tags = {
    Name = "terraform-example-kafka"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port   = 2181
    to_port     = 2181
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port   = 9092
    to_port     = 9092
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6 cidr blocks = [":/:0"]
  }
}
```

```

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-example-instance"
}

```

Where the script **user_data = "\${file("creation.sh")}"** is a local file that could contain some work that you need to perform.

For instance, installing and starting ZooKeeper and Kafka, automatically!

```

#!/bin/bash
echo "Starting..."
cd
sudo wget https://dlcdn.apache.org/zookeeper/zookeeper-3.9.3/apache-zookeeper-3.9.3-bin.tar.gz
sudo tar -zxf apache-zookeeper-3.9.3-bin.tar.gz
sudo mv apache-zookeeper-3.9.3-bin /usr/local/zookeeper
sudo mkdir -p /var/lib/zookeeper
echo "tickTime=2000
dataDir=/var/lib/zookeeper
clientPort=2181" > /usr/local/zookeeper/conf/zoo.cfg

sudo yum -y install java-17-amazon-corretto-devel.x86_64
sudo /usr/local/zookeeper/bin/zkServer.sh start

sudo wget https://downloads.apache.org/kafka/3.9.0/kafka_2.13-3.9.0.tgz
sudo tar -zxf kafka_2.13-3.9.0.tgz
sudo mv kafka_2.13-3.9.0 /usr/local/kafka
sudo mkdir /tmp/kafka-logs

sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon /usr/local/kafka/config/server.properties
echo "Finished."

```

The **egress** tag allows the EC2 instance to outbound to another machine in internet.

Recall that a provisioner is the last action taken. That is why the binary files for Zookeeper and Kafka are being downloaded instead of being received by a provisioner.

If needed, the logging of the EC2 instance provisioning can be consulted at `/var/log/cloud-init-output.log`

F. Change the listener of a Kafka broker deployed with terraform

F.1. Use the previous terraform file for deploying kafka broker in AWS.

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access_key  = "YOUR_ACCESS_KEY"
  secret_key  = "YOUR_SECRET_KEY"
  token       = "YOUR_TOKEN"
}

resource "aws_instance" "exampleKAFKAConfigured" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.small"
  vpc_security_group_ids = [aws_security_group.instance.id]
  key_name     = "vockey"

  user_data = "${file("creation.sh")}"

  user_data_replace_on_change = true

  tags = {
    Name = "terraform-example-kafka2"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port   = 2181
    to_port     = 2181
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port   = 9092
    to_port     = 9092
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-example-instance"
}
```

F.2. Then, the script for provisioning the kafka broker need to be adapted with the new configuration for the listener at file /usr/local/kafka/config/server.properties. For that the name of the EC2 instance need to be known. The instance metadata service available at <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-metadata.html> could be used for that purpose.

```
#!/bin/bash

echo "Starting..."
cd
sudo wget https://dlcdn.apache.org/zookeeper/zookeeper-3.9.3/apache-zookeeper-3.9.3-bin.tar.gz
sudo tar -zxf apache-zookeeper-3.9.3-bin.tar.gz
sudo mv apache-zookeeper-3.9.3-bin /usr/local/zookeeper
sudo mkdir -p /var/lib/zookeeper
echo "tickTime=2000
dataDir=/var/lib/zookeeper
clientPort=2181" > /usr/local/zookeeper/conf/zoo.cfg

sudo yum -y install java-17-amazon-corretto-devel.x86_64
sudo /usr/local/zookeeper/bin/zkServer.sh start

sudo wget https://downloads.apache.org/kafka/3.9.0/kafka_2.13-3.9.0.tgz
sudo tar -zxf kafka_2.13-3.9.0.tgz
sudo mv kafka_2.13-3.9.0 /usr/local/kafka
sudo mkdir /tmp/kafka-logs

ip=`curl http://169.254.169.254/latest/meta-data/public-hostname`
sudo sed -i "s/#listeners=PLAINTEXT:\:\/\/:9092/listeners=PLAINTEXT:\:\/\/$ip:9092/g" /usr/local/kafka/config/server.properties

# due to AWS network stablishment process, check if 60 seconds is enough for your situation
(sleep 120 && sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon /usr/local/kafka/config/server.properties )&
echo "Finished."
```

G. Adding more Kafka brokers to cluster, and configuring each one individually

Solving this problem in small solutions:

- **First solution:** simply add the count element to your aws instance. When you execute the command “terraform apply”, those correspondingly number of similar instances will be created. It’s a fast and easy way of provisioning multiple EC2 instances. However, all the internal detailed configurations of the installed software need to be done manually EC2 instance by EC2 instance: which at scale could be unfeasible.

```
resource "aws_instance" "exampleCLUSTER" {
  ami                  = "ami-045269a1f5c90a6a0"
  instance_type        = "t2.small"

  count = 4

  vpc_security_group_ids  = [aws_security_group.instance.id]
  key_name                = "vokey"

  user_data = "${file("creation.sh")}"
  user_data_replace_on_change = true

  tags = {
    Name = "terraform-example-kafka"
  }
}
```

- **Second solution:** being able to provision multiple EC2 instances, and at the same time configure all the internal detailed configurations of each one of the EC2 instances. It has the advantage of accelerating the manual configurations by automation, and at scale. The drawback is the complexity and the lack of global configuration, i.e., configuration related with all the EC2 instances for all the EC2 instances.

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access_key  = "YOUR_ACCESS_KEY"
  secret_key  = "YOUR_SECRET_KEY"
  token       = "YOUR_TOKEN"
}

variable "nBroker" {
  description = "number of brokers"
  type        = number
  default     = 3
}

resource "aws_instance" "exampleCluster" {
  ami                  = "ami-045269a1f5c90a6a0"
  instance_type        = "t2.small"
  count               = var.nBroker
  vpc_security_group_ids  = [aws_security_group.instance.id]
  key_name             = "vokey"

  user_data      = base64encode(templatefile("creation.sh", {
    idBroker = "${count.index}"
    totalBrokers = var.nBroker
  }))
}
```

```

user_data_replace_on_change = true

tags = {
  Name = "terraform-example-kafka.${count.index}"
}
}

output "publicdnslist" {
  value = "${formatlist("%v", aws_instance.exampleCluster.*.public dns)}"
}

resource "aws_security_group" "instance" {
  name = var.security_group_name
  ingress {
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
  ingress {
    from_port   = 2181
    to_port     = 2181
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
  ingress {
    from_port   = 2888
    to_port     = 2888
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
  ingress {
    from_port   = 3888
    to_port     = 3888
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
  ingress {
    from_port   = 9092
    to_port     = 9092
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
  egress {
    from_port     = 0
    to_port       = 0
    protocol      = "-1"
    cidr_blocks   = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-example-instances"
}

```

```

#!/bin/bash

echo "Starting..."
cd
sudo wget https://dlcdn.apache.org/zookeeper/zookeeper-3.9.3/apache-zookeeper-3.9.3-bin.tar.gz
sudo tar -zxf apache-zookeeper-3.9.3-bin.tar.gz
sudo mv apache-zookeeper-3.9.3-bin /usr/local/zookeeper
sudo mkdir -p /var/lib/zookeeper
echo "tickTime=2000
dataDir=/var/lib/zookeeper
clientPort=2181
maxClientCnxns=60
initLimit=10
syncLimit=5" > /usr/local/zookeeper/conf/zoo.cfg

sudo yum -y install java-17-amazon-corretto-devel.x86_64
echo ${idBroker+1} > /var/lib/zookeeper/myid
sudo /usr/local/zookeeper/bin/zkServer.sh start

```

```

sudo wget https://downloads.apache.org/kafka/3.9.0/kafka_2.13-3.9.0.tgz
sudo tar -zxf kafka_2.13-3.9.0.tgz
sudo mv kafka_2.13-3.9.0 /usr/local/kafka
sudo mkdir /tmp/kafka-logs

ip=`curl http://169.254.169.254/latest/meta-data/public-hostname`
sudo sed -i "s/#listeners=PLAINTEXT:\:\/\/:9092/listeners=PLAINTEXT:\:\/\/$ip:9092/g"
/usr/local/kafka/config/server.properties

sudo sed -i "s/broker.id=0/broker.id=${idBroker+1}/g" /usr/local/kafka/config/server.properties

sudo sed -i "s/offsets.topic.replication.factor=1/offsets.topic.replication.factor=${totalBrokers}/g"
/usr/local/kafka/config/server.properties
sudo sed -i "s/transaction.state.log.replication.factor=1/transaction.state.log.replication.factor=${totalBrokers}/g"
/usr/local/kafka/config/server.properties
sudo sed -i "s/transaction.state.log.min_isr=1/transaction.state.log.min_isr=${totalBrokers}/g"
/usr/local/kafka/config/server.properties

# due to AWS network establishment process, check if 30 seconds is enough for your situation
(sleep 30 && sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon
/usr/local/kafka/config/server.properties )&

echo "Finished."

```

As explained in tutorial P4 - Distributed Kafka service, two global configurations are still missing in all the EC2 instances. Login in each machine and setup the following configuration:

At /usr/local/kafka/config/server.properties

```

#zookeeper connectivity (one per EC2 VM of this cluster)
zookeeper.connect=ec2-54-90-57-82.compute-1.amazonaws.com:2181,ec2-54-173-171-63.compute-1.amazonaws.com:2181,ec2-54-236-47-54.compute-1.amazonaws.com:2181

```

And at /usr/local/zookeeper/conf/zoo.cfg

```

server.1=ec2-54-90-57-82.compute-1.amazonaws.com:2888:3888
server.2=ec2-54-173-171-63.compute-1.amazonaws.com:2888:3888
server.3=ec2-54-236-47-54.compute-1.amazonaws.com:2888:3888

```

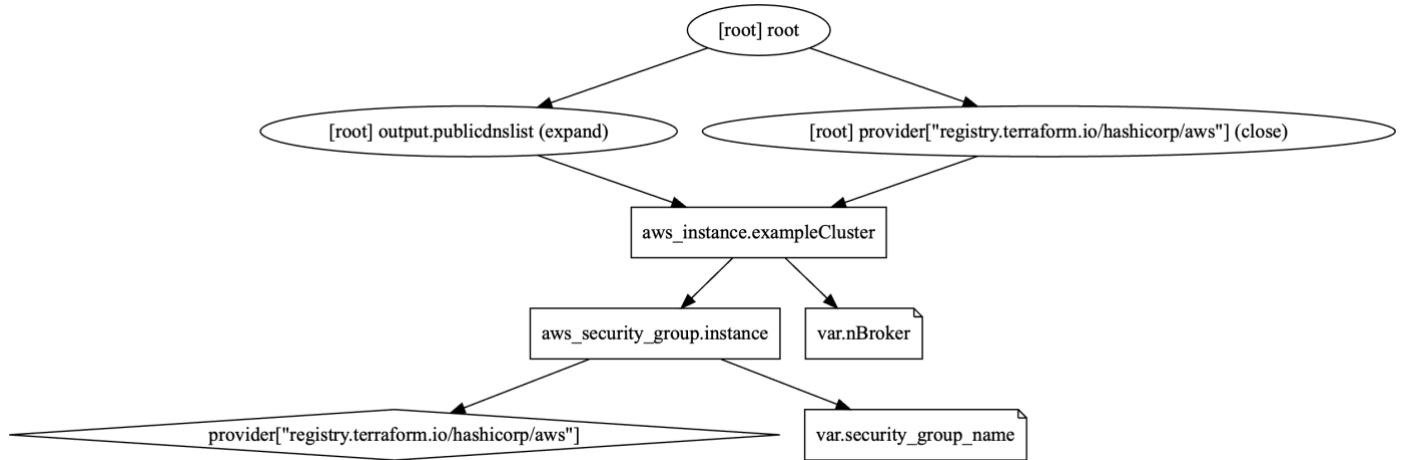
H. How to show the dependencies in terraform?

To show the dependencies of your deployment environment use the following command:

```
terraform graph -draw-cycles
```

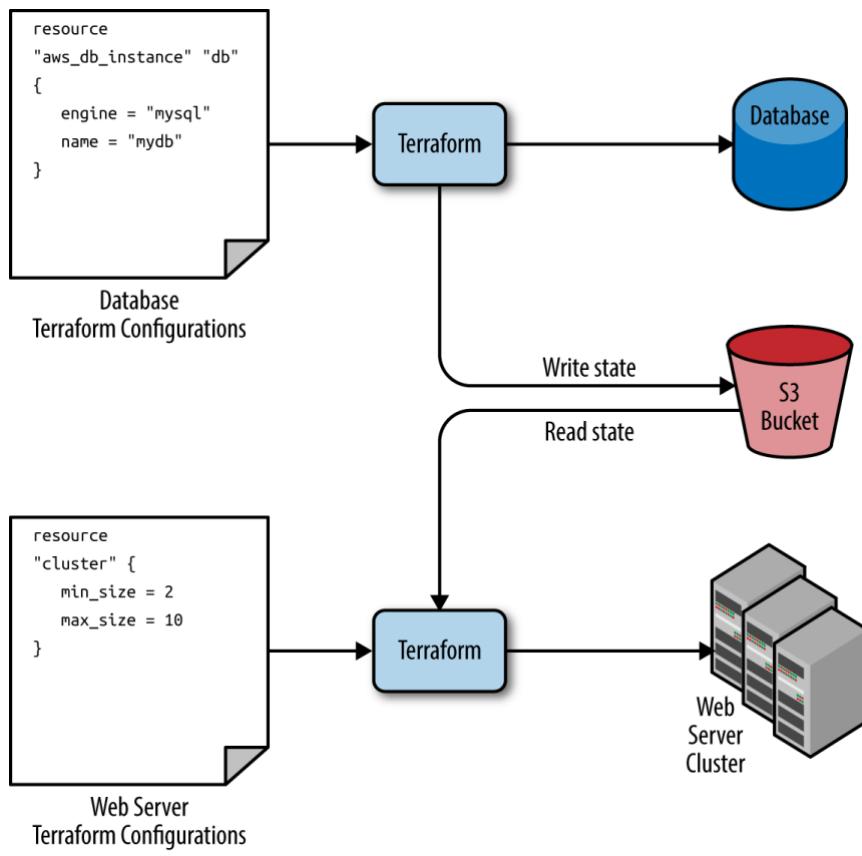
Then copy+paste the output for the graphviz, e.g., on <https://dreampuf.github.io/GraphvizOnline>.

Here is an example of you obtain:



I. Sharing terraform state between resources using AWS S3

The terraform state sharing allows the provisioning of an environment where the posterior resources depend on the formers, as depicted in the next figure. It could be done using a cloud storage or a local storage. S3 is an AWS service to store data.



Source: Brikman, Y. (2022). *Terraform: Up and Running*. " O'Reilly Media, Inc."

Notice that to write the state in the storing an output should be defined in the terraform script file, i.e., to store locally the state of terraform:

```
output "publicdnslist" {
  value = "${formatlist("%v", aws_instance.exampleCluster.*.public_dns)}"
}
```

Then, you can list all the outputs, e.g.:

```
terraform state list
```

returning:

```
aws_instance.exampleCluster[0]
aws_instance.exampleCluster[1]
aws_instance.exampleCluster[2]
aws_security_group.instance
```

Or just show one of them:

```
terraform state show 'aws_instance.exampleCluster[0]'
```

In Terraform, the state is a representation of the configuration of your infrastructure that is stored locally or remotely. Accessing the state locally involves using the `terraform state` command.

I.1. Choose the AWS S3 service, then, click on “create bucket”:

The screenshot shows the AWS S3 Buckets page. At the top, there's a green success message: "Successfully created bucket 'terraform-s3-2025-01-15'". Below it, there's an "Account snapshot - updated every 24 hours" section. The main area shows a table of "General purpose buckets" with one item. The table includes columns for Name, AWS Region, IAM Access Analyzer, and Creation date. A "Create bucket" button is visible at the top right of the table area.

Give it a unique name:

The screenshot shows the "Create bucket" wizard. It starts with the "General configuration" step, where the user has chosen "General purpose" as the bucket type and entered "terraform-s3-2025-01-15" as the bucket name. The next steps are "Object Ownership" (set to "ACLS disabled (recommended)"), "Block Public Access settings for this bucket" (with "Block all public access" checked), and "Bucket Versioning" (set to "Disable"). The wizard is currently on the second step.

I.2. Terraform state will be stored in the S3 cloud environment. For that, firstly, create in a separate directory, a terraform file to create an AWS RDS (e.g.: mySQL cloud database) and store remotely the name of the DB.

```

terraform {
  backend "s3" {
    bucket = "terraform-s3-2025-01-15"
    key    = "stage/data-stores/mysql/terraform.tfstate"
    region = "us-east-1"
    access_key  = "YOUR_ACCESS_KEY"
    secret_key  = "YOUR_SECRET_KEY"
    token      = "YOUR_TOKEN"
  }
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region = "us-east-1"
  access_key  = "YOUR_ACCESS_KEY"
  secret_key  = "YOUR_SECRET_KEY"
  token      = "YOUR_TOKEN"
}

variable "db_username" {
  description = "The username for the database"
  type        = string
  sensitive   = true
  default     = "teste"
}

variable "db_password" {
  description = "The password for the database"
  type        = string
  sensitive   = true
  default     = "testeteste"
}

variable "db_name" {
  description = "The name to use for the database"
  type        = string
  default     = "example_database_stage"
}

resource "aws_db_instance" "example" {
  identifier_prefix  = "terraform-up-and-running"
  engine             = "mysql"
  allocated_storage  = 20
  instance_class     = "db.t4g.micro"
  skip_final_snapshot = true
  publicly_accessible = true
  vpc_security_group_ids = [aws_security_group.rds.id]
  db_name            = var.db_name

  username = var.db_username
  password = var.db_password
}

output "address" {
  value      = aws_db_instance.example.address
  description = "Connect to the database at this endpoint"
}

output "port" {
  value      = aws_db_instance.example.port
  description = "The port the database is listening on"
}

resource "aws_security_group" "rds" {
  name = var.security_group_name
  ingress {
    from_port  = 3306
    to_port    = 3306
    protocol   = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
  egress {
    from_port  = 0
  }
}

```

```

        to_port      = 0
        protocol     = "-1"
        cidr_blocks  = ["0.0.0.0/0"]
        ipv6_cidr_blocks = [":/:0"]
    }
}

variable "security_group_name" {
    description = "The name of the security group"
    type        = string
    default     = "terraform-rds-instance"
}

```

Execute the usual command:

```
terraform init && terraform apply
```

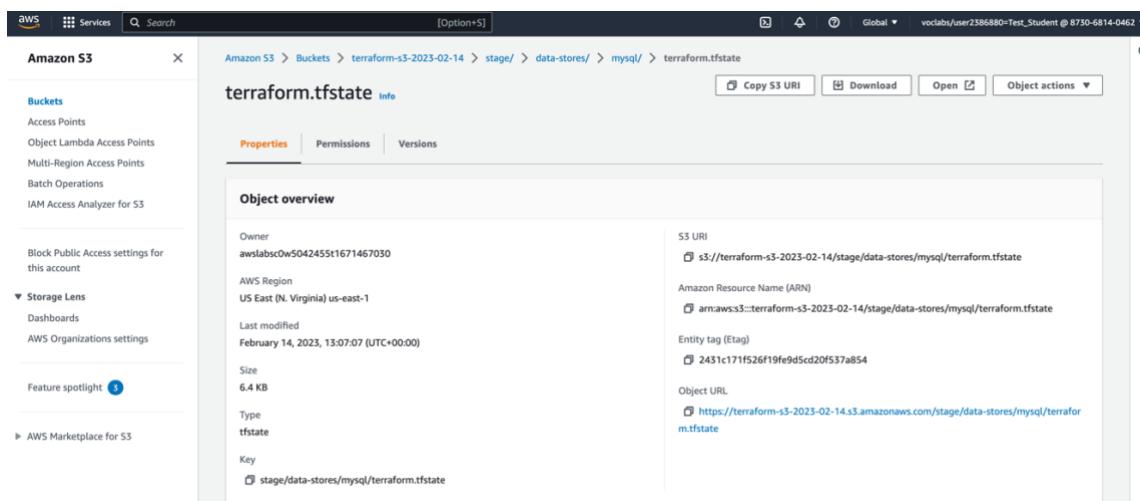
You can now check the outputs were written correctly to terraform state using the following commands:

```
terraform state list
```

```
terraform state show 'aws_db_instance.example'
```

```
terraform output
```

Moreover, if you navigate until:



You'll see that the terraform state of the **recently created mysql database** is stored there.

I.3. Now, in another directory, create a terraform file that can create an AWS EC2 instance writing the previous database address and port to a text file inside the EC2. The address and port are read from the state file of the previous step. With this mechanism you can create resources that are configured with parameters from the previous created resources. Facilitating the automation provisioning process!

```

terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source = "hashicorp/aws"
    }
  }
}

```

```

        version = "~> 4.0"
    }
}
}

data "terraform_remote_state" "db" {
  backend = "s3"
  config = {
    bucket = "terraform-s3-2025-01-15"
    key = "stage/data-stores/mysql/terraform.tfstate"
    region = "us-east-1"
    access_key = "YOUR_ACCESS_KEY"
    secret_key = "YOUR_SECRET_KEY"
    token = "YOUR_TOKEN"
  }
}

provider "aws" {
  region = "us-east-1"
  access_key = "YOUR_ACCESS_KEY"
  secret_key = "YOUR_SECRET_KEY"
  token = "YOUR_TOKEN"
}

resource "aws_instance" "exampleCluster" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.small"
  vpc_security_group_ids = [aws_security_group.instance.id]
  key_name      = "vokey"

  user_data     = base64encode(templatefile("creation.sh", {
    address = data.terraform_remote_state.db.outputs.address
    port   = data.terraform_remote_state.db.outputs.port
  }))

  user_data replace_on_change = true

  tags = {
    Name = "terraform-example-sharestate"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name
  ingress {
    from_port  = 22
    to_port    = 22
    protocol   = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
  egress {
    from_port     = 0
    to_port       = 0
    protocol      = "-1"
    cidr_blocks   = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-example-instance2"
}

```

And creation.sh with the following source code:

```

#!/bin/bash
cd

echo ${address} > address.txt
echo ${port} > port.txt

```

Execute the usual command:

```
terraform init && terraform apply
```

Login in the newly created EC2 instance and verify in the [root directory](#) that two new files are now created with the content provided by DB configuration.

J. Creating Lambda function and API Gateway trigger using S3

J.1. To create a lambda function, written in JAVA 8, that you previously compiled locally, you need to transfer the .jar file to S3, then create the AWS lambda function resource, and finally, add an API Gateway trigger for that lambda function.

```
terraform {
  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.15.0"
    }
    random = {
      source  = "hashicorp/random"
      version = "~> 3.1.0"
    }
    archive = {
      source  = "hashicorp/archive"
      version = "~> 2.2.0"
    }
  }
  required_version = "~> 1.2"
}

provider "aws" {
  region = "us-east-1"
  access_key  =
  secret_key  =
  token =
}

## S3 PART
resource "aws_s3_bucket" "lambda_bucket" {
  bucket = "terraform-lambda-2025-01-16"
}

resource "aws_s3_object" "lambda_hello_world" {
  bucket = aws_s3_bucket.lambda_bucket.id
  key    = "lambda-java-example-0.0.1-SNAPSHOT.jar"
  source = "${path.module}/lambda-java-example-0.0.1-SNAPSHOT.jar"
}

## LAMBDA FUNCTION PART
resource "aws_lambda_function" "hello_world" {
  function_name = "HelloWorld"
  s3_bucket    = aws_s3_bucket.lambda_bucket.id
  s3_key       = aws_s3_object.lambda_hello_world.key
  runtime      = "java11"
  handler     = "example.Hello::handleRequest"
  role        = "arn:aws:iam::8730681408695462:role/LabRole"
}

resource "aws_cloudwatch_log_group" "hello_world" {
  name = "/aws/lambda/${aws_lambda_function.hello_world.function_name}"
  retention_in_days = 30
}

### API GATEWAY PART
resource "aws_api_gateway_rest_api" "example" {
  name          = "ServerlessExample"
  description   = "Terraform Serverless Application Example"
}

resource "aws_api_gateway_resource" "proxy" {
  rest_api_id = "${aws_api_gateway_rest_api.example.id}"
  parent_id   = "${aws_api_gateway_rest_api.example.root_resource_id}"
  #path_part   = "{proxy+}"
  path_part   = "helloworldpath"
}

resource "aws_api_gateway_method" "proxy" {
  rest_api_id = "${aws_api_gateway_rest_api.example.id}"
  resource_id = "${aws_api_gateway_resource.proxy.id}"
  http_method = "ANY"
  authorization = "NONE"
}

resource "aws_api_gateway_integration" "lambda" {
```

```

rest_api_id = "${aws_api_gateway_rest_api.example.id}"
resource_id = "${aws_api_gateway_method.proxy.resource_id}"
http method = "${aws_api_gateway_method.proxy.http_method}"

integration_http_method = "POST"
type                  = "AWS_PROXY"
uri                   = "${aws_lambda_function.hello_world.invoke_arn}"
}

resource "aws_api_gateway_deployment" "example" {
depends_on = [
  aws_api_gateway_integration.lambda,
]
rest_api_id = "${aws_api_gateway_rest_api.example.id}"
stage_name  = "test"
}

resource "aws_lambda_permission" "apigw" {
statement_id  = "AllowAPIGatewayInvoke"
action        = "lambda:InvokeFunction"
function_name = "${aws_lambda_function.hello_world.function_name}"
principal     = "apigateway.amazonaws.com"

# The /*/* portion grants access from any method on any resource
# within the API Gateway "REST API".
source_arn = "${aws_api_gateway_rest_api.example.execution_arn}/*/*"
}

# Output value definitions
output "lambda_bucket_name" {
description = "Name of the S3 bucket used to store function code."
value = aws_s3_bucket.lambda_bucket.id
}

output "function_name" {
description = "Name of the Lambda function."
value = aws_lambda_function.hello_world.function_name
}

output "base_url" {
value = "${aws_api_gateway_deployment.example.invoke_url}"
}

```

Where the field role of resource "aws_lambda_function":

```
role = "arn:aws:iam::8730681408695462:role/LabRole"
```

could be consulted on your AWS account, in the IAM service, or replacing your AWSAccountId number directly in the bold part of the field:

The screenshot shows the AWS IAM interface. On the left, the 'Identity and Access Management (IAM)' section is open, showing various management options like User groups, Users, Roles, Policies, Identity providers, and Account settings. The 'Access management' section is expanded, showing 'Access analyzer', 'Archive rules', 'Analyzers', 'Settings', 'Credential report', 'Organization activity', and 'Service control policies (SCPs)'. The 'Access reports' section is also visible.

In the center, the 'LabRole' configuration page is displayed. It shows the role's summary, creation date (January 31, 2023, 10:13 (JTC)), last activity (30 minutes ago), and maximum session duration (1 hour). The 'Permissions' tab is selected, listing several permissions policies:

- c7424ba151403004900991+e73068140462-VocLabPolicy1-1HHQK35E9QNF
- c7424ba151403004900991+e73068140462-VocLabPolicy2-130ET5NP9BNS
- c7424ba151403004900991+e73068140462-VocLabPolicy3-OJUDRLJ3HC
- AmazonEKSClusterPolicy
- AmazonECKWorkerNodePolicy
- AmazonEC2ContainerRegistryReadOnly
- AmazonSSMManagedInstanceCore

The 'Permissions boundary - (not set)' section indicates no boundary is set. Below it, there's a note about generating a policy based on CloudTrail events, with a 'Generate policy' button.

On the right, the 'Cloud Access' panel is shown. It displays session details: 03:11, Start Lab, End Lab, AWS Details, Readme, Reset, and a link to 'Close'. It also shows the cost used: \$2.9 of \$100. The AWS CLI section contains a command to copy and paste into the credentials file:

```
[default]
aws_access_key_id=ASIA4WRXKS6XGI4LT0ZN
aws_secret_access_key=R1A9iIRd7NaJSIf6icPhjc2DWlMvLk/8J9RIsaOs
aws_session_token=FwoGZXIvYXdzEKn//////////wEaDLxuRLTrQu8nHgJLSK9AU
v5gh/aw1139HcF3bKO42EFh02KnfELEJyXhoCpski8h8h0X63tjIb6nHs3+nrekHo
2/3w7rT9KP060mP/4S2ZWPpb1BSqzExgztf6WL1xln.ukX4+70Qfh8WogSvCEWTl9p9B
SQiPX199a0w/CRT5D4jaRVdvenV5IFu/NX2HrPt0Jg8tIu/JTJ4S4Tu/20uIMMvLD
2ENWgiP+GglKb5vX1leofZYWhrtWe++KRY9UgfUFVjfxtyil7L0FBjItkPtZrLOEkl
xaV+HQj1BsQMpz93ab2FqV9C2tk3L1xk8MPPz8+rNypIO5RJY
```

The 'Cloud Labs' section shows the remaining session time (03:14:46), session start time (2023-02-15T07:11:32-0800), and session end time (2023-02-15T11:11:32-0800). It also displays accumulated lab time (1 day 02:45:00 / 1605 minutes).

The 'No running instance' section shows SSH key, AWS SSO, and AWS URL download buttons. The AWS Account ID is listed as 873068140462, and the Region is us-east-1.

J.2. To test your lambda function, externally from any machine, you need to connect to AWS API Gateway, as explained in the Lambda-AWS tutorial:

```
curl -i -H "Content-Type: Application/json" --data "@body.json" -X POST https://<yourURL>/test/helloworldpath
```

And the result obtained is the following:

```
HTTP/2 200
content-type: application/json
content-length: 43
date: Wed, 15 Feb 2023 16:36:51 GMT
x-amzn-requestid: f2027a23-803d-42d3-8537-e2fe256d7480
x-amz-apigw-id: AY6FiH8XoAMFbvA=
x-custom-header: my custom header value
x-amzn-trace-id: Root=1-63ed0a23-652880af6750972f7523767e; Sampled=0
x-cache: Miss from cloudfront
via: 1.1 cb4f40303e252a22c4df5918669814ac.cloudfront.net (CloudFront)
x-amz-cf-pop: LIS50-C1
x-amz-cf-id: CAQtv416vFuEzDsJQv7tj56smEc7s4ZkcEffM3VobsGSKnB-1H2y4g==

{"message": "Hello Integracao Empresarial!"}
```


K. Kong and Konga deployed in AWS EC2 using Terraform

K.1. To install Kong you will need to follow the indications as explained in the Kong tutorial. They are now adapted to be hyperautomated using Terraform.

Where the .tf file can include the following instructions:

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region     = "us-east-1"
  access_key = ""
  secret_key = ""
  token     = ""
}

resource "aws_instance" "exampleInstallKong" {
  ami           = "ami-045269a1f5c90a6a0"
  instance type = "t2.small"
  vpc security group ids = [aws security group.instance.id]
  key name     = "vockey"

  user_data = "${file("deploy.sh")}"

  user_data_replace_on_change = true

  tags = {
    Name = "terraform-example-Kong"
  }
}

resource "aws security group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }

  egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-kong-instance"
}

output "address" {
  value      = aws_instance.exampleInstallKong.public_dns
  description = "Connect to the database at this endpoint"
}
```

And the correspondingly script with the following:

```

#!/bin/bash
echo "Starting..."

sudo yum install -y docker
sudo service docker start
sudo docker network create kong-net

sudo docker run -d --name kong-database \
--network=kong-net \
-p 5432:5432 \
-e "POSTGRES_USER=kong" \
-e "POSTGRES_DB=kong" \
-e "POSTGRES_PASSWORD=kongpass" \
postgres:13

sudo docker run --rm --network=kong-net \
-e "KONG_DATABASE=postgres" \
-e "KONG_PG_HOST=kong-database" \
-e "KONG_PG_PASSWORD=kongpass" \
-e "KONG_PASSWORD=test" \
kong/kong-gateway:3.9.0.0 kong migrations bootstrap

sudo docker run -d --name kong-gateway \
--network=kong-net \
-e "KONG_DATABASE=postgres" \
-e "KONG_PG_HOST=kong-database" \
-e "KONG_PG_USER=kong" \
-e "KONG_PG_PASSWORD=kongpass" \
-e "KONG_PROXY_ACCESS_LOG=/dev/stdout" \
-e "KONG_ADMIN_ACCESS_LOG=/dev/stdout" \
-e "KONG_PROXY_ERROR_LOG=/dev/stderr" \
-e "KONG_ADMIN_ERROR_LOG=/dev/stderr" \
-e "KONG_ADMIN_LISTEN=0.0.0.0:8001, 0.0.0.0:8444 ssl" \
-e "KONG_ADMIN_GUI_URL=http://localhost:8002" \
-e KONG_LICENSE_DATA \
-p 8000:8000 \
-p 8443:8443 \
-p 8001:8001 \
-p 8002:8002 \
-p 8445:8445 \
-p 8003:8003 \
-p 8004:8004 \
-p 127.0.0.1:8444:8444 \
kong/kong-gateway:3.9.0.0

echo "Finished."

```

K.2. To install the UI Konga again you can follow the indications as explained in the Kong tutorial. They are now adapted to be hyperautomated using Terraform.

Where the .tf file can include the following instructions:

```

terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }

  provider "aws" {
    region      = "us-east-1"
    access_key  =
    secret_key  =
    token       =
  }

  resource "aws_instance" "exampleInstallKonga" {
    ami           = "ami-045269a1f5c90a6a0"
    instance_type = "t2.small"
  }
}

```

```

vpc_security_group_ids  = [aws_security_group.instance.id]
key_name                = "vockey"

user_data = "${file("deploy.sh")}"

user_data_replace_on_change = true

tags = {
  Name = "terraform-example-Konga"
}
}

resource "aws_security_group" "instance" {
  name = var.security_group_name
  ingress {
    from_port    = 0
    to_port      = 0
    protocol     = "-1"
    cidr_blocks  = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }
  egress {
    from_port      = 0
    to_port        = 0
    protocol       = "-1"
    cidr_blocks   = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-konga-instance"
}

output "address" {
  value      = aws_instance.exampleInstallKonga.public_dns
  description = "Connect to the database at this endpoint"
}

```

And the correspondingly script with the following:

```

#!/bin/bash
echo "Starting..."

sudo yum install -y docker
sudo service docker start
sudo docker pull pantsel/konga
sudo docker run -d --name konga -p 1337:1337 pantsel/konga
echo "Finished."

```

L. Camunda Engine deployed in AWS EC2 using Terraform

L.1. To install Camunda Engine you will need to follow the indications as explained in the Camunda tutorial. They are now adapted to be hyperautomated using Terraform.

Where the .tf file can include the following instructions:

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region     = "us-east-1"
  access_key = ""
  secret_key = ""
  token     = ""
}

resource "aws_instance" "exampleInstallCamundaEngine" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.small"
  vpc_security_group_ids = [aws_security_group.instance.id]
  key_name      = "vockey"

  user_data = "${file("deploy.sh")}"

  user_data_replace_on_change = true

  tags = {
    Name = "terraform-example-Camunda"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }

  egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-Camunda-instance"
}

output "address" {
  value      = aws_instance.exampleInstallCamundaEngine.public_dns
  description = "Connect to the database at this endpoint"
}
```

And the correspondingly script with the following:

```
#!/bin/bash
echo "Starting..."
sudo yum update -y
sudo yum install -y docker
sudo service docker start
sudo docker pull camunda/camunda-bpm-platform:latest
sudo docker run -d --name camunda -p 8080:8080 camunda/camunda-bpm-platform:latest
echo "Finished."
```

M. Quarkus project deployed in AWS EC2 using Terraform

M.1. To deploy a previous created and compiled Quarkus project you will need to have the image previously pushed to git, Then, use the following files for reference. Where the .tf file can include the following instructions:

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access_key  = "YOUR ACES KEY"
  secret_key  = "YOUR SECRET KEY"
  token       = "YOUR TOKEN"
}

resource "aws_instance" "exampleDeployQuarkus" {
  ami           = "ami-0e7290665643979b5" # Amazon Linux ARM AMI built by Amazon Web Services - FOR
  DOCKER image COMPATIBILITY if compiled previously on ARM
  instance_type        = "t4g.nano"
  vpc_security_group_ids = [aws_security_group.instance.id]
  key_name          = "vockey"

  user_data = "${file("quarkus.sh")}"

  user_data_replace_on_change = true

  tags = {
    Name = "terraform-deploy-QuarkusProject"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }

  egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-Quarkus-instance"
}

output "address" {
  value      = aws_instance.exampleDeployQuarkus.public_dns
  description = "Address of the Quarkus EC2 machine"
}
```

And the correspondingly script with the following:

```
#!/bin/bash
echo "Starting..."

sudo yum install -y docker
sudo service docker start
sudo docker login -u "YOUR DOCKER USERNAME" -p "YOUR DOCKER PASSWORD"
sudo docker pull YOUR DOCKER USERNAME/tryout1:1.0.0-SNAPSHOT
sudo docker run -d --name tryout2 -p 9000:9000 YOUR DOCKER USERNAME/tryout1:1.0.0-SNAPSHOT
echo "Finished."
```

Docker images built for ARM architecture may not work on AMD architecture machines. This is because ARM and AMD processors use different instruction sets and have different hardware architectures.

For this example, the AMI identifier and instance type were extracted from the launching instance AWS GUI.

The screenshot shows the 'Launch an instance' wizard in the AWS EC2 console. The 'Summary' section is highlighted with a red box, containing the following information:

- Software Image (AMI)**: Amazon Linux 2 LTS Arm64 Kernel 5.10 AMI 2.0.20250108.0 arm64 HVM gp2
- Virtual server type (instance type)**: t4g.nano
- Firewall (security group)**: New security group

Below the summary, there's a note about the Free tier:

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month, 30 GiB of EBS storage, 2 million I/Os, 1 GB of snapshots, and 100 GB of bandwidth to the internet.

The 'Name and tags' section shows a name 'e.g. My Web Server' and an 'Add additional tags' button.

The 'Application and OS Images (Amazon Machine Image)' section shows a search bar and a list of recent AMIs: Amazon Linux, macOS, Ubuntu, Windows, Red Hat, SUSE Linux, and Debian.

The 'Amazon Machine Image (AMI)' section shows the selected AMI details: Amazon Linux 2 AMI (HVM) - Kernel 5.10, SSD Volume Type, ami-0454652560c7f5c55 (64-bit (x86)) / ami-0e7290665643979b5 (64-bit (Arm)).

The 'Architecture' dropdown is set to '64-bit (Arm)'. The 'AMI ID' is 'ami-0e7290665643979b5', 'Username' is 'ec2-user', and 'Verified provider' is checked.

The 'Instance type' section shows 't4g.nano' selected, with details: Family: t4g, 2 vCPU, 0.5 GiB Memory, Current generation: true, On-Demand Ubuntu Pro base pricing: 0.0077 USD per Hour, On-Demand Linux base pricing: 0.0042 USD per Hour, On-Demand SUSE base pricing: 0.0042 USD per Hour. It also mentions 'All generations' and 'Compare instance types'.

N. Ollama deployed in AWS EC2 using Terraform

N.1. To install Ollama in a AWS EC2 instance using Terraform follows the following the .tf file to create the needed resources:

```
terraform {
  required_version = ">= 1.0.0, < 2.0.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 4.0"
    }
  }
}

provider "aws" {
  region      = "us-east-1"
  access_key  = "YOUR AWS ACCESS KEY"
  secret_key  = "YOUR AWS SECRET KEY"
  token       = "YOUR AWS TOKEN"
}

resource "aws_instance" "exampleOllamaConfiguration" {
  ami           = "ami-045269a1f5c90a6a0"
  instance_type = "t2.medium"
  count         = 1
  vpc_security_group_ids = [aws_security_group.instance.id]

  root_block_device {
    volume_size = 24 # In Gb
  }

  key_name  = "vockey"
  user_data = "${file("creation.sh")}"

  user_data_replace_on_change = true

  tags = {
    Name = "terraform-example-ollama"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name

  ingress {
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    from_port   = 11434
    to_port     = 11434
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port     = 0
    to_port       = 0
    protocol     = "-1"
    cidr_blocks  = ["0.0.0.0/0"]
    ipv6_cidr_blocks = [":/:0"]
  }
}

variable "security_group_name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-ollama-example-instance"
}

output "address" {
  value      = aws_instance.exampleOllamaConfiguration[*].public_dns
  description = "Address of the Kafka EC2 machine with ollama"
}
```

N.2. And the following bash script to be executed after EC2 resource creation by the following script **creation.sh**:

```
#!/bin/bash
cd
sudo yum update -y

sudo curl -fsSL https://ollama.com/install.sh | sh

export HOME=$HOME:/usr/local/bin
sudo sed -i "s/\[Install\]/Environment=\"$OLLAMA_HOST=0.0.0.0:11434\"\n\[Install\]/g"
/etc/systemd/system/ollama.service

sudo systemctl enable ollama
sudo systemctl start ollama

ollama pull llama3.2
```

Hint 1: Additionally, if inside the EC2 instance you would like to check the models loaded use the following command:

```
[ec2-user@ip-172-31-23-142 ~]$ ollama list
NAME           ID          SIZE      MODIFIED
llama3.2:1b    baf6a787fdff   1.3 GB    7 minutes ago
```

Hint 2: Additionally, if inside the EC2 instance you would like to check the processors loaded use the following command:

```
[ec2-user@ip-172-31-23-142 ~]$ ollama ps
NAME           ID          SIZE      PROCESSOR      UNTIL
llama3.2:1b    baf6a787fdff   2.2 GB    100% CPU      4 minutes from now
```

Hint 3: Additionally, a request example can be tested with the following remote curl command:

```
curl http://<EC2 DNS NAME>:11434/api/generate -d '{
  "model": "llama3.2",
  "prompt": "Why is the sky blue?",
  "stream": false
}'
```

The response obtained is similar with the following:

```
{"model":"llama3.2","created_at":"2025-02-12T15:18:11.983219976Z","response":"The sky appears blue because of a phenomenon called scattering, which occurs when sunlight interacts with the tiny molecules of gases in the Earth's atmosphere.\n\nHere's a simplified explanation:\n1. **Sunlight**: When the sun shines, it emits white light, which contains all the colors of the visible spectrum (red, orange, yellow, green, blue, indigo, and violet).\n2. **Atmospheric scattering**: As sunlight enters the Earth's atmosphere, it encounters tiny molecules of gases such as nitrogen (N2) and oxygen (O2). These molecules are much smaller than the wavelength of light.\n3. **Scattering**: When sunlight hits these tiny molecules, it scatters in all directions. However, not all wavelengths scatter equally. The shorter (blue) wavelengths are scattered more than the longer (red) wavelengths by the smaller gas molecules.\n4. **Blue light dominates**: As a result of this scattering effect, the blue light is dispersed throughout the atmosphere, while the red and orange light continue to travel in a straight line.\n5. **Our eyes perceive the sky as blue**: From our perspective on the surface of the Earth, we see the scattered blue light coming from all directions, which gives the sky its blue appearance.\n\nThis phenomenon is more pronounced during the daytime when the sun is overhead, and the scattering effect is greater. During sunrise and sunset, the sun's rays have to travel through more of the atmosphere, scattering off even more molecules and creating the warm hues we see.\n\nSo, in short, the sky appears blue because of the scattering of sunlight by tiny molecules in the Earth's atmosphere, which favors shorter wavelengths (like blue light) over longer wavelengths (like red light).","done":true,"done_reason":"stop","context":[]}]
```

```
n":3846465248,"prompt_eval_count":31,"prompt_eval_duration":3108000000,"eval_count":344,"eval_duration":65784000000}  
%
```

References

URLs with other resources

- Terraform Up & Running, Writing Infrastructure as Code, Yevgeniy Brikman, 3rd edition (2022).
<https://www.terraformupandrunning.com/>
 - The github with all the source code from the book is at: <https://github.com/brikis98/terraform-up-and-running-code>
- <https://developer.hashicorp.com/terraform/downloads>
- <https://developer.hashicorp.com/terraform/tutorials/aws-get-started>
- <https://stackoverflow.com/questions/41596412/how-to-use-terraform-output-as-input-variable-of-another-terraform-template>
- <https://developer.hashicorp.com/terraform/tutorials/aws/lambda-api-gateway>
- <https://docs.aws.amazon.com/lambda/latest/dg/lambda-runtimes.html>
- <https://registry.terraform.io/providers/hashicorp/aws/2.34.0/docs/guides/serverless-with-aws-lambda-and-api-gateway>
- Ollama API documentation: <https://github.com/ollama/ollama/blob/main/docs/api.md>
- Ollama FAQs documentation: <https://github.com/ollama/ollama/blob/main/docs/faq.md>
- Ollama readme page <https://github.com/ollama/ollama?tab=readme-ov-file>

The goal of this document is to show how to use the Kafka in a cluster environment.

Consider that all the consumers, producers, and topics management indicated in this tutorial are executed, by the scripts delivered in the Kafka distribution, in a command line, but outside the EC2 AWS instance with your Kafka installation.

For that end, create an additional new EC2 AWS instance with Kafka installed without starting its services. Then you can create as many sessions as you need.

The following contents is presented in this document.

Contents

A.	Listing all topics	2
B.	Creating one topic with multiple partitions	2
C.	Adding partitions for a specific topic	2
D.	Deleting one topic	2
E.	Listing the count of partitions for a specific topic	2
F.	Producing messages to a specific topic with multiple partitions	2
G.	Consuming messages from a topic with multiple partitions	2
H.	Listing the consumer groups.....	2
I.	Listing the consumer-ids from a specific consumer group	2
J.	Defining the group-id in a specific consumer	3
K.	Creating two brokers in the same EC2 VM	3
L.	Checking that cluster of brokers is launched.....	3
M.	Producing messages to a cluster of brokers	4
N.	Consuming messages from a cluster of brokers.....	4
O.	Creating N brokers in different EC2 VMs.....	5
P.	Broker failure test	6
	Appendix: URLs with other resources	8

A. Listing all topics

```
sudo /usr/local/kafka/bin/kafka-topics.sh --list --bootstrap-server <YourIP_or_DNS>:9092

CustomerCountry
EXERCISE3
__consumer_offsets
novo_topico_automatico1
novo_topico_automatico2
novo_topico_automatico3
novotopico
test
```

B. Creating one topic with multiple partitions

```
sudo /usr/local/kafka/bin/kafka-topics.sh --create --bootstrap-server <YourIP_or_DNS>:9092 -replication-factor 1 --partitions 8 --topic clicks

Created topic "clicks".
```

C. Adding partitions for a specific topic

```
sudo /usr/local/kafka/bin/kafka-topics.sh --alter --bootstrap-server <YourIP_or_DNS>:9092 --partitions 10 --topic clicks
```

WARNING: If partitions are increased for a topic that has a key, the partition logic or ordering of the messages will be affected
Adding partitions succeeded!

D. Deleting one topic

```
sudo /usr/local/kafka/bin/kafka-topics.sh --delete --bootstrap-server <YourIP_or_DNS>:9092 --topic clicks

Topic clicks is marked for deletion.
Note: This will have no impact if delete.topic.enable is not set to true.
```

E. Listing the count of partitions for a specific topic

```
/usr/local/kafka/bin/kafka-topics.sh --describe --bootstrap-server <YourIP_or_DNS>:9092 --topic clicks

Topic:clicks    PartitionCount:8    ReplicationFactor:1    Configs:
          Topic: clicks    Partition: 0    Leader: 0        Replicas: 0    Isr: 0
          Topic: clicks    Partition: 1    Leader: 0        Replicas: 0    Isr: 0
          Topic: clicks    Partition: 2    Leader: 0        Replicas: 0    Isr: 0
          Topic: clicks    Partition: 3    Leader: 0        Replicas: 0    Isr: 0
          Topic: clicks    Partition: 4    Leader: 0        Replicas: 0    Isr: 0
          Topic: clicks    Partition: 5    Leader: 0        Replicas: 0    Isr: 0
          Topic: clicks    Partition: 6    Leader: 0        Replicas: 0    Isr: 0
          Topic: clicks    Partition: 7    Leader: 0        Replicas: 0    Isr: 0
```

F. Producing messages to a specific topic with multiple partitions

Remember that:

```
/usr/local/kafka/bin/kafka-console-producer.sh --broker-list <YourIP_or_DNS>:9092 --topic clicks
```

G. Consuming messages from a topic with multiple partitions

Remember that:

```
/usr/local/kafka/bin/kafka-console-consumer.sh --bootstrap-server <YourIP_or_DNS>:9092 --topic clicks --from-beginning
```

H. Listing the consumer groups

```
sudo /usr/local/kafka/bin/kafka-consumer-groups.sh --bootstrap-server <YourIP_or_DNS>:9092 -list

console-consumer-2817
console-consumer-57329
console-consumer-28356
```

I. Listing the consumer-ids from a specific consumer group

```
sudo /usr/local/kafka/bin/kafka-consumer-groups.sh --bootstrap-server <YourIP_or_DNS>:9092 -describe -group console-consumer-28356

TOPIC      PARTITION  CURRENT-OFFSET  LOG-END-OFFSET  LAG      CONSUMER-ID                                     HOST          CLIENT-ID
clicks      1          -              0              -      consumer-1-4924ed20-7fc8-46e4-909a-35c71d5a79fc  /194.210.61.144  consumer-1
clicks      4          -              0              -      consumer-1-4924ed20-7fc8-46e4-909a-35c71d5a79fc  /194.210.61.144  consumer-1
clicks      7          -              0              -      consumer-1-4924ed20-7fc8-46e4-909a-35c71d5a79fc  /194.210.61.144  consumer-1
clicks      0          -              1              -      consumer-1-4924ed20-7fc8-46e4-909a-35c71d5a79fc  /194.210.61.144  consumer-1
clicks      5          -              0              -      consumer-1-4924ed20-7fc8-46e4-909a-35c71d5a79fc  /194.210.61.144  consumer-1
clicks      2          -              1              -      consumer-1-4924ed20-7fc8-46e4-909a-35c71d5a79fc  /194.210.61.144  consumer-1
clicks      3          -              1              -      consumer-1-4924ed20-7fc8-46e4-909a-35c71d5a79fc  /194.210.61.144  consumer-1
clicks      6          -              0              -      consumer-1-4924ed20-7fc8-46e4-909a-35c71d5a79fc  /194.210.61.144  consumer-1
```

J. Defining the group-id in a specific consumer

In command line:

```
sudo /usr/local/kafka/bin/kafka-console-consumer.sh --bootstrap-server <YourIP_or_DNS>:9092 --topic clicks --from-beginning --group group-id-teste
```

In Java:

Define the property:

```
props.put("group.id", "group-id-teste");
```

as explained in the section C. of P3-Kafka-ClientsJAVA document.

K. Creating two brokers in the same EC2 VM

1. Stop the kafka server if running.
2. Create the two configuration files

```
cp /usr/local/kafka/config/server.properties /usr/local/kafka/config/server-1.properties  
cp /usr/local/kafka/config/server.properties /usr/local/kafka/config/server-2.properties
```

3. Edit these new files with the following set of properties (if in the same AWS EC2 instance, otherwise the port and log directory could be the same):

```
config/server-0.properties:  
broker.id=0  
listeners=PLAINTEXT://<YourIP_or_DNS>:9093  
offsets.topic.replication.factor=2  
transaction.state.log.replication.factor=2  
transaction.state.log.min_isr=2  
log.dir=/tmp/kafka-logs-0
```

```
config/server-1.properties:  
broker.id=1  
listeners=PLAINTEXT://<YourIP_or_DNS>:9094  
offsets.topic.replication.factor=2  
transaction.state.log.replication.factor=2  
transaction.state.log.min_isr=2  
log.dir=/tmp/kafka-logs-1
```

4. Start Kafka broker 1 with the command referring to server config file #1

```
sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon /usr/local/kafka/config/server-0.properties
```

5. Start Kafka broker 2 with the command referring to server config file #2

```
sudo /usr/local/kafka/bin/kafka-server-start.sh -daemon /usr/local/kafka/config/server-1.properties
```

6. Open the in-bound ports 9093 and 9094 in the AWS EC2 console

7. Create a topic with replication factor=2 and multiple partitions

```
sudo /usr/local/kafka/bin/kafka-topics.sh --create --bootstrap-server <YourIP_or_DNS>:9092 -replication-factor 2 --partitions 2 --topic testing-cluster
```

Created topic "testing-cluster".

L. Checking that cluster of brokers is launched

Executing the following command:

```
ps -ef |grep java |grep server
```

You should receive something similar with (one PID for each broker using each configuration):

```
root      31619      1  4 08:43 pts/0    00:00:08 java -Xmx1G -Xms1G -server -XX:+UseG1GC -XX:MaxGCPauseMillis=20 -  
XX:InitiatingHeapOccupancyPercent=35 -XX:+ExplicitGCInvokesConcurrent -Djava.awt.headless=true -Xloggc:/usr/local/kafka/bin/..../logs/kafkaServer-  
gc.log -verbose:gc -XX:+PrintGCDetails -XX:+PrintGCDetailedStatistics -XX:+PrintGCDateStamps -XX:+PrintGCTimeStamps -XX:+UseGCLogFileRotation -XX:NumberOfGCLogFiles=10 -  
XX:GCLogFileMaxSize=100M -Dcom.sun.management.jmxremote -Dcom.sun.management.jmxremote.authenticate=false -Dcom.sun.management.jmxremote.ssl=false -  
Dkafka.logs.dir=/usr/local/kafka/bin/..../logs -Dlog4j.configuration=file:/usr/local/kafka/bin/..../config/log4j.properties -cp  
/usr/local/kafka/bin/..../libs/activation-1.1.1.jar:/usr/local/kafka/bin/..../libs/aopalliance-repackaged-2.5.0-  
b42.jar:/usr/local/kafka/bin/..../libs/argparse4j-0.7.0.jar:/usr/local/kafka/bin/..../libs/commons-lang3-  
0.5.0.jar:/usr/local/kafka/bin/..../libs/commons-logging-1.1.3.jar:/usr/local/kafka/bin/..../libs/connect-basic-auth-extension-2.1.0.jar:/usr/local/kafka/bin/..../libs/connect-file-  
2.1.0.jar:/usr/local/kafka/bin/..../libs/connect-json-2.1.0.jar:/usr/local/kafka/bin/..../libs/connect-runtime-  
2.1.0.jar:/usr/local/kafka/bin/..../libs/connect-transforms-2.1.0.jar:/usr/local/kafka/bin/..../libs/guava-20.0.jar:/usr/local/kafka/bin/..../libs/hk2-  
api-2.5.0-b42.jar:/usr/local/kafka/bin/..../libs/hk2-locator-2.5.0-b42.jar:/usr/local/kafka/bin/..../libs/hk2-utils-2.5.0-
```

```

b42.jar:/usr/local/kafka/bin/../libs/jackson-annotations-2.9.7.jar:/usr/local/kafka/bin/../libs/jackson-core-2.9.7.jar:/usr/local/kafka/bin/../libs/jackson-databind-2.9.7.jar:/usr/local/kafka/bin/../libs/jackson-jaxrs-base-2.9.7.jar:/usr/local/kafka/bin/../libs/jackson-jaxrs-json-provider-2.9.7.jar:/usr/local/kafka/bin/../libs/jackson-module-jaxb-annotations-2.9.7.jar:/usr/local/kafka/bin/../libs/javassist-3.22.0-CR2.jar:/usr/local/kafka/bin/../libs/javax.annotation-api-1.2.jar:/usr/local/kafka/bin/../libs/javax.inject-1.jar:/usr/local/kafka/bin/../libs/javax.inject-2.5.0-b42.jar:/usr/local/kafka/bin/../libs/javax.servlet-api-3.1.0.jar:/usr/local/kafka/bin/../libs/javax.ws.rs-api-2.1.1.jar:/usr/local/kafka/bin/../libs/javax.ws.rs-api-2.1.jar:/usr/local/kafka/bin/../libs/jaxb-api-2.3.0.jar:/usr/local/jersey-client-2.27.jar:/usr/local/kafka/bin/../libs/jersey-common-2.27.jar:/usr/local/kafka/bin/../libs/jersey-container-servlet-2.27.jar:/usr/local/kafka/bin/../libs/jersey-container-servlet-core-2.27.jar:/usr/local/kafka/bin/../libs/jersey-hk2-2.27.jar:/usr/local/kafka/bin/../libs/jersey-media-jaxb-2.27.jar:/usr/local/kafka/bin/../libs/jersey-server-2.27.jar:/usr/local/kafka/bin/../libs/jetty-client-9.4.12.v20180830.jar:/usr/local/kafka/bin/../libs/jetty-continuation-9.4.12.v20180830.jar:/usr/local/kafka/bin/../libs/jetty-http-9.4.12.v20180830.jar:/usr/local/kafka/bin/../libs/jetty-io-9.4.12.v20180830.jar:/usr/local/kafka/bin/../libs/jetty-security-9.4.12.v20180830.jar:/usr/local/kafka/bin/../libs/jetty-server-9.4.12.v20180830.jar:/usr/local/kafka/bin/../libs/jetty-servlet-9.4.12.v20180830.jar:/usr/local/kafka/bin/../libs/jetty-servlets-9.4.12.v20180830.jar:/usr/local/kafka/bin/../libs/jetty-util-9.4.12.v20180830.jar:/usr/local/kafka/bin/../libs/jopt-simple-5.0.4.jar:/usr/local/kafka/bin/../libs/kafka_2.12-2.1.0.jar:/usr/local/kafka/bin/../libs/kafka-appender-2.1.0.jar:/usr/local/kafka/bin/../libs/kafka-clients-2.1.0.jar:/usr/local/kafka/bin/../libs/kafka-log4j-appender-2.1.0.jar:/usr/local/kafka/bin/../libs/kafka-streams-2.1.0.jar:/usr/local/kafka/bin/../libs/kafka-streams-examples-2.1.0.jar:/usr/local/kafka/bin/../libs/kafka-streams-scala_2.12-2.1.0.jar:/usr/local/kafka/bin/../libs/kafka-streams-test-utils-2.1.0.jar:/usr/local/kafka/bin/../libs/kafka-tools-2.1.0.jar:/usr/local/kafka/bin/../libs/kafka-bin-1.2.17.jar:/usr/local/kafka/bin/..libs/lz4-java-1.5.0.jar:/usr/local/kafka/bin/../libs/maven-artifact-3.5.4.jar:/usr/local/kafka/bin/../libs/plexus-utils-2.2.0.jar:/usr/local/kafka/bin/../libs/osgi-resource-locator-1.0.1.jar:/usr/local/kafka/bin/../libs/reflections-0.9.11.jar:/usr/local/kafka/bin/../libs/rocksdbjni-5.14.2.jar:/usr/local/kafka/bin/../libs/scala-library-2.12.7.jar:/usr/local/kafka/bin/../libs/scala-logging-2.12-3.9.0.jar:/usr/local/kafka/bin/../libs/scala-reflect-2.12.7.jar:/usr/local/kafka/bin/..libs/slf4j-api-1.7.25.jar:/usr/local/kafka/bin/..libs/slf4j-log4j12-1.7.25.jar:/usr/local/kafka/bin/..libs/snappy-java-1.1.7.2.jar:/usr/local/kafka/bin/..libs/validation-api-1.1.0.Final.jar:/usr/local/kafka/bin/..libs/zkclient-0.10.jar:/usr/local/kafka/bin/..libs/zookeeper-3.4.13.jar:/usr/local/kafka/bin/..libs/zstd-jni-1.3.5-4.jar kafka.Kafka /usr/local/kafka/config/server-0.properties

root 31959 1 3 08:43 pts/0 00:00:06 java -Xmx1G -Xms1G -server -XX:+UseG1GC -XX:MaxGCPauseMillis=20 -XX:InitiatingHeapOccupancyPercent=35 -XX:+ExplicitGCInvokesConcurrent -Djava.awt.headless=true -Xloggc:/usr/local/kafka/bin/..logs/kafkaServer-gc.log -verbose:gc -XX:+PrintGCDetails -XX:+PrintGCDateStamps -XX:+PrintGCTimeStamps -XX:+UseGCLogFileRotation -XX:NumberOfGCLogFiles=10 -XX:GCLogFileSize=100M -Dcom.sun.management.jmxremote -Dcom.sun.management.jmxremote.authenticate=false -Dcom.sun.management.jmxremote.ssl=false -Dkafka.logs.dir=/usr/local/kafka/bin/..logs -Dlog4j.configuration=file:/usr/local/kafka/bin/..config/log4j.properties -cp /usr/local/kafka/bin/..libs/activation-1.1.1.jar:/usr/local/kafka/bin/..libs/aopalliance-repackaged-2.5.0-b42.jar:/usr/local/kafka/bin/..libs/argparse4j-0.7.0.jar:/usr/local/kafka/bin/..libs/audience-annotations-0.5.0.jar:/usr/local/kafka/bin/..libs/commons-lang3-3.5.jar:/usr/local/kafka/bin/..libs/compileScala.mapping:/usr/local/kafka/bin/..libs/compileScala.mapping.asc:/usr/local/kafka/bin/..libs/connect-api-2.1.0.jar:/usr/local/kafka/bin/..libs/connect-basic-auth-extension-2.1.0.jar:/usr/local/kafka/bin/..libs/connect-file-2.1.0.jar:/usr/local/kafka/bin/..libs/connect-json-2.1.0.jar:/usr/local/kafka/bin/..libs/connect-runtime-2.1.0.jar:/usr/local/kafka/bin/..libs/connect-transforms-2.1.0.jar:/usr/local/kafka/bin/..libs/guava-20.0.jar:/usr/local/kafka/bin/..libs/hk2-api-2.5.0-b42.jar:/usr/local/kafka/bin/..libs/hk2-locator-2.5.0-b42.jar:/usr/local/kafka/bin/..libs/hk2-utils-2.5.0-b42.jar:/usr/local/kafka/bin/..libs/jackson-annotations-2.9.7.jar:/usr/local/kafka/bin/..libs/jackson-core-2.9.7.jar:/usr/local/kafka/bin/..libs/jackson-databind-2.9.7.jar:/usr/local/kafka/bin/..libs/jackson-jaxrs-base-2.9.7.jar:/usr/local/kafka/bin/..libs/jackson-jaxrs-json-provider-2.9.7.jar:/usr/local/kafka/bin/..libs/jackson-module-jaxb-annotations-2.9.7.jar:/usr/local/kafka/bin/..libs/javassist-3.22.0-CR2.jar:/usr/local/kafka/bin/..libs/javax.annotation-api-1.2.jar:/usr/local/kafka/bin/..libs/javax.inject-1.jar:/usr/local/kafka/bin/..libs/javax.ws.rs-api-2.1.1.jar:/usr/local/kafka/bin/..libs/javax.ws.rs-api-2.1.jar:/usr/local/kafka/bin/..libs/jaxb-api-2.3.0.jar:/usr/local/kafka/bin/..libs/jersey-client-2.27.jar:/usr/local/kafka/bin/..libs/jersey-common-2.27.jar:/usr/local/kafka/bin/..libs/jersey-container-servlet-2.27.jar:/usr/local/kafka/bin/..libs/jersey-container-servlet-core-2.27.jar:/usr/local/kafka/bin/..libs/jersey-hk2-2.27.jar:/usr/local/kafka/bin/..libs/jersey-media-jaxb-2.27.jar:/usr/local/kafka/bin/..libs/jersey-server-2.27.jar:/usr/local/kafka/bin/..libs/jetty-client-9.4.12.v20180830.jar:/usr/local/kafka/bin/..libs/jetty-continuation-9.4.12.v20180830.jar:/usr/local/kafka/bin/..libs/jetty-http-9.4.12.v20180830.jar:/usr/local/kafka/bin/..libs/jetty-io-9.4.12.v20180830.jar:/usr/local/kafka/bin/..libs/jetty-security-9.4.12.v20180830.jar:/usr/local/kafka/bin/..libs/jetty-server-9.4.12.v20180830.jar:/usr/local/kafka/bin/..libs/jetty-servlet-9.4.12.v20180830.jar:/usr/local/kafka/bin/..libs/jetty-servlets-9.4.12.v20180830.jar:/usr/local/kafka/bin/..libs/jetty-util-9.4.12.v20180830.jar:/usr/local/kafka/bin/..libs/jopt-simple-5.0.4.jar:/usr/local/kafka/bin/..libs/kafka_2.12-2.1.0-sources.jar:/usr/local/kafka/bin/..libs/kafka-clients-2.1.0.jar:/usr/local/kafka/bin/..libs/kafka-log4j-appender-2.1.0.jar:/usr/local/kafka/bin/..libs/kafka-streams-2.1.0.jar:/usr/local/kafka/bin/..libs/kafka-streams-examples-2.1.0.jar:/usr/local/kafka/bin/..libs/kafka-streams-scala_2.12-2.1.0.jar:/usr/local/kafka/bin/..libs/kafka-tools-2.1.0.jar:/usr/local/kafka/bin/..libs/lz4-java-1.5.0.jar:/usr/local/kafka/bin/..libs/maven-artifact-3.5.4.jar:/usr/local/kafka/bin/..libs/plexus-utils-2.2.0.jar:/usr/local/kafka/bin/..libs/osgi-resource-locator-1.0.1.jar:/usr/local/kafka/bin/..libs/reflections-0.9.11.jar:/usr/local/kafka/bin/..libs/rocksdbjni-5.14.2.jar:/usr/local/kafka/bin/..libs/scala-library-2.12.7.jar:/usr/local/kafka/bin/..libs/scala-logging-2.12-3.9.0.jar:/usr/local/kafka/bin/..libs/scala-reflect-2.12.7.jar:/usr/local/kafka/bin/..libs/slf4j-api-1.7.25.jar:/usr/local/kafka/bin/..libs/slf4j-log4j12-1.7.25.jar:/usr/local/kafka/bin/..libs/snappy-java-1.1.7.2.jar:/usr/local/kafka/bin/..libs/validation-api-1.1.0.Final.jar:/usr/local/kafka/bin/..libs/zkclient-0.10.jar:/usr/local/kafka/bin/..libs/zookeeper-3.4.13.jar:/usr/local/kafka/bin/..libs/zstd-jni-1.3.5-4.jar kafka.Kafka /usr/local/kafka/config/server-1.properties

```

The topic can also be described:

```

sudo /usr/local/kafka/bin/kafka-topics.sh --describe --topic testing-cluster --bootstrap-server <Your_IP_or_DNS>:9092

Topic:testing-cluster PartitionCount:2 ReplicationFactor:2 Configs:
Topic: testing-cluster Partition: 0 Leader: 1 Replicas: 1,0 Isr: 1,0
Topic: testing-cluster Partition: 1 Leader: 0 Replicas: 0,1 Isr: 0,1

```

M. Producing messages to a cluster of brokers

Remember that:

```
/usr/local/kafka/bin/kafka-console-producer.sh --broker-list <YourIP_or_DNS>:9093,<YourIP_or_DNS>:9094 --topic testing-cluster
```

N. Consuming messages from a cluster of brokers

Remember that:

```
/usr/local/kafka/bin/kafka-console-consumer.sh --bootstrap-server <YourIP_or_DNS>:9093,<YourIP_or_DNS>:9094 --topic testing-cluster
```

O. Creating N brokers in different EC2 VMs

1. Create multiple EC2 VMs adding the following inbound rules for ports: 2888 and 3888

Hint: to facilitate the creation of multiple instances you can create an image from one instance. Then, when you launch new instances select your new image. With this approach all the kafka installation procedure is performed only once.

2. Install zookeeper and Kafka in each EC2 VM
3. If previously started, then stop the Zookeeper and Kafka servers in each EC2 VM
4. In each EC2 VM change Kafka server configuration file in `/usr/local/kafka/config/server.properties` accordingly:

```
# the broker ID starting on 0 (this is for broker 2 within a 3 brokers cluster)
broker.id=2

# your EC2 VM name to be accessed from outside
listeners=PLAINTEXT://ec2-54-236-47-54.compute-1.amazonaws.com:9092

# replication factor accordingly with the number of brokers
offsets.topic.replication.factor=3
transaction.state.log.replication.factor=3
transaction.state.log.min.isr=3

#zookeeper connectivity (one per EC2 VM of this cluster)
zookeeper.connect=ec2-54-90-57-82.compute-1.amazonaws.com:2181,ec2-54-173-171-63.compute-1.amazonaws.com:2181,ec2-54-236-47-54.compute-1.amazonaws.com:2181
```

5. In each EC2 VM change the Zookeeper server configuration file in `/usr/local/zookeeper/conf/zoo.cfg` accordingly:

```
tickTime=2000
dataDir=/var/lib/zookeeper
clientPort=2181
maxClientCnxns=60
initLimit=10
syncLimit=5
server.1=ec2-54-90-57-82.compute-1.amazonaws.com:2888:3888
server.2=ec2-54-173-171-63.compute-1.amazonaws.com:2888:3888
server.3=ec2-54-236-47-54.compute-1.amazonaws.com:2888:3888
```

6. In each EC2 VM create a file with the ID of the Zookeeper server in `/var/lib/zookeeper/myid` containing only the ID of each specific Zookeeper server, example:

```
[ec2-user@ip-172-31-50-87 conf]$ cd /var/lib/zookeeper/
[ec2-user@ip-172-31-50-87 zookeeper]$ ls -ltr
total 8
-rw-r--r-- 1 root root 2 Feb 2 14:45 myid
-rw-r--r-- 1 root root 4 Feb 2 16:14 zookeeper_server.pid
drwxr-xr-x 2 root root 273 Feb 2 16:19 version-2
[ec2-user@ip-172-31-50-87 zookeeper]$ more myid
3
```

7. In each EC2 VM start the Kafka and Zookeeper servers.
8. Check the available topics:

```
sudo /usr/local/kafka/bin/kafka-topics.sh --list --bootstrap-server ec2-54-90-57-82.compute-1.amazonaws.com:9092
```

9. Create a new topic, for example in a 3 broker cluster:

```
sudo /usr/local/kafka/bin/kafka-topics.sh --create --bootstrap-server ec2-54-173-171-63.compute-1.amazonaws.com:9092 -replication-factor 3 --partitions 6 --topic PURCHASE
```

10. Check how the topic has been created:

```
/usr/local/kafka/bin/kafka-topics.sh --describe --bootstrap-server ec2-54-236-47-54.compute-1.amazonaws.com:9092 --topic PURCHASE
```

11. Start a new consumer (in a different machine, VM or session) for the new topic using the predefined cluster:

```
/usr/local/kafka/bin/kafka-console-producer.sh --broker-list ec2-54-90-57-82.compute-1.amazonaws.com:9092,ec2-54-173-171-63.compute-1.amazonaws.com:9092,ec2-54-236-47-54.compute-1.amazonaws.com:9092 --topic PURCHASE
```

12. Start a new producer (in a different machine, VM or session) for the new topic using the predefined cluster and start sending messaging. Verify if they are consumed correctly:

```
/usr/local/kafka/bin/kafka-console-producer.sh --broker-list ec2-54-90-57-82.compute-1.amazonaws.com:9092,ec2-54-173-171-63.compute-1.amazonaws.com:9092,ec2-54-236-47-54.compute-1.amazonaws.com:9092 --topic PURCHASE
```

13. Try to stop one Kafka server from the cluster and recheck the topic configuration with:

```
/usr/local/kafka/bin/kafka-topics.sh --describe --bootstrap-server ec2-54-236-47-54.compute-1.amazonaws.com:9092 --topic PURCHASE
```

14. Restart the Kafka server and execute the kafka-leader-election to refresh the topic' cluster configuration:

```
/usr/local/kafka/bin/kafka-leader-election.sh --bootstrap-server ec2-54-90-57-82.compute-1.amazonaws.com:9092,ec2-54-173-171-63.compute-1.amazonaws.com:9092,ec2-54-236-47-54.compute-1.amazonaws.com:9092 --election-type preferred --all-topic-partitions
```

Hints: if needed the Kafka server logs are located at: /usr/local/kafka/logs/kafkaServer.out where you can see the execution logging.

If you are using already used kafka servers, to make them as a cluster, you may need to remove all old data. Simply remove the /tmp/kafka-logs directory, and recreate it empty.

P. Broker failure test

- Launch the cluster with 2 brokers
- Create a new topic with 4 partitions

```
sudo /usr/local/kafka/bin/kafka-topics.sh --create --bootstrap-server <YourIP_or_DNS>:9092 -replication-factor 2 --partitions 4 --topic events
```

Created topic "events".

- Check who are the leaders:

```
/usr/local/kafka/bin/kafka-topics.sh --describe --bootstrap-server <YourIP_or_DNS>:9092 --topic events
```

Topic:events	PartitionCount:4	ReplicationFactor:2	Configs:
Topic: events	Partition: 0	Leader: 0	Replicas: 0,1 Isr: 0,1
Topic: events	Partition: 1	Leader: 1	Replicas: 1,0 Isr: 1,0
Topic: events	Partition: 2	Leader: 0	Replicas: 0,1 Isr: 0,1
Topic: events	Partition: 3	Leader: 1	Replicas: 1,0 Isr: 1,0

- Produce message to the topic

```
/usr/local/kafka/bin/kafka-console-producer.sh --broker-list <YourIP_or_DNS>:9093, <YourIP_or_DNS>:9094 --topic events
```

- Consume message from the topic <start two different consumer sessions>

```
/usr/local/kafka/bin/kafka-console-consumer.sh --bootstrap-server <YourIP_or_DNS>:9093, <YourIP_or_DNS>:9094 --topic events --group group-id-teste-cluster
```

- Check that messages are being balanced between the two consumers
- Stop broker id 0 with the following command:

```
sudo kill `ps aux | grep java | grep server-0.properties | tr -s " " | cut -d " " -f2`
```

- Check the change in the leaders of the cluster

```
/usr/local/kafka/bin/kafka-topics.sh --describe --bootstrap-server <YourIP_or_DNS>:9092 --topic events
```

```
Topic:events    PartitionCount:4      ReplicationFactor:2      Configs:  
  Topic: events    Partition: 0        Leader: 1            Replicas: 0,1    Isr: 1  
  Topic: events    Partition: 1        Leader: 1            Replicas: 1,0    Isr: 1  
  Topic: events    Partition: 2        Leader: 1            Replicas: 0,1    Isr: 1  
  Topic: events    Partition: 3        Leader: 1            Replicas: 1,0    Isr: 1
```

- Check that messages are still being balanced between the two consumers

Appendix: URLs with other resources

- <https://github.com/confluentinc/confluent-kafka-python/issues/187>
- <http://cloudurable.com/blog/kafka-tutorial-kafka-failover-kafka-cluster/index.html>
- if kafka brokers IDs are corrupted: <https://stackoverflow.com/questions/59592518/kafka-broker-doesnt-find-cluster-id-and-creates-new-one-after-docker-restart>

The goal of this document is to introduce the fundamentals of stream processing using Kafka. Three examples taken from: **Narkhede, N., Shapira, G., & Palino, T. (2017). Kafka: the definitive guide: real-time data and stream processing at scale.** " O'Reilly Media, Inc.", and are used to illustrate (i) how to develop the source code and (ii) how to test it.

First example is a simple word count example that is used to demonstrate the map/filter pattern and simple aggregates.

Second example contains a calculation of different statistics on stock market trades, which will allow us to demonstrate window aggregations.

Third example demonstrates the streaming joins.

The following contents is presented in this document.

Contents

A.	Installing git in AWS EC2	2
B.	Installing mvn in AWS EC2	3
C.	Testing the Word count example	4
D.	Testing the Stock Market Statistics example.....	7
E.	Testing the click Stream Enrichment example	10
F.	Other tools and references.....	13

A. Installing git in AWS EC2

In command line:

```
sudo yum install git
```

Check installation:

```
git --version  
git version 2.47.1 (or later)
```

B. Installing mvn in AWS EC2

1. Create an EC2 Amazon linux image instance
2. Check that JDK 8 is available in the new created instance with the following commands:

```
[ec2-user@ip---- ~]$ java -version
openjdk version "17.0.14" 2025-01-21 LTS
OpenJDK Runtime Environment Corretto-17.0.14.7.1 (build 17.0.14+7-LTS)
OpenJDK 64-Bit Server VM Corretto-17.0.14.7.1 (build 17.0.14+7-LTS, mixed mode, sharing)
[ec2-user@ip-172-31-85-55 ~]$ javac -version
javac 17.0.14
```

Hint 1: if not available, use the following commands or similar:

```
[ec2-user@ip- ~]$ yum search java-17
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
=====
N/S matched: java-17
=====
java-17-amazon-corretto.x86_64 : Amazon Corretto development environment
java-17-amazon-corretto-debugsymbols.x86_64 : Amazon Corretto 17 zipped debug symbols
java-17-amazon-corretto-devel.x86_64 : Amazon Corretto 17 development tools
java-17-amazon-corretto-headless.x86_64 : Amazon Corretto headless development environment
java-17-amazon-corretto-javadoc.x86_64 : Amazon Corretto 17 API documentation
java-17-amazon-corretto-jmods.x86_64 : Amazon Corretto 17 jmods

[ec2-user@ip- ~]$ sudo yum -y install java-17-amazon-corretto-devel.x86_64
```

3. Execute the following commands in the EC2 instance:

```
[ec2-user@ip- - - - ~]$ sudo wget https://repo.maven.apache.org/maven2/org/apache/maven/apache-maven/3.9.9/apache-maven-3.9.9-bin.tar.gz
[ec2-user@ip- - - - ~]$ tar xzvf apache-maven-3.9.9-bin.tar.gz
[ec2-user@ip- - - - ~]$ cd apache-maven-3.9.9/bin
[ec2-user@ip- - - - bin]$ export PATH=/home/ec2-user/apache-maven-3.9.9/bin:$PATH

[ec2-user@ip- - - - ~]$ mvn -version
Maven home: /home/ec2-user/apache-maven-3.9.9
Java version: 17.0.14, vendor: Amazon.com Inc., runtime: /usr/lib/jvm/java-17-amazon-corretto.x86_64
Default locale: en_US, platform encoding: ANSI_X3.4-1968
OS name: "linux", version: "5.10.235-227.919.amzn2.x86_64", arch: "amd64", family: "unix"
```

Hint 2: To automate the maven path with instance login update the export command in the file:

.bash_profile

C. Testing the Word count example

C.1. Create working directory

```
mkdir WordCount
```

C.2. Change to working directory

```
cd WordCount
```

C.3. Clone the source code from git

```
[ec2-user@ip- WordCount]$ git clone https://github.com/gwenshap/kafka-streams-wordcount
Cloning into 'kafka-streams-wordcount'...
remote: Enumerating objects: 57, done.
remote: Total 57 (delta 0), reused 0 (delta 0), pack-reused 57
Unpacking objects: 100% (57/57), done.
```

C.4. Change directory

```
cd kafka-streams-wordcount/
```

C.5. Change pom.xml from:

```
<dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-streams</artifactId>
    <version>0.10.3.0-SNAPSHOT</version>
</dependency>
```

To a most recent library¹:

```
<dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-streams</artifactId>
    <version>3.2.0</version>
</dependency>
```

C.6. Study the following source code that is also located in:

```
/home/ec2-user/WordCount/kafka-streams-wordcount/src/main/java/com/shapira/examplesstreams/wordcount/
WordCountExample.java
```

```
1 package com.shapira.examplesstreams.wordcount;

2 import org.apache.kafka.clients.CommonClientConfigs;
3 import org.apache.kafka.clients.consumer.ConsumerConfig;
4 import org.apache.kafka.common.serialization.Serdes;
5 import org.apache.kafka.streams.KafkaStreams;
6 import org.apache.kafka.streams.KeyValue;
7 import org.apache.kafka.streams.StreamsConfig;
8 import org.apache.kafka.streams.kstream.KStream;
9 import org.apache.kafka.streams.kstream.KStreamBuilder;

10 import java.util.Arrays;
11 import java.util.Properties;
12 import java.util.regex.Pattern;

13 public class WordCountExample {

14     public static void main(String[] args) throws Exception{

15         Properties props = new Properties();
16         props.put(StreamsConfig.APPLICATION_ID_CONFIG, "wordcount");
17         props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
18         props.put(StreamsConfig.KEY_SERDE_CLASS_CONFIG, Serdes.String().getClass().getName());
19         props.put(StreamsConfig.VALUE_SERDE_CLASS_CONFIG, Serdes.String().getClass().getName());
```

¹ The most recent libraries should be consulted at <https://mvnrepository.com/>

```

20 // setting offset reset to earliest so that we can re-run the demo code with the same pre-loaded data
21 // Note: To re-run the demo, you need to use the offset reset tool:
22 // https://cwiki.apache.org/confluence/display/KAFKA/Kafka+Streams+Application+Reset+Tool
23 props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

24 // work-around for an issue around timing of creating internal topics
25 // Fixed in Kafka 0.10.2.0
26 // don't use in large production apps - this increases network load
27 // props.put(CommonClientConfigs.METADATA_MAX_AGE_CONFIG, 500);

28 KStreamBuilder builder = new KStreamBuilder();

29 KStream<String, String> source = builder.stream("wordcount-input");

30 final Pattern pattern = Pattern.compile("\\\\W+");
31 KStream counts = source.flatMapValues(value-> Arrays.asList(pattern.split(value.toLowerCase())))
32 .map((key, value) -> new KeyValue<Object, Object>(value, value))
33 .filter((key, value) -> (!value.equals("the")))
34 .groupByKey()
35 .count("CountStore").mapValues(value->Long.toString(value)).toStream();
36 counts.to("wordcount-output");

37 KafkaStreams streams = new KafkaStreams(builder, props);

38 // This is for reset to work. Don't use in production - it causes the app to re-load the state from Kafka
39 // on every start
39 streams.cleanUp();

40 streams.start();

41 // usually the stream application would be running forever,
42 // in this example we just let it run for some time and stop since the input data is finite.
43 Thread.sleep(5000L);

44 streams.close();

45 }
46 }

```

C.7. Change the location of your kafka server in the java source code:

```
/home/ec2-user/WordCount/kafka-streams-wordcount/src/main/java/com/shapira/examplesstreams/wordcount/
WordCountExample.java
```

In the following source code line:

```
props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "<YOUR_DNS_NAME>:9092");
```

C.8. Build the project with mvn package, this will generate an uber-jar with the streams app and all its dependencies.

```
mvn package
```

C.9. Create a wordcount-input topic:

```
/usr/local/kafka/bin/kafka-topics.sh --bootstrap-server <YOUR DNS NAME>:9092 --create --topic wordcount-input
--partitions 1 --replication-factor 1
```

C.10. Produce some text to the topic. Don't forget to repeat words (so we can count higher than 1) and to use the word "the", so we can filter it.

```
/usr/local/kafka/bin/kafka-console-producer.sh --broker-list <YOUR_DNS_NAME>:9092 --topic wordcount-input
```

C.11. Run the app:

```
java -cp target/uber-kafka-streams-wordcount-1.0-SNAPSHOT.jar
com.shapira.examplesstreams.wordcount.WordCountExample
```

C.12. View the results obtained:

```
/usr/local/kafka/bin/kafka-console-consumer.sh --topic wordcount-output --from-beginning --bootstrap-server <YOUR_DNS_NAME>:9092 --property print.key=true
```

You will receive something similar with:

```
um      1  
este    2  
outro   1  
testte  1teste 1  
tets    1  
test    4  
      1  
a      1
```

C.13. If you want to reset state and re-run the application (maybe with some changes?) on existing input topic, you can stop the kafka producer and consumer and then:

Reset internal topics (used for shuffle the topic and state-stores):

```
/usr/local/kafka/bin/kafka-streams-application-reset.sh --application-id wordcount --bootstrap-servers <YOUR_DNS_NAME>:9092 --input-topics wordcount-input
```

(optional) Delete the output topic:

```
/usr/local/kafka/bin/kafka-topics.sh --bootstrap-server <YOUR_DNS_NAME>:9092 --delete --topic wordcount-output
```

D. Testing the Stock Market Statistics example

D.1. Create working directory

```
mkdir StockMarketStatistics
```

D.2. Change to working directory

```
cd StockMarketStatistics
```

D.3. Clone the source code from git

```
[ec2-user@ip- StockMarketStatistics]$ git clone https://github.com/gwenshap/kafka-streams-stockstats
Cloning into 'kafka-streams-stockstats'...
remote: Enumerating objects: 173, done.
remote: Total 173 (delta 0), reused 0 (delta 0), pack-reused 173
Receiving objects: 100% (173/173), 27.77 KiB | 3.08 MiB/s, done.
Resolving deltas: 100% (45/45), done.
```

D.4. Change pom.xml from:

```
<properties>
    <kafka.version>2.8.0</kafka.version>
    <confluent.version>6.1.0</confluent.version>
    <avro.version>1.8.2</avro.version>
    <project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>
</properties>
```

To a most recent library²:

```
<properties>
    <kafka.version>3.2.0</kafka.version>
    <confluent.version>6.1.0</confluent.version>
    <avro.version>1.8.2</avro.version>
    <project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>
</properties>
```

D.5. Study the following source code that is also located in:

```
/home/ec2-user/StockMarketStatistics/kafka-streams-stockstats/src/main/java/com/Shapira/examples/streams/
stockstats/StockStatsExample.java
```

```
1 package com.shapira.examples.streams.stockstats;
2
3 import com.shapira.examples.streams.stockstats.serde.JsonDeserializer;
4 import com.shapira.examples.streams.stockstats.serde.JsonSerializer;
5 import com.shapira.examples.streams.stockstats.serde.WrapperSerde;
6 import com.shapira.examples.streams.stockstats.model.Trade;
7 import com.shapira.examples.streams.stockstats.model.TradeStats;
8 import org.apache.kafka.clients.admin.AdminClient;
9 import org.apache.kafka.clients.admin.DescribeClusterResult;
10 import org.apache.kafka.clients.consumer.ConsumerConfig;
11 import org.apache.kafka.common.serialization.Serdes;
12 import org.apache.kafka.common.utils.Bytes;
13 import org.apache.kafka.streams.KafkaStreams;
14 import org.apache.kafka.streams.StreamsConfig;
15 import org.apache.kafka.streams.Topology;
16 import org.apache.kafka.streams.kstream.KStream;
17 import org.apache.kafka.streams.StreamsBuilder;
18 import org.apache.kafka.streams.kstream.Materialized;
19 import org.apache.kafka.streams.kstream.Produced;
20 import org.apache.kafka.streams.kstream.TimeWindows;
21 import org.apache.kafka.streams.kstream.Windowed;
22 import org.apache.kafka.streams.kstream.WindowedSerdes;
23 import org.apache.kafka.streams.state.WindowStore;
```

² The most recent libraries should be consulted at <https://mvnrepository.com/>

```

23 import java.util.Properties;
24 /**
25 Input is a stream of trades
26 Output is two streams: One with minimum and avg "ASK" price for every 10 seconds window
27 Another with the top-3 stocks with lowest minimum ask every minute
28 */
29 public class StockStatsExample {
30
31     public static void main(String[] args) throws Exception {
32
33         Properties props;
34         if (args.length==1)
35             props = LoadConfigs.loadConfig(args[0]);
36         else
37             props = LoadConfigs.loadConfig();
38
39         props.put(StreamsConfig.APPLICATION_ID_CONFIG, "stockstat-2");
40         props.put(StreamsConfig.DEFAULT KEY SERDE CLASS CONFIG, Serdes.String().getClass().getName());
41         props.put(StreamsConfig.DEFAULT VALUE SERDE CLASS CONFIG, TradeSerde.class.getName());
42
43         // setting offset reset to earliest so that we can re-run the demo code with the same pre-loaded data
44         // Note: To re-run the demo, you need to use the offset reset tool:
45         // https://cwiki.apache.org/confluence/display/KAFKA/Kafka+Streams+Application+Reset+Tool
46         props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");
47
48         // creating an AdminClient and checking the number of brokers in the cluster, so I'll know how many
49         // replicas we want...
50
51         AdminClient ac = AdminClient.create(props);
52         DescribeClusterResult dcr = ac.describeCluster();
53         int clusterSize = dcr.nodes().get().size();
54
55         if (clusterSize<3)
56             props.put("replication.factor",clusterSize);
57         else
58             props.put("replication.factor",3);
59
60         StreamsBuilder builder = new StreamsBuilder();
61
62         KStream<String, Trade> source = builder.stream(Constants.STOCK_TOPIC);
63
64         KStream<Windowed<String>, TradeStats> stats = source
65             .groupByKey()
66             .windowedBy(TimeWindows.of(5000).advanceBy(1000))
67             .<TradeStats>aggregate(() -> new TradeStats(), (k, v, tradestats) -> tradestats.add(v),
68             Materialized.<String, TradeStats, WindowStore<Bytes, byte[]>>as("trade-aggregates")
69             .withValueSerde(new TradeStatsSerde())
70             .toStream()
71             .mapValues((trade) -> trade.computeAvgPrice());
72
73         stats.to("stockstats-output", Produced.keySerde(WindowedSerdes.timeWindowedSerdeFrom(String.class)));
74
75         Topology topology = builder.build();
76
77         KafkaStreams streams = new KafkaStreams(topology, props);
78
79         System.out.println(topology.describe());
80
81         streams.cleanUp();
82
83         streams.start();
84
85         // Add shutdown hook to respond to SIGTERM and gracefully close Kafka Streams
86         Runtime.getRuntime().addShutdownHook(new Thread(streams::close));
87
88     }
89
90     static public final class TradeSerde extends WrapperSerde<Trade> {
91         public TradeSerde() {
92             super(new JsonSerializer<Trade>(), new JsonDeserializer<Trade>(Trade.class));
93         }
94     }
95
96     static public final class TradeStatsSerde extends WrapperSerde<TradeStats> {
97         public TradeStatsSerde() {
98             super(new JsonSerializer<TradeStats>(), new JsonDeserializer<TradeStats>(TradeStats.class));
99         }
100    }
101}

```

D.6. Change directory

```
cd kafka-streams-stockstats
```

D.7. Build the project with mvn package, this will generate an uber-jar with the streams app and all its dependencies.

```
mvn package
```

D.8. Change to working directory

```
cd kafka-streams-stockstats
```

D.9. Create a stocks input topic and output topic:

```
/usr/local/kafka/bin/kafka-topics.sh --bootstrap-server localhost:9092 --create --topic stocks --partitions 1 --replication-factor 1
```

```
/usr/local/kafka/bin/kafka-topics.sh --bootstrap-server localhost:9092 --create --topic stockstats-output --partitions 1 --replication-factor 1
```

D.10. Create the configuration file, next to you JAVA project, with the command:

```
cat > configfile << EOF
```

and then write the following content to the file directly in the command line:

```
bootstrap.servers = PLAINTEXT://<YOUR DNS NAME>:9092
EOF
```

D.11. Generate some trades so we can analyze them. Start running the trades producer and stop it with ctrl-c when you think there's enough data:

```
java -cp target/uber-kafka-streams-stockstats-1.1-SNAPSHOT.jar -DLOGLEVEL=INFO
com.shapira.examples.streams.stockstats.StockGenProducer configfile
```

D.12. Run the streams app:

```
java -cp target/uber-kafka-streams-stockstats-1.1-SNAPSHOT.jar -DLOGLEVEL=INFO
com.shapira.examples.streams.stockstats.StockStatsExample configfile
```

D.13. Check the results:

```
/usr/local/kafka/bin/kafka-console-consumer.sh --topic stockstats-output --from-beginning --bootstrap-server
<YOUR_DNS_NAME>:9092 --property print.key=true
```

D.14. If you want to reset state and re-run the application (maybe with some changes?) on existing input topic, you can stop the kafka producer and consumer and then:

Reset internal topics (used for shuffle the topic and state-stores):

```
/usr/local/kafka/bin/kafka-streams-application-reset.sh --application-id wordcount --bootstrap-servers
<YOUR_DNS_NAME>:9092 --input-topics stocks
```

(optional) Delete the output topic:

```
/usr/local/kafka/bin/kafka-topics.sh --bootstrap-server <YOUR_DNS_NAME>:9092 --delete --topic stockstats-
output
```

E. Testing the click Stream Enrichment example

E.1. Proceed accordingly with the previous examples:

```
[ec2-user@ip- ~]$ mkdir StreamEnrichment  
[ec2-user@ip- ~]$ cd StreamEnrichment/  
[ec2-user@ip- StreamEnrichment]$ git clone https://github.com/gwenshap/kafka-clickstream-enrich  
Cloning into 'kafka-clickstream-enrich'...  
remote: Enumerating objects: 65, done.  
remote: Total 65 (delta 0), reused 0 (delta 0), pack-reused 65  
Unpacking objects: 100% (65/65), done.  
[ec2-user@ip- StreamEnrichment]$ cd kafka-clickstream-enrich
```

E.2. Study the following source code that is also located in:

```
/home/ec2-user/StreamEnrichment/kafka-clickstream-enrich/src/main/java/com/shapira/examplesstreams/  
clickstreamenrich/ClickstreamEnrichment.java
```

```
1 package com.shapira.examplesstreams.clickstreamenrich;  
2 import com.shapira.examplesstreams.clickstreamenrich.model.PageView;  
3 import com.shapira.examplesstreams.clickstreamenrich.model.Search;  
4 import com.shapira.examplesstreams.clickstreamenrich.model.UserActivity;  
5 import com.shapira.examplesstreams.clickstreamenrich.model.UserProfile;  
6 import com.shapira.examplesstreams.clickstreamenrich.serde.JsonDeserializer;  
7 import com.shapira.examplesstreams.clickstreamenrich.serde.JsonSerializer;  
8 import com.shapira.examplesstreams.clickstreamenrich.serde.WrapperSerde;  
9 import org.apache.kafka.clients.consumer.ConsumerConfig;  
10 import org.apache.kafka.common.serialization.Serdes;  
11 import org.apache.kafka.streams.KafkaStreams;  
12 import org.apache.kafka.streams.StreamsConfig;  
13 import org.apache.kafka.streams.kstream.JoinWindows;  
14 import org.apache.kafka.streams.kstream.KStream;  
15 import org.apache.kafka.streams.kstream.KStreamBuilder;  
16 import org.apache.kafka.streams.kstream.KTable;  
17  
18 public class ClickstreamEnrichment {  
19  
20     public static void main(String[] args) throws Exception {  
21  
22         Properties props = new Properties();  
23         props.put(StreamsConfig.APPLICATION_ID_CONFIG, "clicks");  
24         props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, Constants.BROKER);  
25         // Since each step in the stream will involve different objects, we can't use default Serde  
26  
27         // setting offset reset to earliest so that we can re-run the demo code with the same pre-loaded data  
28         // Note: To re-run the demo, you need to use the offset reset tool:  
29         // https://cwiki.apache.org/confluence/display/KAFKA/Kafka+Streams+Application+Reset+Tool  
30         props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");  
31  
32         // work-around for an issue around timing of creating internal topics  
33         // this was resolved in 0.10.2.0 and above  
34         // don't use in large production apps - this increases network load  
35         // props.put(CommonClientConfigs.METADATA_MAX_AGE_CONFIG, 500);  
36  
37         KStreamBuilder builder = new KStreamBuilder();  
38  
39         KStream<Integer, PageView> views = builder.stream(Serdes.Integer(), new PageViewSerde(),  
40             Constants.PAGE_VIEW_TOPIC);  
41         KTable<Integer, UserProfile> profiles = builder.table(Serdes.Integer(), new ProfileSerde(),  
42             Constants.USER_PROFILE_TOPIC, "profile-store");  
43         KStream<Integer, Search> searches = builder.stream(Serdes.Integer(), new SearchSerde(),  
44             Constants.SEARCH_TOPIC);  
45  
46         KStream<Integer, UserActivity> viewsWithProfile = views.leftJoin(profiles,  
47             (page, profile) -> {  
48                 if (profile != null)
```

```

20     return new UserActivity(profile.getUserId(), profile.getUserName(), profile.getZipcode(),
21         profile.getInterests(), "", page.getPage());
22     else
23   });
24 
25     KStream<Integer, UserActivity> userActivityKStream = viewsWithProfile.leftJoin(searches,
26     (userActivity, search) -> {
27       if (search != null)
28         userActivity.updateSearch(search.getSearchTerms());
29       else
30         userActivity.updateSearch("");
31     },
32     JoinWindows.of(1000), Serdes.Integer(), new UserActivitySerde(), new SearchSerde());
33 
34     userActivityKStream.to(Serdes.Integer(), new UserActivitySerde(), Constants.USER_ACTIVITY_TOPIC);
35 
36   KafkaStreams streams = new KafkaStreams(builder, props);
37 
38   streams.cleanUp();
39 
40   streams.start();
41 
42   // usually the stream application would be running forever,
43   // in this example we just let it run for some time and stop since the input data is finite.
44   Thread.sleep(60000L);
45 
46   streams.close();
47 
48 }
49 
50 static public final class PageViewSerde extends WrapperSerde<PageView> {
51   public PageViewSerde() {
52     super(new JsonSerializer<PageView>(), new JsonDeserializer<PageView>(PageView.class));
53   }
54 }
55 
56 static public final class ProfileSerde extends WrapperSerde<UserProfile> {
57   public ProfileSerde() {
58     super(new JsonSerializer<UserProfile>(), new JsonDeserializer<UserProfile>(UserProfile.class));
59   }
60 }
61 
62 static public final class SearchSerde extends WrapperSerde<Search> {
63   public SearchSerde() {
64     super(new JsonSerializer<Search>(), new JsonDeserializer<Search>(Search.class));
65   }
66 }
67 
68 static public final class UserActivitySerde extends WrapperSerde<UserActivity> {
69   public UserActivitySerde() {
70     super(new JsonSerializer<UserActivity>(), new JsonDeserializer<UserActivity>(UserActivity.class));
71   }
72 }

```

E.3.Change pom.xml from:

```

<dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-clients</artifactId>
    <version>0.10.3.0-SNAPSHOT</version>
</dependency>
<dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-streams</artifactId>
    <version>0.10.3.0-SNAPSHOT</version>
</dependency>

```

To a most recent library:

```

<dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-clients</artifactId>

```

```
<version>3.2.0</version>
</dependency>
<dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-streams</artifactId>
    <version>3.2.0</version>
</dependency>
```

E.4. Change the location of your kafka server in the java source code:

```
/home/ec2-user/StreamEnrichment/kafka-clickstream-
enrich/src/main/java/com/shapira/examplesstreams/clickstreamenrich/Constants.java
```

In the line 8:

```
public static final String BROKER = "<YOUR_DNS_NAME>:9092";
```

E.5. Build the project with mvn package:

```
mvn package
```

E.6. Generate some clicks, searches, and profiles. Run the generator. It should take about 5 seconds to run. Don't worry about complete lack of output:

```
java -cp target/uber-kafka-clickstream-enrich-1.0-SNAPSHOT.jar
com.shapira.examples.streams.clickstreamenrich.GenerateData
```

E.7. Run the streams app:

```
java -cp target/uber-kafka-clickstream-enrich-1.0-SNAPSHOT.jar
com.shapira.examples.streams.clickstreamenrich.ClickstreamEnrichment
```

E.8. Check the results:

```
/usr/local/kafka/bin/kafka-console-consumer.sh --topic clicks.user.activity --from-beginning --bootstrap-
server <YOUR_DNS_NAME>:9092 --property print.key=true
```

E.9. If you want to reset state and re-run the application (maybe with some changes?) on existing input topic, you can stop the kafka producer and consumer and then:

Reset internal topics (used for shuffle the topic and state-stores):

```
/usr/local/kafka/bin/kafka-streams-application-reset.sh --application-id clicks --bootstrap-servers
<YOUR_DNS_NAME>:9092 --input-topics clicks.user.profile,clicks.pages.views,clicks.search
```

(optional) Delete the output topic:

```
/usr/local/kafka/bin/kafka-topics.sh --bootstrap-server <YOUR_DNS_NAME>:9092 --delete --topic
clicks.user.activity
```

F. Other tools and references

- If you prefer, install emacs in AWS EC2 using the following command:

```
sudo yum install emacs
```

- Where to download apache maven <https://maven.apache.org/download.cgi>
- How to install apache maven <https://maven.apache.org/install.html>

The goal of this document is to show a database available in a cloud environment. The project should be built from these indications.

The usage of a cloud database requires an Amazon AWS account.

Scalability:

RDS virtualization could be scaled, creating different SCHEMAS, or DATABASES instances, for different Application services.

The following contents is presented in this document.

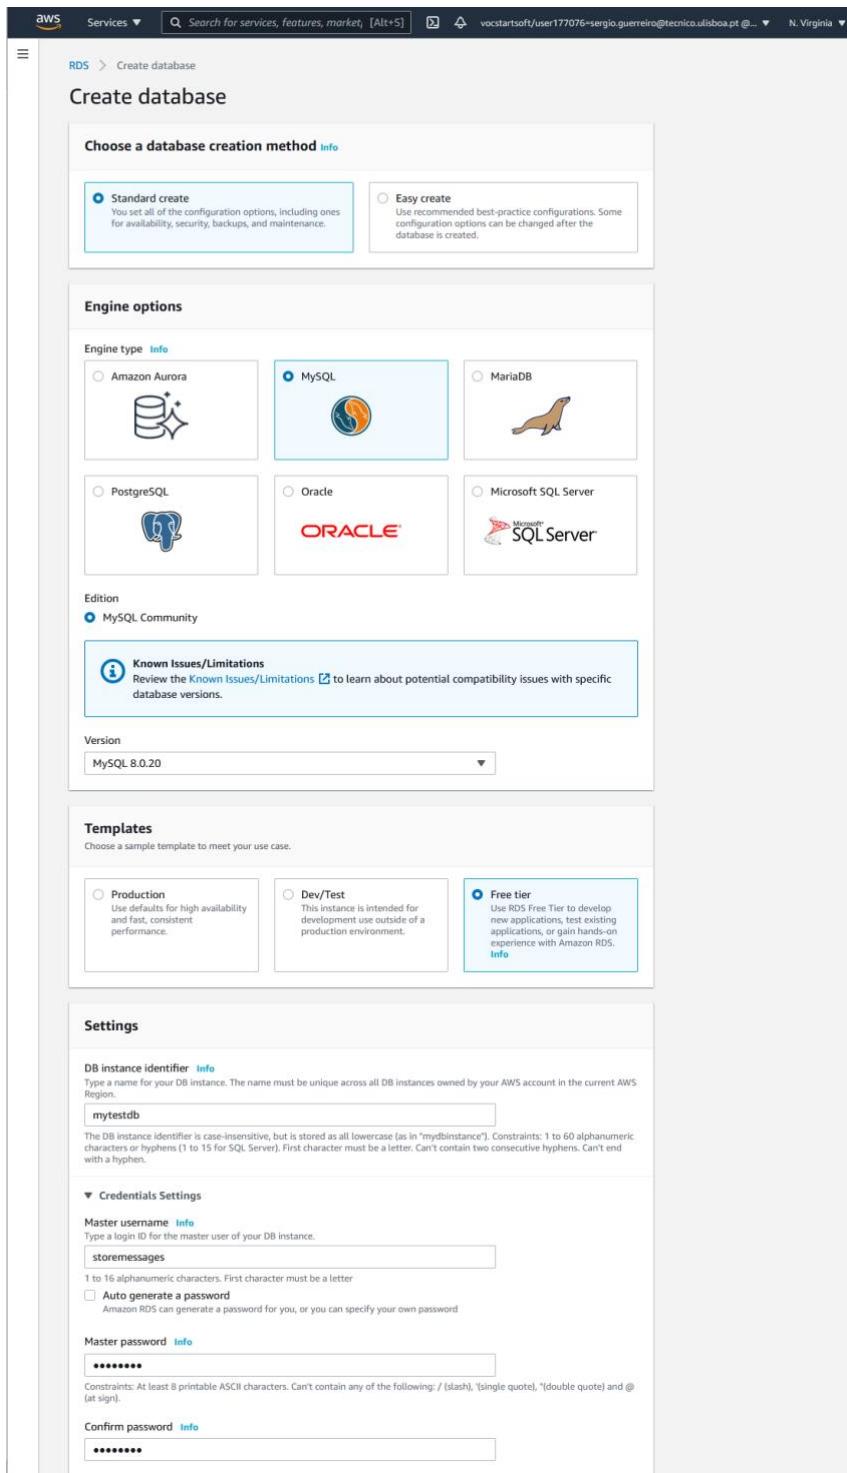
Contents

- A. Create an AWS RDS Database, accessing it, and inserting data: example of.....2

A. Create an AWS RDS Database, accessing it, and inserting data: example of

The goal of this section is to exemplify how to create your own cloud database in AWS, to access its contents through a database workbench, and then creating a JAVA microservice to consume Kafka messages that are then inserted in the cloud database.

1. Go to your AWS account console and click on Database RDS -> My SQL and select the following configurations in step 2:



DB instance size

DB instance class [Info](#)
 Choose a DB instance class that meets your processing power and memory requirements. The DB instance class options below are limited to those supported by the engine you selected above.

- Standard classes (includes m classes)
- Memory Optimized classes (includes r and x classes)
- Burstable classes (includes t classes)

db.t2.micro
 1 vCPUs 1 GiB RAM Not EBS Optimized

New instance classes are available for specific engine versions. [Info](#)

Include previous generation classes

Storage

Storage type [Info](#)
 General Purpose (SSD)

Allocated storage
 20 GiB
 (Minimum: 20 GiB, Maximum: 16 384 GiB) Higher allocated storage may improve IOPS performance.

Storage autoscaling [Info](#)
 Provides dynamic scaling support for your database's storage based on your application's needs.

Enable storage autoscaling
 Enabling this feature will allow the storage to increase once the specified threshold is exceeded.

Maximum storage threshold [Info](#)
 Charges will apply when your database autoscales to the specified threshold
 1000 GiB
 Minimum: 21 GiB, Maximum: 16 384 GiB

Availability & durability

Multi-AZ deployment [Info](#)

- Do not create a standby instance
- Create a standby instance (recommended for production usage)
 Creates a standby in a different Availability Zone (AZ) to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups.

Connectivity

Virtual private cloud (VPC) [Info](#)
 VPC that defines the virtual networking environment for this DB instance.
Default VPC (vpc-3413dd49)

Only VPCs with a corresponding DB subnet group are listed.

Subnet group [Info](#)
 DB subnet group that defines which subnets and IP ranges the DB instance can use in the VPC you selected.
default

Public access [Info](#)

- Yes
 Amazon EC2 instances and devices outside the VPC can connect to your database. Choose one or more VPC security groups that specify which EC2 instances and devices inside the VPC can connect to the database.
- No
 RDS will not assign a public IP address to the database. Only Amazon EC2 instances and devices inside the VPC can connect to your database.

VPC security group
 Choose a VPC security group to allow access to your database. Ensure that the security group rules allow the appropriate incoming traffic.

Choose existing
 Choose existing VPC security groups

Create new
 Create new VPC security group

Existing VPC security groups
 Choose VPC security groups
 launch-wizard-5 X default X

Availability Zone [Info](#)
 No preference

Additional configuration

Database authentication

Database authentication options [Info](#)

- Password authentication
 Authenticates using database passwords.
- Password and IAM database authentication
 Authenticates using the database password and user credentials through AWS IAM users and roles.
- Password and Kerberos authentication (not available for this version)
 Choose a directory in which you want to allow authorized users to authenticate with this DB instance using Kerberos Authentication.

▼ Additional configuration
Database options, backup enabled, backtrack disabled, Enhanced Monitoring disabled, maintenance, CloudWatch Logs, delete protection disabled

Database options

Initial database name [Info](#)
 If you do not specify a database name, Amazon RDS does not create a database.

DB parameter group [Info](#)

Option group [Info](#)

Backup

Creates a point-in-time snapshot of your database

Enable automatic backups
Enabling backups will automatically create backups of your database during a certain time window.

⚠ Please note that automated backups are currently supported for InnoDB storage engine only. If you are using MyISAM, refer to details [here](#).

Backup retention period [Info](#)
Choose the number of days that RDS should retain automatic backups for this instance.

Backup window [Info](#)
Select the period for which you want automated backups of the database to be created by Amazon RDS.
 Select window
 No preference

Copy tags to snapshots

Monitoring

Enable Enhanced monitoring
Enabling Enhanced monitoring metrics are useful when you want to see how different processes or threads use the CPU

Log exports

Select the log types to publish to Amazon CloudWatch Logs

Error log
 General log
 Slow query log

IAM role

The following service-linked role is used for publishing logs to CloudWatch Logs.

ⓘ Ensure that general, slow query, and audit logs are turned on. Error logs are enabled by default. [Learn more](#)

Maintenance

Auto minor version upgrade [Info](#)

Enable auto minor version upgrade
Enabling auto minor version upgrade will automatically upgrade to new minor versions as they are released. The automatic upgrades occur during the maintenance window for the database.

Maintenance window [Info](#)
Select the period you want pending modifications or maintenance applied to the database by Amazon RDS.
 Select window
 No preference

Deletion protection

Enable deletion protection
Protects the database from being deleted accidentally. While this option is enabled, you can't delete the database.

Hint: Later, check if your VPC configuration allows connections from anywhere in the internet. Choose the VPC Security groups and edit the configuration similarly with:

The screenshot shows the AWS EC2 Services dashboard with the 'Resource Groups' tab selected. On the left, a sidebar lists various EC2 features like Instances, AMIs, and Security Groups. The main pane displays a table of security groups. One row is selected, showing details for a group named 'default' associated with VPC ID 'vpc-c798bcb'. Below the table, a detailed view for 'Security Group: sg-671f0033' is shown, with the 'Inbound' tab selected. It lists two rules: one for 'All traffic' on 'All' protocol and port ranges from '0.0.0.0/0' and another identical rule.

2. You can now check if the database is started.
3. Use your MySQL database client (for example MySQL workbench), connect to your RDS AWS database and execute the following scripts:

```
DROP DATABASE IF EXISTS collectedMSGs;
CREATE DATABASE IF NOT EXISTS collectedMSGs;
```

```
USE collectedMSGs;
DROP TABLE IF EXISTS Message;

CREATE TABLE Message( offset INTEGER PRIMARY KEY,
    groupId VARCHAR(100) NOT NULL,
    contentKey VARCHAR(100),
    contentValue VARCHAR(1000));
```

The goal of this document is to present the details related with the QUARKUS microservices framework.

The examples address the following aspects: environment set-up, creation of projects and adding elements to it, accessing databases, and controlling the database transaction, using messages in JSON format, how to use KAFKA, how to provide an OpenAPI interface, and finally, deploying in docker and deploying in AWS.

The following contents is presented in this document.

Contents

A.	Setup steps for initial setup of QUARKUS environment.....	2
B.	Create a Simple QUARKUS project with starter code.....	3
C.	Add a new JAVA class to and already existing QUARKUS project.....	4
D.	QUARKUS exposed by OpenAPI interface	5
E.	QUARKUS accessing database with all CRUD Operations	6
F.	Deploying a QUARKS project with docker locally	9
G.	QUARKUS with Kafka messaging embedded (in development environment)	10
H.	QUARKUS with Kafka messaging not embedded	16
I.	Consuming a Kafka Topic that changes dynamically.....	17
J.	Creating and Removing Kafka Topics dynamically.....	20
K.	Deploying QUARKUS in AWS EC2 environment with github and terraform!.....	23
	References	26

A. Setup steps for initial setup of QUARKUS environment

A.1. Clone this **repository** <https://github.com/quarkusio/quarkus-quickstarts.git> using GitHub desktop or web.

A.2. Install **Maven 3.9.9** from <https://archive.apache.org/dist/maven/maven-3/3.9.9/binaries/>
(for detail on installing in different operating systems:

<https://mkyong.com/maven/install-maven-on-mac-osx/>
<https://mkyong.com/maven/how-to-install-maven-in-windows/>)

A.3. Install **Docker** from <https://www.docker.com/get-started> and launch it.

A.4. Install **Mandrel** (for linux only) or **GraalVM** (other platforms) with **JDK 17** from
<https://github.com/graalvm/mandrel> or <https://www.oracle.com/java/technologies/downloads/#graalvmjava17>
respectively

(for details on installing in different operating systems

<https://github.com/graalvm/homebrew-tap>
<https://www.graalvm.org/latest/getting-started/linux/>
<https://www.graalvm.org/latest/getting-started/windows/>)

A.5. Check java version with "javac -version" in command line, and change JDK if needed following these instructions:
https://mkyong.com/java/how-to-set-java_home-environment-variable-on-mac-os-x/
https://mkyong.com/java/how-to-set-java_home-on-windows-10/

A.6. In the repository base directory of step A.1. execute the following to assure that your environment is ok:

```
mvn clean install
```

A.7. (Optional) Install **mysql server** locally (optionally, for testing purposes of QUARKUS accessing data reactively) or, later, use AWS RDS.

A.8. (Optional) Install VS Code (<https://code.visualstudio.com/>) or other preferred IDE.

B. Create a Simple QUARKUS project with starter code

B.1. Navigate to the QUARKUS project creation website at <https://code.quarkus.io> and use the following configuration:

The screenshot shows the Quarkus project creation interface. At the top, it says "QUARKUS 3.17 io.quarkus.platform". Below that, the title "CONFIGURE YOUR APPLICATION" is displayed. There are four main configuration sections: "Group" set to "org.ie", "Version" set to "1.0.0-SNAPSHOT", "Artifact" set to "tryout1", and "Build Tool" set to "Maven". Underneath these, there are dropdowns for "Java Version" (set to 17) and "Starter Code" (set to Yes). On the right side, there is a "CLOSE" button, a notification icon with "0" notifications, and a prominent blue "Generate your application" button. At the bottom, there is a search bar with the placeholder "Search across all available extensions: rest, hibernate, web, data...".

B.2. Generate, download the zip file and extract it to your local folder project

B.3. Import your newly created project to VS Code

B.4. Execute the command

```
./mvnw quarkus:dev
```

on command line or calling maven inside VS Code

B.5. Go to URL <http://localhost:8080/hello> and check the message that is produced.

B.6. Open the source code on your IDE, e.g., VS Code and navigate to `GreetingResource.java` file. Study the source code and change the endpoint at line:

```
@Path("/hello")
```

To

```
@Path("/newapi")
```

B.7. Execute the command

```
./mvnw quarkus:dev
```

on command line or calling maven inside VS Code

B.8. Go to URL <http://localhost:8080/newapi> and <http://localhost:8080/welcome> and check the result that is produced.

C. Add a new JAVA class to and already existing QUARKUS project

C.1.Create a new class in your project with the name: **SecondClass.java**, in the same location as previous GreetingResource.java.

C.2.Update the new file SecondClass.java accordingly with this source code:

```
package org.ie;

import jakarta.ws.rs.GET;
import jakarta.ws.rs.Path;
import jakarta.ws.rs.Produces;
import jakarta.ws.rs.core.MediaType;

@Path("/api2")
public class SecondClass {

    @GET
    @Produces(MediaType.TEXT_PLAIN)
    public String hello() {
        return "Hello to the second API";
    }
}
```

Hint: you may also add a test class to automate the code verification

C.3.Execute the command

```
./mvnw quarkus:dev
```

on command line or calling maven inside VS Code

C.4.Go to URL <http://localhost:8080/newapi> and <http://localhost:8080/api2> and check the result that is produced.

C.5.If needed, you can change the port, adding the following configuration in **application.properties** of your project (resources directory of your Quarkus project) :

```
quarkus.http.port=9000
```

D. QUARKUS exposed by OpenAPI interface

The aim of this example is to provide QUARKUS APIs with OpenAPI (<https://quarkus.io/guides/openapi-swaggerui>).

- D.1. Execute the following command to add the openapi extension:

```
./mvnw quarkus:add-extension -Dextensions='quarkus-smallrye-openapi'
```

- D.2. Change the end point for swagger-ui, adding the following configuration in application.properties:

```
quarkus.swagger-ui.path=swagger-ui  
quarkus.swagger-ui.always-include=true
```

- D.3. Execute the command on command line or calling maven inside VS Code:

```
./mvnw quarkus:dev
```

- D.4. Navigate to <http://localhost:9000/q/swagger-ui/>, a similar result will be presented to you:

The screenshot shows the OpenAPI UI interface for the tryout1 API. At the top, it displays the title "tryout1 API" with version "1.0.0-SNAPSHOT" and "OAS3". Below the title, there's a navigation bar with tabs for "Swagger UI" and "Explore". The main content area is titled "Second Class". It shows a single endpoint: "GET /api2". Under "Parameters", it says "No parameters". Under "Responses", it lists a 200 OK response with "text/plain" as the media type and "string" as the example value. Below this, there's a section for "Greeting Resource" with a "GET /newapi" endpoint.

E. QUARKUS accessing database with all CRUD Operations

E.1. A ready-to-use MySQL server to try out this example.

```
docker run -it --name mysqlInstance -e MYSQL_ROOT_PASSWORD=password -p 3306:3306 -d mysql:latest
```

E.2. Create a MySQL table to try out the examples with the following command in MYQSL Workbench;

```
CREATE DATABASE quarkus_test;
```

E.3. Navigate to the QUARKUS project creation website at <https://code.quarkus.io> and use the following configuration:

The screenshot shows the Quarkus Project Creation interface. At the top, it displays "QUARKUS 3.17 | io.quarkus.platform". Below that, the title "CONFIGURE YOUR APPLICATION" is visible. On the left, there are input fields for "Group" (org.acme), "Artifact" (reactive-mysl-client-quickstart), "Build Tool" (Maven), "Version" (1.0.0-SNAPSHOT), "Java Version" (17), and "Starter Code" (No). To the right of these fields is a "CLOSE" button. Below the input fields is a search bar with the query "resteasy jackson". Underneath the search bar, it says "Extensions found by origin: 3 in platform" and "1 in other". A checkbox for "REST Jackson [quarkus-rest-jackson]" is checked, with a note below stating: "Jackson serialization support for Quarkus REST. This extension is not compatible with the quarkus-resteasy extension, or any of the RESTEasy Jackson dependencies it provides. If you want to use both, you will need to exclude them from your project dependencies." On the right side of the interface, a sidebar titled "Generate your application (⟳ + ↻)" is open, showing "Selected Extensions" which include "SmallRye OpenAPI", "Reactive MySQL client", and "REST Jackson". It also shows "Transitive extensions (6)".

E.4. Generate, download the zip file and extract it to your local folder project

E.5. Import your newly created project to VS Code

E.6. Add the following configuration in src/main/resources/application.properties of your project

```
quarkus.datasource.db-kind=mysql
quarkus.datasource.username=root
quarkus.datasource.password=<your MySQL root password>
quarkus.datasource.reactive.url=mysql://localhost:3306/quarkus_test

quarkus.swagger-ui.path=swagger-ui
quarkus.swagger-ui.always-include=true
```

E.7. Add the file Fruit.java to src/main/java/Fruit.java

```
import io.smallrye.mutiny.Multi;
import io.smallrye.mutiny.Uni;
import io.vertx.mutiny.mysqlclient.MySQLPool;
import io.vertx.mutiny.sqlclient.Row;
import io.vertx.mutiny.sqlclient.RowSet;
import io.vertx.mutiny.sqlclient.Tuple;

public class Fruit {

    public Long id;
    public String name;

    public Fruit() {
    }

    public Fruit(String name) {
        this.name = name;
    }

    public Fruit(Long id, String name) {
        this.id = id;
        this.name = name;
    }
}
```

```

private static Fruit from(Row row) {
    return new Fruit(row.getLong("id"), row.getString("name"));
}

public static Multi<Fruit> findAll(MySQLPool client) {
    return client.query("SELECT id, name FROM fruits ORDER BY name ASC").execute()
        .onItem().transformToMulti(set -> Multi.createFrom().iterable(set))
        .onItem().transform(Fruit::from);
}

public static Uni<Fruit> findById(MySQLPool client, Long id) {
    return client.preparedQuery("SELECT id, name FROM fruits WHERE id = ?").execute(Tuple.of(id))
        .onItem().transform(RowSet::iterator)
        .onItem().transform(iterator -> iterator.hasNext() ? from(iterator.next()) : null);
}

public Uni<Boolean> save(MySQLPool client) {
    return client.preparedQuery("INSERT INTO fruits(name) VALUES (?)").execute(Tuple.of(name))
        .onItem().transform(pgRowSet -> pgRowSet.rowCount() == 1);
}

public static Uni<Boolean> delete(MySQLPool client, Long id) {
    return client.preparedQuery("DELETE FROM fruits WHERE id = ?").execute(Tuple.of(id))
        .onItem().transform(pgRowSet -> pgRowSet.rowCount() == 1);
}

public static Uni<Boolean> update(MySQLPool client, Long id, String name) {
    return client.preparedQuery("UPDATE fruits SET name = ? WHERE id = ?").execute(Tuple.of(name,id))
        .onItem().transform(pgRowSet -> pgRowSet.rowCount() == 1);
}
}

```

E.8. Add the file FruitResource.java to src/main/java/FruitResource.java

```

import java.net.URI;
import jakarta.enterprise.event.Observes;
import jakarta.inject.Inject;
import jakarta.ws.rs.*;
import org.eclipse.microprofile.config.inject.ConfigProperty;

import io.quarkus.runtime.StartupEvent;
import io.smallrye.mutiny.Multi;
import io.smallrye.mutiny.Uni;
import jakarta.ws.rs.core.Response;
import jakarta.ws.rs.core.Response.ResponseBuilder;
import jakarta.ws.rs.core.MediaType;

@Path("fruits")
public class FruitResource {

    @Inject
    io.vertx.mysqlclient.MySQLPool client;

    @Inject
    @ConfigProperty(name = "myapp.schema.create", defaultValue = "true")
    boolean schemaCreate;

    void config(@Observes StartupEvent ev) {
        if (schemaCreate) {
            initdb();
        }
    }

    private void initdb() {
        client.query("DROP TABLE IF EXISTS fruits").execute()
            .flatMap(r -> client.query("CREATE TABLE fruits (id SERIAL PRIMARY KEY, name TEXT NOT NULL)").execute())
            .flatMap(r -> client.query("INSERT INTO fruits (name) VALUES ('Orange')").execute())
            .flatMap(r -> client.query("INSERT INTO fruits (name) VALUES ('Pear')").execute())
            .flatMap(r -> client.query("INSERT INTO fruits (name) VALUES ('Apple')").execute())
            .flatMap(r -> client.query("INSERT INTO fruits (name) VALUES ('other')").execute())
            .await().indefinitely();
    }

    @GET
    public Multi<Fruit> get() {

```

```

        return Fruit.findAll(client);
    }

    @GET
    @Path("/{id}")
    public Uni<Response> getSingle(Long id) {
        return Fruit.findById(client, id)
            .onItem().transform(fruit -> fruit != null ? Response.ok(fruit) :
Response.status(Response.Status.NOT_FOUND))
            .onItem().transform(ResponseBuilder::build);
    }

    @POST
    public Uni<Response> create(Fruit fruit) {
        return fruit.save(client)
            .onItem().transform(id -> URI.create("/fruits/" + id))
            .onItem().transform(uri -> Response.created(uri).build());
    }

    @DELETE
    @Path("/{id}")
    public Uni<Response> delete(Long id) {
        return Fruit.delete(client, id)
            .onItem().transform(deleted -> deleted ? Response.Status.NO_CONTENT :
Response.Status.NOT_FOUND)
            .onItem().transform(status -> Response.status(status).build());
    }

    @PUT
    @Path("/{id}/{name}")
    public Uni<Response> update(Long id, String name) {
        return Fruit.update(client, id, name)
            .onItem().transform(updated -> updated ? Response.Status.NO_CONTENT :
Response.Status.NOT_FOUND)
            .onItem().transform(status -> Response.status(status).build());
    }
}

```

E.9. Compile and execute the project:

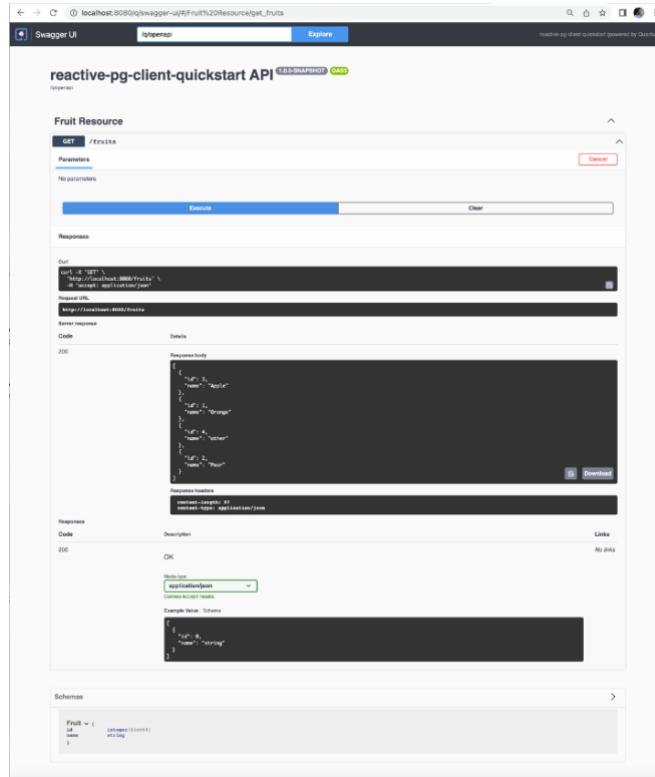
```

cd reactive-mysql-CRUD-client-quickstart
./mvnw quarkus:dev

```

E.10. Navigate to http://localhost:8080/q/swagger-ui/ and execute the all the CRUD APIs.

Try to create, update, and delete an item.



F. Deploying a QUARKS project with docker locally

F.1. Execute the following command in your project directory:

```
./mvnw quarkus:add-extension -Dextensions="quarkus-container-image-docker"
```

F.2. Add the following configuration to application.properties in order to include all the dependencies to execute in a different environment from dev:

```
quarkus.container-image.build=true
```

F.3. Build docker image

Hint 1: if you are not logged in on docker use the "docker login" command

Hint 2: have in mind that all the network requirements for the docker image need to be met beforehand.

To facilitate start, e.g., with a simple tryout project as in section B of this document.

```
./mvnw clean package
```

F.4. Check that image has been created successfully.

```
% docker image ls -a
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
sergioguerreiro/tryout1  1.0.0-SNAPSHOT  eaf911ae3716  15 hours ago  492MB
```

F.5. Create container. Have in mind that the port available and the name of the project depends on your previous settings.

```
docker run -d --name tryout1 -p 9000:9000 sergioguerreiro/tryout1:1.0.0-SNAPSHOT
```

Test the container opening the following URL in your browser:

```
# open browser with url: http://localhost:9000/q/swagger-ui/
```

Hint 1: Later if you would like to list all the available container use the following command:

```
> docker container ls --all
```

```
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              NAMES
2d4cc67021e3      sergioguerreiro/tryout1:1.0.0-SNAPSHOT   "/usr/local/s2i/run"
8080/tcp, 8443/tcp, 8778/tcp, 0.0.0.0:9000->9000/tcp           2 minutes ago    Up 2 minutes     tryout1
```

Hint 2: Later if you would like to stop the container use the following command:

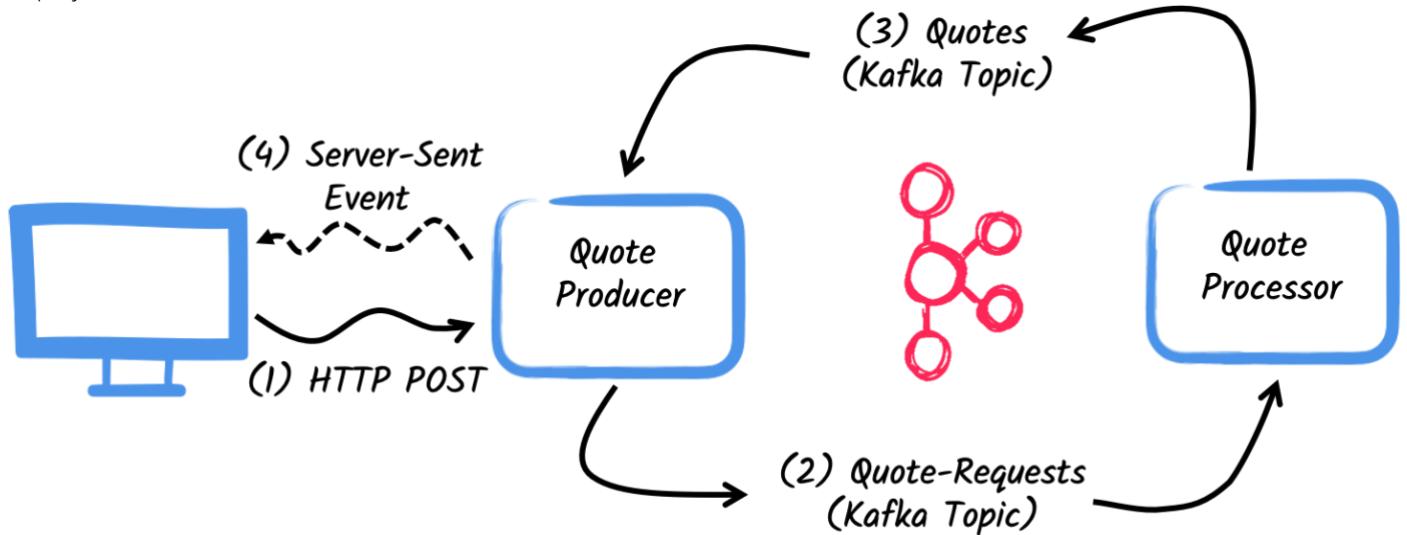
```
docker container stop <containerID_hash>
```

Hint 3: Later if you would like to restart the container use the following command:

```
docker container start <containerID_hash>
```

G. QUARKUS with Kafka messaging embedded (in development environment)

The following example aims at implementing the following messaging architecture. Kafka is embedded in QUARKUS deployment.



G.1. Create an empty project (KAFKA message producer) at <https://code.quarkus.io/> with the following dependencies.

Group: org.acme
Artifact: kafka-dynamic-producer
Build Tool: Maven
Version: 1.0.0-SNAPSHOT
Java Version: 17
Starter Code: No

Selected Extensions:

- SmallRye OpenAPI
- Messaging - Kafka Connector STARTER-CODE
- REST Jackson

G.2.Create an empty project (KAFKA message consumer) at <https://code.quarkus.io/> with the following dependencies.

Group: org.acme
Artifact: kafka-dynamic-consumer
Build Tool: Maven
Version: 1.0.0-SNAPSHOT
Java Version: 17
Starter Code: No

Selected Extensions:

- SmallRye OpenAPI
- Messaging - Kafka Connector STARTER-CODE

G.3.In project kafka-dynamic-producer create the following JAVA classes at
src/main/java/org/acme/kafka/model/Quote.java

```

package org.acme.kafka.model;

public class Quote {
    public String id;
    public int price;
}
  
```

```

/**
 * Default constructor required for Jackson serializer
 */
public Quote() { }

public Quote(String id, int price) {
    this.id = id;
    this.price = price;
}

@Override
public String toString() {
    return "Quote{" +
        "id='" + id + '\'' +
        ", price=" + price +
        '}';
}
}

```

Add to kafka-dynamic-producer the `src/main/java/org/acme/kafka/producer/QuotesResource.java`

```

package org.acme.kafka.producer;

import java.util.UUID;

import jakarta.ws.rs.GET;
import jakarta.ws.rs.POST;
import jakarta.ws.rs.Path;
import jakarta.ws.rs.Produces;
import jakarta.ws.rs.core.MediaType;

import org.acme.kafka.model.Quote;
import org.eclipse.microprofile.reactive.messaging.Channel;
import org.eclipse.microprofile.reactive.messaging.Emitter;

import io.smallrye.mutiny.Multi;

@Path("/quotes")
public class QuotesResource {

    @Channel("quote-requests")
    Emitter<String> quoteRequestEmitter;

    @Channel("quotes")
    Multi<Quote> quotes;
    /**
     * Endpoint to generate a new quote request id and send it to "quote-requests" Kafka topic using the
     * emitter.
     */
    @POST
    @Path("/request")
    @Produces(MediaType.TEXT_PLAIN)
    public String createRequest() {
        UUID uuid = UUID.randomUUID();
        quoteRequestEmitter.send(uuid.toString());
        return uuid.toString();
    }

    /**
     * Endpoint retrieving the "quotes" Kafka topic and sending the items to a server sent event.
     */
    @GET
    @Produces(MediaType.SERVER_SENT_EVENTS)
    public Multi<Quote> stream() {
        return quotes;
    }
}

```

Add to kafka-dynamic-producer the following `index.html` file to `/src/main/resources/META-INF/resources/`

```

<!DOCTYPE html>
<html lang="en">
<head>
    <meta charset="UTF-8">
    <title>kafka-quickstart - 1.0.0-SNAPSHOT</title>
    <style>
        h1, h2, h3, h4, h5, h6 {
            margin-bottom: 0.5rem;
            font-weight: 400;
        }
    </style>

```

```

        line-height: 1.5;
    }

h1 {
    font-size: 2.5rem;
}

h2 {
    font-size: 2rem;
}

h3 {
    font-size: 1.75rem;
}

h4 {
    font-size: 1.5rem;
}

h5 {
    font-size: 1.25rem;
}

h6 {
    font-size: 1rem;
}

.lead {
    font-weight: 300;
    font-size: 2rem;
}

.banner {
    font-size: 2.7rem;
    margin: 0;
    padding: 2rem 1rem;
    background-color: #00A1E2;
    color: white;
}

body {
    margin: 0;
    font-family: -apple-system, system-ui, "Segoe UI", Roboto, "Helvetica Neue", Arial, sans-serif,
    "Apple Color Emoji", "Segoe UI Emoji", "Segoe UI Symbol", "Noto Color Emoji";
}

code {
    font-family: SFMono-Regular, Menlo, Monaco, Consolas, "Liberation Mono", "Courier New",
    monospace;
    font-size: 87.5%;
    color: #e83e8c;
    word-break: break-word;
}

.left-column {
    padding: .75rem;
    max-width: 75%;
    min-width: 55%;
}

.right-column {
    padding: .75rem;
    max-width: 25%;
}

.container {
    display: flex;
    width: 100%;
}

li {
    margin: 0.75rem;
}

.right-section {
    margin-left: 1rem;
    padding-left: 0.5rem;
}

.right-section h3 {
    padding-top: 0;
    font-weight: 200;
}

```

```

        }

    .right-section ul {
        border-left: 0.3rem solid #00A1E2;
        list-style-type: none;
        padding-left: 0;
    }


```

</style>

</head>

<body>

<div class="banner lead">
Your new Cloud-Native application is ready!
</div>

<div class="container">
 <div class="left-column">
 <p class="lead"> Congratulations, you have created a new Quarkus application.</p>

 <h2>Why do you see this?</h2>

 <p>This page is served by Quarkus. The source is in
 <code>src/main/resources/META-INF/resources/index.html</code>.</p>

 <h2>What can I do from here?</h2>

 <p>If not already done, run the application in dev mode using: <code>mvn quarkus:dev</code>.
</p>

 Add REST resources, Servlets, functions and other services in <code>src/main/java</code>.
 Your static assets are located in <code>src/main/resources/META-INF/resources</code>.
 Configure your application in <code>src/main/resources/application.properties</code>.

 <h2>How do I get rid of this page?</h2>
 <p>Just delete the <code>src/main/resources/META-INF/resources/index.html</code> file.</p>
</div>

<div class="right-column">
 <div class="right-section">
 <h3>Application</h3>

 GroupId: org.acme.quarkus.sample
 ArtifactId: kafka-quickstart
 Version: 1.0.0-SNAPSHOT
 Quarkus Version: 999-SNAPSHOT

 </div>
 <div class="right-section">
 <h3>Next steps</h3>

 Setup your IDE
 Getting started
 Quarkus Web Site

 </div>
</div>
</div>

</body>

</html>

Add to kafka-dynamic-producer the following quotes.html file to /src/main/resources/META-INF/resources/

```

<!DOCTYPE html>
<html lang="en">
<head>
    <meta charset="UTF-8">
    <title>Prices</title>

    <link rel="stylesheet" type="text/css"
        href="https://cdnjs.cloudflare.com/ajax/libs/patternfly/3.24.0/css/patternfly.min.css">
    <link rel="stylesheet" type="text/css"
        href="https://cdnjs.cloudflare.com/ajax/libs/patternfly/3.24.0/css/patternfly-additions.min.css">

```

</head>

<body>

<div class="container">
 <div class="card">

```

<div class="card-body">
    <h2 class="card-title">Quotes</h2>
    <button class="btn btn-info" id="request-quote">Request Quote</button>
    <div class="quotes"></div>
</div>
</div>
</body>
<script src="https://code.jquery.com/jquery-3.6.0.min.js"></script>
<script>
    $("#request-quote").click((event) => {
        fetch("/quotes/request", {method: "POST"})
            .then(res => res.text())
            .then(qid => {
                var row = $(`<h4 class='col-md-12' id='${qid}'>Quote # <i>${qid}</i> | 
<strong>Pending</strong></h4>`);
                $(".quotes").prepend(row);
            });
    });

    var source = new EventSource("/quotes");
    source.onmessage = (event) => {
        var json = JSON.parse(event.data);
        `#${json.id}`).html((index, html) => {
            return html.replace("Pending", `\$xA0${json.price}`);
        });
    };
</script>
</html>

```

G.4. Add the following configuration in src/main/resources/application.properties of your consumer project:

```

%dev.quarkus.http.port=8081

# Go bad to the first records, if it's out first access
mp.messaging.incoming.requests.auto.offset.reset=earliest

# Set the Kafka topic, as it's not the channel name
mp.messaging.incoming.requests.topic=quote-requests

```

Hint: For now, the producer project do not need any configuration in the application.properties.

G.5. In project kafka-dynamic-consumer create the following JAVA classes at
src/main/java/org/acme/kafka/model/Quote.java

```

package org.acme.kafka.model;

public class Quote {

    public String id;
    public int price;

    /**
     * Default constructor required for Jackson serializer
     */
    public Quote() { }

    public Quote(String id, int price) {
        this.id = id;
        this.price = price;
    }

    @Override
    public String toString() {
        return "Quote{" +
            "id='" + id + '\'' +
            ", price=" + price +
            '}';
    }
}

```

Add to kafka-dynamic-consumer the src/main/java/org/acme/kafka/processor/QuotesProcessor.java

```

package org.acme.kafka.processor;

```

```

import java.util.Random;
import jakarta.enterprise.context.ApplicationScoped;
import org.acme.kafka.model.Quote;
import org.eclipse.microprofile.reactive.messaging.Incoming;
import org.eclipse.microprofile.reactive.messaging.Outgoing;
import io.smallrye.reactive.messaging.annotations.Blocking;
/**
 * A bean consuming data from the "quote-requests" Kafka topic (mapped to "requests" channel) and giving out
 * a random quote.
 * The result is pushed to the "quotes" Kafka topic.
 */
@ApplicationScoped
public class QuotesProcessor {

    private Random random = new Random();

    @Incoming("requests")
    @Outgoing("quotes")
    @Blocking
    public Quote process(String quoteRequest) throws InterruptedException {
        // simulate some hard working task
        Thread.sleep(200);
        return new Quote(quoteRequest, random.nextInt(100));
    }
}

```

G.6. Execute in a separate command line:

```

cd kafka-quickstart-consumer
./mvnw quarkus:dev

```

G.7. Execute in a separate command line:

```

cd kafka-quickstart-producer
./mvnw quarkus:dev

```

G.8. Navigate to <http://localhost:8080/quotes.html>

Then, click “Request Quote” and check in the consumer window that the message has been printed.

The screenshot shows a web browser window with the following details:

- Address Bar:** localhost:8080/quotes.html
- Page Title:** Quotes
- Buttons:** Back, Forward, Stop, Refresh, and a search icon.
- Content:**
 - Request Quote** button (highlighted in blue)
 - Four quote entries:
 - Quote # 12e911a6-07d2-4b62-9bbe-17f2e31e5cc0 | \$ 92
 - Quote # b94bb1fa-ed39-4178-a01f-5a71c78b6bfb | \$ 66
 - Quote # 66bb9f9b-a7d6-4700-8dd4-49102f15d04d | \$ 26
 - Quote # 80192582-fd64-4fc5-a790-b4036a55effe | \$ 71

H. QUARKUS with Kafka messaging not embedded

Using the previous project of “G. QUARKUS with Kafka messaging embedded (in development environment)” execute the following configuration changes:

- H.1. Add the following configuration in src/main/resources/application.properties of your consumer project. You are adding a remote integration with a previously installed Kafka server (or broker) in AWS EC2 at port 9092.

```
kafka.bootstrap.servers=<YOUR EC2 NAME>:9092  
mp.messaging.incoming.requests.topic=quote-requests  
mp.messaging.incoming.requests.auto.offset.reset=earliest
```

- H.2. Then, add the following configuration in src/main/resources/application.properties of your producer project. You are adding a remote integration with a previously installed Kafka server (or broker) in AWS EC2 at port 9092.

```
kafka.bootstrap.servers=<YOUR EC2 NAME>:9092
```

- H.3. Execute in a separate command line

```
cd kafka-quickstart-consumer  
./mvnw quarkus:dev
```

- H.4. Execute in a separate command line

```
cd kafka-quickstart-producer  
./mvnw quarkus:dev
```

- H.5. Navigate to <http://localhost:8080/quotes.html>

I. Consuming a Kafka Topic that changes dynamically

In Quarkus Kafka, the @Incoming annotation defines a method that acts as a Kafka consumer. You are not allowed to change it dynamically at runtime, as the topic is determined at compile time. A simple workaround solution is to use the old JAVA way of dealing with consuming Kafka messages. For exemplification consider a REST API exposing a POST where the topic is received, then a new JAVA Thread is started containing the Kafka consuming source code for that topic.

- I.1. Create an empty project (KAFKA message consumer) at <https://code.quarkus.io/> with the following dependencies.

The screenshot shows the Quarkus Code Editor interface. At the top, there's a header with the Quarkus logo, version 3.17, and the URL io.quarkus.platform. To the right are links for 'Back to quarkus.io' and 'Available with Enterprise Support'. Below the header, the main title is 'CONFIGURE YOUR APPLICATION'. On the left, there are input fields for 'Group' (org.acme), 'Artifact' (kafka-dynamic), 'Build Tool' (Maven), 'Version' (1.0.0-SNAPSHOT), 'Java Version' (17), and 'Starter Code' (No). There's also a 'CLOSE' button. In the center, there's a search bar with the text 'resteasy' and a dropdown menu below it. On the right, a sidebar titled 'Selected Extensions' lists three items: 'SmallRye OpenAPI', 'Messaging - Kafka Connector', and 'REST Jackson', each with a small icon and a 'STARTER-CODE' badge. A blue button at the bottom right says 'Generate your application (⟳ + ↲)'.

And then, add the following files to your project.

- I.2. The KafkaProvisioningResource.java to src/main/java/org/acme/KafkaProvisioningResource.java

```
package org.acme;

import jakarta.ws.rs.POST;
import jakarta.ws.rs.Path;
import jakarta.ws.rs.Produces;
import jakarta.ws.rs.core.MediaType;
import org.acme.model.Topic;
import org.eclipse.microprofile.config.inject.ConfigProperty;

@Path("/createTopic")
public class KafkaProvisioningResource {

    @ConfigProperty(name = "kafka.bootstrap.servers")
    String kafka_servers;

    @POST
    @Produces(MediaType.TEXT_PLAIN)
    public String ProvisioningConsumer(Topic topic) {
        Thread worker = new DynamicTopicConsumer(topic.TopicName , kafka_servers );
        worker.start();
        return "New worker started";
    }
}
```

- I.3. The DynamicTopicConsumer.java to src/main/java/org/acme/DynamicTopicConsumer.java

```
package org.acme;

import org.apache.kafka.clients.consumer.Consumer;
import org.apache.kafka.clients.consumer.ConsumerRecords;
import org.apache.kafka.clients.consumer.ConsumerRecord;
import org.apache.kafka.clients.consumer.KafkaConsumer;
import java.time.Duration;
import java.util.Collections;
import java.util.Properties;

public class DynamicTopicConsumer extends Thread {
    private String kafka_servers;
    private String topic;

    public DynamicTopicConsumer(String topic_received , String kafka_servers_received)
```

```

{
    topic = topic_received;
    kafka_servers = kafka_servers received;
}

public void run()
{
    try
    {
        Properties properties = new Properties();
        properties.put("bootstrap.servers", kafka_servers);
        properties.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        properties.put("value.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
        properties.put("group.id", "your-group-id");

        try (Consumer<String, String> consumer = new KafkaConsumer<>(properties)) {
            consumer.subscribe(Collections.singletonList(topic));

            while (true) {
                ConsumerRecords<String, String> records = consumer.poll(Duration.ofMillis(100));
                for (ConsumerRecord<String, String> record : records)
                {
                    System.out.printf("topic = %s, partition = %s, offset = %d, key = %s, value = %s\n",
                        record.topic(), record.partition(), record.offset(),
                        record.key(), record.value());
                }
            }
        }
    }
    catch (Exception e)
    {
        System.out.println("Exception is caught");
    }
}
}

```

I.4. The Topic.java to src/main/java/org/acme/Topic.java

```

package org.acme.model;

public class Topic {
    public String TopicName;
    public Topic() { }
    public Topic(String topicName) { TopicName = topicName; }
    @Override
    public String toString()
    { return "Topic [TopicName=" + TopicName + "]"; }
}

```

I.5. Add the following configuration in src/main/resources/application.properties

```

quarkus.swagger-ui.path=swagger-ui
quarkus.swagger-ui.always-include=true

kafka.bootstrap.servers=<YOUR KAFKA NAME>:9092

```

I.6. Execute in a separate command line

```
./mvnw quarkus:dev
```

I.7. Navigate to <http://localhost:8080/q/swagger-ui/> and execute the POST API. Then, check that the new kafka topic is being consumed, sending some new messages to that topic. At this moment, you can change dynamically the topic that is being consumed.

localhost:8080/q/swagger-ui/#/Kafka%20Provisioning%20Resource/post_createTopic

kafka-dynamic API 1.0.0-SNAPSHOT OAS 3.0

/q/openapi Explore

Kafka Provisioning Resource

POST /createTopic

Parameters

No parameters

Request body

application/json

```
{ "TopicName": "teste" }
```

Cancel Reset

Execute Clear

Responses

Curl

```
curl -X 'POST' \
'http://localhost:8080/createTopic' \
-H 'accept: text/plain' \
-H 'Content-Type: application/json' \
-d '{' \
"TopicName": "teste" \
'}
```

Request URL

<http://localhost:8080/createTopic>

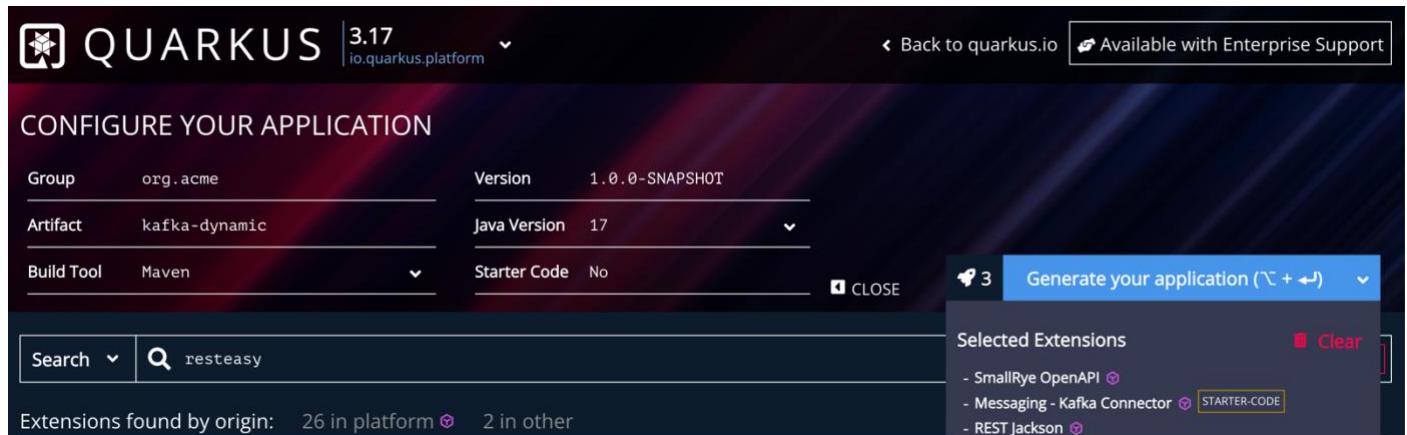
Server response

Code	Details	Links
200	<p>Response body</p> <pre>New worker started</pre> <p>Download</p> <p>Response headers</p> <pre>content-length: 18 content-type: text/plain;charset=UTF-8</pre>	No links
Responses		
Code	Description	Links
200	OK	No links
	Media type	
	text/plain	
	Controls Accept header.	

J. Creating and Removing Kafka Topics dynamically

Before starting the kafka message consumption it is required to create the specific kafka topic. Using a microservice approach, consider a REST API exposing a POST where the topic, to be created or deleted, is received, then the library kafka-client provides APIs to perform the desired operation on a kafka cluster.

- J.1. Create an empty project (KAFKA message consumer) at <https://code.quarkus.io/> with the following dependencies.



The screenshot shows the Quarkus application configuration interface. At the top, it displays 'QUARKUS | 3.17 io.quarkus.platform'. Below that, the title 'CONFIGURE YOUR APPLICATION' is shown. There are four main configuration sections: 'Group' set to 'org.acme', 'Version' set to '1.0.0-SNAPSHOT', 'Artifact' set to 'kafka-dynamic', and 'Build Tool' set to 'Maven'. Underneath these, there's a search bar containing 'resteasy'. To the right, a 'Selected Extensions' sidebar lists extensions: 'SmallRye OpenAPI', 'Messaging - Kafka Connector' (which is highlighted with a yellow border and labeled 'STARTER-CODE'), and 'REST Jackson'. At the bottom left, it says 'Extensions found by origin: 26 in platform 2 in other'. On the far right, there's a button to 'Generate your application'.

And then, add the following files to your project.

- J.2. The KafkaProvisioningResource.java to src/main/java/org/acme/KafkaProvisioningResource.java

```
package org.acme;

import jakarta.ws.rs.POST;
import jakarta.ws.rs.Path;
import jakarta.ws.rs.Produces;
import jakarta.ws.rs.core.MediaType;
import java.util.Arrays;
import java.util.Properties;
import java.util.concurrent.ExecutionException;
import org.acme.model.Topic;
import org.eclipse.microprofile.config.inject.ConfigProperty;
import org.apache.kafka.clients.admin.AdminClient;
import org.apache.kafka.clients.admin.CreateTopicsResult;
import org.apache.kafka.clients.admin.NewTopic;
import org.apache.kafka.clients.admin.DeleteTopicsResult;
import org.apache.kafka.common.errors.TopicExistsException;
import org.apache.kafka.common.errors.UnknownTopicOrPartitionException;

@Path("/TopicManagement")
public class KafkaProvisioningResource {

    @ConfigProperty(name = "kafka.bootstrap.servers")
    String kafka_servers;

    @POST
    @Path("/createTopic")
    @Produces(MediaType.TEXT_PLAIN)
    public String CreateNewTopic(Topic topic) {
        Properties properties = new Properties();
        properties.put("bootstrap.servers", kafka_servers);
        properties.put("connections.max.idle.ms", 10000);
        properties.put("request.timeout.ms", 5000);
        try (AdminClient client = AdminClient.create(properties)) {
            CreateTopicsResult result = client.createTopics(Arrays.asList(
                new NewTopic(topic.TopicName, 1, (short) 1) ));
            try { result.all().get(); }
            catch ( org.apache.kafka.common.errors.TopicExistsException exc )
            { throw exc; }
            catch ( ExecutionException | InterruptedException e )
            { throw new IllegalStateException(e); }
        }
        return ("New Topic created = " + topic);
    }
}
```

```

@POST
@Path("/DeleteTopic")
@Produces(MediaType.TEXT_PLAIN)
public String RemoveOneTopic(Topic topic) {
    Properties properties = new Properties();
    properties.put("bootstrap.servers", kafka_servers);
    properties.put("connections.max.idle.ms", 10000);
    properties.put("request.timeout.ms", 5000);
    try (AdminClient client = AdminClient.create(properties))
    {
        DeleteTopicsResult result = client.deleteTopics(Arrays.asList( topic.TopicName ));
        try { result.all().get(); } catch ( org.apache.kafka.common.errors.UnknownTopicOrPartitionException exc ) { throw exc; }
        catch ( ExecutionException | InterruptedException e ) { throw new IllegalStateException(e); }
    }
    return ("New Topic deleted = " + topic);
}

```

J.3. The Topic.java to src/main/java/org/acme/Topic.java

```

package org.acme.model;

public class Topic {
    public String TopicName;
    public Topic() { }
    public Topic(String topicName) { TopicName = topicName; }
    @Override
    public String toString()
    { return "Topic [TopicName=" + TopicName + "]"; }
}

```

J.4. Add the following configuration in src/main/resources/application.properties

```

quarkus.swagger-ui.path=swagger-ui
quarkus.swagger-ui.always-include=true

kafka.bootstrap.servers=<YOUR KAFKA NAME>:9092

```

J.5. Add the following dependency in pom.xml

```

<dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-clients</artifactId>
</dependency>

```

J.6. Execute in a separate command line

```

./mvnw quarkus:dev

```

Navigate to <http://localhost:8080/q/swagger-ui/> and execute the POST APIs. Then, check that the new kafka topic is created or deleted, listing the topics in that kafka cluster.

localhost:8080/q/swagger-ui/#/Kafka%20Provisioning%20Resource/post_TopicManagement_RemoveTopic

Swagger
powered by SMARTBEAR

/q/openapi Explore

kafka-dynamic API 1.0.0-SNAPSHOT OAS 3.0

/q/openapi

Kafka Provisioning Resource

POST /TopicManagement/RemoveTopic

Parameters

No parameters

Request body

application/json

```
{
  "TopicName": "novotopico"
}
```

Cancel Reset

Execute Clear

Responses

Curl

```
curl -X 'POST' \
  'http://localhost:8080/TopicManagement/RemoveTopic' \
  -H 'accept: text/plain' \
  -H 'Content-Type: application/json' \
  -d '{
    "TopicName": "novotopico"
}'
```

Request URL

http://localhost:8080/TopicManagement/RemoveTopic

Server response

Code	Details	Links
404 Undocumented	Error: Not Found Response headers content-length: 0	No links
Responses		
Code	Description	Links
200	OK Media type text/plain Controls Accept header. Example Value Schema string	No links

POST /TopicManagement/createTopic

K. Deploying QUARKUS in AWS EC2 environment with github and terraform!

K.1.Create your new Quarkus project as before.

K.2.Execute the following command in your project directory to add the docker dependencies:

```
./mvnw quarkus:add-extension -Dextensions="quarkus-container-image-docker"
```

K.3.Add the following configuration to **application.properties** in order to include all the dependencies to execute in a different environment from dev:

```
quarkus.swagger-ui.path=swagger-ui  
quarkus.swagger-ui.always-include=true  
  
quarkus.container-image.build=true  
quarkus.container-image.push=true
```

Also, add your Git Hub accountID in the **application.properties** of your project, to identify it correctly and push it directly:

```
quarkus.container-image.group=YOUR DOCKER USERNAME
```

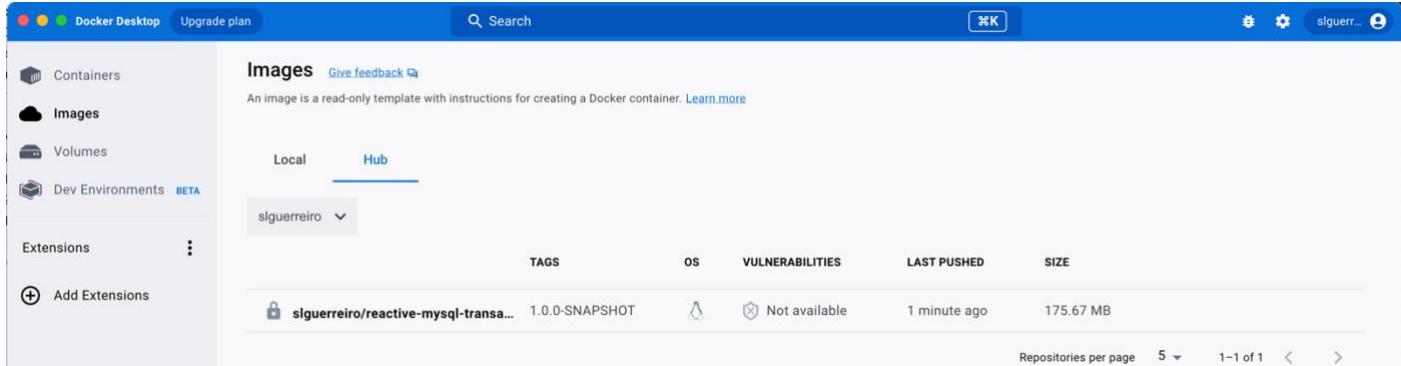
K.4.Login in your account (you can create one in <http://hub.docker.com> if needed)

```
docker login -u "your docker username" -p "your password"
```

K.5. Build docker image and push it directly to you github account:

```
./mvnw clean package
```

K.6.Check that image has been created successfully on github or on the docker dashboard (optional):



K.7.Now, to deploy your Quarkus project in EC2, using terraform, and based on the image previously pushed to git, create a .tf file with the following instructions. Note that the AMI to create should be compatible with the operating systems where you previously compiled the docker image.

```
terraform {  
    required_version = ">= 1.0.0, < 2.0.0"  
  
    required_providers {  
        aws = {  
            source  = "hashicorp/aws"  
            version = "~> 4.0"  
        }  
    }  
  
    provider "aws" {  
        region      = "us-east-1"  
        access_key  = "YOUR ACES KEY"  
        secret_key  = "YOUR SECRET KEY"  
        token       = "YOUR TOKEN"  
    }  
}
```

```

resource "aws_instance" "exampleDeployQuarkus" {
  ami
    = "ami-08cf815cff6ee258a" # Amazon Linux ARM AMI built by Amazon Web Services - FOR
  DOCKER image COMPATIBILITY if compiled previously on ARM
  instance type
    = "t4g.nano"
  vpc_security_group_ids
    = [aws_security_group.instance.id]
  key_name
    = "vokey"

  user_data = "${file("quarkus.sh")}"
  user_data_replace_on_change = true

  tags = {
    Name = "terraform-deploy-QuarkusProject"
  }
}

resource "aws_security_group" "instance" {
  name = var.security_group_name
  ingress {
    from_port    = 0
    to_port      = 0
    protocol     = "-1"
    cidr_blocks = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }
  egress {
    from_port      = 0
    to_port        = 0
    protocol       = "-1"
    cidr_blocks   = ["0.0.0.0/0"]
    ipv6_cidr_blocks = ["::/0"]
  }
}

variable "security group name" {
  description = "The name of the security group"
  type        = string
  default     = "terraform-Quarkus-instance"
}

output "address" {
  value      = aws_instance.exampleDeployQuarkus.public_dns
  description = "Address of the Quarkus EC2 machine"
}

```

And the correspondingly script "quarkus.sh" with the following:

```

#!/bin/bash
echo "Starting..."

sudo yum install -y docker

sudo service docker start

sudo docker login -u "YOUR DOCKER USERNAME" -p "YOUR DOCKER PASSWORD"

sudo docker pull YOUR DOCKER USERNAME/tryout1:1.0.0-SNAPSHOT

sudo docker run -d --name tryout2 -p 9000:9000 YOUR DOCKER USERNAME/YOUR ARTIFACT ID:1.0.0-SNAPSHOT

echo "Finished."

```

Docker images built for ARM architecture may not work on AMD architecture machines. This is because ARM and AMD processors use different instruction sets and have different hardware architectures.

For this example, the AMI identifier and instance type were extracted from the launching instance AWS GUI.

EC2 > Instances > Launch an instance

Launch an instance Info

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

Name and tags Info

Name
e.g. My Web Server

Application and OS Images (Amazon Machine Image) Info

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below.

Search our full catalog including 1000s of application and OS images

Recents

Amazon Linux macOS Ubuntu Windows Red Hat SI
Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

Amazon Linux 2 AMI (HVM) - Kernel 5.10, SSD Volume Type
Free tier eligible

ami-0dfcb1ef8550277af (64-bit (x86)) / ami-0cd7323ab3e63805f (64-bit (Arm))
Virtualization: hvm ENA enabled: true Root device type: ebs

Description
Amazon Linux 2 LTS Arm64 Kernel 5.10 AMI 2.0.20230207.0 arm64 HVM gp2

Architecture 64-bit (Arm)

Instance type Info

c6g.medium
Family: c6g 1 vCPU 2 GiB Memory
On-Demand Linux pricing: 0.034 USD per Hour
On-Demand RHEL pricing: 0.094 USD per Hour
On-Demand SUSE pricing: 0.065 USD per Hour

Summary

Number of instances Info
1

Software image (AMI)
Amazon Linux 2 LTS Arm64 Kernel 5.10 AMI 2.0.20230207.0 arm64 HVM gp2 ami-0cd7323ab3e63805f

Virtual server type (instance type)
c6g.medium

Firewall (security group)
New security group

Storage (volumes)
1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 30 GiB of EBS storage, 2 million I/Os, 1 GiB of snapshots, and 100 GiB of bandwidth to the internet.

References

Full official documentation:

- <https://quarkus.io/guides/>
- <https://quarkus.io/guides/container-image#building>
- <https://dev.mysql.com/downloads/>
- <https://quarkus.io/guides/reactive-sql-clients>
- <https://quarkus.io/guides/kafka-reactive-getting-started>
- <https://quarkus.io/guides/amazon-lambda>
- <https://aws.amazon.com/pt/blogs/aws-brasil/integracao-do-aws-lambda-com-o-quarkus/>
- <https://dev.to/marcuspaulo/tutorial-publish-a-quarkus-application-in-kubernetes-minikube-and-dockerhub-36nd>
- <https://stackoverflow.com/questions/23935141/how-to-copy-docker-images-from-one-host-to-another-without-using-a-repository>
- https://quarkus.io/guides/container-image#quarkus-container-image-docker_quarkus.docker.executable-name

The goal of this document is to present the details related with: Installation and execution of the Camunda platform, Designing the BPMN 2.0 models using the main constructors of the language, Deploy those previous modelled BPMN 2.0 models, Testing the models, and Operating the platform.

The examples address the following aspects: how to integrate a BPMN process model with an external worker written in JAVA, how to add a user task to the BPMN process model, how to add gateways to the BPMN process model, how to implement DMN rules in the BPMN process model, and how to invoke a cloud REST API.

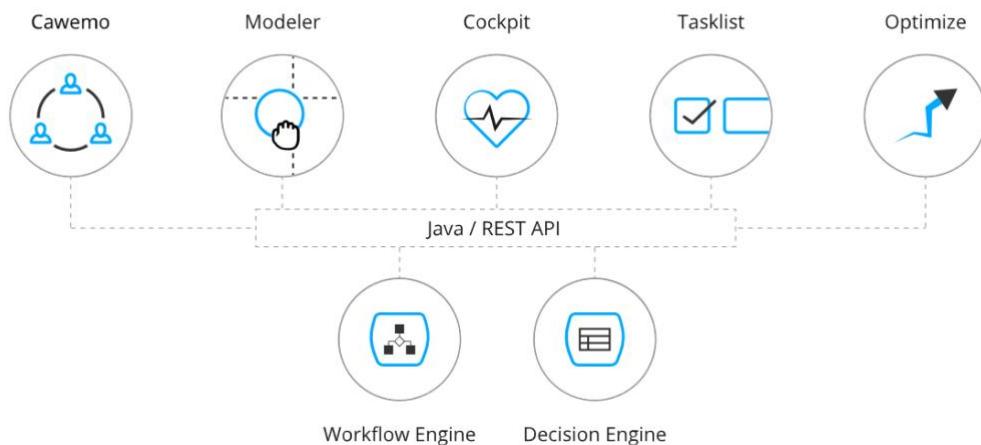
The following contents is presented in this document.

Contents

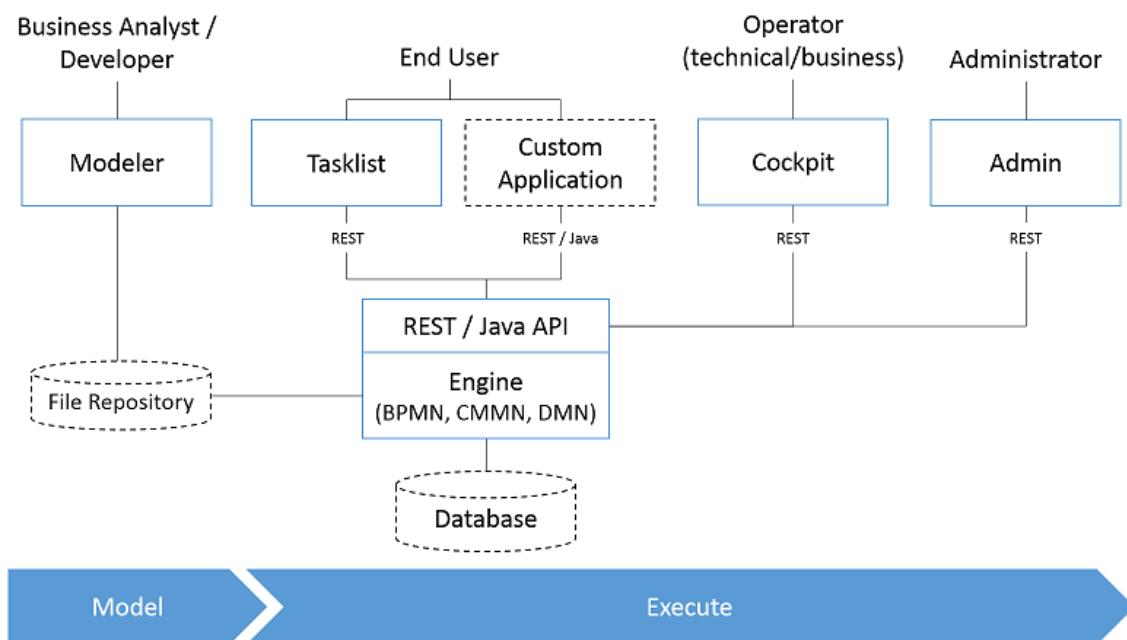
A. Camunda brief overview	2
B. Installation and operation of Camunda Engine in AWS EC2	3
C. Add a User Task to the Process	4
D. Yet, another option using Camunda forms	10
E. Add Gateways to the Process	12
F. Business rule enforcement	17
G. Integrate with a REST API installed remotely using Camunda http-connector – REQUEST ONLY	23
H. Integrate with a REST API installed remotely using Camunda http-connector – REQUEST with variable	25
I. Integrate with a REST API installed remotely using Camunda http-connector – REQUEST & RESPONSE	28
J. Integrate with a REST API installed remotely using Camunda http-connector – REQUEST & RESPONSE & Headers fields	33
K. Using an intermediate message catch event inside a process	34
References	36

A. Camunda brief overview

Camunda encompasses the following software modules.



The Community edition that we are going to use have a basic cockpit and optimize component is also not available. The main stakeholders involved in Camunda are the following:



B. Installation and operation of Camunda Engine in AWS EC2

C.1. Create an EC2 new instance of the type: Amazon Linux 2 AMI (HVM), SSD Volume Type. The exact versions may change with time. And define the inbound rules to allow the 8080 and 22 accessed from anywhere.

C.2. Install the most recent version of camunda-modeler (**version 7 recommended**) using the download public available at: <https://camunda.com/download/modeler/>

C.3. Access the new EC2 instance and execute the following commands:

```
sudo yum update -y
sudo yum install -y docker
sudo service docker start
sudo usermod -aG docker ec2-user
docker ps
exit
```

C.4. Then, access again your EC2 instance and install and execute the Camunda engine with the following command:

```
docker pull camunda/camunda-bpm-platform:latest
docker run -d --name camunda -p 8080:8080 camunda/camunda-bpm-platform:latest
```

Test the installed Camunda engine opening the following URL in your browser (the default user/password is **demo/demo**):

```
# open browser with url: http://YOUR AWS EC2 INSTANCE:8080/camunda-welcome/index.html
```

Hint 1: Later if you would like to list all the available container use the following command:

```
PS C:\Users\Sérgio Guerreiro> docker container ls --all
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              STATUS
PORTS      NAMES
44509250aa0c        camunda/camunda-bpm-platform:latest    "/sbin/tini -- ./cam..."   11 days ago        Exited
```

Hint 2: Later if you would like to stop the container use the following command:

```
docker container stop <containerID_hash>
```

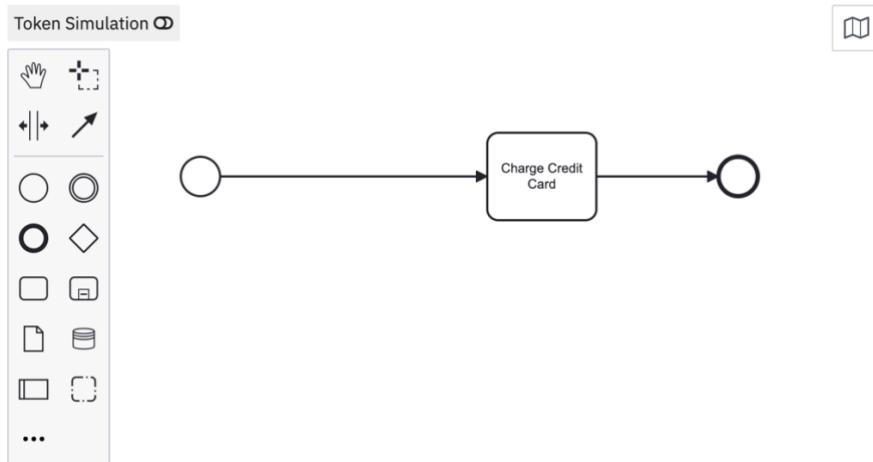
Hint 3: Later if you would like to restart the container use the following command:

```
docker container start <containerID_hash>
```

C. Add a User Task to the Process

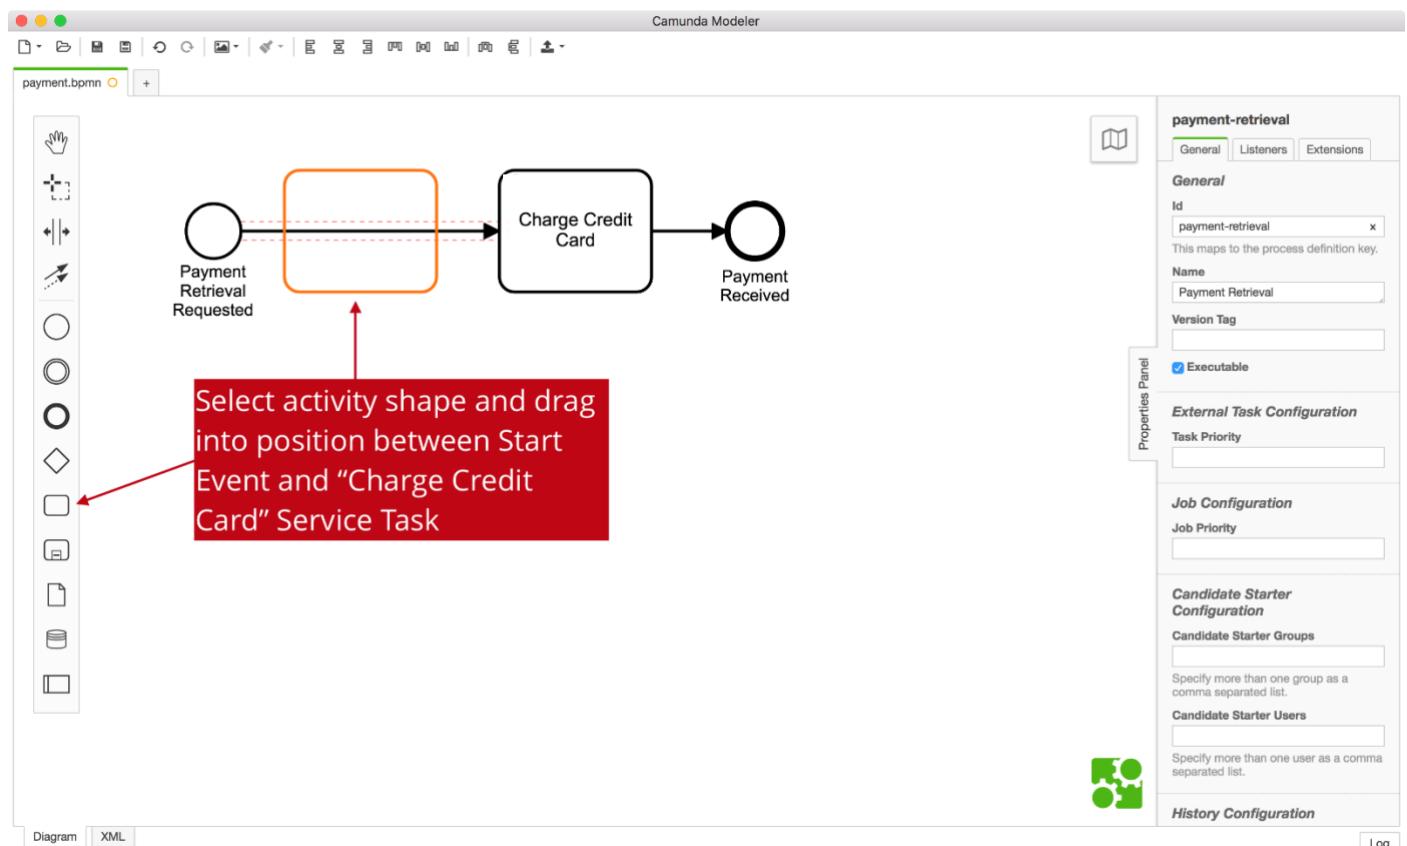
This example is based in the Camunda documentation available at <https://docs.camunda.org/get-started/quick-start/user-task/>. The goal of the example is to show how to add a human task in the middle of a process execution. After deploying, go to the tasklist (all tasks) identify the “Approve payment”, select the claim, and check approved and then complete. Then, check that instances and tasklist are not available anymore.

- D.1. Add an abstract task with the name “Charge Credit Card” between a start and an end event.

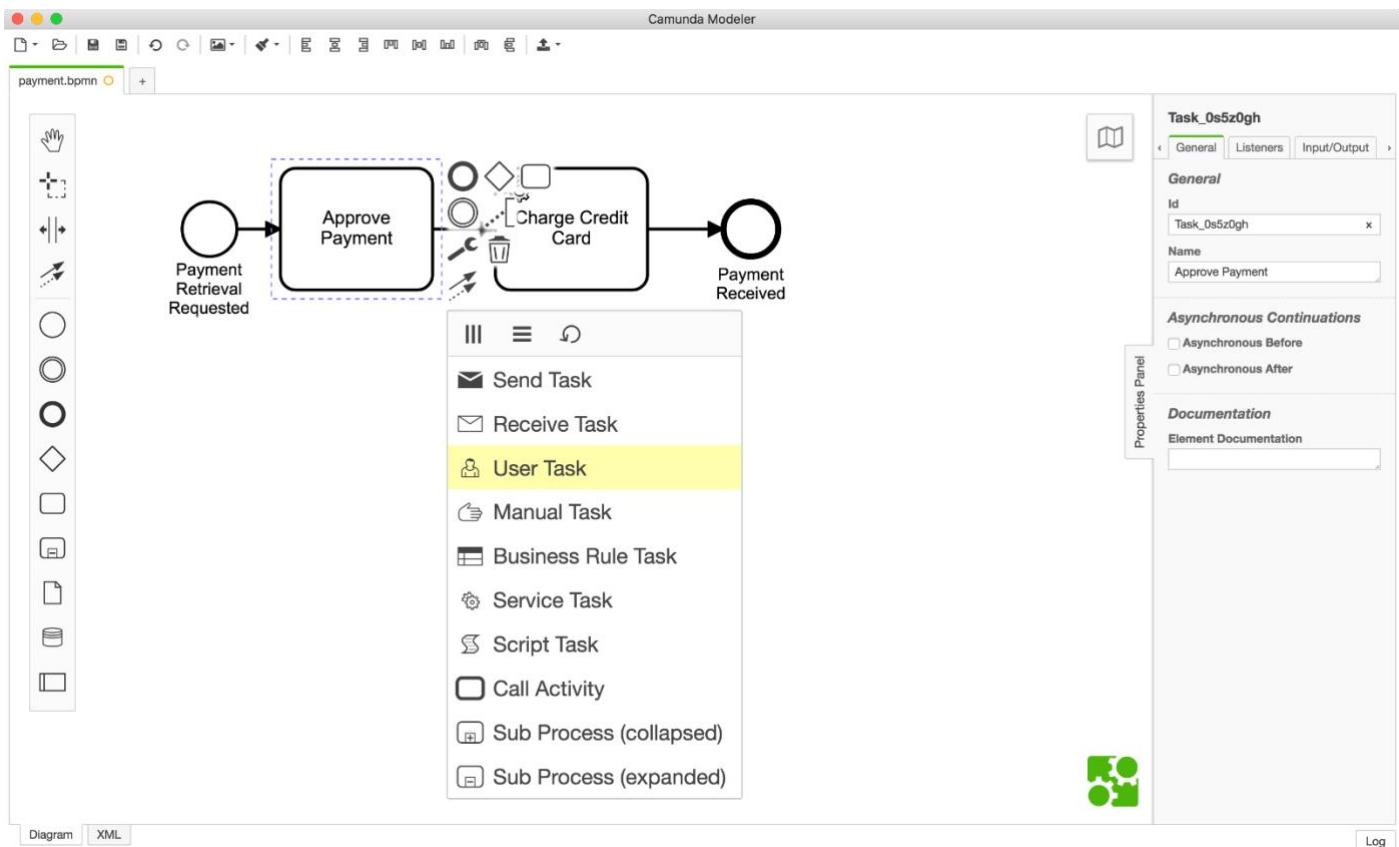


- D.2. Add a User Task

Now, we want to modify our process so that we can involve humans. To do so, open the process in the Camunda Modeler. Next, from the Modeler’s left-hand menu, select the activity shape (rectangle) and drag it into position between the Start Event and the “Charge Credit Card” Service Task. Name it *Approve Payment*.



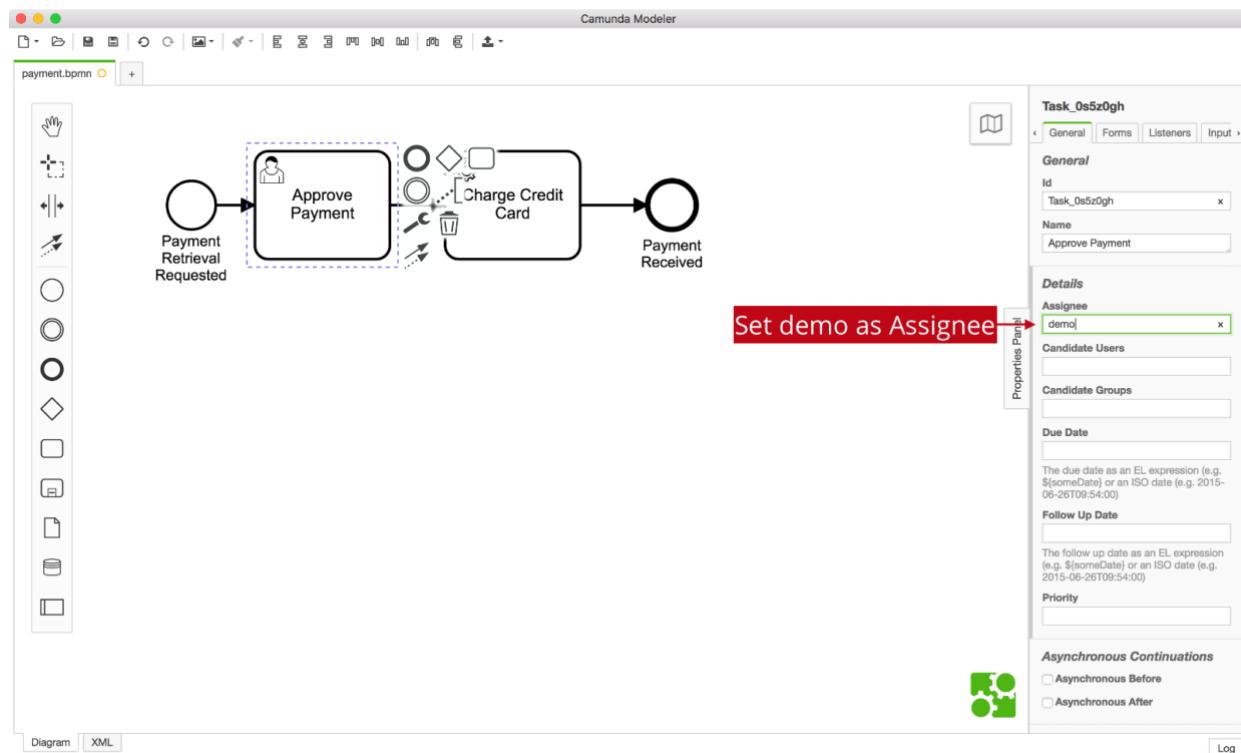
Change the activity type to *User Task* by clicking on it and using the wrench button menu.



D.3. Configure a User Task

Next, open the properties view. If the properties view is not already visible, click on the "Properties Panel" label on the right-hand side of the Modeler canvas.

Select the User Task on the canvas. This will update the selection in the properties view. Scroll to the property named *Assignee*. Type *demo*.



D.4. Configure a basic form in the User Task

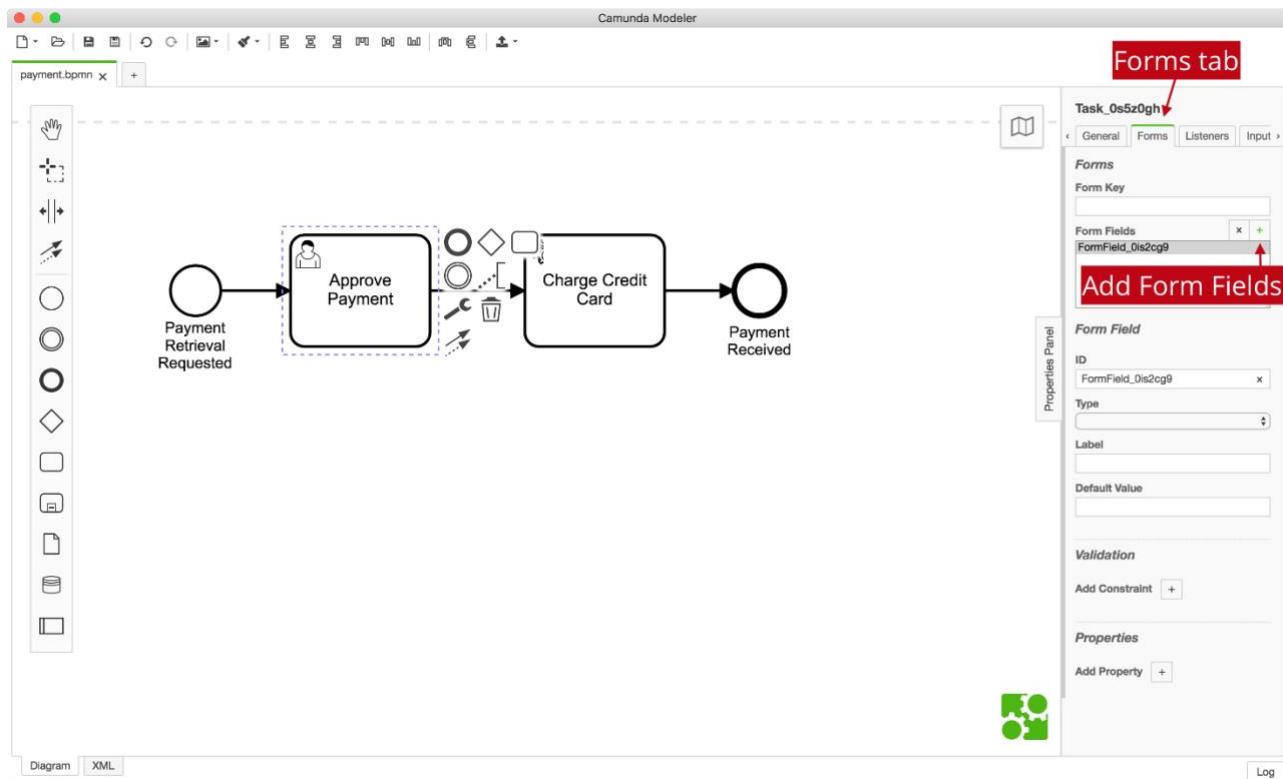
This step will also take place in the properties panel. If the panel is not already visible, click on the "Properties Panel" label on the right-hand side of the Modeler canvas.

Select the User Task on the canvas. This will update the selection in the properties view.

Click on the Tab **Forms** in the properties panel.

(version 5.8.0) -> In the type of Forms choose "Generated Task Forms"

Add three form fields by clicking on the plus button:

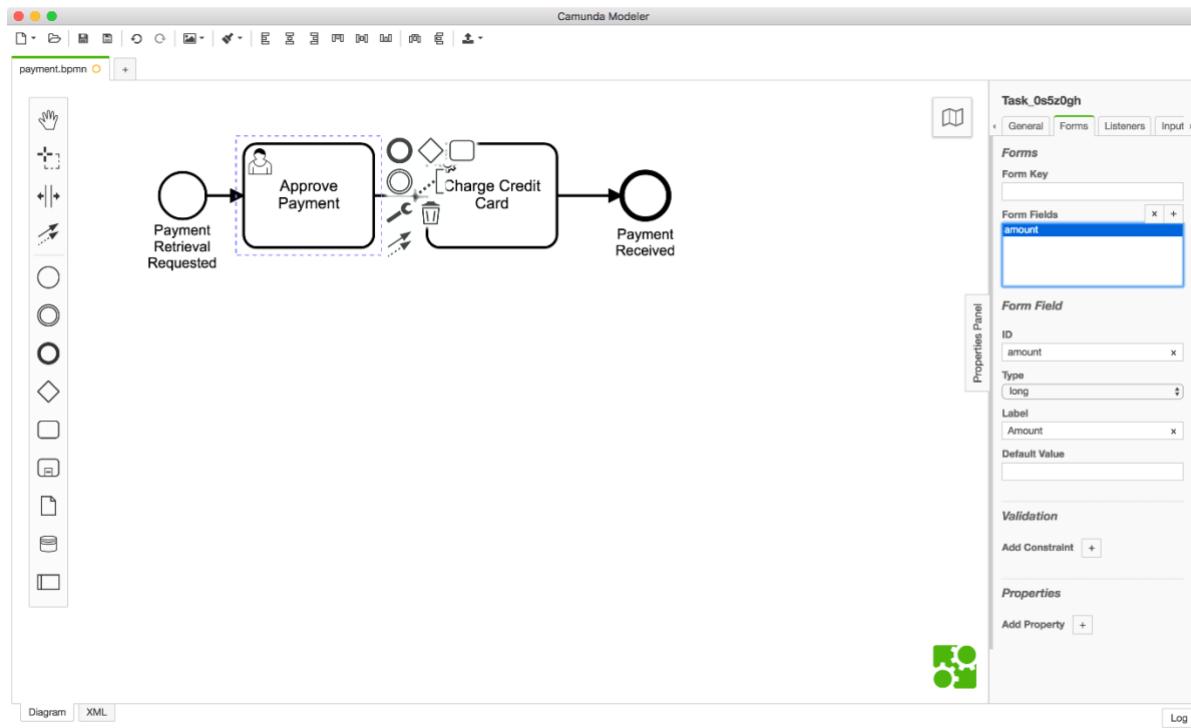


Field 1:

ID: amount

Type: long

Label: Amount

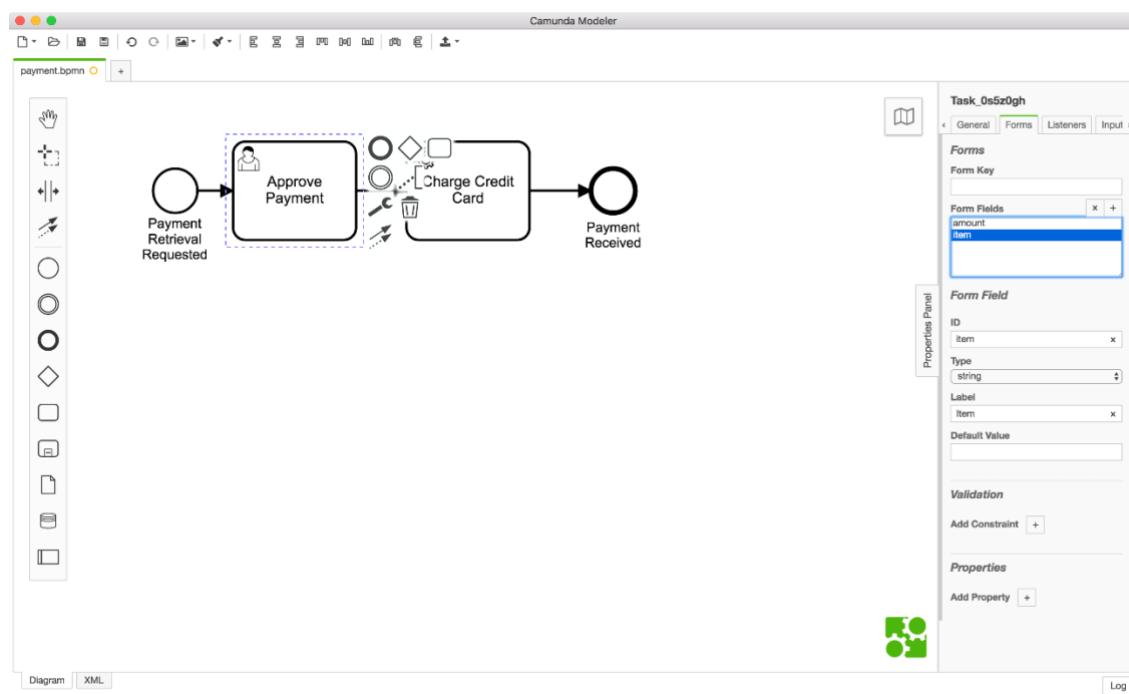


Field 2:

ID: item

Type: string

Label: Item

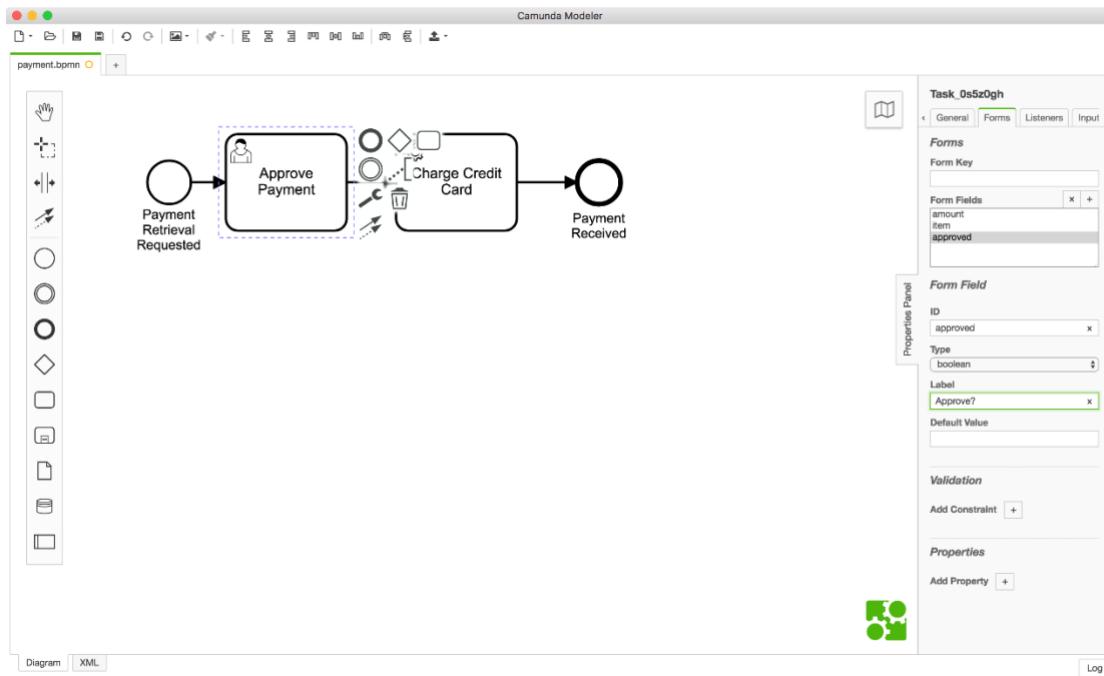


Field 3:

ID: approved

Type: boolean

Label: Approved?

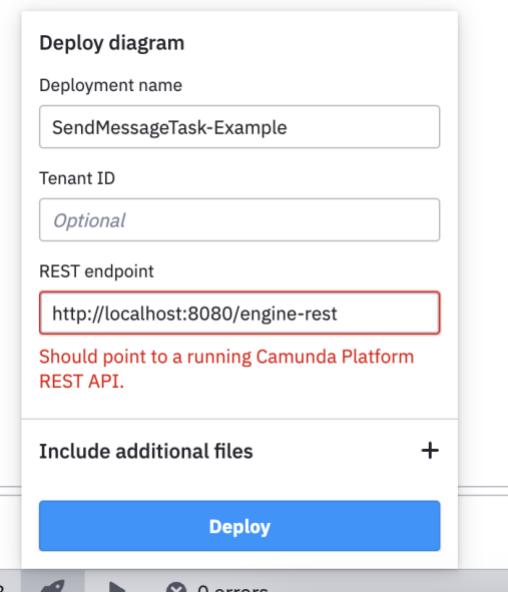


D.5. Deploy the Process

Use the Deploy Button in the Camunda Modeler to deploy the updated process to Camunda.



Replacing the localhost for the AWS EC2 DNS_name where you deployed the camunda-engine.

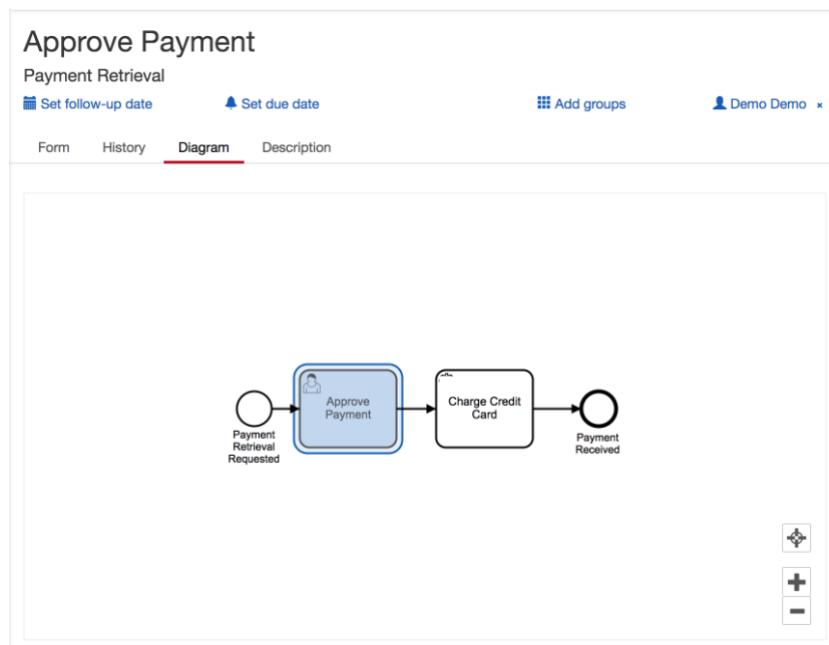


D.6. Work on the Task

Go to Camunda Tasklist and log in with the credentials "demo / demo". Click on the button to start a process instance. This opens a dialog where you can select *Payment Retrieval* from the list. Now you can set variables for the process instance using a generic form.

The generic form can be used whenever you have not added a dedicated form for a User Task or a Start Event. Click on the *Add a variable* button to create a new row. Fill in the form as shown in the screenshot. When you're done, click *Start*.

You should now see the *Approve Payment* task in your Tasklist. Select the task and click on the *Diagram* tab. This displays the process diagram highlighting the User Task that's waiting to be worked on.

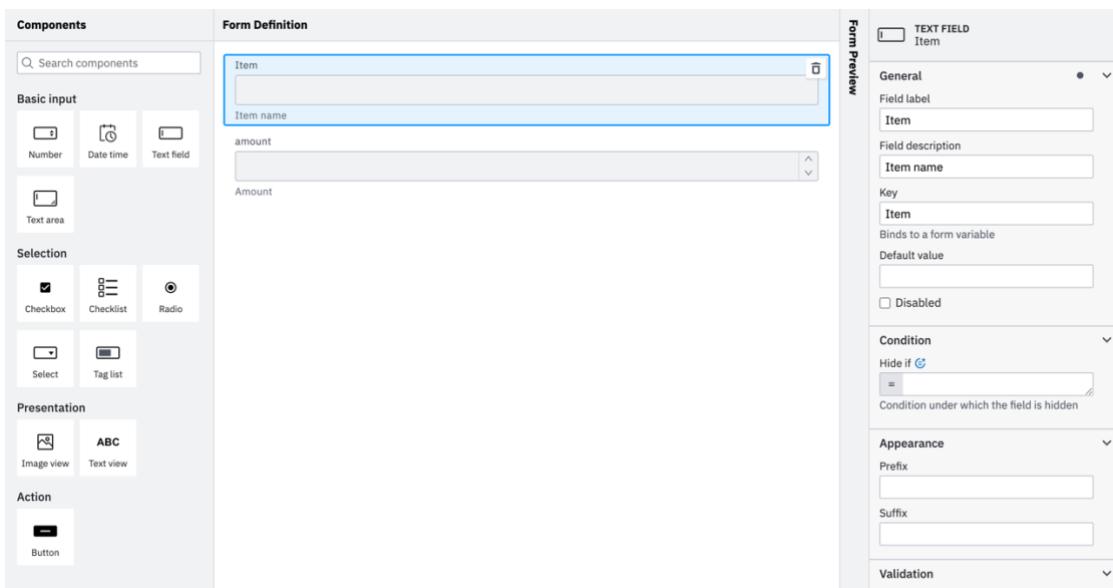


To work on the task, select the *Form* tab. Because we defined the variables in the Form Tab in the Camunda Modeler, the Tasklist has automatically generated form fields for us.

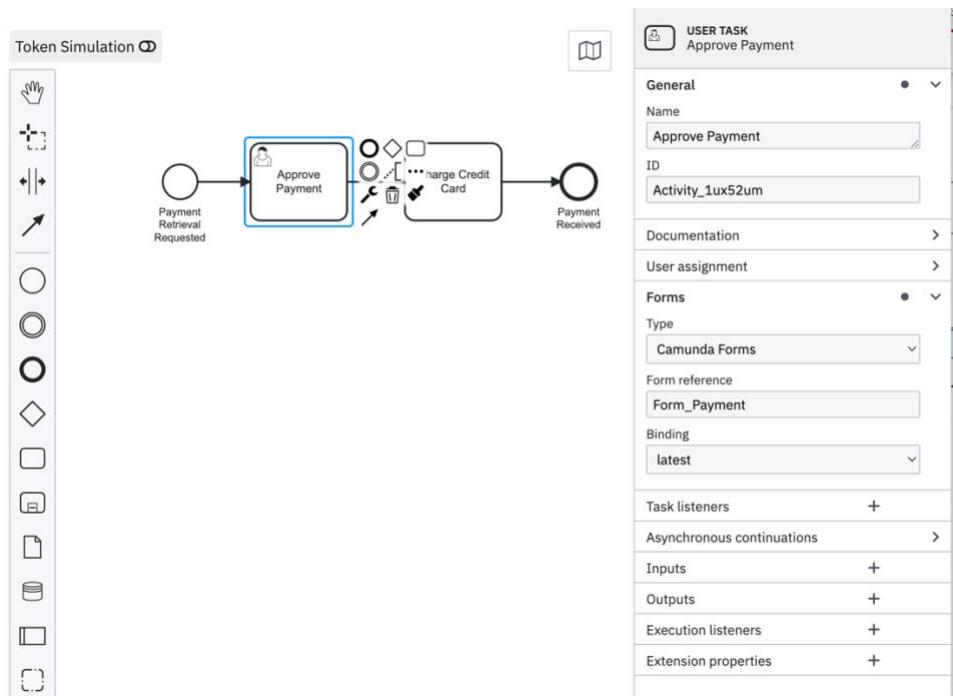
D. Yet, another option using Camunda forms

Instead of hard coding the form directly in the user task you can create the form separately, deploy it, and then refer to it in your process. However, remind that Camunda is a process engine, not an optimized platform for user interaction. When intensive user interactions are required consider the usage of solutions such as <https://github.com/formio/formio>.

- E.1. Create a new form, using the option: File > New File > Form (Camunda Platform 7)
- E.2. Define an ID for the entire form, i.e., Form_Payment
- E.3. Add 1 text field and 1 number, similar with the following figure:



- E.4. Deploy the form.
- E.5. Change the configuration of the previous user task accordingly with the following figure:



- E.6. Deploy the process
- E.7. Start a process instance, using the **Start instance** Button in the Camunda Modeler to deploy the updated process to Camunda.



- E.8. Check in the cockpit that a pending instance is waiting for a user task:

The screenshot shows the Camunda cockpit's "Process Instances" tab. At the top, there are tabs for "Process Instances", "Incidents", "Called Process Definitions", and "Job Definitions". The "Process Instances" tab is selected. Below the tabs, there is a table with two rows of data. The columns are "State", "ID", "Start Time", and "Business Key". The first row has a green checkmark in the "State" column, and the ID is "3dae78d0-b3aa-11ed-8bae-0242ac110002". The start time is "2023-02-23T18:45:27" and the business key is "default". The second row also has a green checkmark in the "State" column, and the ID is "f1139329-b3a9-11ed-8bae-0242ac110002". The start time is "2023-02-23T18:43:19" and the business key is "default".

- E.9. Choose your process instance, and then change to user tasks tab, and select one pending action in the task list. You can now see your form!

The screenshot shows the Camunda Tasklist interface. On the left, there is a sidebar with a "My Tasks (4)" section and a list of other tasks: "My Group Tasks", "Accounting", "John's Tasks", "Mary's Tasks", "Peter's Tasks", and "All Tasks". The main area shows a table of tasks. One task, "Approve Payment" from "Process_1823xr6", is highlighted and expanded. This task was created 3 minutes ago and has a due date of 50. The task details show "Demo Demo" and "50". The task is assigned to "Demo Demo". The "Form" tab is selected in the task details view. The form fields include "Item" (with a placeholder "Item name"), "amount" (with a placeholder "Amount"), and "Invoice Number" (with a value "PSACE-5342").

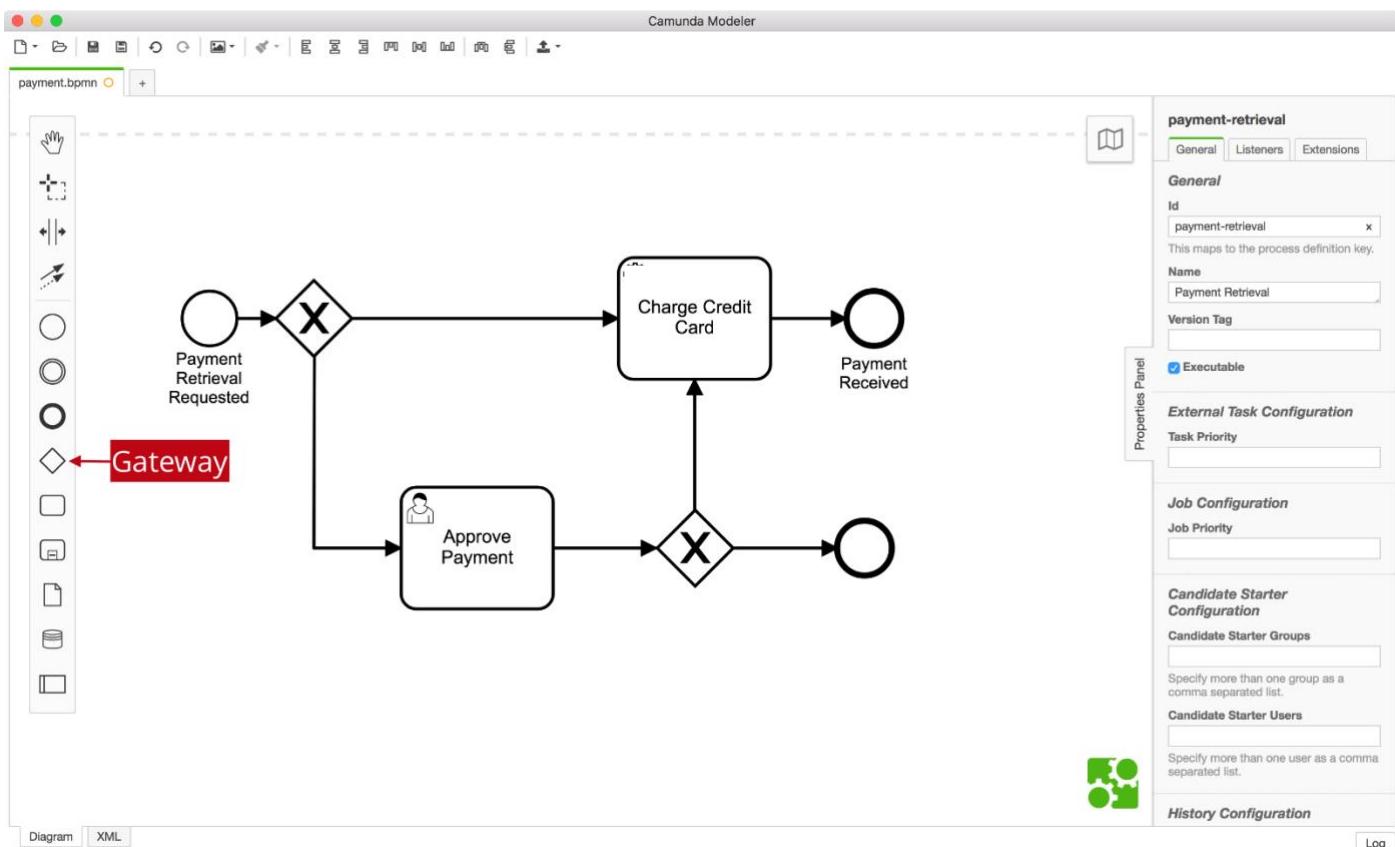
E. Add Gateways to the Process

This example is based in the Camunda documentation available at <https://docs.camunda.org/get-started/quick-start/gateway/>. The goal of the example is to implement decisions in the flow of the business process, using the BPMN gateway element.

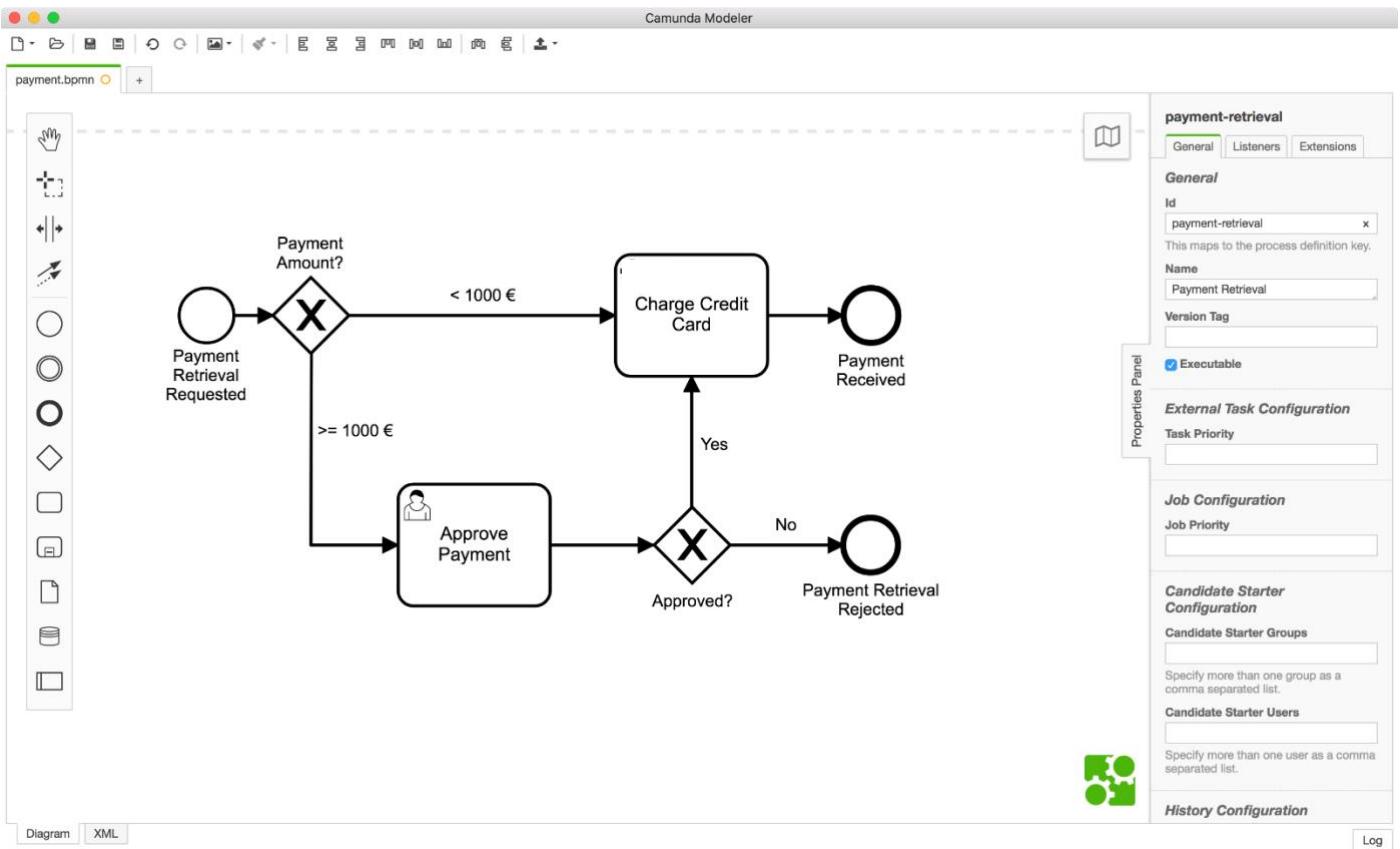
F.1. Add Two Gateways

We want to modify our process so that it is more dynamic. To do so, open the process in the Camunda Modeler.

Next, from the Modeler's left-hand menu, select the gateway shape (diamond) and drag it into position between the Start Event and the Service Task. Move the User Task down and add another Gateway after it. Lastly, adjust the Sequence Flows so that the model looks like this:

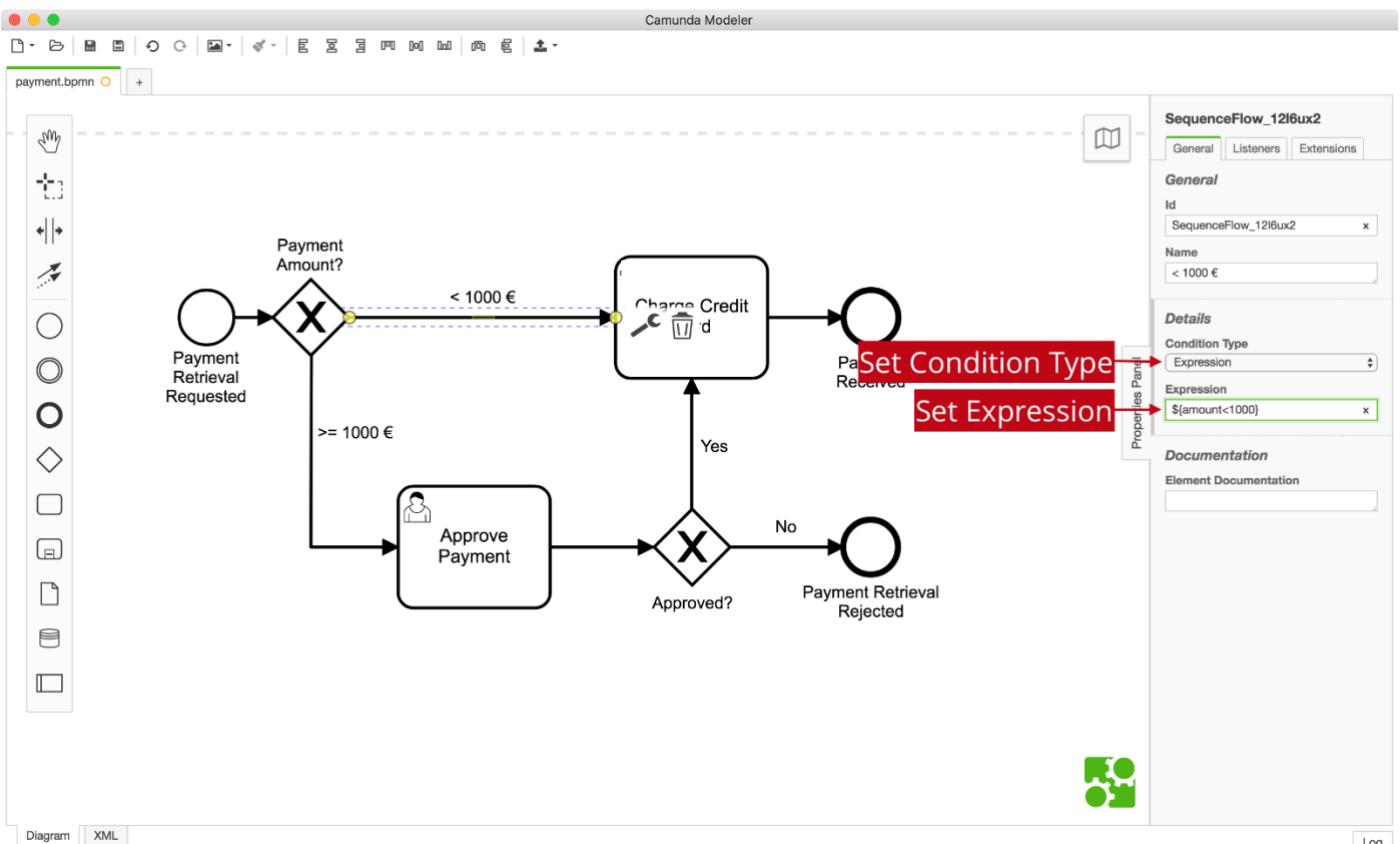


Now also name the new elements accordingly:



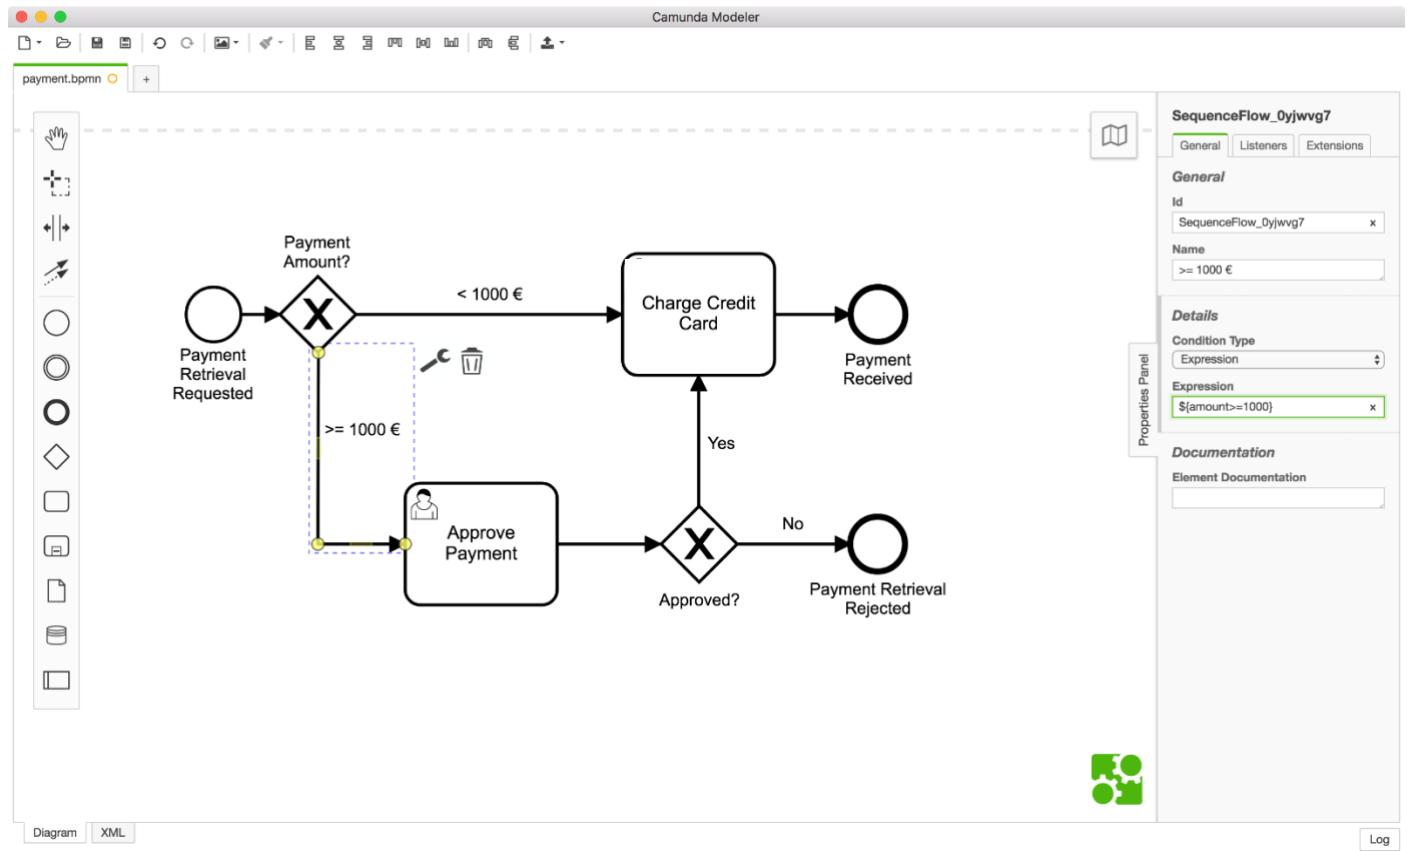
F.2. Configure the Gateways

Next, open the properties panel and select the $< 1000 \text{ €}$ Sequence Flow after the Gateway on the canvas. This will update the selection in the properties panel. Scroll to the property named **Condition Type** and change it to **Expression**. Then input `#{amount<1000}` as the Expression. We are using the [Java Unified Expression Language](#) to evaluate the Gateway.

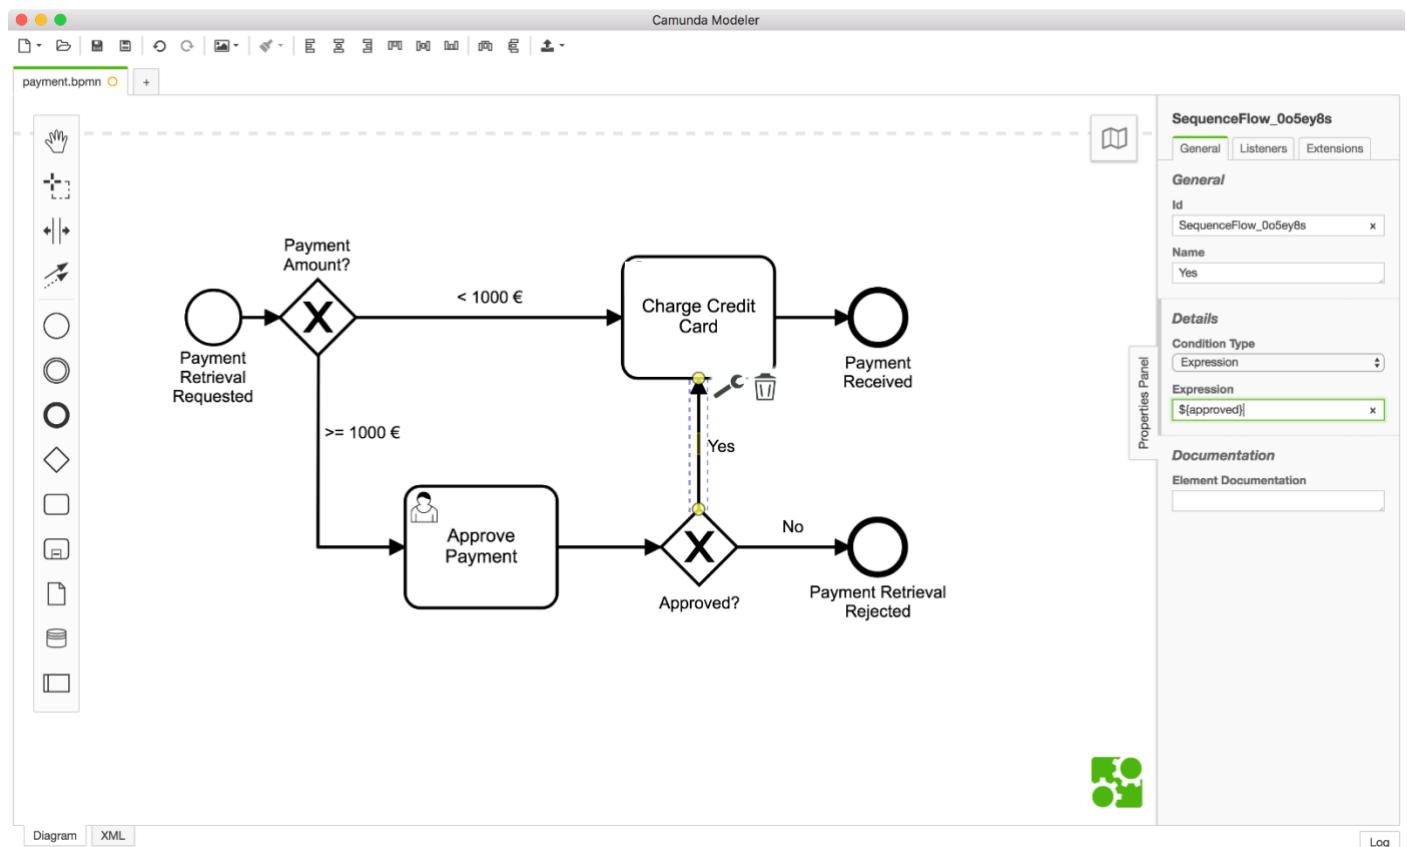


Next, change the Expressions for the other Sequence Flows, too.

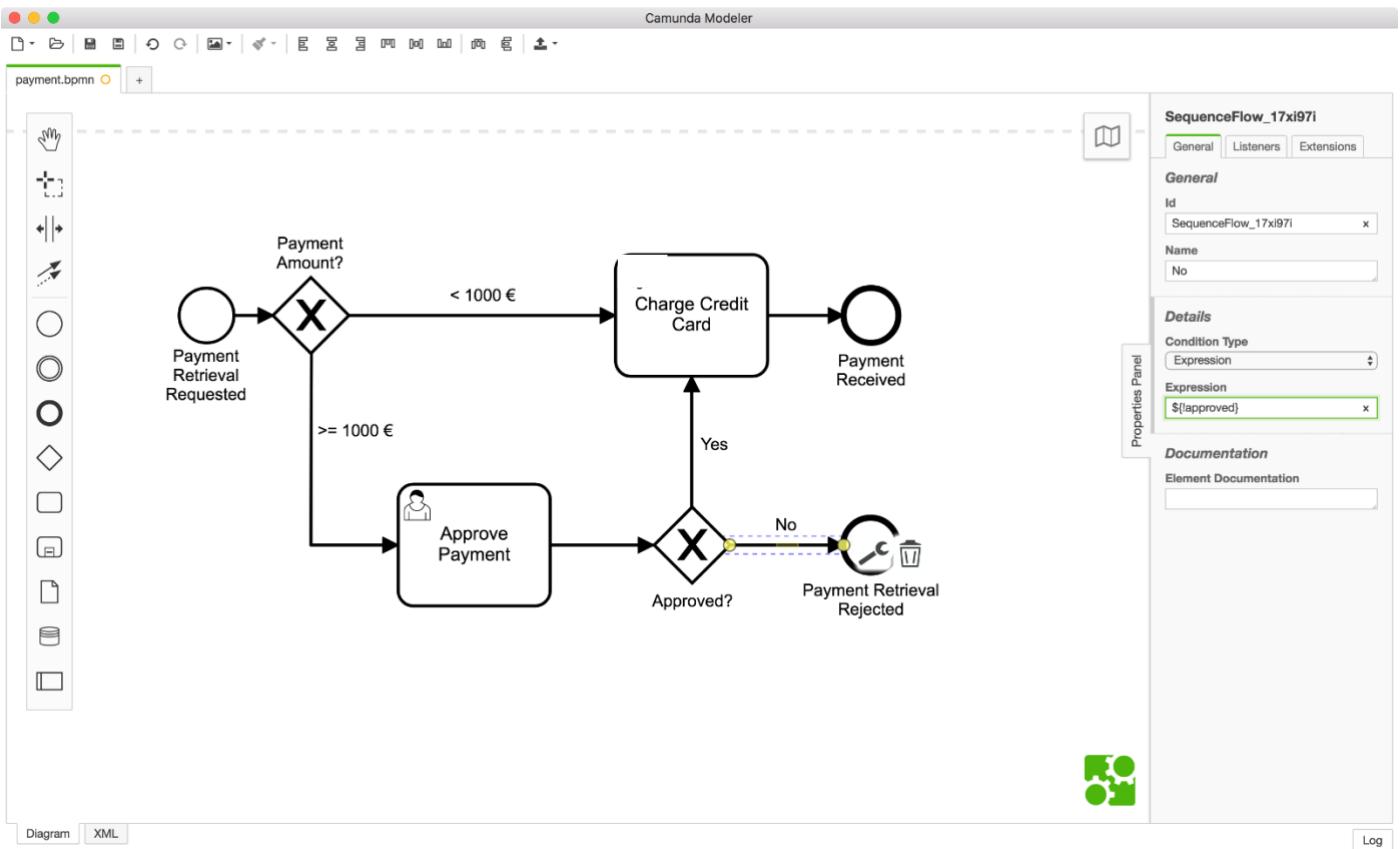
For the $\geq 1000 \text{ €}$ Sequence Flow, use the Expression `#{amount} >= 1000`:



For the Yes Sequence Flow, use the Expression `#{approved}`:



For the No Sequence Flow, use the Expression `#{!approved}`:



F.3. Deploy the Process

F.4. Work on the Task

Go to Tasklist and log in with the credentials “demo / demo”. Click on the “Start process” button to start a process instance for the *Payment Retrieval* Process. Next, set variables for the process instance using the generic form as we learned in the *User Tasks* section.

Add a varia... +	Name	Type	Value
Rem... x	amount	Long	555
Rem... x	item	String	item-xyz

Fill in the form as shown in the screenshot and make sure you use an amount that is larger or equal to 1000 in order to see the User Task *Approve Payment*. When you are done, click *Start*.

You should see the *Approve Payment* task when you click on *All Tasks*. In this quick start, we’re logged into Tasklist as an admin user, and so we can see all tasks associated with our processes. However, it’s possible to create [filters in Tasklist](#) to determine which users can see which tasks based on [user authorization](#) as well as other criteria.

To work on the task, select the *Form* tab and check the *approved* checkbox so that our payment retrieval gets approved. We should see that our worker prints something to the console.

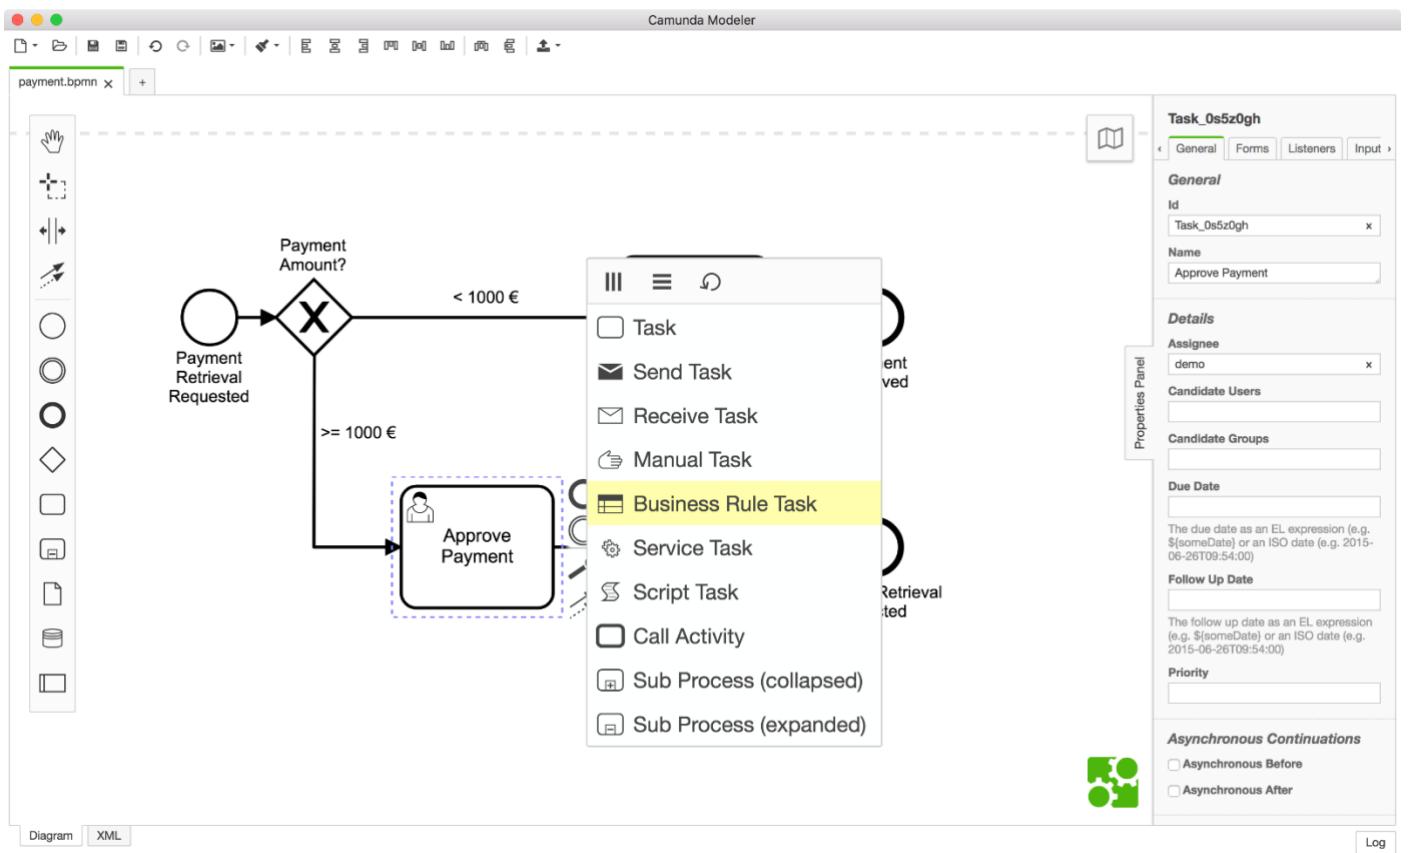
Next, repeat the same steps, and this time, reject the payment. You should also create one instance with an amount less than 1000 to confirm that the first gateway works correctly.

F. Business rule enforcement

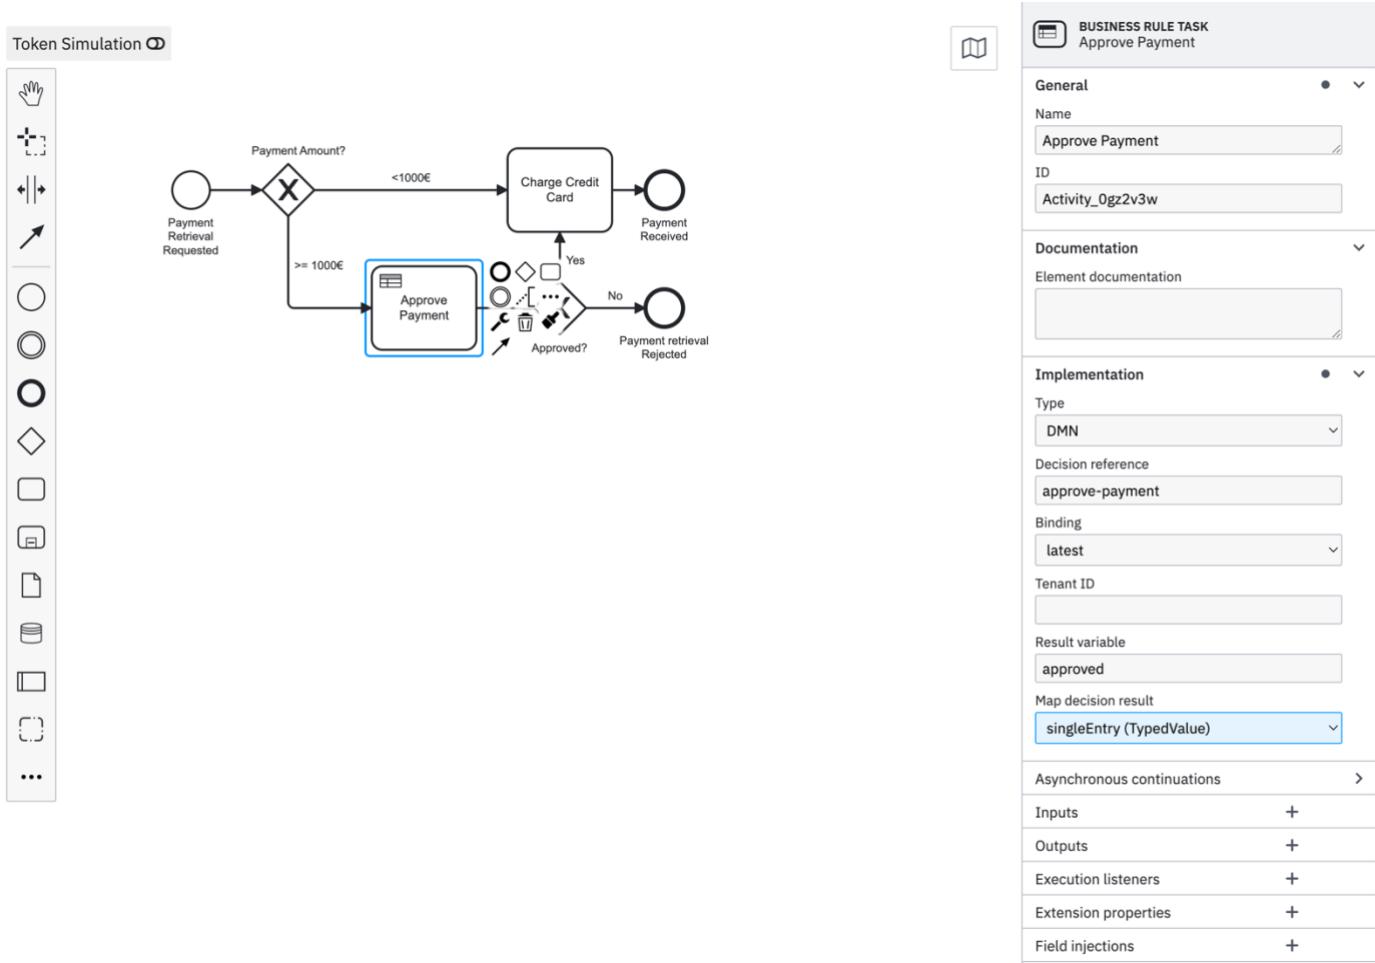
This example is based in the Camunda documentation available at <https://docs.camunda.org/get-started/quick-start/decision-automation/>. The goal of this example is to implement business rules using the DMN 1.1 standard.

G.1. Add a Business Rule Task to the Process

Use the Camunda Modeler to open the Payment Retrieval process then click on the Approve Payment Task. Change the activity type to *Business Rule Task* in the wrench button menu.



Next, link the Business Rule Task to a DMN table by changing `Implementation` to DMN and `Decision Ref` to `approve-payment` in the properties panel. In order to retrieve the result of the evaluation and save it automatically as a process instance variable in our process, we also need to change the `Result Variable` to `approved` and use `singleEntry` as the `Map Decision Result` in the properties panel.



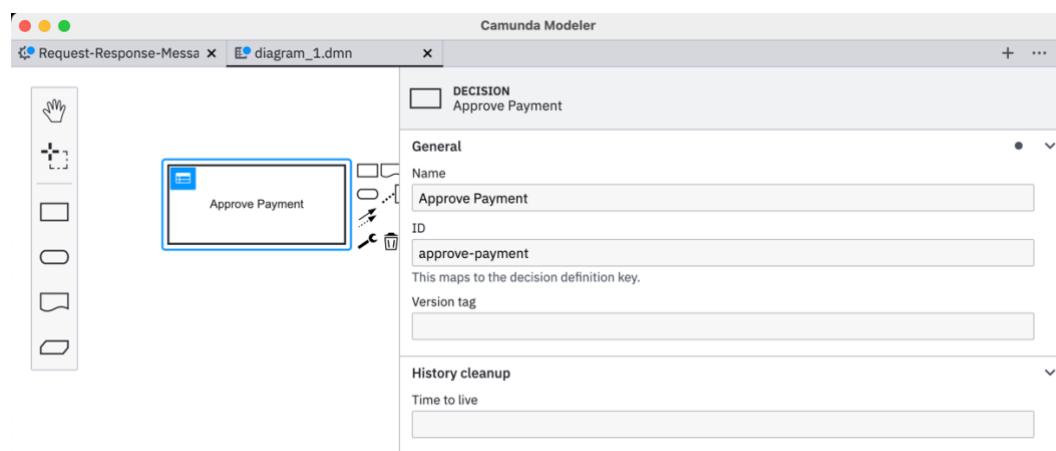
Save your changes and deploy the updated process using the Deploy Button in the Camunda Modeler.

G.2. Create a DMN table using the Camunda Modeler

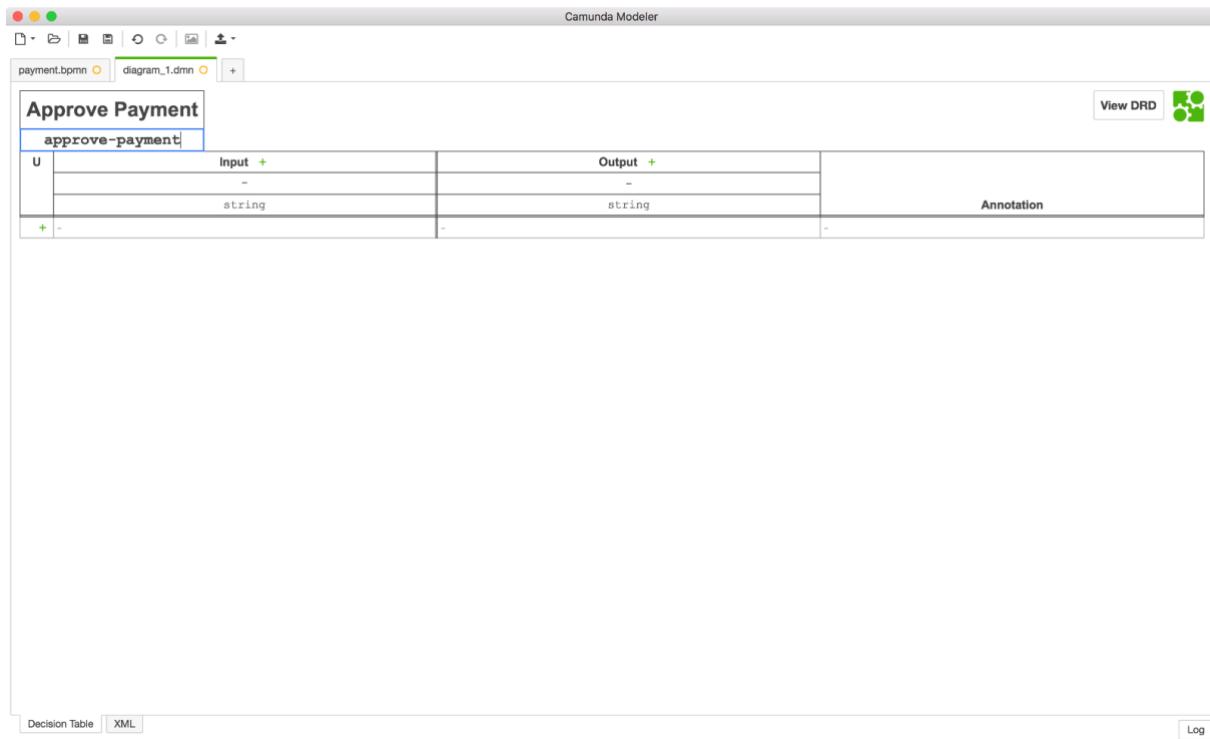
Create a new DMN table by clicking *File > New File > DMN Diagram (Camunda Platform 7)*.

G.3. Specify the DMN table

First, give the DMN table the name *Approve Payment* and the ID *approve-payment*, and then double click on it, as following.



The DMN table ID will match the *Decision Ref* in your BPMN process.



Next, specify the input expressions for the DMN table. In this example, we'll decide whether a payment is approved based on the item name. Your rules can also make use of the FEEL Expression Language, JUEL or Script. If you like, you can [read more about Expressions in the DMN Engine](#).

For the input column, use the following configuration :

The screenshot shows the Camunda Modeler interface with the "dmn_diagram.dmn" tab selected. The "Approve Payment" table is open, and the "Input" column is being configured. The "Expression" dropdown is set to "feel". The "Input Variable" dropdown is set to "Item". The "Type" dropdown is set to "string". The "Annotations" section is visible on the right.

Next, set up the output column, as following:

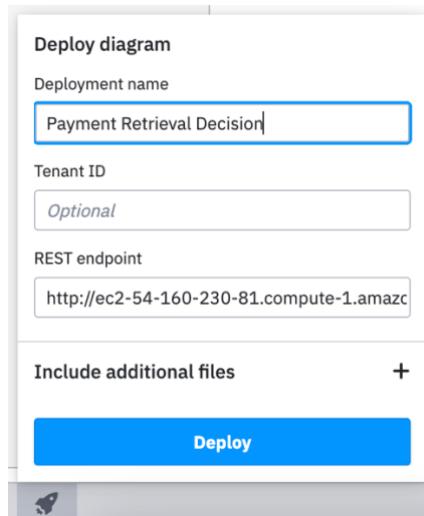
Approve Payment		Hit Policy:	Unique	
	When	Item	Output	Annotations
1	"item-xyz"	string	approved	
2	not("item-xyz")		-	
+	-			

Let's create some rules by clicking on the plus icon on the left side of the DMN table. After setup, your DMN table should look like this:

Approve Payment		Hit Policy:	Unique		
	When	Item	Then	Type	Annotations
1	"item-xyz"	string	true	boolean	
2	not("item-xyz")		false		
+	-				

G.4. Deploy the DMN table

To deploy the Decision Table, click on the Deploy button in the Camunda Modeler, give it Deployment Name “**Payment Retrieval Decision**”, then hit the Deploy button.



G.5. Verify the Deployment with Cockpit

Now, use Cockpit to see if the decision table was successfully deployed. Log in with the credentials *demo / demo*. Navigate to the “Decision definitions” section. Your decision table *Approve Payment* should be listed as deployed decision definition.

The screenshot shows the Camunda Cockpit interface with the 'Decisions' tab selected. The main content area displays the message '3 decision definitions deployed'. Below this, there is a table with three rows, each representing a deployed decision definition:

Name
Approve Payment
Assign Approver Group
Invoice Classification

G.6. Inspect using Cockpit and Tasklist

Next, use Tasklist to start two new Process Instances and verify that depending on your input, the Process Instance will be routed differently.

Click on the **Start process** button to start a process instance and choose the *Payment* process. Use the generic form to add the variables as follows:

The screenshot shows the Camunda Tasklist interface with the 'Start process' dialog open. The dialog has the following fields:

- Business Key:** An empty input field.
- Add a varia... +**: A button to add more variables.
- Name**: The name of the variable.
- Type**: The type of the variable.
- Value**: The value of the variable.

Name	Type	Value
item	String	item-xyz
amount	Long	1200

At the bottom of the dialog, there are 'Back', 'Close', and 'Start' buttons.

Hit the Start Instance button.

Next, click again on the **Start process** button to start another process instance and choose the *Payment* process. Use the generic form to add the variables as follows:

Camunda Tasklist

Create a filter

Keyboard Shortcuts Create task Start process Demo Demo

My Tasks (3)

My Group Tasks

Accounting

John's Task

Mary's Task

Peter's Task

All Tasks

Start process

Business Key

Add a variable +

Name	Type	Value
item	String	item-zzz
amount	Long	1200

[Back](#) [Close](#) [Start](#)

Assign Reviewer

You will see that depending on the input, the worker will either charge or not charge the credit card. You can also verify that the DMN tables were evaluated by using Camunda Cockpit. Log in with the credentials *demo / demo*. Navigate to the "Decisions" section and click on Approve Payment. Check the different Decision Instances that were evaluated by clicking on the ID in the table. A single DMN table that was executed could look like this in Camunda Cockpit:

Camunda Cockpit

Processes Decisions Human Tasks More Demo Demo

Dashboard » Decisions » Approve Payment » 5005d608-7ba9-11e9-8d9e-0242ac110002

Instance ID: 5005d608-7ba9-11e9-8d9e-0242ac110002

Definition Version: 1

Definition ID: approve-payment:1:3c3d99ac-7ba9-1...

Definition Key: approve-payment

Definition Name: Approve Payment

Tenant ID: null

Deployment ID: 3c3bc4e9-7ba9-11e9-8d9e-0242ac110002

Process Instance ID: 50049e7d-7ba9-11e9-8d9e-0242ac110002

Case Instance ID: null

Decision Requirements Definition: null

Approve Payment

approve-payment		Annotation	
U	Input		Output
	Item = item-xyz	Approved	
	string	boolean	
1	"item-xyz"	true = true	-
2	not("item-xyz")	false	-

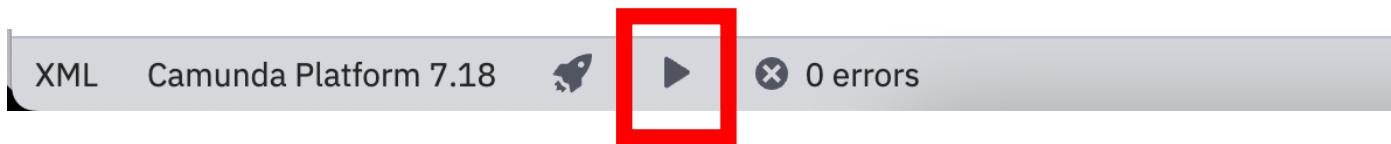
Inputs Outputs

Name	Type	Value
Item	String	item-xyz

G. Integrate with a REST API installed remotely using Camunda http-connector – REQUEST ONLY

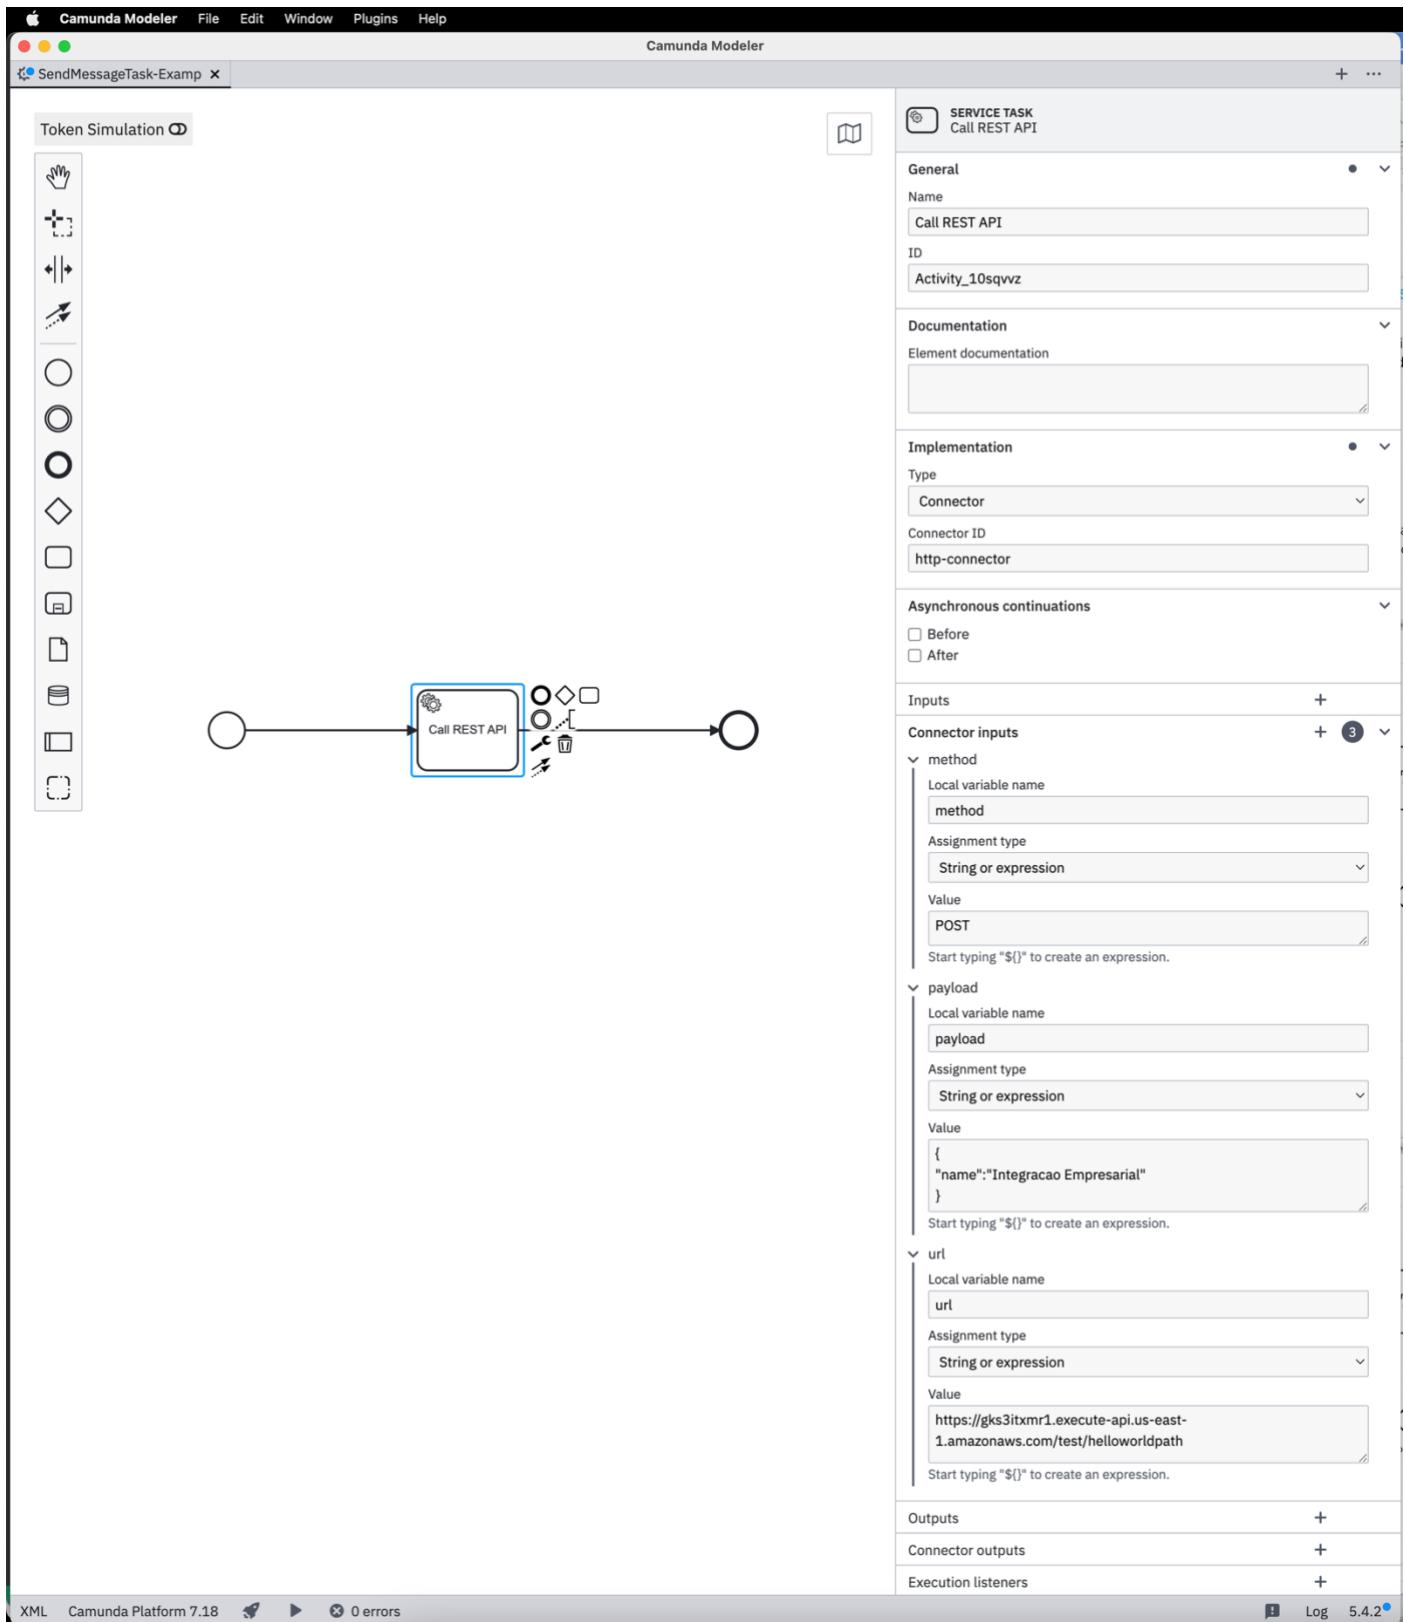
This example shows how to invoke a REST API from a **BPMN service task using the Camunda connector if indications for installation in section B were followed**. You need to create the service task element in the BPMN model as explained previously in this tutorial. Then, to invoke a REST API during the BPMN execution the following configuration in the Camunda modeller is required.

- H.1. Configure your task as a service task and choose the “**Connector**” in the implementation details, the connector ID should be “**http-connector**”.
- H.2. Then, add three “**Connector inputs**”:
 - a. “**method**”: of type String or expression with value = POST
 - b. “**payload**”: of type String or expression with value containing the JSON to be send with the request
 - c. “**url**”: of type String or expression with value = endpoint previously available, e.g., Kong API or AWS Gateway API, or other
- H.3. Deploy the BPMN model
- H.4. Start a process instance, using the **Start instance** Button in the Camunda Modeler to deploy the updated process to Camunda.



- H.5. Verify in the REST API endpoint that it was invoked. For instance, in AWS CloudWatch.

The previous configuration is also depicted in the following figure:



H. Integrate with a REST API installed remotely using Camunda http-connector – REQUEST with variable

This example shows how to invoke a REST API from a **BPMN service task using the Camunda connector if indications for installation in section B were followed** and that includes some data previously processed in the BPMN process flow into the REST API Request. You need to create the service task element in the BPMN model as explain previously in this tutorial. Then, to invoke a REST API during the BPMN execution the following configuration in the Camunda modeller is required.

If you need to add any process variable to the web service call you may add it, using the payload, as follows (*CourseName* is used to only for demonstration purposes).

Firstly, a user task is created to allow the input of the course name as text description from the end user. For that, you need to add a Task and choose the “**User Task**” option. Then, assign a user to the task. In this example is used the “demo” user. After that, choose a form of type “Generated Task Forms”, and create a Form field:

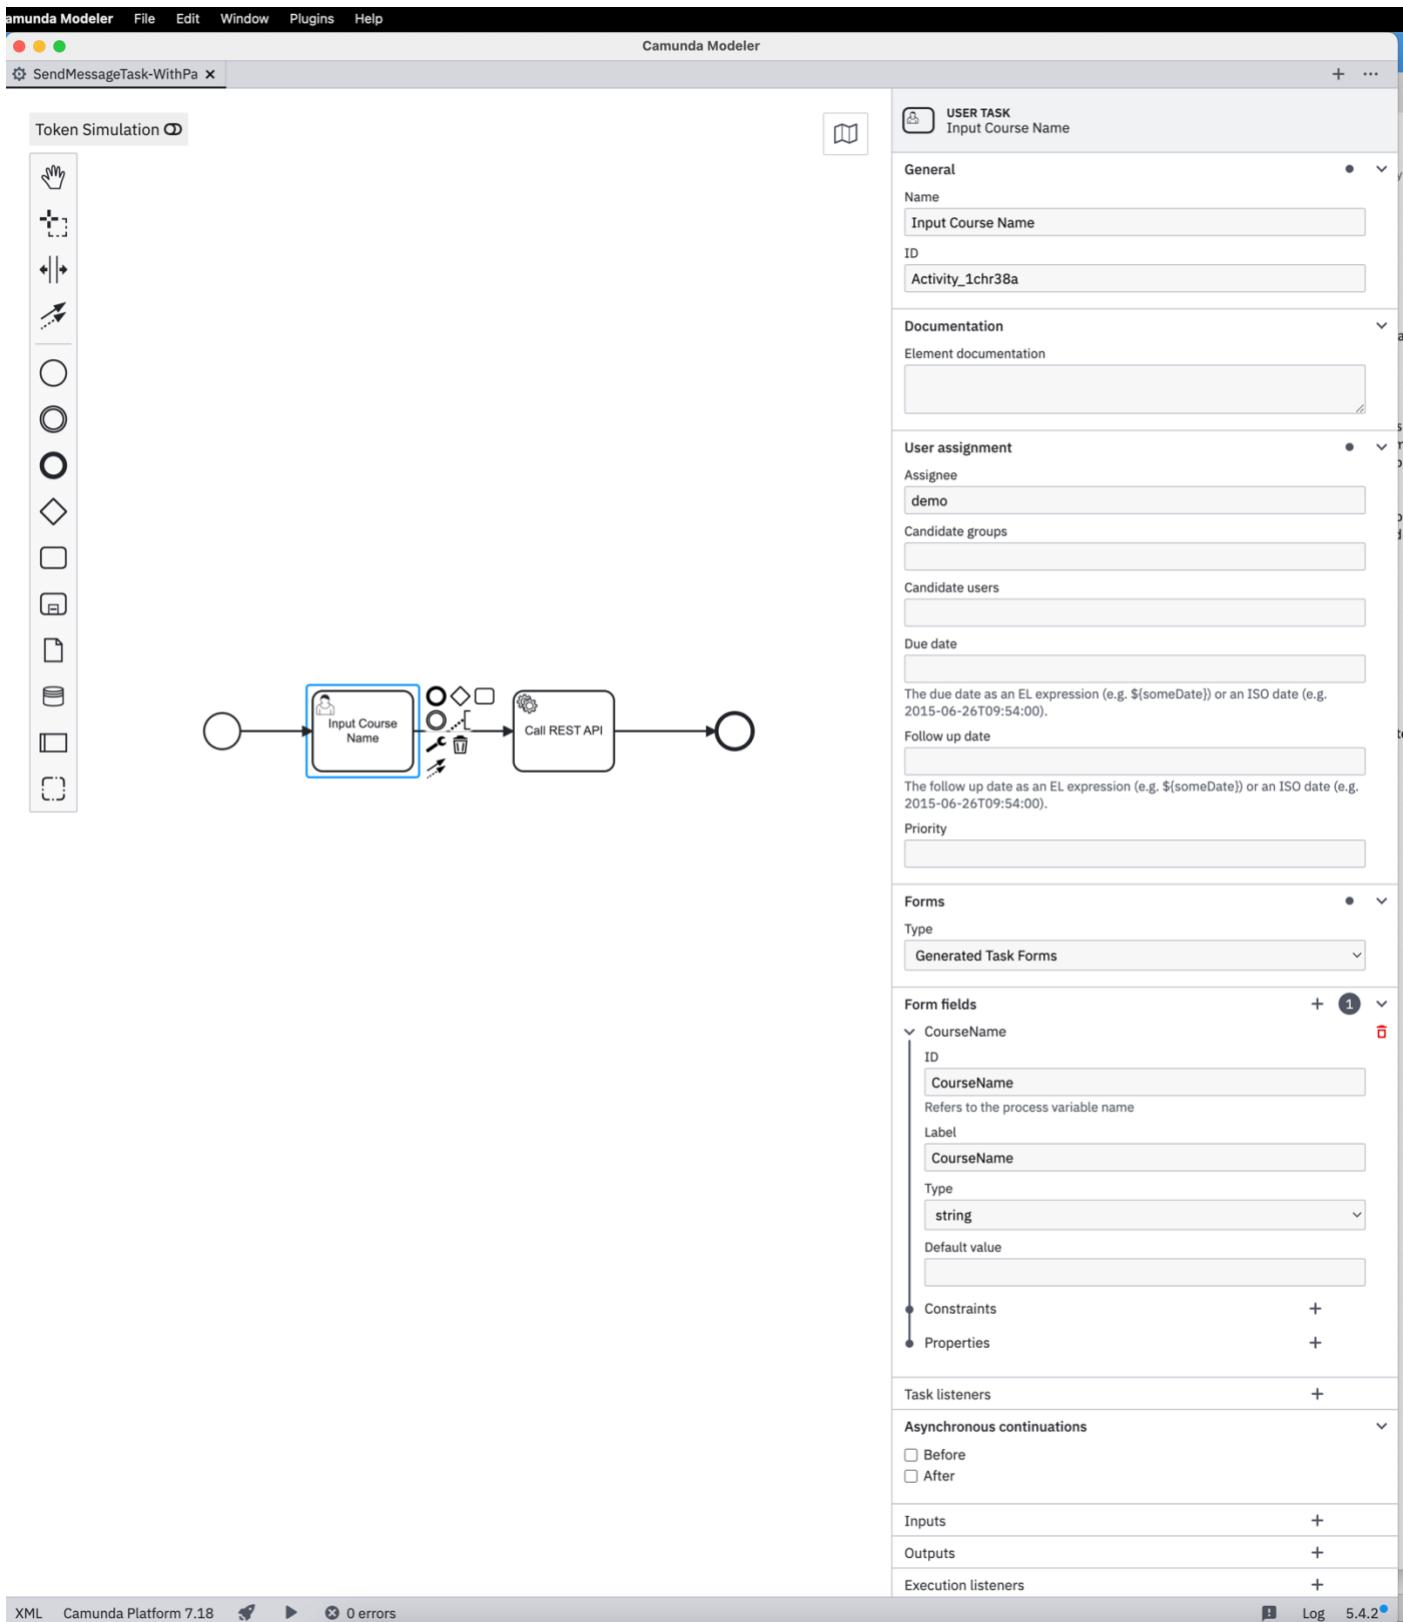
- ID = CourseName
- Label = CourseName
- Type = string

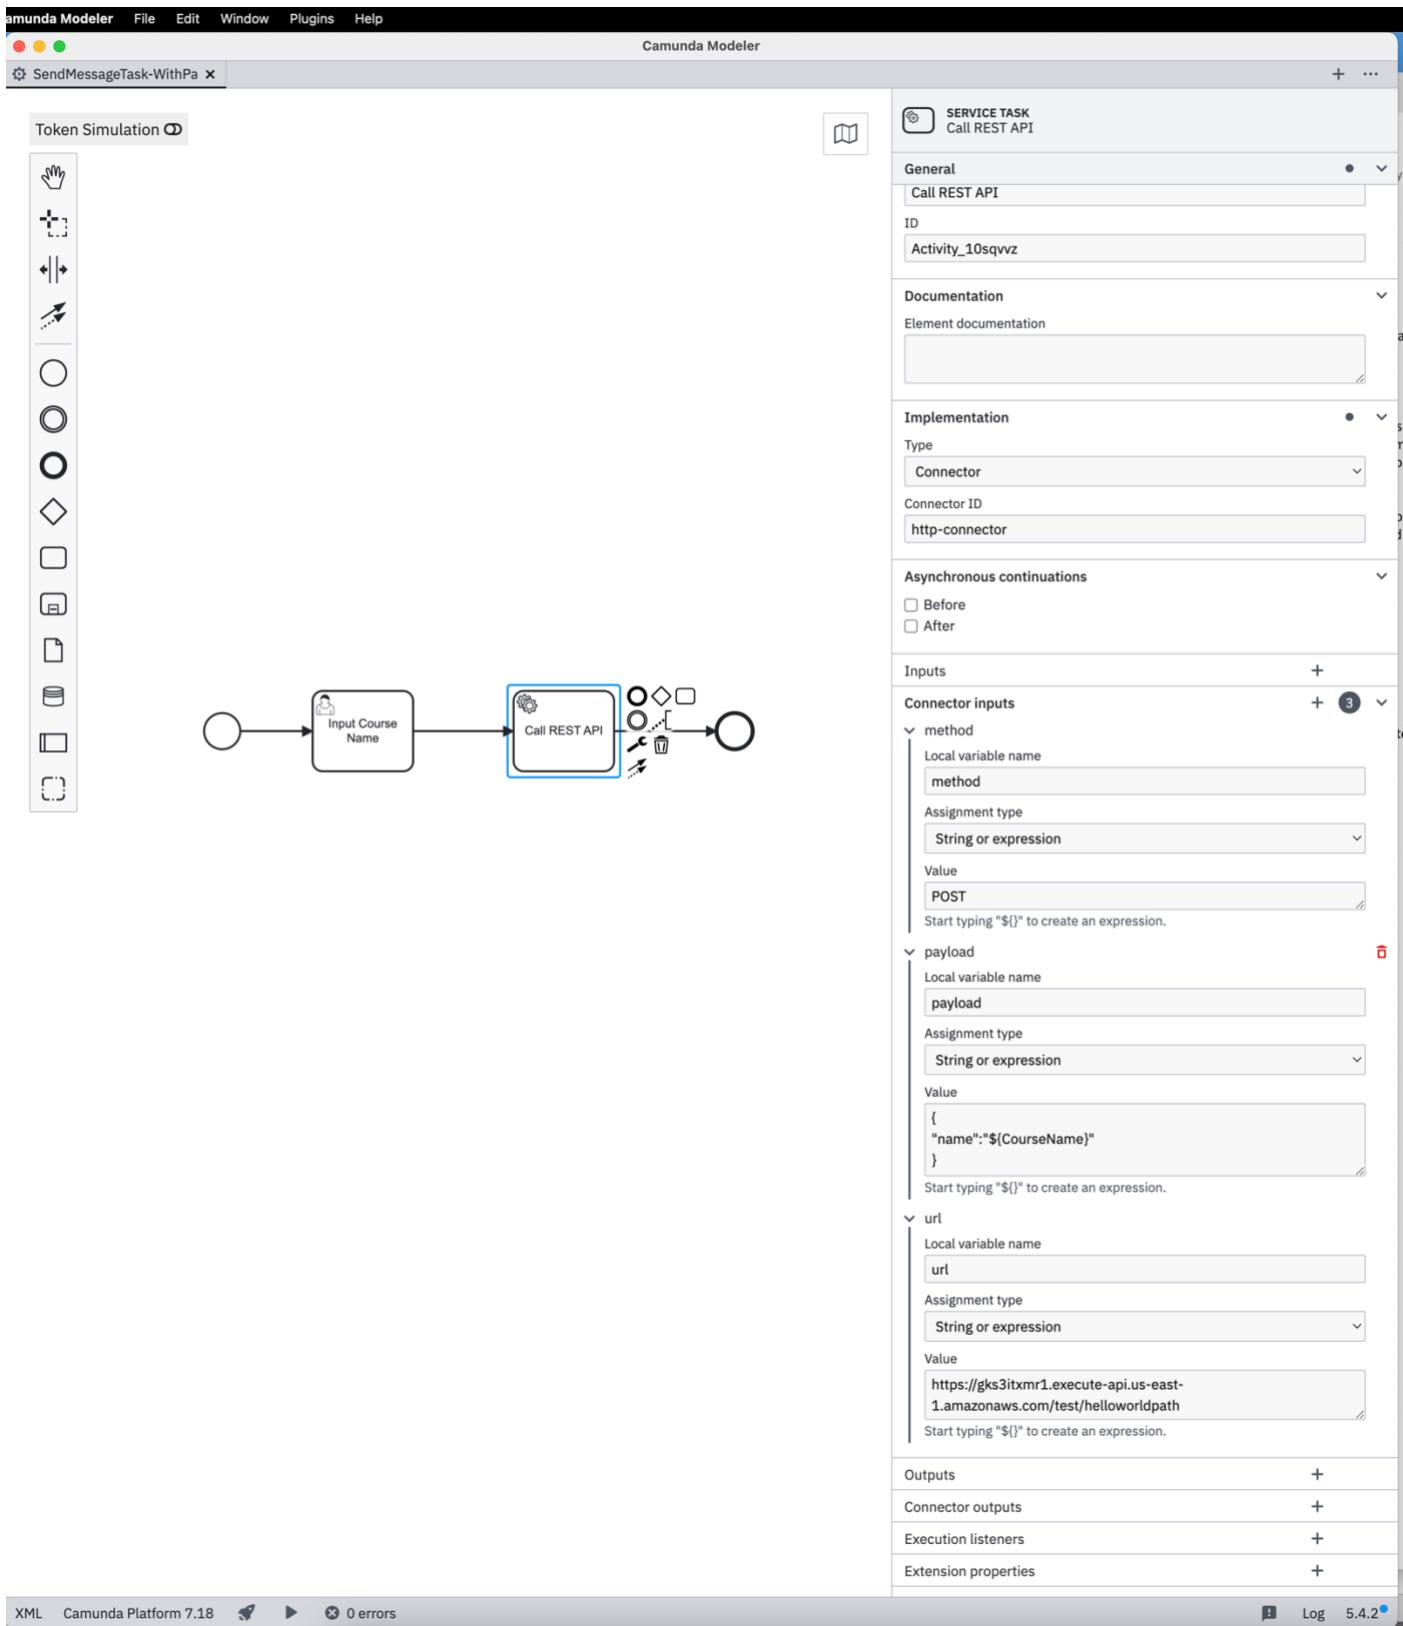
With this configuration you are now accepting the coursename from the user and allocate a variable that could be used afterwards.

Secondly, choose the “Call REST API” task and change the payload to be dynamically created on the previous **CourseName**, using the following syntax:

```
{  
  "name": "${CourseName}"  
}
```

The previous configuration is also depicted in the following figures:

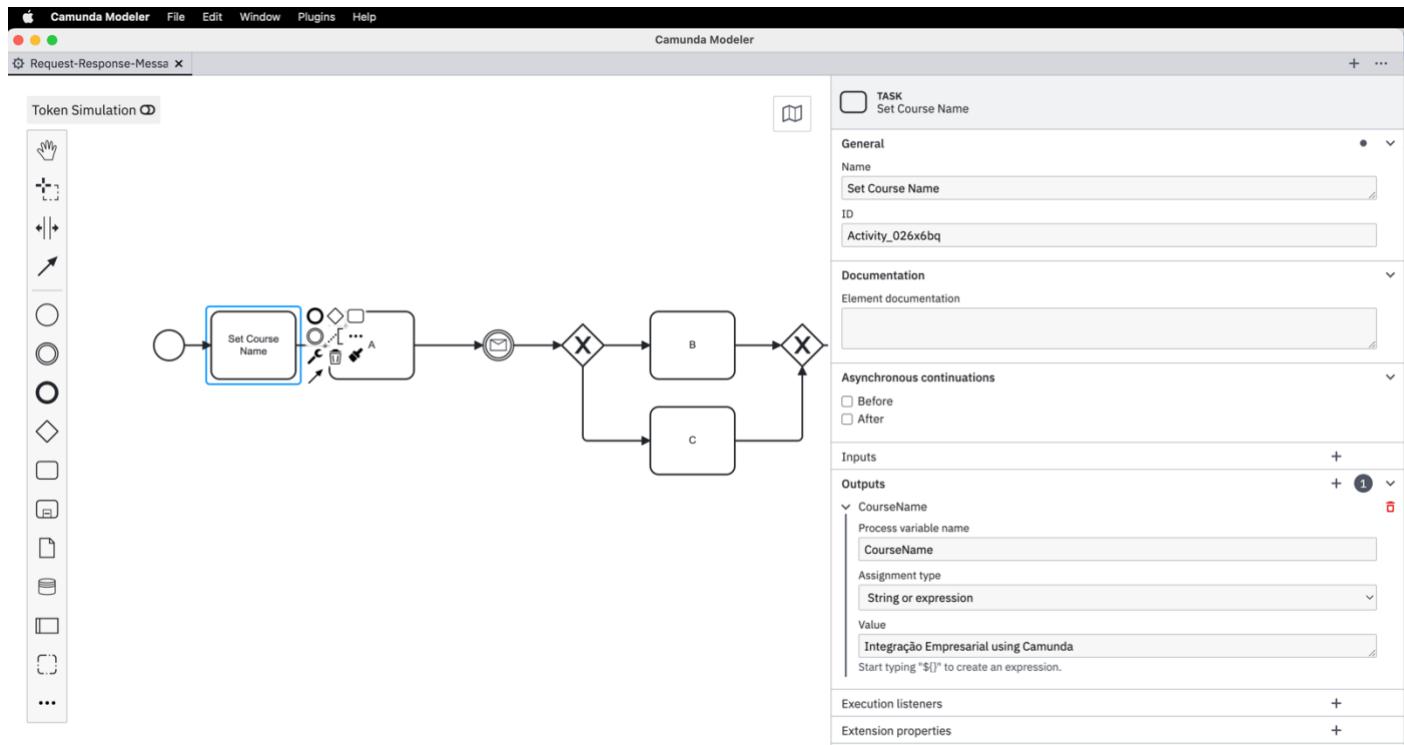




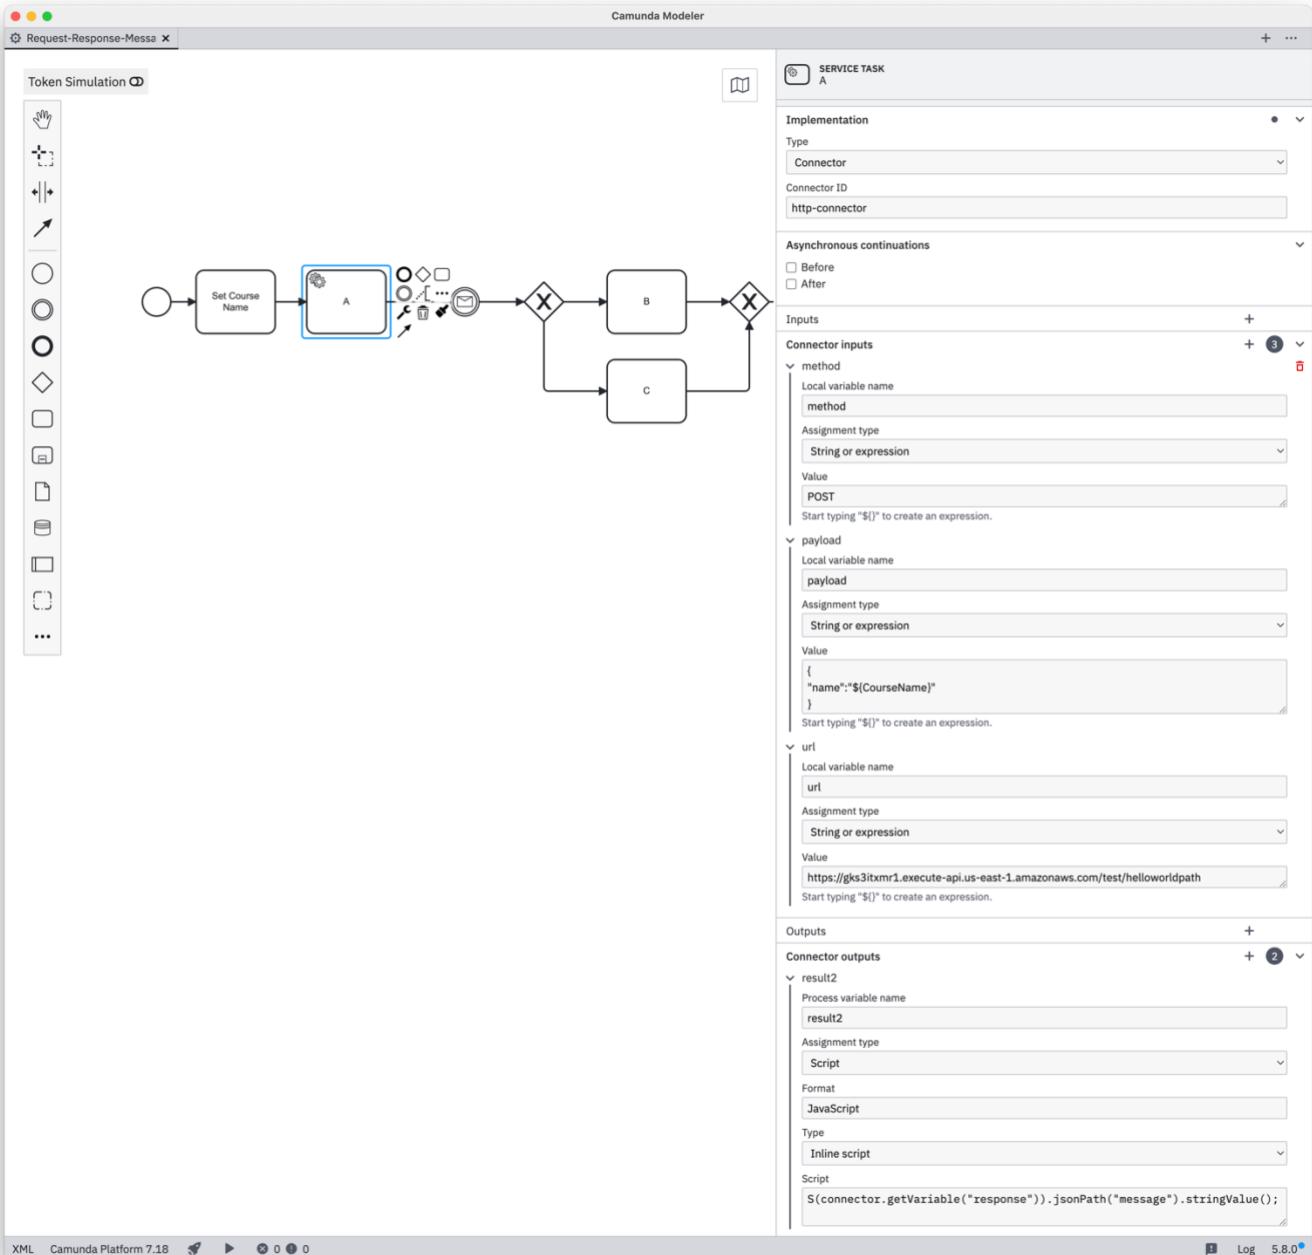
I. Integrate with a REST API installed remotely using Camunda http-connector – REQUEST & RESPONSE

For configuring the parsing of a REST API response, you may use inline *javascript* for parsing the received message, as depicted in the next figures. Many other JSON queries could be programmed, check reference manual at <https://docs.camunda.org/manual/7.4/reference/spin/json/03-querying-json/>.

Firstly, in this example, instead of receiving the course name from the user, we are adding the process variable “**CourseName**” as the output of first task, as depicted in the next figure.



Secondly, the process variable “**CourseName**” is used in the payload of the connector inputs. Then, in the connector output you decode the fields that you like to receive from the response.



Notice that the inline script used in the connector output is reading “**message**” field from response and writing that to the “**result**” process variable:

```
S(connector.getVariable("response")).jsonPath("message").stringValue();
```

If you try to invoke your REST API using the curl command, you see the matching between the response and the name of the field “**message**”.

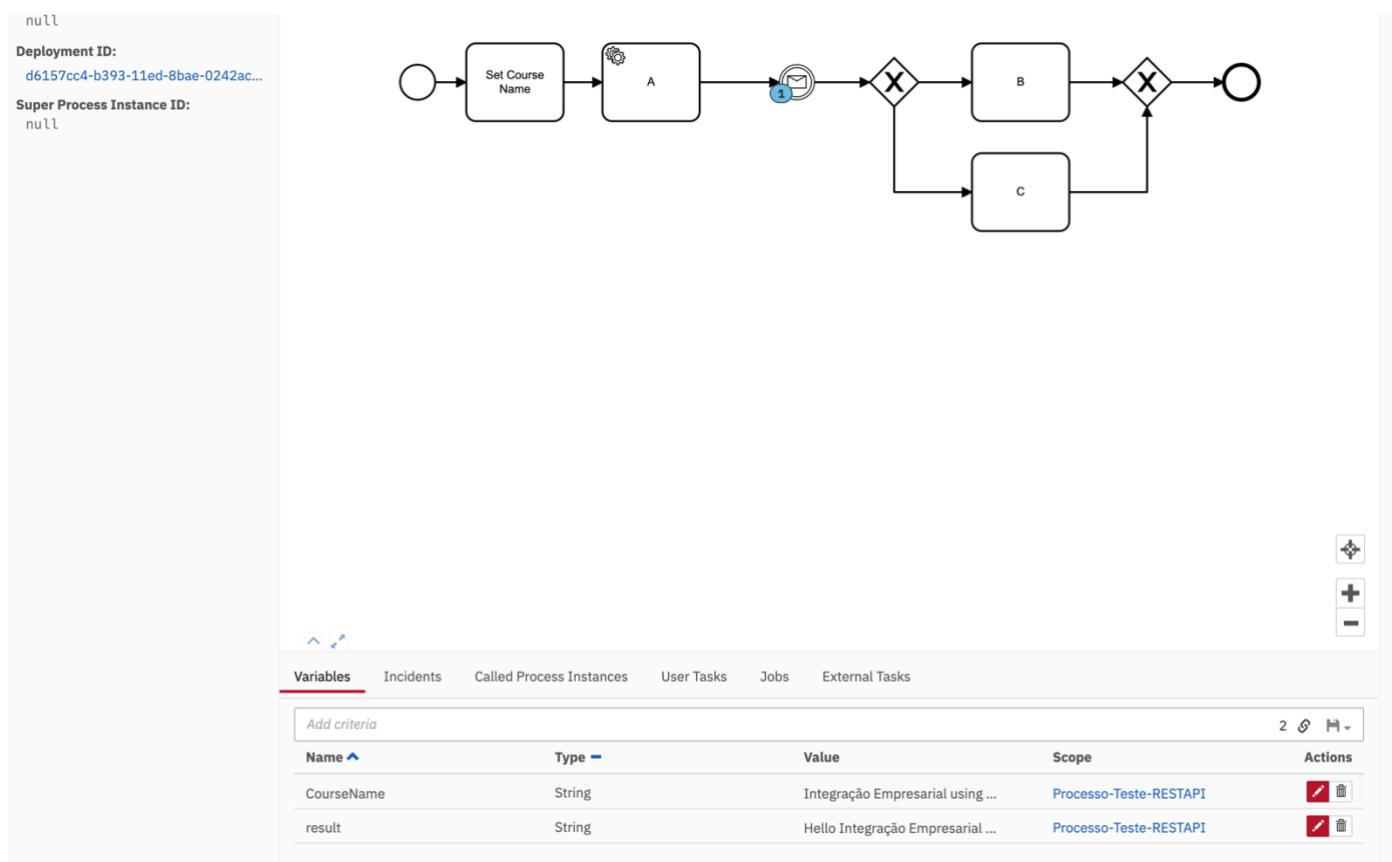
```

Lambda-api-gateway -- zsh -- 127x17
sergioguerreiro@MBP-de-Sergio Lambda-api-gateway % curl -i -H "Content-Type: Application/json" --data "@body.json" -X POST http://gks3itxr1.execute-api.us-east-1.amazonaws.com/test/helloworldpath
HTTP/2 200
content-type: application/json
content-length: 43
date: Thu, 23 Feb 2023 16:01:40 GMT
x-amzn-requestid: b5cae982-52ec-4e19-ae6c-9ea9523ec98e
x-amz-apigw-id: AzMbtHAYIAMTFQ=
x-custom-header: my custom header value
x-amzn-trace-id: Root=1-63f78de4-5650c464514343ac3e382adc;Sampled=0
x-cache: Miss from cloudfront
via: 1.1 307b5e33f74f1fe7c0f94fe6d2fd888.cloudfront.net (CloudFront)
x-amz-cf-pop: LISS0-C1
x-amz-cf-id: dVYjryOSD4-JWzyG6G9qp6sr3CfcapHTmUUtQRdzj8mRN2Q1alVL0Q==

{"message":"Hello Integração Empresarial!"}
sergioguerreiro@MBP-de-Sergio Lambda-api-gateway %

```

Moreover, in the Camunda cockpit, choose the process instance (an intermediate catch message event has been deliberately added to the process to hold it!), and verify that your instance now has the “**result**” process variable.



If you have many variables to retrieve from REST API response you only need to add more connector outputs with different process variable names, as depicted in the following figure.

Connector outputs

- result2**
 - Process variable name: `result2`
 - Assignment type: Script
 - Format: JavaScript
 - Type: Inline script
 - Script: `S(connector.getVariable("response")).jsonPath("message").stringValue();`
- result**
 - Process variable name: `result`
 - Assignment type: Script
 - Format: JavaScript
 - Type: Inline script
 - Script: `S(connector.getVariable("response")).jsonPath("message").stringValue();`

After that, you can check, in Camunda cockpit, that both variables are available.

Variables

Name	Type	Value	Scope	Actions
CourseName	String	Integração Empresarial using ...	Processo-Teste-RESTAPI	
result	String	Hello Integração Empresarial ...	Processo-Teste-RESTAPI	
result2	String	Hello Integração Empresarial ...	Processo-Teste-RESTAPI	

--*--

Hint: If you would like to debug a json message that has been received by a service task, you can use the following source code in the connector output:

The BPMN diagram illustrates a process flow between two systems: Customer/Supplier (TS) and Fruit Supplier - Executer (TS'). The process starts with a customer placing an order. The system retrieves the fruit catalog and checks if the product is available. If yes, it accepts the order; if no, it rejects it. The supplier then updates its fruit stock and sends a confirmation back to the customer.

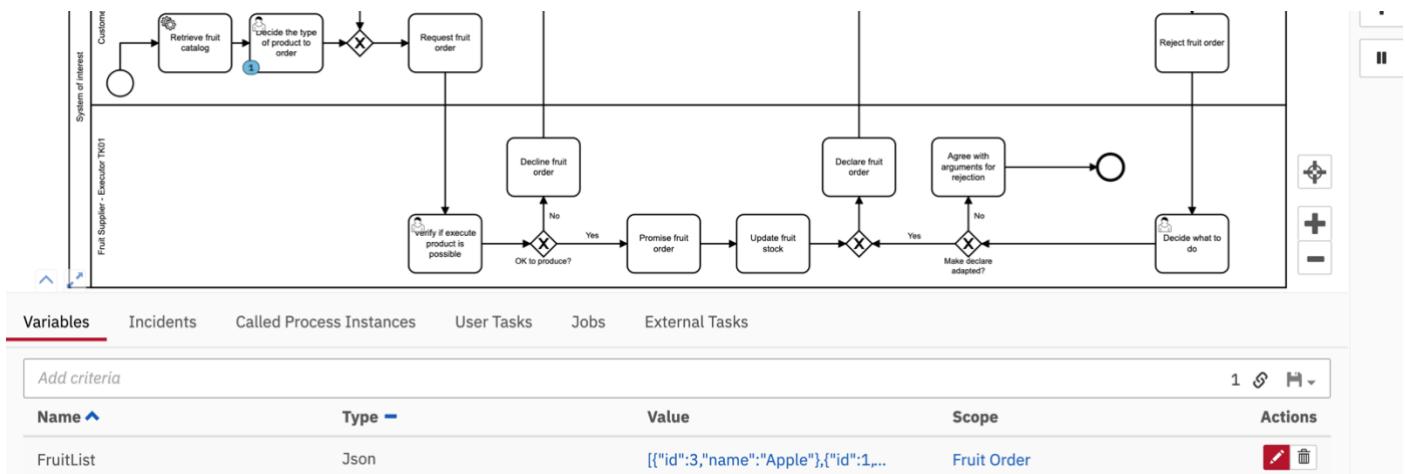
Activity_1y7hoeb

- Documentation
- Implementation
- Asynchronous continuations
- Inputs
- Connector inputs
- Outputs
- Connector outputs

FruitList

- Process variable name: `FruitList`
- Assignment type: Script
- Format: JavaScript
- Type: Inline script
- Script: `S(connector.getVariable("response"));`

And then, start an instance and check in the Camunda cockpit the message has been received correctly.

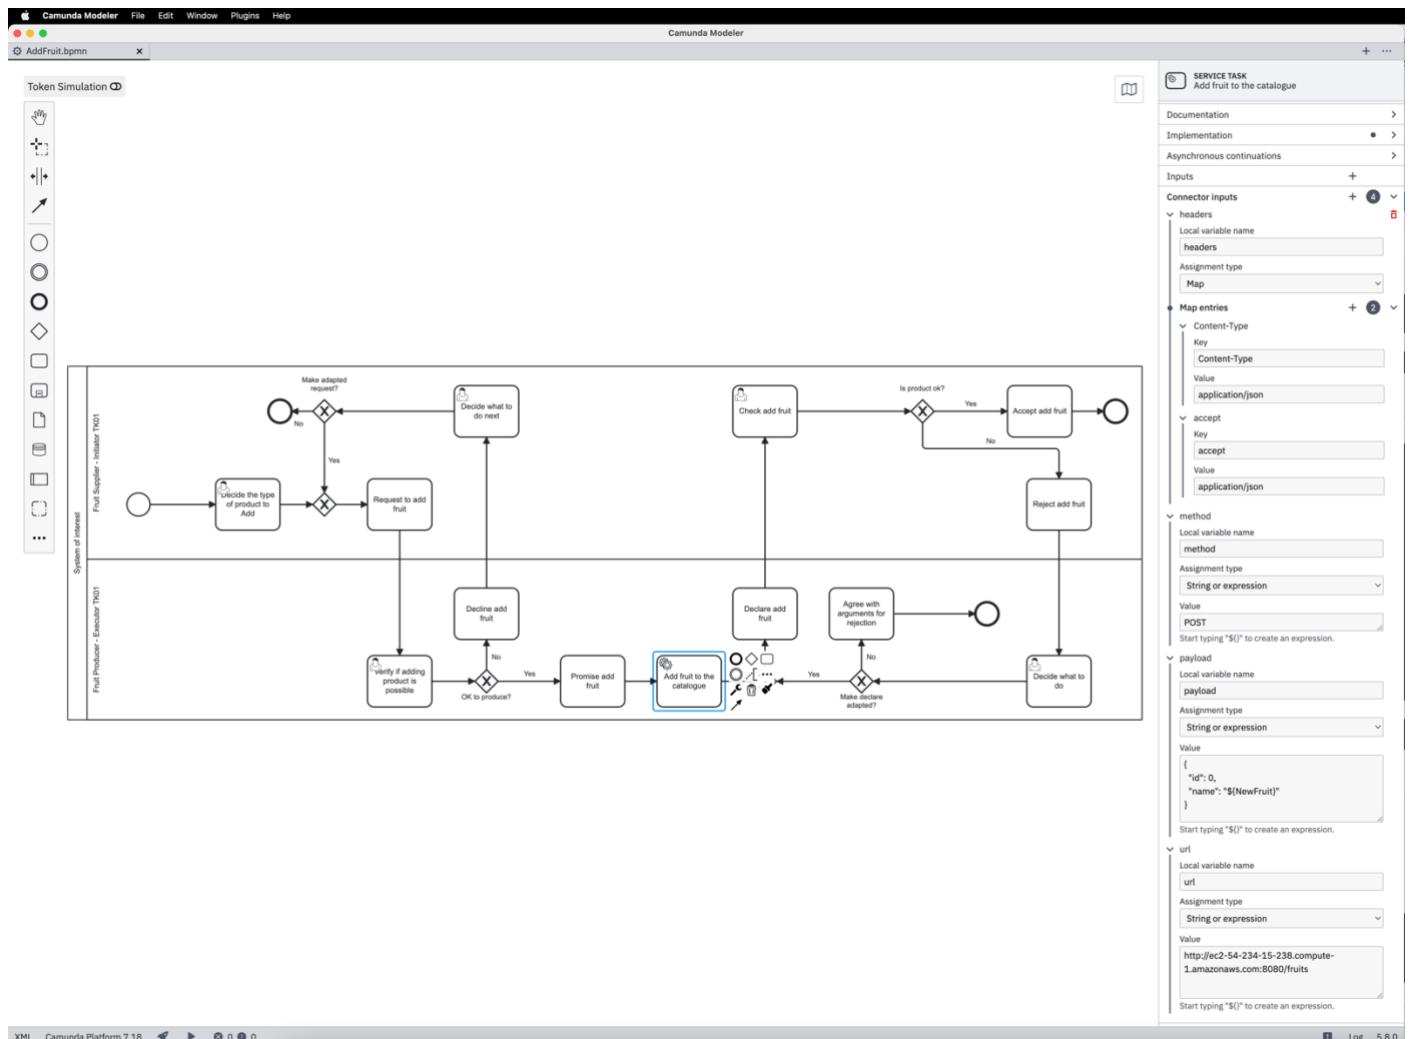


J. Integrate with a REST API installed remotely using Camunda http-connector – REQUEST & RESPONSE & Headers fields

When needed more information can be added to an API REST. For instance, if you need to specify the messages media, Headers can be included in the connector inputs. Considering the following POST request:

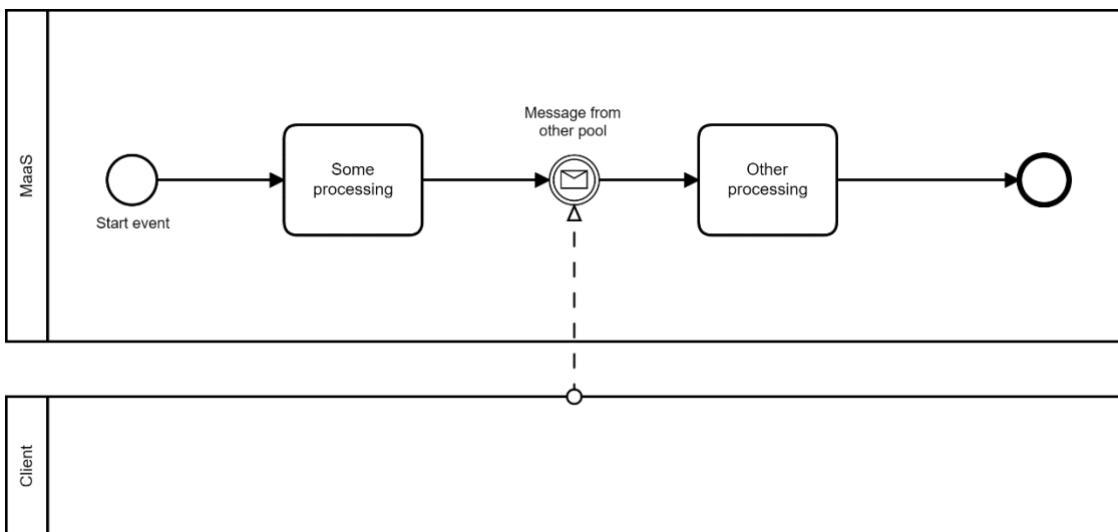
```
curl -X 'POST' \
  'http://ec2-54-234-15-238.compute-1.amazonaws.com:8080/fruits' \
  -H 'accept: application/json' \
  -H 'Content-Type: application/json' \
  -d '{
    "id": 0,
    "name": "sd"
  }'
```

The same can be configured in a Camunda service task accordingly:

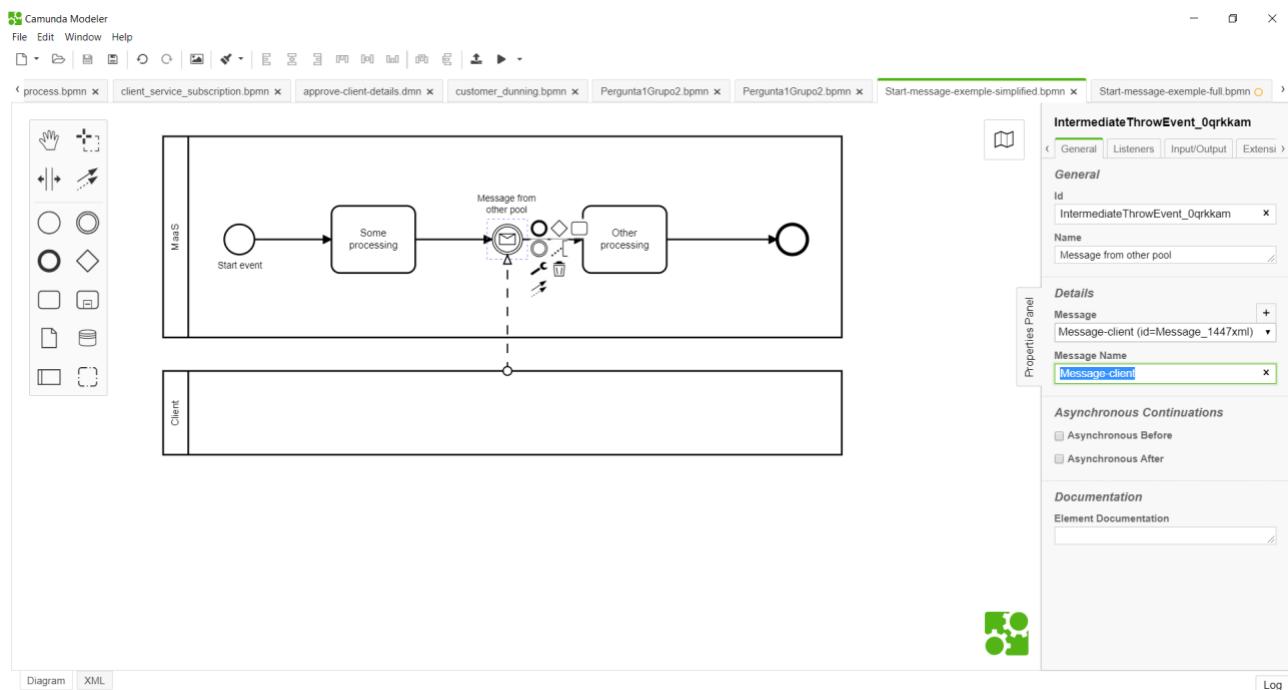


K. Using an intermediate message catch event inside a process.

L.1. Design the following collaboration in Camunda Modeller:



Where the intermediate message catch event should have a configuration similar with the following:



L.2. Deploy the model in your Camunda engine.

L.3. Start a new process instance with new configuration, where the **businessKey** is filled with any unique identifier.

Start Process Instance - Step 2 of 2

Enter details to start a process instance on Camunda Platform. Alternatively, you can start a process instance [via a Rest Client](#).

Business Key
dgfffc123455

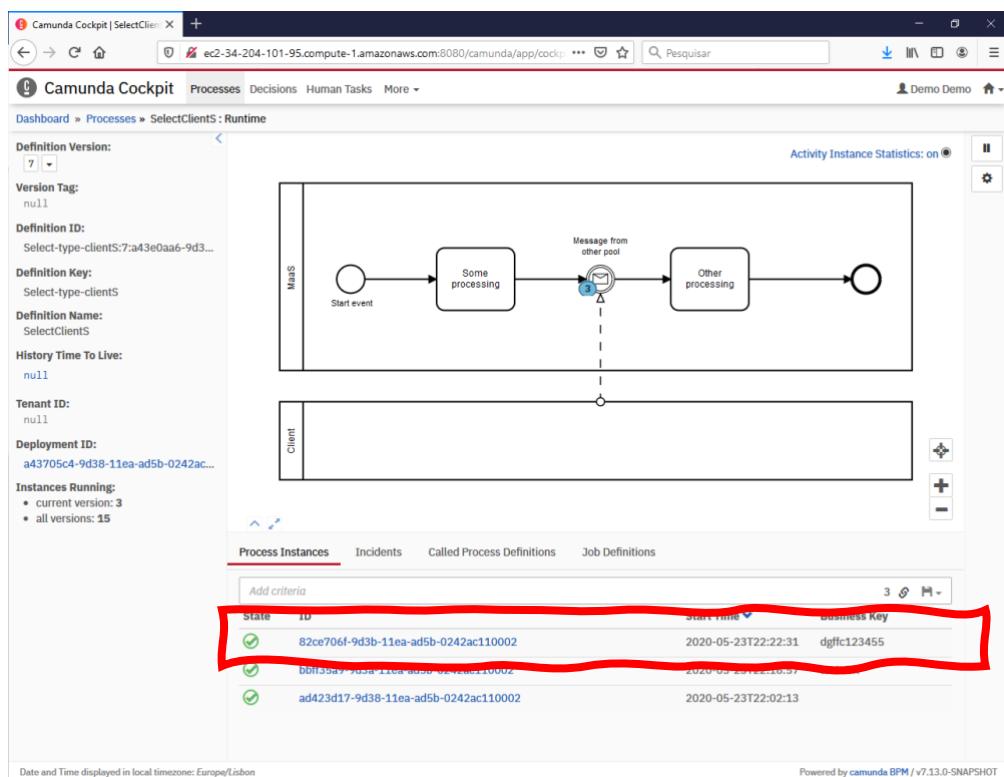
Variables (optional)
A JSON string representing the variables the process instance is started with.

Must be a proper [JSON string](#) representing process instance variables.

Start

▶ × 0 ⏪ 0 ⏩ 0

You can verify the creation of the instance in the camunda cockpit. You will obtain something similar with the following business key as defined:



L.4. Then, the instance is expecting a message. To test the catch of a REST message, use a similar curl command:

```
curl -H "Content-Type: application/json" --data "@body.json" -X POST http://<YOUR CAMUNDA EC2 DNS>:8080/engine-rest/message
```

where bodyStart.json defines the message listener and the correlation with the key for your instance and sending some process variables:

```
{
  "messageName" : "Message-client",
  "businessKey" : "dgfffc123455",
  "processVariables": {"amount": {"value":240,"type":"long"}},
  "resultEnabled" : true
}
```

References

Full official documentation:

- <https://docs.camunda.org/manual/>
- <https://docs.camunda.org/javadoc/camunda-bpm-platform/7.12/>

Testing the basics of the Camunda:

- <https://docs.camunda.org/get-started/quick-start/>
- Videos tutorials: <https://camunda.com/learn/videos/>

URLs with other resources

- <https://camunda.com/products/>
- <https://github.com/camunda/camunda-bpm-examples>
- Send Task - <https://docs.camunda.org/manual/7.7/reference/bpmn20/tasks/send-task/>
- Receive task - <https://docs.camunda.org/manual/7.7/reference/bpmn20/tasks/receive-task/>
- <https://github.com/camunda/camunda-bpm-examples/tree/master/servicetask/rest-service>
- <https://github.com/rob2universe/camunda-http-connector-example>
- <https://docs.camunda.org/manual/7.3/guides/user-guide/#process-engine-connectors>
- <https://docs.camunda.org/manual/7.4/reference/spin/json/03-querying-json/>



The goal of this document is to present the details about installing, coding, deploying and testing a lambda service in Amazon AWS Lambda service.

With AWS Lambda, you run functions to process events. You can send events to your function by invoking it with the Lambda API, or by configuring an AWS service or resource to

Function: is a resource that you can invoke to run your code in AWS Lambda. A function has code that processes events, and a runtime that passes requests and responses between Lambda and the function code. You provide the code, and you can use the provided runtimes or create your own.

Runtime: Lambda runtimes allow functions in different languages to run in the same base execution environment. You configure your function to use a runtime that matches your programming language. The runtime sits in between the Lambda service and your function code, relaying invocation events, context information, and responses between the two. You can use runtimes provided by Lambda or build your own.

Event: is a JSON formatted document that contains data for a function to process. The Lambda runtime converts the event to an object and passes it to your function code. When you invoke a function, you determine the structure and contents of the event. When an AWS service invokes your function, the service defines the event.

Concurrency: is the number of requests that your function is serving at any given time. When your function is invoked, Lambda provisions an instance of it to process the event. When the function code finishes running, it can handle another request. If the function is invoked again while a request is still being processed, another instance is provisioned, increasing the function's concurrency.

Trigger: is a resource or configuration that invokes a Lambda function. This includes AWS services that can be configured to invoke a function, applications that you develop, and event source mappings. An event source mapping is a resource in Lambda that reads items from a stream or queue and invokes a function.

The following contents is presented in this document.

Contents

A. Creating a Lambda function using VSCode IDE (Java)	2
B. Testing a Lambda function internally.....	5
C. Invoking a Lambda function remotely	7
D. Delete a Lambda function	10
E. Checking consumption of a Lambda function	11
Other references.....	12

A. Creating a Lambda function using VSCode IDE (Java)

(Partially based on AWS Lambda Developer Guide, 2025 Amazon Web Services, Inc., public available at <https://docs.aws.amazon.com/lambda/>)

In this step, you start VSCode and create a Maven project. You will add the necessary dependencies and build the project. The build will produce a .jar, which is your deployment package.

A.0. Add a new Java Project...

A.1. Select project Type: Maven

A.2. Select No Archetype

A.3. Write examples as group Id

A.4. Write lambda-java-example as artifact Id

A.5. Select your destination folder

A.6. Add Java class to the project.

a. Open the context (right-click) menu for the src/main/java subdirectory in the project, choose **New**, and then choose **New File**...add the name **Hello.java**

```
package example;

import com.amazonaws.services.lambda.runtime.Context;
import com.amazonaws.services.lambda.runtime.LambdaLogger;
import com.amazonaws.services.lambda.runtime.RequestHandler;
import com.amazonaws.services.lambda.runtime.RequestStreamHandler;

import java.util.ArrayList;
import java.util.LinkedHashMap;
import java.util.Map;

import org.json.simple.JSONArray;
import org.json.simple.JSONObject;
import org.json.simple.JSONValue;
import org.json.simple.parser.JSONParser;
import org.json.simple.parser.ParseException;

import java.io.*;

public class Hello implements RequestStreamHandler{
    public void handleRequest(
        InputStream inputStream,
        OutputStream outputStream,
        Context context)
    {
        LambdaLogger logger = context.getLogger();
        try
        {
            JSONParser parser = new JSONParser();
            BufferedReader reader = new BufferedReader(new InputStreamReader(inputStream));
            JSONObject responseJson = new JSONObject();
            JSONObject event = (JSONObject) parser.parse(reader);
            String name = new String();

            logger.log("input:" + (String) event.toString()+"\n");

            if (event.get("body") != null)
            {
                JSONObject bodyjson = (JSONObject) parser.parse((String) event.get("body"));
                if (bodyjson.get("name") != null)
                    name = (String) bodyjson.get("name");
            }

            JSONObject responseBody = new JSONObject();
            if (name != null)
                responseBody.put("message", "Hello " + name + "!");
            else
                responseBody.put("message", "No name defined. Add {\"name\":\"name\"} in the request.");

            JSONObject headerJson = new JSONObject();
            headerJson.put("x-custom-header", "my custom header value");
            responseJson.put("statusCode", 200);
            responseJson.put("headers", headerJson);
            responseJson.put("body", responseBody.toString());
        }
    }
}
```

```

        OutputStreamWriter writer = new OutputStreamWriter(outputStream, "UTF-8");
        writer.write(responseJson.toString());
        writer.close();
    }
    catch(Exception exc)
    {
        logger.log("Error : " + exc.toString());
    }
}
}

```

A.7. Verify that your **pom.xml** file contains the following configuration:

```

<?xml version="1.0" encoding="UTF-8"?>
<project xmlns="http://maven.apache.org/POM/4.0.0"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://maven.apache.org/POM/4.0.0 http://maven.apache.org/xsd/maven-4.0.0.xsd">
<modelVersion>4.0.0</modelVersion>

<groupId>examples</groupId>
<artifactId>lambdas-java-example</artifactId>
<version>0.0.1-SNAPSHOT</version>
<name>lambdas-java-example</name>

<properties>
    <maven.compiler.source>11</maven.compiler.source>
    <maven.compiler.target>11</maven.compiler.target>
</properties>

<dependencies>
    <dependency>
        <groupId>com.amazonaws</groupId>
        <artifactId>aws-lambda-java-core</artifactId>
        <version>1.2.2</version>
    </dependency>
    <dependency>
        <groupId>com.googlecode.json-simple</groupId>
        <artifactId>json-simple</artifactId>
        <version>1.1.1</version>
    </dependency>
</dependencies>
<build>
    <plugins>
        <plugin>
            <groupId>org.apache.maven.plugins</groupId>
            <artifactId>maven-shade-plugin</artifactId>
            <version>3.4.1</version>
        </plugin>
    </plugins>
</build>
</project>

```

A.8. Configure **JDK11** to be available on your project root command line with Apache Maven 3.9.9.

A.9. Execute the packaging procedure with

```
mvn package shade:shade
```

A.10. You will find the following file in the project target folder:

```
lambda-java-example-0.0.1-SNAPSHOT.jar.
```

- A.11. **Login** to AWS Educate, then choose the classroom, and press AWS console. Expand “All services” option.
A.12. Search for the AWS Lambda Service (under Compute option)

- A.13. Press Create function
- A.14. Choose “Author from scratch”
- A.15. Define Function name as: PrimeiroMicroServico
- A.16. Choose JAVA_11 for Runtime
- A.17. Choose x86_64 Architecture
- A.18. Change Default execution role > choose Use an existing role> choose LabRole
- A.19. Press Create function
- A.20. Then, in the next interface, press the Upload button:

Function package

Upload

For files larger than 10 MB, consider uploading using Amazon S3.

and choose the previously created jar file:

`lambda-java-example-0.0.1-SNAPSHOT.jar.`

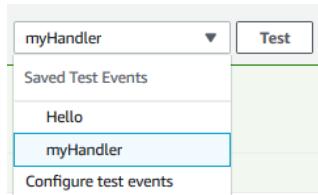
- A.21. Check the Handler field with the name of the handler defined within your java class. You will obtain something similar with this:

Runtime settings Info				
Runtime Java 8 on Amazon Linux 2	Handler Info example.Hello::handleRequest	Architecture Info x86_64		
Runtime management configuration <table border="1"> <tr> <td>Runtime version ARN Info <code>arn:aws:lambda:us-east-1:runtime:dc76617de5625b2968de9b2d31e98df</code></td> <td>Update runtime version Info Auto</td> </tr> </table>			Runtime version ARN Info <code>arn:aws:lambda:us-east-1:runtime:dc76617de5625b2968de9b2d31e98df</code>	Update runtime version Info Auto
Runtime version ARN Info <code>arn:aws:lambda:us-east-1:runtime:dc76617de5625b2968de9b2d31e98df</code>	Update runtime version Info Auto			

- A.22. Your lambda service is deployed!

B. Testing a Lambda function internally

C.1. Choose option "Test" and select create new event



C.2. Define a name for "Event name" field

C.3. Define the JSON message accordingly with the test example, with the example of section A, define the message as follows:

A detailed screenshot of the AWS Lambda function configuration interface. At the top, it shows a function overview with a thumbnail of the function code ('t3'), 0 layers, and a 'Description' field. Below this, there are tabs for 'Code', 'Test' (which is selected and highlighted in blue), 'Monitor', 'Configuration', 'Aliases', and 'Versions'. In the 'Test' tab, there's a 'Test event' section with a 'Save' and 'Test' button. It prompts to invoke the function without saving an event and to configure a JSON event. Below this, there are sections for 'Test event action' (with 'Create new event' selected), 'Event name' (set to 'testeinitial'), 'Event sharing settings' (with 'Private' selected), and 'Template - optional' (set to 'hello-world'). At the bottom, there's an 'Event JSON' section containing the following JSON code:

```
1 - {  
2   "name": "Integração Empresarial"  
3 }
```

A 'Format JSON' button is located to the right of the JSON code.

C.4. Press Test button

C.5. You should receive a response similar with:

Execution result: succeeded ([logs](#)) X

▼ Details
The area below shows the last 4 KB of the execution log.

```
{  
  "headers": {  
    "x-custom-header": "my custom header value"  
  },  
  "body": "{\"message\":\"Hello !\"}",  
  "statusCode": 200  
}
```

Summary

Code SHA-256	Request ID
QRVeqKESj4gQ5hUN/cI0ITI4A2hWvuHQ0Mgy9kQzqw4=	5937a1a3-f28a-42e9-a706-4f71622ba768
Init duration	Duration
240.74 ms	74.82 ms
Billed duration	Resources configured
75 ms	512 MB
Max memory used	
59 MB	

Log output

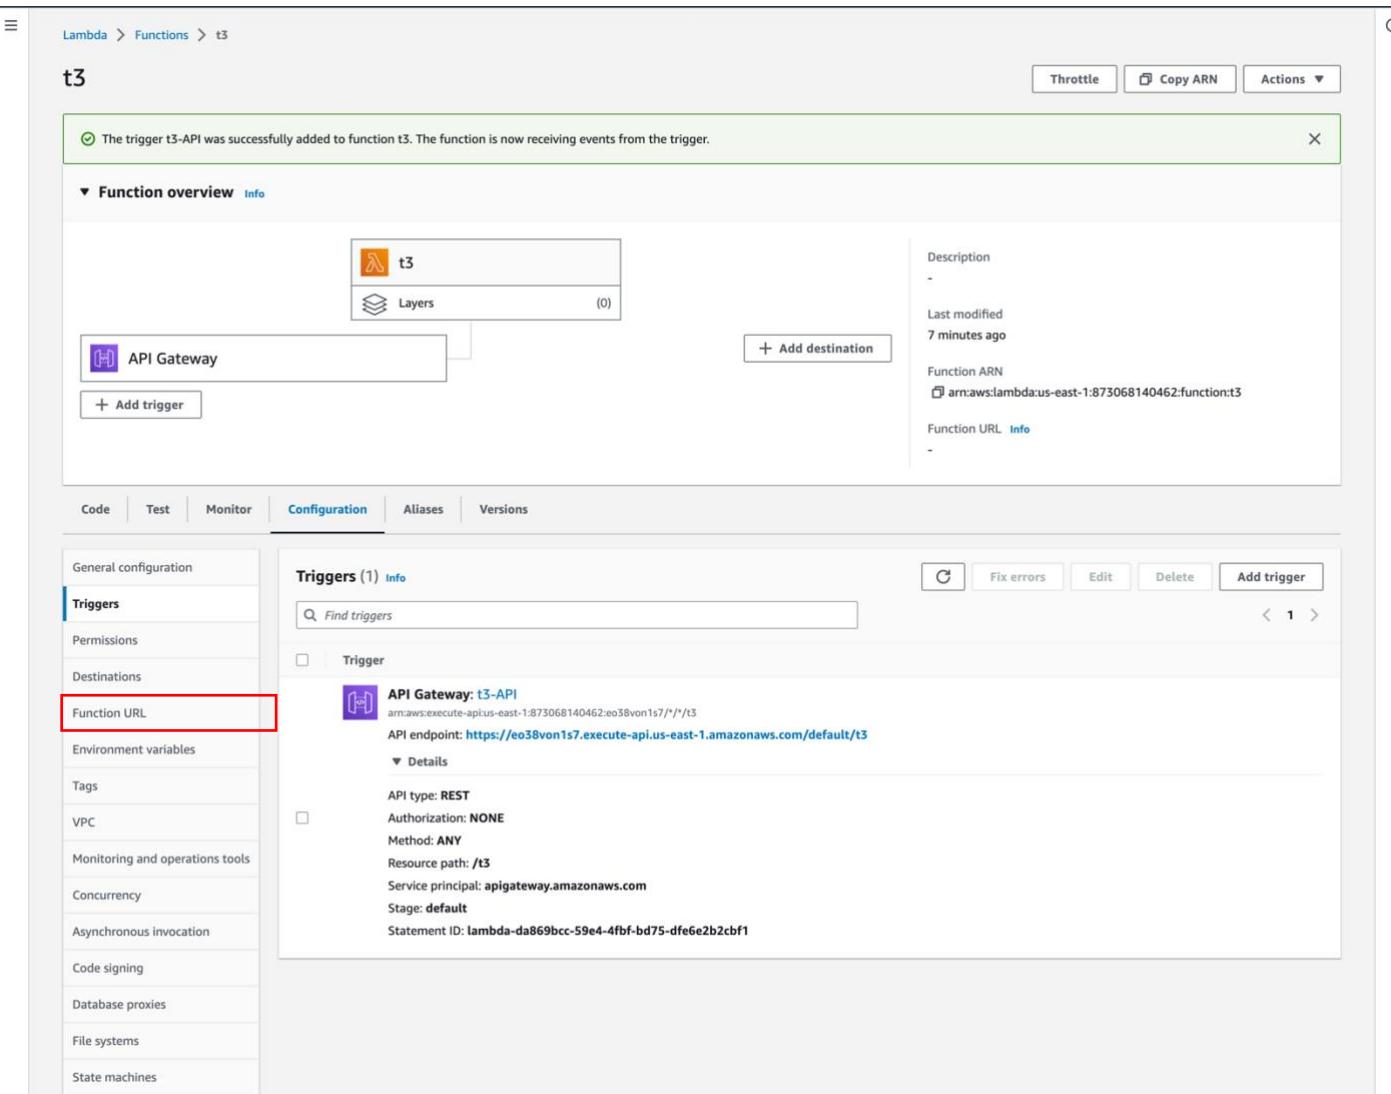
The section below shows the logging calls in your code. [Click here](#) to view the corresponding CloudWatch log group.

```
START RequestId: 5937a1a3-f28a-42e9-a706-4f71622ba768 Version: $LATEST  
input:{name:"Integração Empresarial"}  
END RequestId: 5937a1a3-f28a-42e9-a706-4f71622ba768  
REPORT RequestId: 5937a1a3-f28a-42e9-a706-4f71622ba768 Duration: 74.82 ms Billed Duration: 75 ms Memory Size: 512 MB Max Memory Used: 59 MB Init Duration: 240.74 ms
```

C. Invoking a Lambda function remotely

To allow a Lambda function to be invoked remotely the API Gateway need to be assigned to the trigger of the function

- D.1. Access your Lambda function in AWS console as described in the previous sections
- D.2. Press the Add trigger button: 
- D.3. In the trigger confirmation drop-down menu choose "API Gateway"
- D.4. Then, in the API drop-down menu choose "Create a new API"
- D.5. Choose REST API option
- D.6. Choose the following option from Security drop down box: Open
- D.7. Press Add button 
- D.8. You will obtain something similar with the following image:



The screenshot shows the AWS Lambda Functions interface for a function named 't3'. At the top, there's a success message: 'The trigger t3-API was successfully added to function t3. The function is now receiving events from the trigger.' Below this, the 'Function overview' section displays the function name 't3', its layers '(0)', and an 'API Gateway' trigger. The 'Configuration' tab is selected, showing the 'Triggers' section with one entry: 'API Gateway: t3-API'. This entry includes details like the ARN, API endpoint (<https://eo38von1s7.execute-api.us-east-1.amazonaws.com/default/t3>), and method details (Method: ANY, Resource path: /t3). A red box highlights the 'Function URL' link in the left sidebar.

Click in the "t3-API" link, or similar, as identified in the red box of the image, to select the AWS API Gateway service of your lambdaservice. A similar interface is obtained:

The screenshot shows the AWS API Gateway interface. On the left, the navigation pane is open with the 'APIs' section expanded, showing 'API: teste-API' with 'Resources' selected. The main content area displays the 'Resources' for the '/teste' endpoint. The URL path is shown as '/teste' with 'ANY' selected. To the right, a flow diagram illustrates the request handling process: Client → Method request → Integration request → Lambda integration. Below the diagram, tabs for 'Method request', 'Integration request', 'Integration response', 'Method response', and 'Test' are present, with 'Test' being the one highlighted by a red box. The 'Method request settings' section shows 'Authorization' set to 'NONE' and 'API key required' set to 'False'. At the top right, there are 'API actions' and 'Deploy API' buttons.

D.9. Select the TEST option as identified in the red box of the previous image and obtain the following:

This screenshot shows the 'Test method' configuration page for the '/teste' endpoint. It includes fields for 'Method type' (set to 'PUT'), 'Query strings' (containing 'param1=value1¶m2=value2'), 'Headers' (containing 'header1:value1' and 'header2:value2'), and a 'Client certificate' section. The 'Request body' field contains the JSON object: { "name": "IE" }. At the bottom, a large orange 'Test' button is highlighted with a red box.

D.10. Alternatively, create the file body.json with the following contents:

```
{
  "name": "Integracao Empresarial"
}
```

And execute the following command in another ssh or putty session (in linux like environment) where is available as in step C.8. (e.g.: <https://uwajykv2j.execute-api.us-east-1.amazonaws.com/default/PrimeiroMicroServico>):

```
curl -i -H "Content-Type: Application/json" --data "@body.json" -X POST <LAMBDA_URL>
```

It is expected that you receive something similar with:

```
$ curl -i -H "Content-Type: Application/json" --data "@body.json" -X POST https://2p3fp5acxj.execute-api.us-east-1.amazonaws.com/default/PrimeiroMicroServico

% Total    % Received % Xferd  Average Speed   Time      Time      Time  Current
          Dload  Upload   Total  Spent   Left  Speed
100    79  100    43  100     36     43    36  0:00:01  --:--:--  0:00:01   117
HTTP/1.1 200 OK
Date: Tue, 07 Jan 2020 22:26:07 GMT
Content-Type: application/json
Content-Length: 43
Connection: keep-alive
x-amzn-RequestId: a072be98-cefe-4847-8558-2360b6106e30
x-amz-apigw-id: F83L8FjKIAMFzrQ=
x-custom-header: my custom header value
X-Amzn-Trace-Id: Root=1-5e15057f-3a08bd2b9dd68b3d22208a89; Sampled=0

{"message":"Hello Integracao Empresarial!"}
```

The trace of the request can be verified in the “Monitor” tab of the Lambda function. Choosing the desired LogStream (left part of the following figure) will bring you to the right part of the figure with all the details of the invocation (inside AWS CloudWatch).

The left screenshot shows the AWS Lambda console. Under the 'HelloWorld' function, the 'Logs' tab is selected. It displays the CloudWatch Logs section with a table of recent invocations. One invocation is highlighted with a duration of 32.3 ms and a memory usage of 128 MB. The right screenshot shows the CloudWatch Log events page for the same invocation. It displays the log stream with detailed log entries, including the Lambda function code and environment variables.

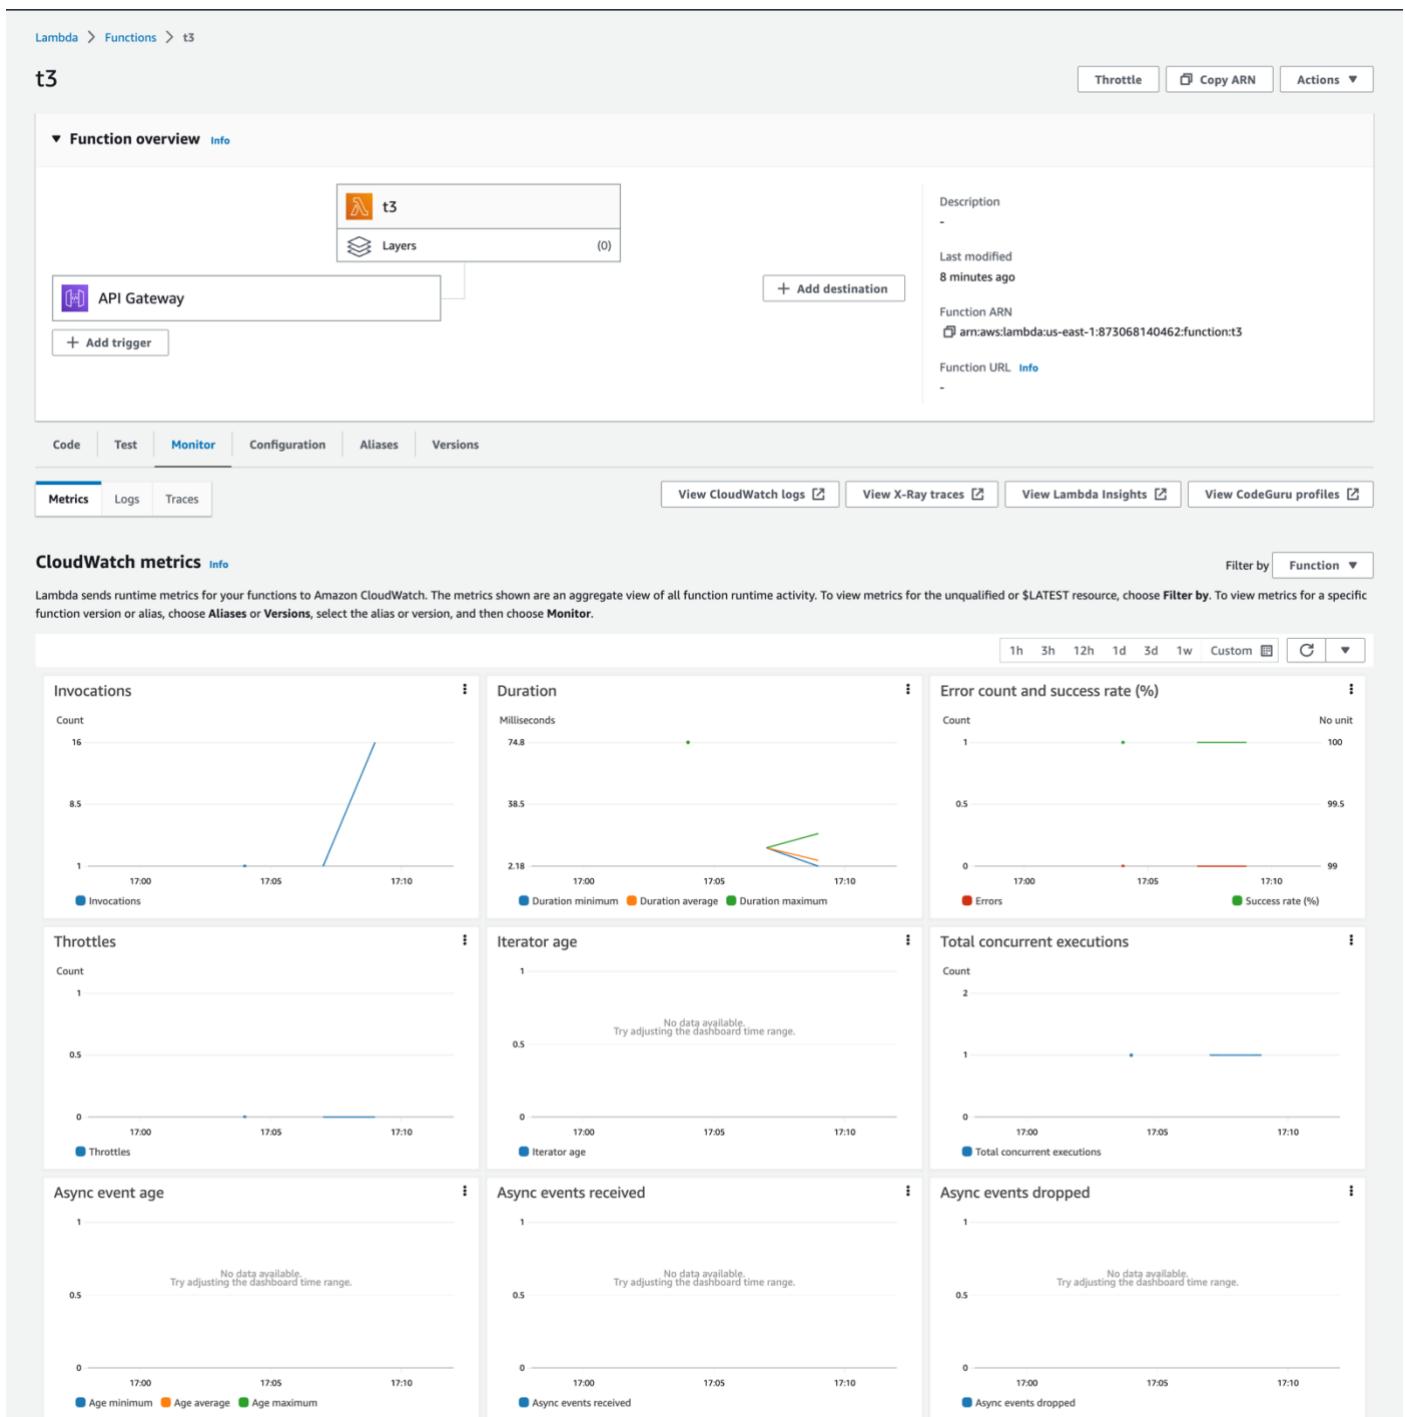
D. Delete a Lambda function

- E.1. Choose the Lambda > Functions options, then choose the function to delete and then chose Delete from Action drop-down menu.

Functions (1)					
Function name	Description	Runtime	Code size	Last modified	Actions
PrimeiraFuncao		Java 8	9.6 kB	14 minutes ago	View details Test Delete

E. Checking consumption of a Lambda function

F.1. Choose the Lambda > Dashboard to view all the detail regarding each of the invocations to each of the lambda service. For a specific Function, use the option Monitoring available in all Functions.



Other references

- AWS Lambda Concepts <https://docs.aws.amazon.com/lambda/latest/dg/welcome.html>, available in PDF.
- Amazon API Gateway Developer Guide
https://docs.aws.amazon.com/pt_br/apigateway/latest/developerguide/welcome.html, available in PDF.
- Set up an Integration Response in API Gateway,
<https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-integration-settings-integration-response.html>
- <https://www.mkyong.com/java/json-simple-example-read-and-write-json/>
- https://www.tutorialspoint.com/json_simple/json_simple_quick_guide.htm
- <https://www.baeldung.com/aws-lambda-api-gateway>
- <https://github.com/eugenp/tutorials/blob/master/aws-lambda/src/main/java/com/baeldung/lambda/apigateway/APIDemoHandler.java>
- Amazon Relational Database Service User Guide -
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/rds-ug.pdf#CHAP_GettingStarted

The goal of this document is to demonstrate how to install and operate the Kong open source service. Kong is a platform aiming the integration between all the micro services (λ and JAVA), and the BPMN processes developed in Camunda. Moreover, a Kong management platform, named Konga, is introduced.

The following contents is presented in this document.

Contents

A. Installing and starting the Kong Open Source Service as a docker in EC2	2
B. Configuring Kong to be used remotely	3
C. Creating a service and a route in Kong.....	4
D. Invoking a service	6
E. Connecting to a remote microservice.....	7
F. Other example of connecting to a JAVA microservice	8
G. Endpoints with parameters.....	9
H. Listing all the available services, routes and routes associated with services in Kong	10
I. Install Konga Open Source platform.....	12
Other references.....	14

A. Installing and starting the Kong Open Source Service as a docker in EC2

A.1.Create an EC2 new instance of the type: Amazon Linux 2 AMI (HVM), SSD Volume Type. The exact versions may change with time. And define the inbound rules to allow the 8080, 8002 and 22 accessed from anywhere.

A.2.Access the new EC2 instance and execute the following commands:

```
sudo yum install -y docker  
sudo service docker start  
sudo docker network create kong-net
```

```
sudo docker run -d --name kong-database \  
--network=kong-net \  
-p 5432:5432 \  
-e "POSTGRES_USER=kong" \  
-e "POSTGRES_DB=kong" \  
-e "POSTGRES_PASSWORD=kongpass" \  
postgres:13
```

```
sudo docker run --rm --network=kong-net \  
-e "KONG_DATABASE=postgres" \  
-e "KONG_PG_HOST=kong-database" \  
-e "KONG_PG_PASSWORD=kongpass" \  
-e "KONG_PASSWORD=test" \  
kong/kong-gateway:3.9.0.0 kong migrations bootstrap
```

```
sudo docker run -d --name kong-gateway \  
--network=kong-net \  
-e "KONG_DATABASE=postgres" \  
-e "KONG_PG_HOST=kong-database" \  
-e "KONG_PG_USER=kong" \  
-e "KONG_PG_PASSWORD=kongpass" \  
-e "KONG_PROXY_ACCESS_LOG=/dev/stdout" \  
-e "KONG_ADMIN_ACCESS_LOG=/dev/stdout" \  
-e "KONG_PROXY_ERROR_LOG=/dev/stderr" \  
-e "KONG_ADMIN_ERROR_LOG=/dev/stderr" \  
-e "KONG_ADMIN_LISTEN=0.0.0.0:8001, 0.0.0.0:8444 ssl" \  
-e "KONG_ADMIN_GUI_URL=http://localhost:8002" \  
-e KONG_LICENSE_DATA \  
-p 8000:8000 \  
-p 8443:8443 \  
-p 8001:8001 \  
-p 8002:8002 \  
-p 8445:8445 \  
-p 8003:8003 \  
-p 8004:8004 \  
-p 127.0.0.1:8444:8444 \  
kong/kong-gateway:3.9.0.0
```

A.3.Verify your installation in EC2 instance:

```
curl -i -X GET --url http://localhost:8001/services
```

You should receive a 200 status code.

B. Configuring Kong to be used remotely

B.1. Choose your EC2 instance. Click on security-group (for example):

And change the configuration to allow all the inbound traffic:

B.2. Test with the following command from your local computer:

```
curl -i http://<YOURIP>:8000/
```

You should receive a similar response with the following:

```
% Total    % Received % Xferd  Average Speed   Time     Time      Time  Current
          Dload  Upload   Total Spent  Left Speed
100  48  100  48    0      0  48      0  0:00:01 --:--:-- 0:00:01  192
HTTP/1.1 404 Not Found
Date: Tue, 14 Jan 2020 15:34:55 GMT
Content-Type: application/json; charset=utf-8
Connection: keep-alive
Content-Length: 48
X-Kong-Response-Latency: 0
Server: kong/2.0.0rc2

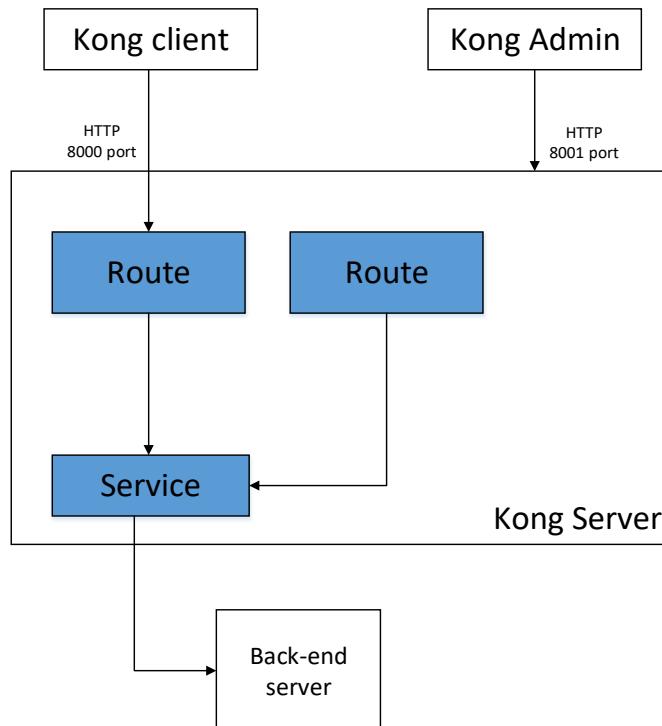
{"message": "no Route matched with those values"}
```

C. Creating a service and a route in Kong

From Kong definitions:

- a Service; that is the name Kong uses to refer to the upstream APIs and microservices it manages.
- and a Route specify how (and if) requests are sent to their Services after they reach Kong. A single Service can have many Routes.

In Kong the ports are organized as follows: Port 8001 – is used for managing APIs and Port 8000 – for service invocation. This is depicted by the following figure:



To try-out use the following commands:

- C.1. In your Kong AWS EC2 instance (for security reasons is safer to do it locally) issue the following cURL request to add your first Service (pointing to the Mockbin API) to Kong:

```
curl -i -X POST \
--url http://localhost:8001/services/ \
--data 'name=example-service' \
--data 'url=http://mockbin.org'
```

You should receive a response similar to:

```
HTTP/1.1 201 Created
Content-Type: application/json
Connection: keep-alive

{
  "host": "mockbin.org",
  "created_at": 1519130509,
  "connect_timeout": 60000,
  "id": "92956672-f5ea-4e9a-b096-667bf55bc40c",
  "protocol": "http",
  "name": "example-service",
  "read_timeout": 60000,
  "port": 80,
  "path": null,
  "updated_at": 1519130509,
  "retries": 5,
  "write_timeout": 60000
}
```

C.2. Then, add a Route for the previous Service:

```
curl -i -X POST \
--url http://localhost:8001/services/example-service/routes \
--data 'hosts[]="example.com"
```

The answer should be similar to:

```
HTTP/1.1 201 Created
Content-Type: application/json
Connection: keep-alive

{
    "created_at":1519131139,
    "strip_path":true,
    "hosts":[
        "example.com"
    ],
    "preserve_host":false,
    "regex_priority":0,
    "updated_at":1519131139,
    "paths":null,
    "service":{
        "id":"79d7ee6e-9fc7-4b95-aa3b-61d2e17e7516"
    },
    "methods":null,
    "protocols":[
        "http",
        "https"
    ],
    "id":"f9ce2ed7-c06e-4e16-bd5d-3a82daef3f9d"
}
```

The service is now ready to be invoked. Continue to the next section of this document.

D. Invoking a service

Issue the following cURL request to verify that Kong is properly forwarding requests to your Service. Note that by default Kong handles proxy requests on port :8000. The following request indicates which is the route that you would like to call:

```
$ curl -i -X GET \  
--url http://localhost:8000/ \  
--header 'Host: example.com'
```

A successful response means Kong is now forwarding requests made to http://localhost:8000 to the url we configured in previous section and is forwarding the response back to us. Kong knows to do this through the header defined in the above curl request:

Host: <given host>

hint: If you are testing the invocation remotely from the AWS EC2 Kong instance, then you should replace the localhost by the AWS EC2 DNS name.

E. Connecting to a remote microservice

E.1. DIRECTLY TO LAMBDA BACKEND (as previously introduced in correspondingly tutorial):

```
curl -i -H "Content-Type: Application/json"\n--data "@body.json"\n-X POST https://2p3fp5acxj.execute-api.us-east-1.amazonaws.com/default/PrimeiroMicroServico
```

E.2. CREATE SERVICE (in the AWS EC2 Kong instance):

```
curl -i -X POST \
--url http://localhost:8001/services/ \
--data 'name=invoke-lambda-service' \
--data 'url=https://2p3fp5acxj.execute-api.us-east-1.amazonaws.com/default/PrimeiroMicroServico'\
```

E.3. CREATE ROUTE (in the AWS EC2 Kong instance):

```
curl -i -X POST \
--url http://localhost:8001/services/invoke-lambda-service/routes \
--data 'hosts[] = disciplinas.com'
```

E.4. INVOKE REMOTELY (from your local computer), considering that your JSON file should be available

```
curl -i -X POST \
-H "Content-Type: Application/json" \
--url http://ec2-18-212-198-84.compute-1.amazonaws.com:8000/ \
--header 'Host: disciplinas.com' \
--data '@body.json'
```

You are expected to receive something similar with:

```
% Total    % Received % Xferd  Average Speed   Time     Time      Time  Current\n          Dload  Upload   Total Spent  Left Speed\n100  79  100    43  100     36      43     36  0:00:01  --::--  0:00:01  220\nHTTP/1.1 200 OK\nContent-Type: application/json\nContent-Length: 43\nConnection: keep-alive\nDate: Tue, 14 Jan 2020 17:16:50 GMT\nx-amzn-RequestId: 39bf3f69-af84-4568-96fb-b0277ea9e896\nx-amz-apigw-id: GTOceEKUoAMFc4g=\nx-custom-header: my custom header value\nX-Amzn-Trace-Id: Root=1-5e1df782-21cac32327af61b6064ed89a; Sampled=0\nX-Kong-Upstream-Latency: 28\nX-Kong-Proxy-Latency: 0\nVia: kong/2.0.0rc2\n\n{"message": "Hello Integracao Empresarial!"}
```

F. Other example of connecting to a JAVA microservice

F.1. DIRECTLY TO JAVA WEBSERVICE BACKEND (as previous introduced in correspondingly tutorial):

```
curl -i -X POST -H "Content-Type: text/xml;charset=UTF-8" \
--url http://ec2-52-91-101-30.compute-1.amazonaws.com:9298/calculate \
--data '@calculate.xml'
```

F.2. CREATE SERVICE (in the AWS EC2 Kong instance):

```
curl -i -X POST \
--url http://localhost:8001/services/ \
--data 'name=invoke-java-service' \
--data 'url=http://ec2-52-91-101-30.compute-1.amazonaws.com:9298/calculate'\
```

F.3. CREATE ROUTE (in the AWS EC2 Kong instance):

```
curl -i -X POST \
--url http://localhost:8001/services/invoke-java-service/routes \
--data 'hosts[] = calculate.com'
```

F.4. INVOKE REMOTELY (from your local computer), considering that your XML file should be available

```
curl -i -X POST -H "Content-Type: text/xml;charset=UTF-8" \
--url ec2-3-86-244-114.compute-1.amazonaws.com:8000/ \
--header 'Host: calculate.com' \
--data '@calculate.xml'
```

You are expected to receive something similar with:

```
% Total    % Received % Xferd  Average Speed   Time     Time      Time  Current
          Dload  Upload   Total  Spent   Left  Speed
100  461     0    202  100    259     202    259  0:00:01 --::--  0:00:01  1054
HTTP/1.1 200 OK
Content-Type: text/xml; charset=utf-8
Transfer-Encoding: chunked
Connection: keep-alive
Date: Tue, 14 Jan 2020 21:44:18 GMT
X-Kong-Upstream-Latency: 8
X-Kong-Proxy-Latency: 3
Via: kong/2.0.0rc2

<?xml version="1.0" ?><S:Envelope
xmlns:S="http://schemas.xmlsoap.org/soap/envelope/"><S:Body><ns2:mulResponse
xmlns:ns2="http://Webservice/"><return>182</return></ns2:mulResponse></S:Body></S:Envelope>
```

G. Endpoints with parameters

G.1.To allow the access to endpoints with parameters, e.g.,

```
http://URL_OF_ENDPOINT:8080/fruits/2/bananas
```

where the exact command if you invoke it by curl is:

```
curl -X 'PUT' \
'http://URL_OF_ENDPOINT:8080/fruits/2/bananas' \
-H 'accept: application/json'
```

You need to create a service and route accordingly:

G.2.The Kong service need to be defined only with the base path without the service identification, for instance:

```
curl -i -X POST \
--url "http://URL_OF_KONG:8001/services/" \
--data "name=invoke-quarkus4-service" \
--data "url=http://URL_OF_ENDPOINT:8080/"
```

G.3.After that the Kong route need to be defined using a regular expression, containing the fixed part and the variables with type indication, for instance:

```
curl -i -X POST \
--url "http://URL_OF_KONG:8001/services/invoke-quarkus4-service/routes" \
--data-urlencode "paths[0]=~/fruits/(?<id>\d+)/(?<name>\S+)" \
--data "strip_path=false"
```

id and name are stored for later use by any Kong plug-in. d+ and S+ means, one or more decimal, and one or more alphanumeric. It is here that you customize the path accordingly with your endpoint application.

G.4.After that you can use the Kong route to invoke your endpoint with parameters.

```
curl -v -i -X PUT -H "accept: application/json" \
--url "http://URL_OF_KONG:8000/fruits/2/bananas"
```

H. Listing all the available services, routes and routes associated with services in Kong

Assuming the example in Section C - Creating a service and a route in Kong, the following command retrieved the available services:

```
curl -i http://localhost:8001/services
```

the answer should be similar to:

```
HTTP/1.1 200 OK
Date: Tue, 14 Jan 2020 21:46:40 GMT
Content-Type: application/json; charset=utf-8
Connection: keep-alive
Access-Control-Allow-Origin: *
Server: kong/2.0.0rc2
Content-Length: 1029
X-Kong-Admin-Latency: 1

{"next":null,"data":[{"host":"ec2-52-91-101-30.compute-1.amazonaws.com","created_at":1579036862,"connect_timeout":60000,"id":"2158a1c8-4ae4-48e4-a312-e464b3010104","protocol":"http","name":"invoke-java-service","read_timeout":60000,"port":9298,"path":"/calculate","updated_at":1579036862,"retries":5,"write_timeout":60000,"tags":null,"client_certificate":null}, {"host":"2p3fp5acxj.execute-api.us-east-1.amazonaws.com","created_at":1579033755,"connect_timeout":60000,"id":"783cfce7-9142-4222-9ee0-a2b50a507cd3","protocol":"https","name":"invoke-lambda-service","read_timeout":60000,"port":443,"path":"/default/PrimeiroMicroServico","updated_at":1579033755,"retries":5,"write_timeout":60000,"tags":null,"client_certificate":null}, {"host":"mockbin.org","created_at":1579033726,"connect_timeout":60000,"id":"e9f67cc4-fale-4a23-ade3-d5e692b7368b","protocol":"http","name":"example-service","read_timeout":60000,"port":80,"path":null,"updated_at":1579033726,"retries":5,"write_timeout":60000,"tags":null,"client_certificate":null}]}{}
```

Then, the following command retrieved the available routes:

```
curl -i http://localhost:8001/routes
```

and the answer should be similar to:

```
HTTP/1.1 200 OK
Date: Tue, 14 Jan 2020 21:47:24 GMT
Content-Type: application/json; charset=utf-8
Connection: keep-alive
Access-Control-Allow-Origin: *
Server: kong/2.0.0rc2
Content-Length: 1318
X-Kong-Admin-Latency: 1

{"next":null,"data":[{"id":"38cd55c6-adea-4fc6-86b9-19ca0749e4a3","path_handling":"v0","paths":null,"destinations":null,"headers":null,"protocols":["http","https"]}, {"methods":null,"snis":null,"service":{"id":"783cfce7-9142-4222-9ee0-a2b50a507cd3"}, "name":null,"strip_path":true,"preserve_host":false,"regex_priority":0,"updated_at":1579033770}, {"sources":null,"hosts":["disciplinas.com"], "https_redirect_status_code":426,"tags":null,"created_at":1579033770}, {"id":"8d188dde-c65a-402b-b659-d8c39054087f", "path_handling":"v0", "paths":null, "destinations":null, "headers":null, "protocols":["http","https"]}, {"methods":null,"snis":null,"service":{"id":"2158a1c8-4ae4-48e4-a312-e464b3010104"}, "name":null,"strip_path":true,"preserve_host":false,"regex_priority":0,"updated_at":1579036920}, {"sources":null,"hosts":["calculate.com"], "https_redirect_status_code":426,"tags":null,"created_at":1579036920}, {"id":"cca09713-5cdf-4db1-8ab0-606e2bd76d1c", "path_handling":"v0", "paths":null, "destinations":null, "headers":null, "protocols":["http","https"]}, {"methods":null,"snis":null,"service":{"id":"e9f67cc4-fale-4a23-ade3-d5e692b7368b"}, "name":null,"strip_path":true,"preserve_host":false,"regex_priority":0,"updated_at":1579033736}, {"sources":null,"hosts":["example.com"], "https_redirect_status_code":426,"tags":null,"created_at":1579033736}]}{}
```

Finally, the following command retrieved the available routes for a specific service:

```
curl -i http://localhost:8001/services/example-service/routes
```

and the answer should be similar to:

```
HTTP/1.1 200 OK
Date: Tue, 14 Jan 2020 16:28:47 GMT
Content-Type: application/json; charset=utf-8
Connection: keep-alive
Access-Control-Allow-Origin: *
Server: kong/2.0.0rc2
Content-Length: 452
```

X-Kong-Admin-Latency: 2

```
{"next":null,"data":[{"id":"6b7df64c-9ed9-4d5c-87a3-58b9935d01f4","path_handling":"v0","paths":null,"destinations":null,"headers":null,"protocols":["http","https"],"methods":null,"snis":null,"service":{"id":"af50ee81-66e3-4dbe-b2f6-f25e72c549a1"}, "name":null,"strip_path":true,"preserve_host":false,"regex_priority":0,"updated_at":1579017405,"sources":null,"hosts":[ "example.com"], "https_redirect_status_code":426,"tags":null,"created_at":1579017405}]}{}
```

I. Install Konga Open Source platform

Optionally, for administration purposes of Kong you can install Konga, which is an unofficial app. No affiliated with Kong. It offer an easy management of all the Kong concepts.

- I.1. Pull the following Docker image locally on your PC:

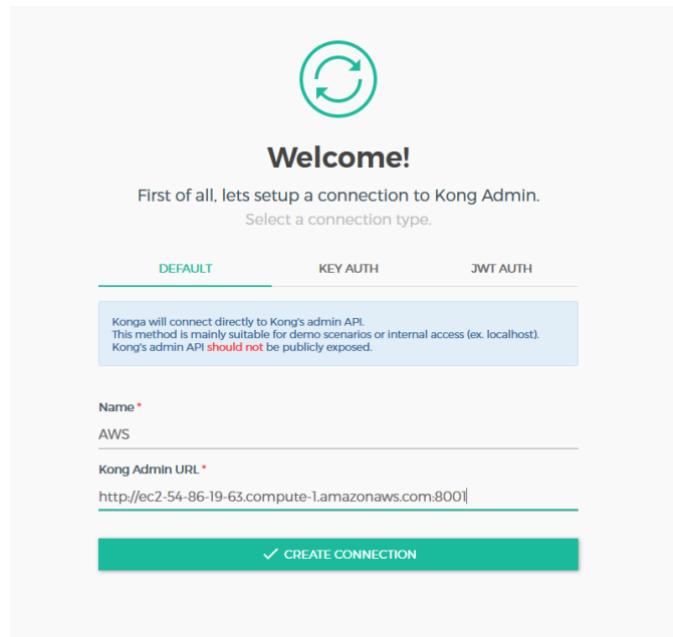
```
docker pull pantsel/konga
```

- I.2. Then, execute it:

```
docker run -d --name konga -p 1337:1337 pantsel/konga
```

- I.3. Open your browser with the following url: http://<ec2_DNS_NAME>:1337 and create an admin account (keep the password!)

- I.4. Then, provide the Kong URL with the 8001 port reference in your browser, accordingly with the following example:



A similar result is obtained, where you can manage the Kong concepts:

localhost:1337#!/dashboard

KONGA

DASHBOARD

API GATEWAY

- INFO
- SERVICES
- ROUTES
- CONSUMERS
- PLUGINS
- UPSTREAMS
- CERTIFICATES

APPLICATION

- USERS
- CONNECTIONS
- SNAPSHOTS
- SETTINGS

CONNECTIONS

ACTIVE	READING	WRITING	WAITING	ACCEPTED	HANDED
2	0	1	1	72	72

Total Requests: 71

NODE INFO

HostName	ip-172-31-84-175.ec2.internal
Tag Line	Welcome to kong
Version	2.0.0rc2
LUA Version	LuaJIT 2.1.0-beta3
Admin listen	[{"0.0.0.0:8001 reuseport backlog=16384"}, {"127.0.0.1:8444 http2 ssl reuseport backlog=16384"}]

TIMERS

DATASTORE INFO

DBMS	postgres
Host	127.0.0.1
Database	kong
User	kong
Port	5432

PLUGINS

reachable

consulation_id pre-functions cors klap-auth loggy hmac.auth zipkin request-size-limiting azure-functions request-transformer oauth2 response-transformer ip-restriction stated jwt proxy-cache basic-auth key-auths http-log datalog tcp-log post-function prometheus acl syslog file-log session udp-log response-size-limiting aws-lambda bot-detection rate-limiting request-termination

KONGA 0.14.7 GitHub Issues Support the project Connected to AWS

Other references

- Kong Inc. <https://konghq.com/>
- https://docs.konghq.com/install/redhat/?_ga=2.90626585.229207644.1612972455-1733373400.1609865717
- Konga <https://pantsel.github.io/konga/>
- https://wiki.postgresql.org/wiki/YUM_Installation
- <https://www.postgresql.org/download/linux/redhat/>
- <https://www.cyberciti.biz/faq/install-and-setup-postgresql-on-rhel-8/>
- <https://docs.konghq.com/1.4.x/admin-api/>
- For deploying Auth2.0 plug/in in Kong read the manual: <https://docs.konghq.com/hub/kong-inc/oauth2/>
- <https://docs.konghq.com/gateway/2.8.x/install-and-run/docker/>